



Architect's Open-Source Guide for a Data Mesh Architecture

Lena Hall
Microsoft

Lena Hall



Director at Microsoft
Azure Engineering

- ✓ Architecture
- ✓ Cloud
- ✓ Data
- ✓ ML/AI



lenadroid



Entry Point

How to Move Beyond a Monolithic Data Lake to a Distributed Data Mesh

<https://martinfowler.com/articles/data-monolith-to-mesh.html>

Data Mesh Principles and Logical Architecture

<https://martinfowler.com/articles/data-mesh-principles.html>

Slack for Data-Mesh-Learning

<https://launchpass.com/data-mesh-learning>

Talk Snapshot

- What is Data Mesh
- When is Data Mesh a Good Idea
- Core Principles and Concepts
- Example: Drone Delivery Service
- Challenges
- OSS and Open Standards



When and Why **Data Mesh**

The background is a solid red color. On the right side, there are several overlapping circles of varying shades of red, creating a layered, abstract effect. The circles are semi-transparent, allowing the underlying red to show through.

Data Mesh is Not For Everyone



Challenges Indicating **Data Mesh** May Be Considered

Drone Delivery Service



WHYs

- Ambiguity in Ownership and Responsibility
- Slow Change due to Coupling to Monolithic System
- Data Engineering Resources Bottleneck



Ideas Composing **Data Mesh** Concept

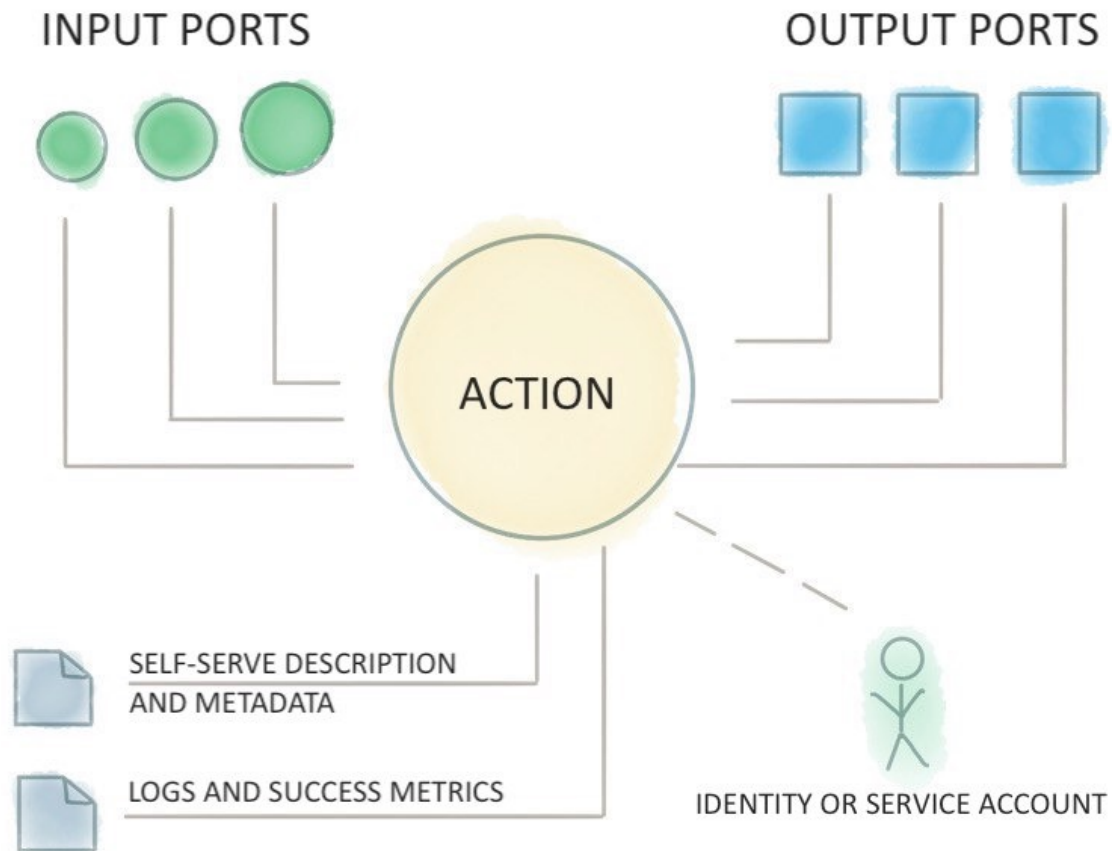
Core Ideas

- ✓ Decentralized teams and data ownership

Core Ideas

- ✓ Decentralized teams and data ownership
- ✓ Data Products powered by Domain Driven Design

High-Level View of a Data Product



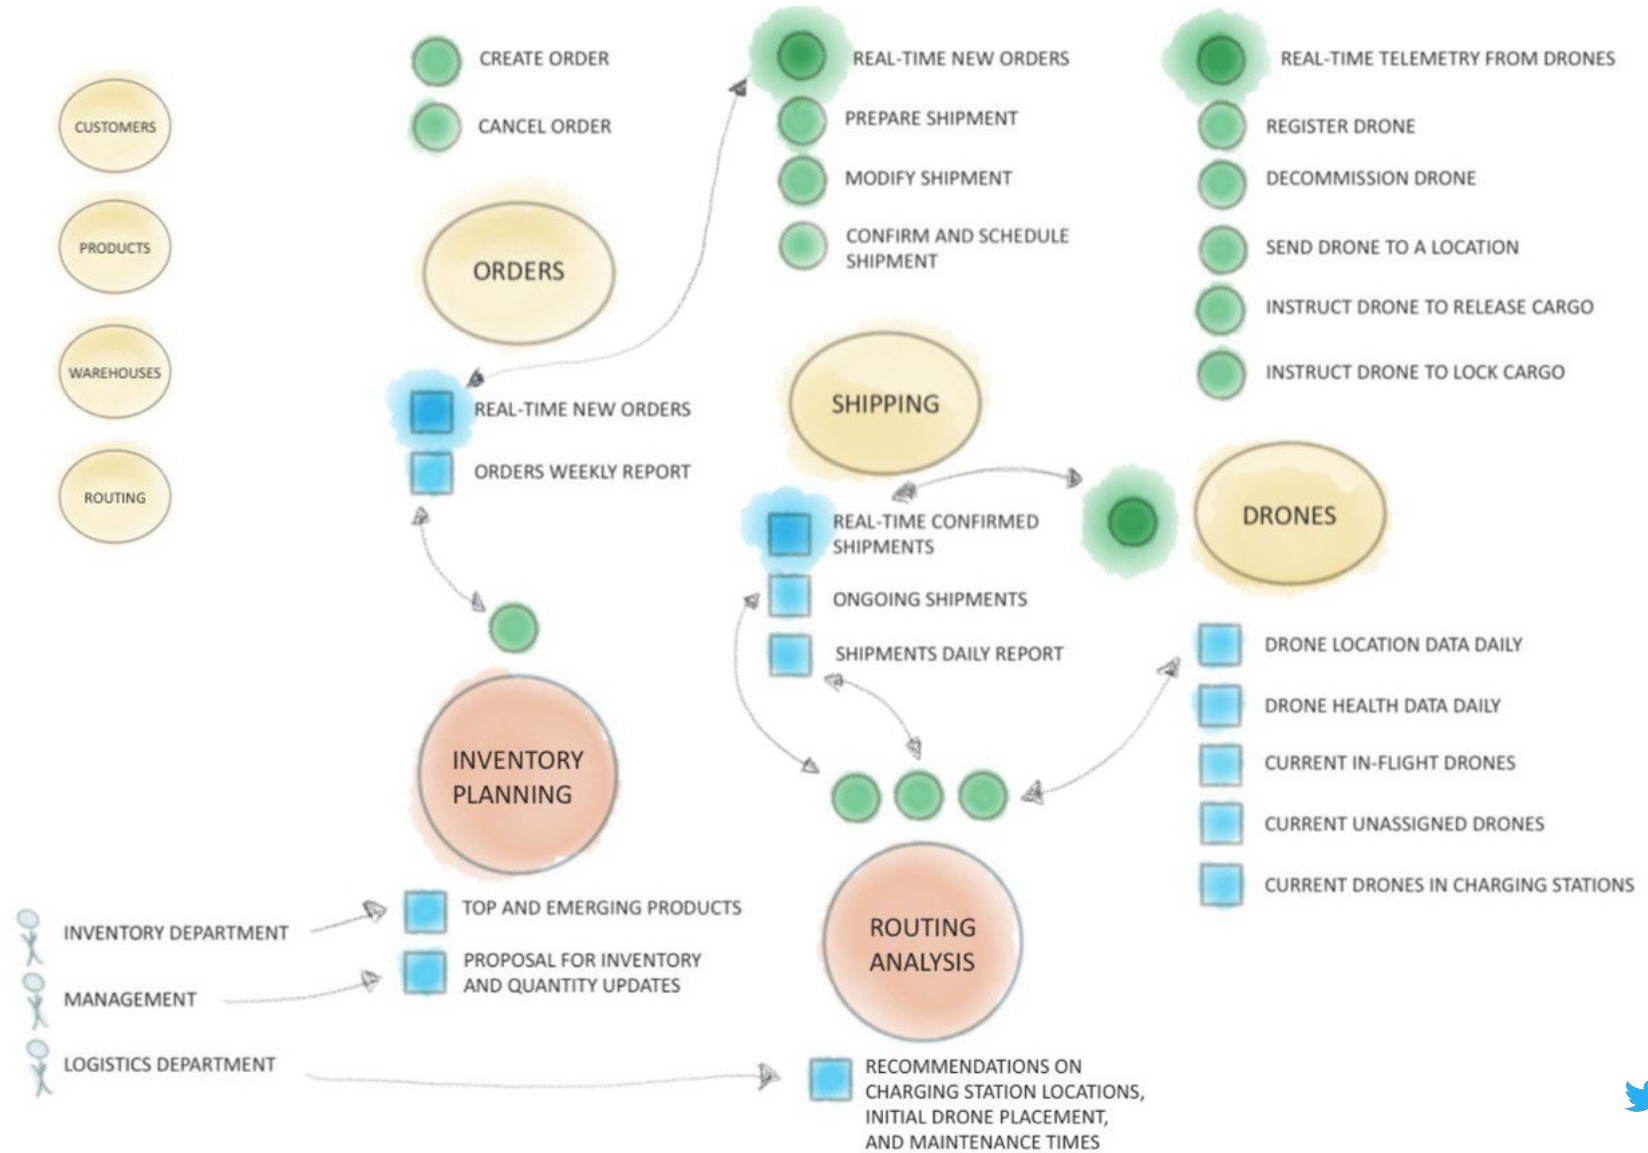
Core Ideas

- ✓ Decentralized teams and data ownership
- ✓ Data Products powered by Domain Driven Design
- ✓ Self-serve Shared Data Infrastructure

Core Ideas

- ✓ Decentralized teams and data ownership
- ✓ Data Products powered by Domain Driven Design
- ✓ Self-serve Shared Data Infrastructure
- ✓ Global Federated Governance

Drone Delivery Service **Data Products**



Core Principles for **Data Products**

Core Principles for Data Products




DISCOVERABLE

Core Principles for Data Products



DISCOVERABLE




SELF-DESCRIBING

Core Principles for Data Products



DISCOVERABLE



SELF-DESCRIBING



ADDRESSABLE


Core Principles for Data Products



DISCOVERABLE



SECURE



SELF-DESCRIBING



ADDRESSABLE


Core Principles for Data Products



DISCOVERABLE



SECURE



SELF-DESCRIBING



TRUSTWORTHY



ADDRESSABLE


Core Principles for Data Products



DISCOVERABLE



SECURE




SELF-DESCRIBING



TRUSTWORTHY



ADDRESSABLE



INTEROPERABLE

Cheat Sheet for Planning Data Products

Input Ports Questions

- **Data Source** - Where is the data coming from? External dataset or another data product?
- **Data Format** - What is the format of the source input?
- **Rate of Updates** - How frequently does the input need to be updated?

Output Ports Questions

- **End-consumers** - Who are the end-users of the data product?
- **Data purpose** - What are they planning to do with the data outputs?
- **Data access** - Who needs to have access? How do they prefer to access the data output?
- **Data address** - How do they prefer to access the data output?
- **Data Format** - What format of the data do they expect?

Identity and Permission Policies Questions

- Which resources can this data product be allowed to access?
- Which data products or users can read which output ports of this data product?
- Are all sensitive resources this data product offers protected according their required privacy standards (e.g. HIPAA, GDPR, PII, CCPA, etc.)
- Is the permissions policy stored and managed in the same package as the data product?

Data Product Action Questions

- What is the action that needs to happen to produce the outcomes for the end-users?
- What are the required adjustments, transformations, filters, updates, or quality improvements to the input data?

Operational Questions

- How can this data product be discovered and how should it be described to other data products that might want to consume it?
- Which metadata and information should it make available to the end-users?
- Where and how should data product versioning be managed during updates to ensure consistency with how the end-users consume it?
- Which SLAs or SLOs does the data product provide?
- Which product success metrics can this data product expose and keep track of? (adoption, usage, quality)
- Is the automation/resource orchestration logic stored in the same package?

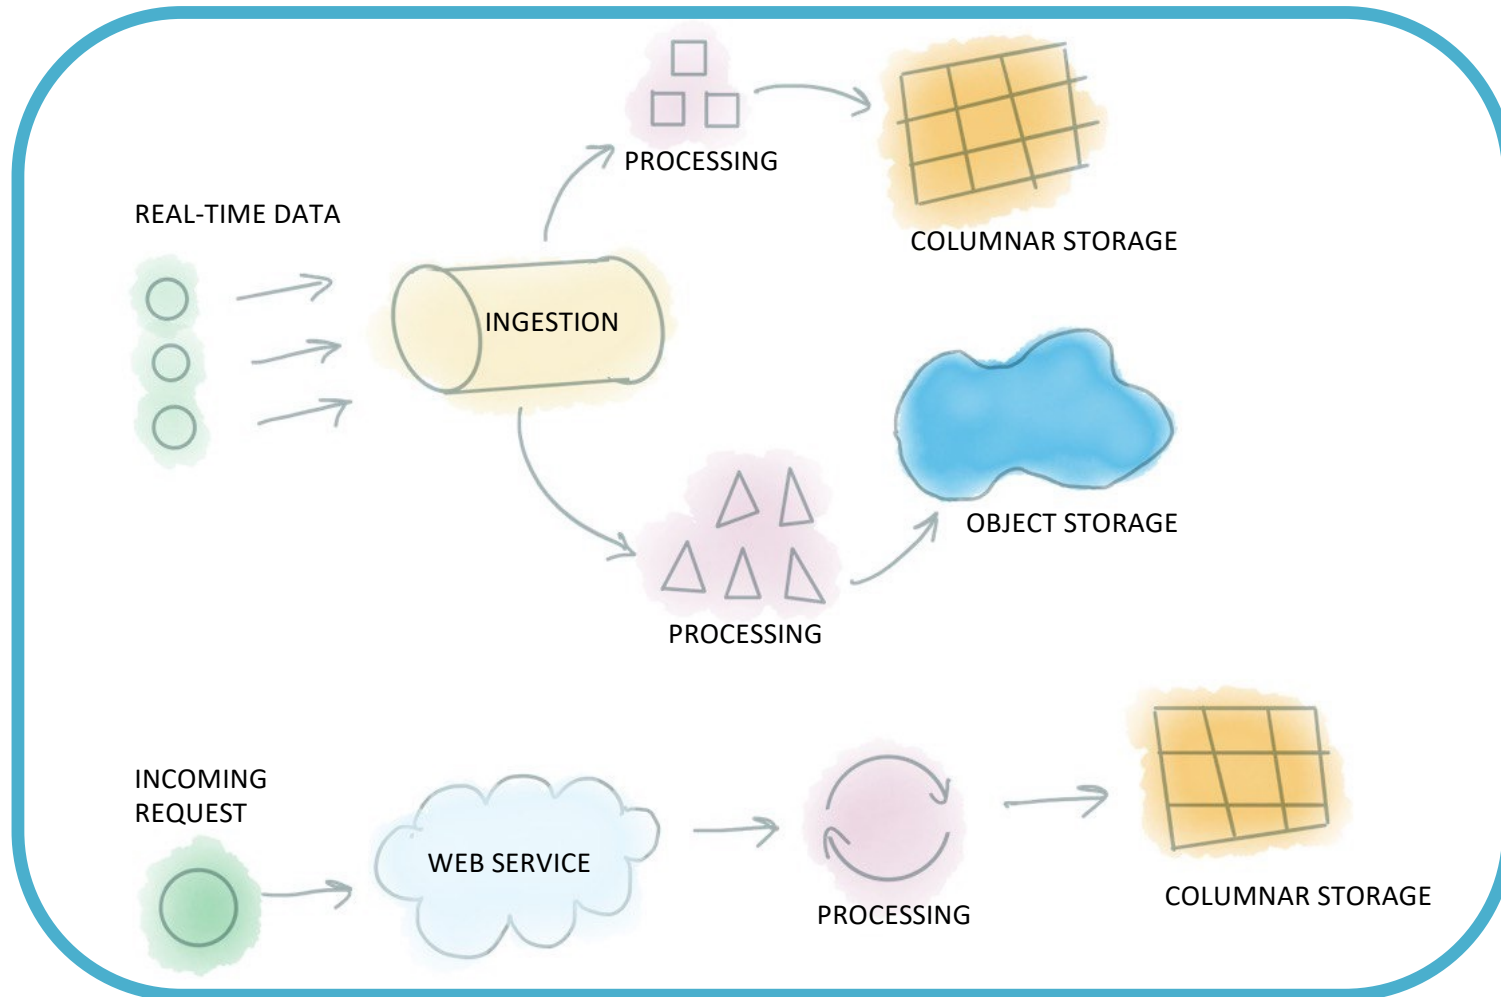
Other Questions

- Is this product not tightly coupled to any other data source, data product, or any other resource that makes him not interoperable?
- Does this data product follow the defined global governance standards and practices defined by the organization?
- Does this data product have any implementation details that could interfere with its portability?

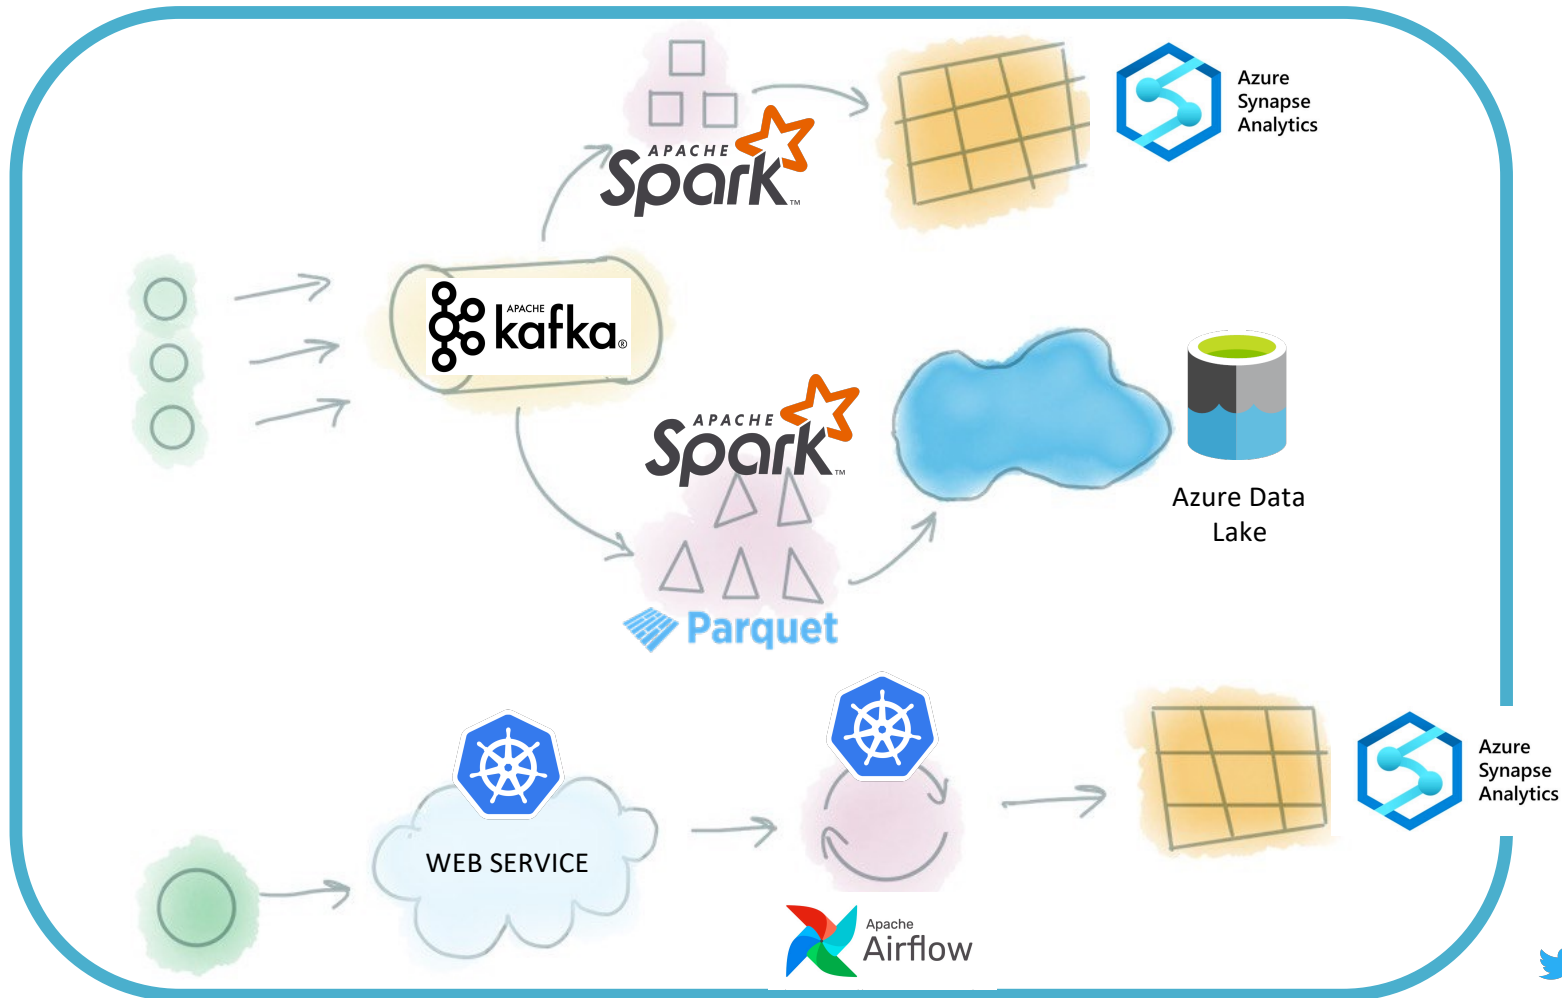


Self-Serve Shared Infrastructure

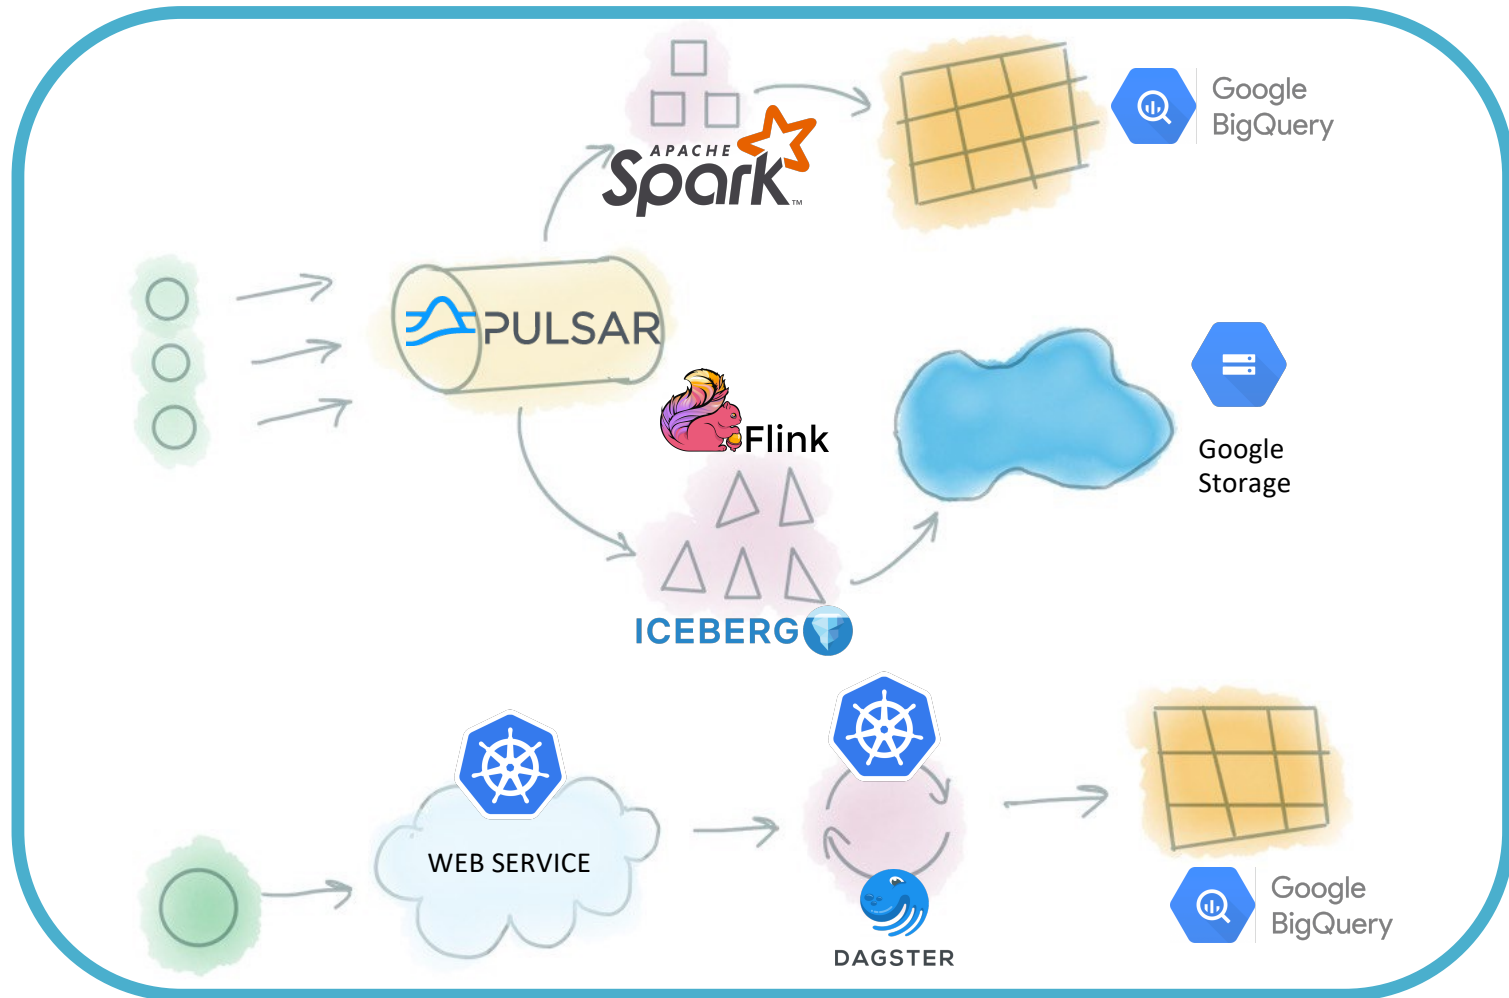
Types of Workloads Within a Data Product



It can look like this



Or, it can look like this




Self-Serve Shared Infrastructure


SHARED PLATFORM FOR
STREAMING INGESTION




SHARED PLATFORM FOR
RAW DATA STORAGE



SHARED PLATFORM FOR
COLUMNAR DATA STORAGE



DATA CATALOGUE



SHARED PLATFORM FOR
CONTAINER WORKLOADS



SHARED PLATFORM FOR
CONTINUOUS DELIVERY



SHARED PLATFORM
FOR OBSERVABILITY



AND MORE...

DEPENDING ON THE ORGANIZATION

Data Mesh



DISCOVERABLE



SELF-DESCRIBING




ADDRESSABLE



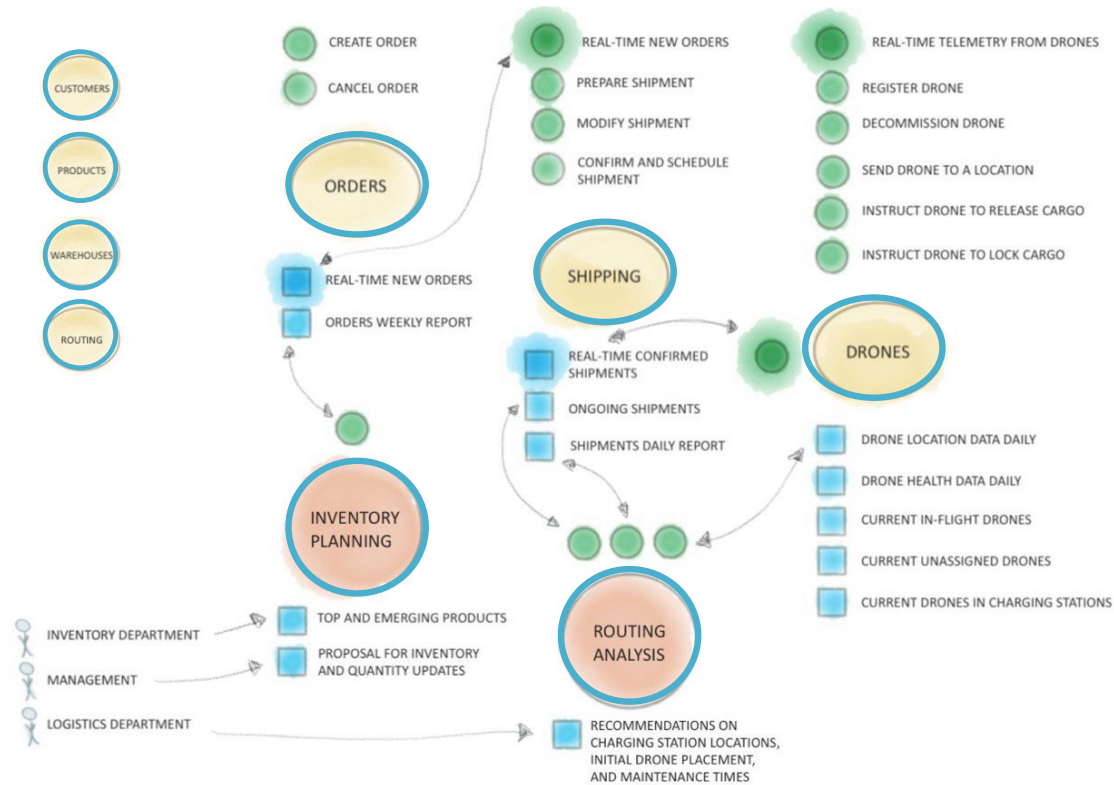
SECURE



TRUSTWORTHY




INTEROPERABLE




SHARED PLATFORM FOR STREAMING INGESTION



SHARED PLATFORM FOR RAW DATA STORAGE



SHARED PLATFORM FOR COLUMNAR DATA STORAGE



SHARED PLATFORM FOR CONTAINER WORKLOADS



SHARED PLATFORM FOR CONTINUOUS DELIVERY



SHARED PLATFORM FOR OBSERVABILITY



DATA CATALOGUE



Wait, What About the OSS Tools for Data Mesh??

Challenges with Data Mesh

Challenges

- Cost questions
- Lack of end-to-end examples
- Efforts to shift from centralized architecture to decentralization-friendly techniques
- Automation required for enabling creating data products
- Underestimating the importance organizational aspects


























































Considerations for **Technology** Choices

Considerations for Technology Choices

- Workload sharing and multi-tenancy
- No-copy data and compute mobility support
- Granularity of access-control
- Richness of automation and extensibility capabilities
- Flexibility and elasticity
- Provider-agnostic/multi-cloud operations support
- Variety of limitations
(quotas, data volume, resource count, etc.)
- **Open Standards, Open Protocols, Open-Source Integrations**



Examples of Data Mesh-friendly Technologies

Data Catalogue, Data Lineage, Data Governance	OSS Data Analytics, Data Processing, Data Querying	Data Ingestion, Streaming	Products for Data Analytics and Processing	Data Visualization and BI Tools	Cross-Platform Concepts and Tools
       	       	     Data Orchestration, Workflows   	     	    	    
Open Formats	Cloud Storage	OSS Storage	Infrastructure Automation	Data Experimentation	Multi and Hybrid Cloud Tools
   	  	 	  	 	  

Data Governance Systems

- Metadata
- Data lineage
- Data schemas
- Data relationships
- Data classification
- Data security
- Data catalog



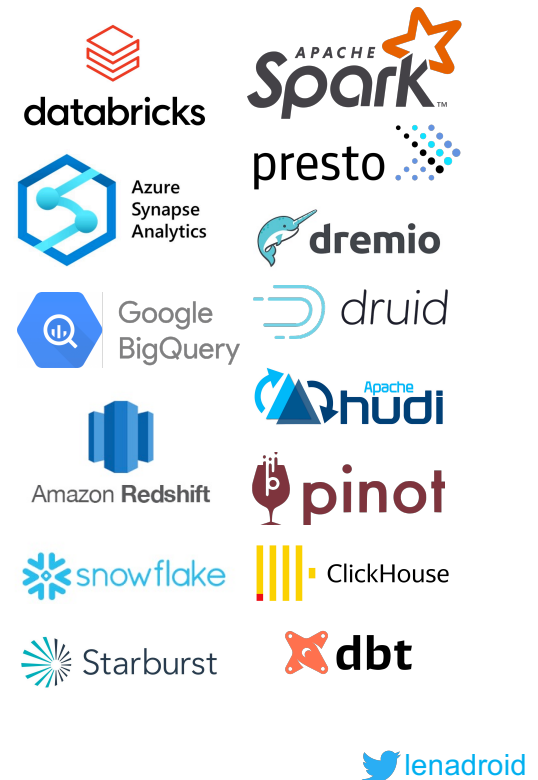
Open Formats

- Open standard
- Atomic updates, serializable isolation, transactions
- Concurrent operations
- Versioning, rollbacks, time-travel
- Schema Evolution
- Scale, Efficiency, Data Volumes
- Compatibility with existing data stores and languages



Data Platforms (Cloud or OSS)

- Separation of storage and compute
- Support for no-copy data sharing
- Bringing compute to data
- Fine-tuned granularity of permissions for access
- Support for automation and resource management
- Open standards and interoperability with other platforms and tools for governance, visualization, analytics, etc.



Multi-Cloud Infrastructure Management

- Terraform

Open-source infrastructure as code software tool that enables you to safely and predictably create, change, and improve infrastructure.

- Pulumi

Open-source infrastructure as code SDK that enables you to create, deploy, and manage infrastructure on any cloud, using your favorite languages.

- Crossplane

Assemble infrastructure from multiple vendors, and expose higher level self-service APIs for application teams to consume, without having to write any code.

Multi-Cloud Workload Portability

- **Azure Arc**

Build cloud-native apps anywhere, at scale. Run Azure services in any Kubernetes environment, whether it's on-premises, multi-cloud, or at the edge

- **Google Anthos**

A modern application management platform that provides a consistent development and operations experience for **cloud** and on-premises environments

Kubernetes Open-Standard Technologies

NOT AN EXHAUSTIVE LIST

- **Open Application Model**

An open standard for defining cloud native apps.

KubeVella - <https://kubevela.io/docs/concepts>

- **Open Policy Agent**

Declarative Policy-as-Code, enables portability, combination with Infra-as-Code.

<https://www.openpolicyagent.org/docs/latest>

- **Service Catalog**

Provision managed services and make them available within a Kubernetes cluster.

<https://kubernetes.io/docs/concepts/extend-kubernetes/service-catalog/>

Benefits Brought by Data Mesh

- Data Quality
- Tailored resource and focus allocation
- Organizational cohesion while allowing flexibility
- Reducing complexity
- Democratizing creating value
- Better understanding of value and innovation opportunities
- Empowering a more consistent and fast change

Important Focus Areas for Technology Providers

- **Open Standards, Open Protocols, Open-Source Integrations**
- Workload sharing and multi-tenancy
- No-copy data and compute mobility support
- Granularity of access-control
- Richness of automation and extensibility capabilities
- Flexibility and elasticity
- Provider-agnostic/multi-cloud operations support
- Variety of limitations
(quotas, data volume, resource count, etc.)

Data Mesh will drive better Interoperability, Open Standards, and Data Quality in the Industry

Thank you!

Follow  [lenadroid](#) for more insights