

Tracing

Kafka

with

 OpenTelemetry

Daniel Kim, Lead Developer Relations Engineer @ New Relic

Daniel Kim

Lead Developer Advocate, New Relic

- Likes spicy food (hotpot)
- Likes spicy talks
- Likes spicy tweets
- Follow me on **@learnwdaniel**



Agenda

- Why trace Kafka
- What is Distributed Tracing?
- What is OpenTelemetry?
- Implementing OpenTelemetry with Kafka
- Live Demo



I decided to do a live demo

Live Demo

Daniel's Awesome Kafka Talk

https://docs.google.com/forms/d/e/1FAIpQLScOYAbpPUVdqSSSFuE...

Daniel's Awesome Kafka Talk Questions

This Google form is linked to a Kafka Producer + Consumer, which will be demoed at Daniel's awesome talk. If you want to see your name in the demo portion of the talk, please ask a question :)

dgkim@ucdavis.edu (not shared) [Switch account](#)

* Required

What's your name? *

Your answer

What is your Question? *

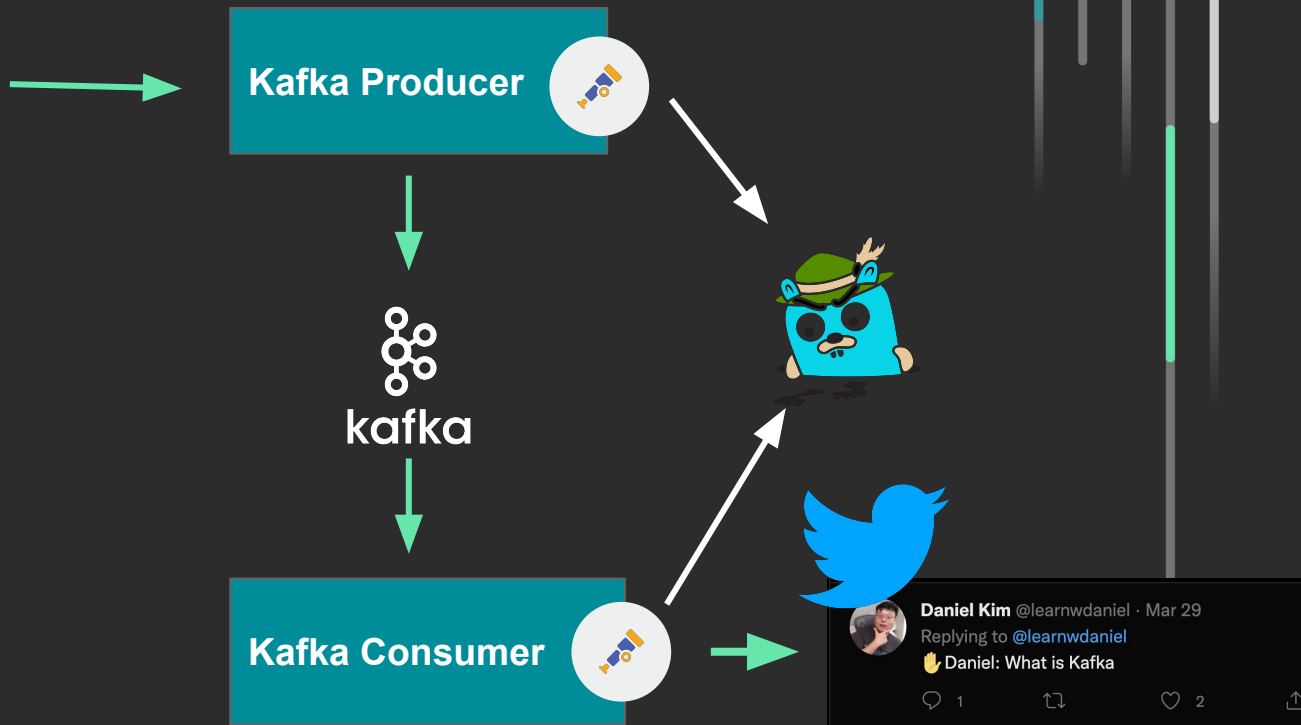
Your answer

Submit Clear form

er submit passwords through Google Forms.

This form was created inside of UC Davis. [Report Abuse](#)

Google Forms




We need test data - so please ask questions!

ask.danielkim.fun

Daniel's Awesome Kafka Talk Questions

This Google form is linked to a Kafka Producer + Consumer, which will be demoed at Daniel's awesome talk. If you want to see your name in the demo portion of the talk, please ask a question :)

 dgkim@ucdavis.edu

* Required

What's your name?


Your answer

What is your Question?

Your answer

Submit

Never submit passwords through Google Forms

 Daniel is in LONDON 🇬🇧🇬🇧
@learnwdaniel

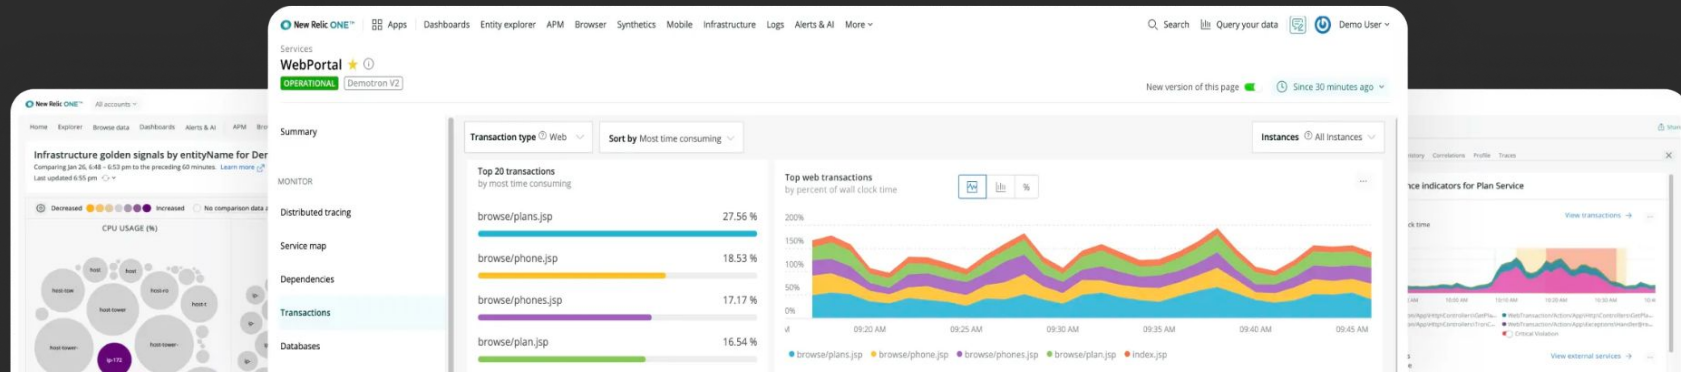
I'm talking all things @opentelemetry at #kafkasummit today! Super excited 🇬🇧🇬🇧




Why

Trace


Kafka?





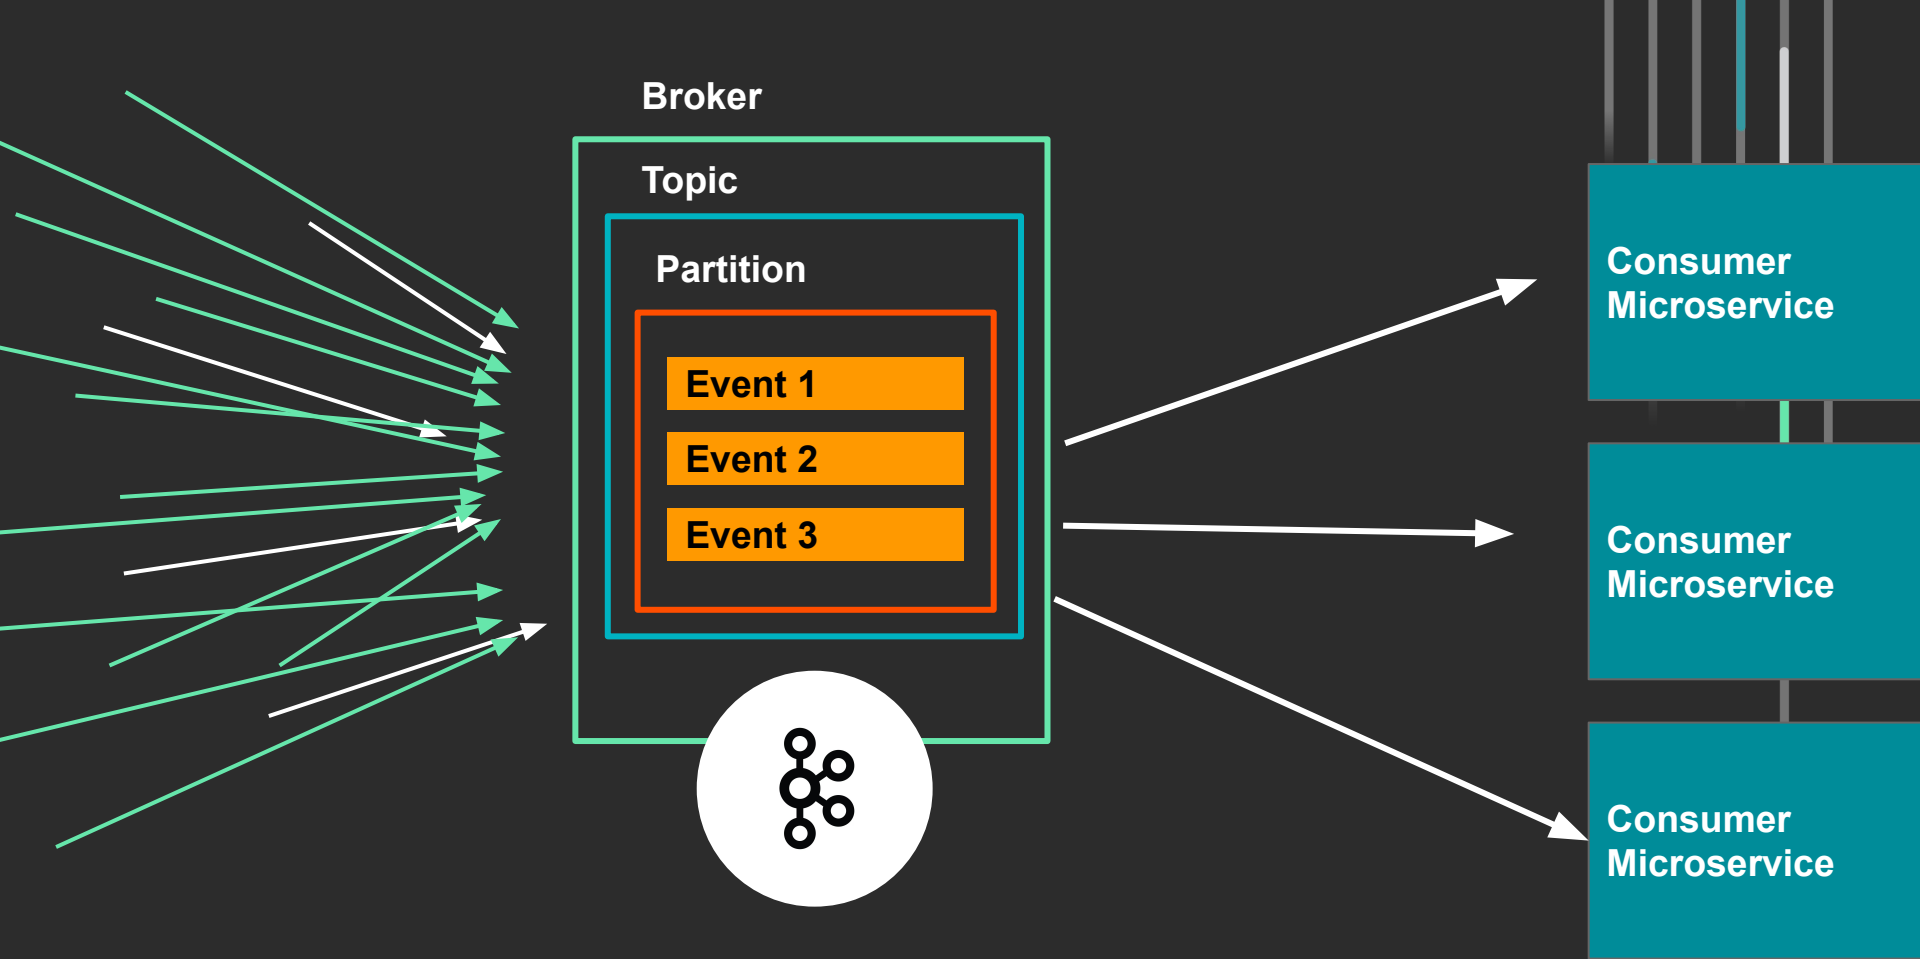
**200 petabytes a month
= 9,000,000 records/sec**

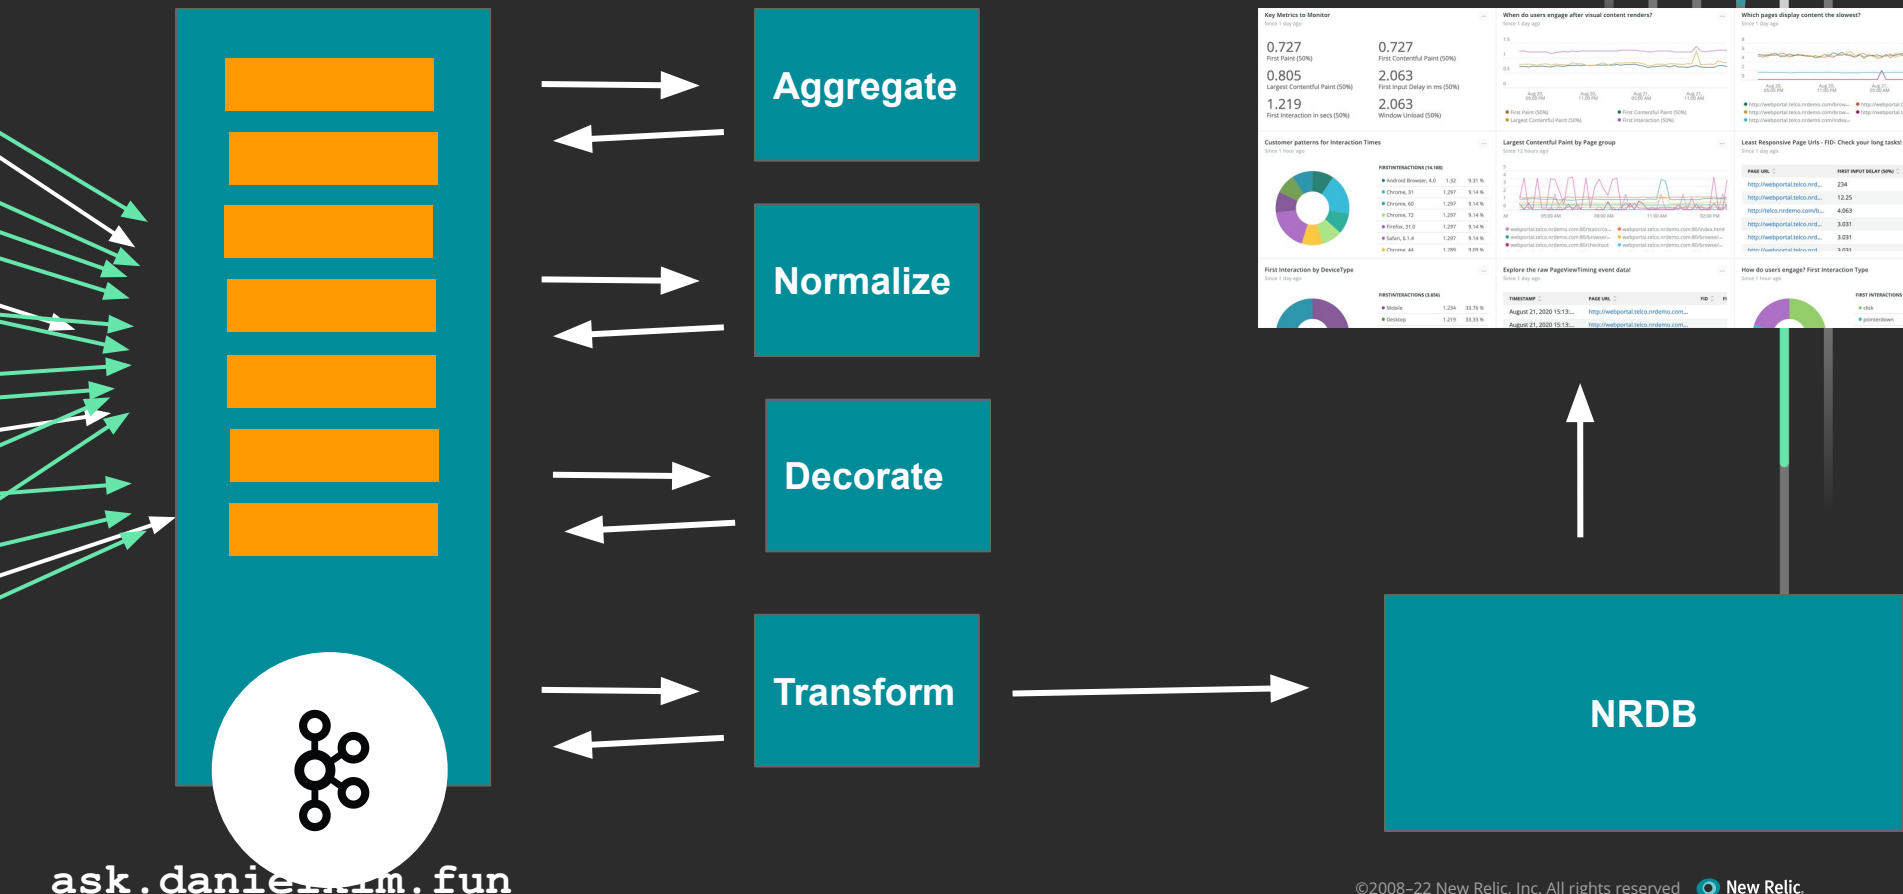
ask.danielkim.fun

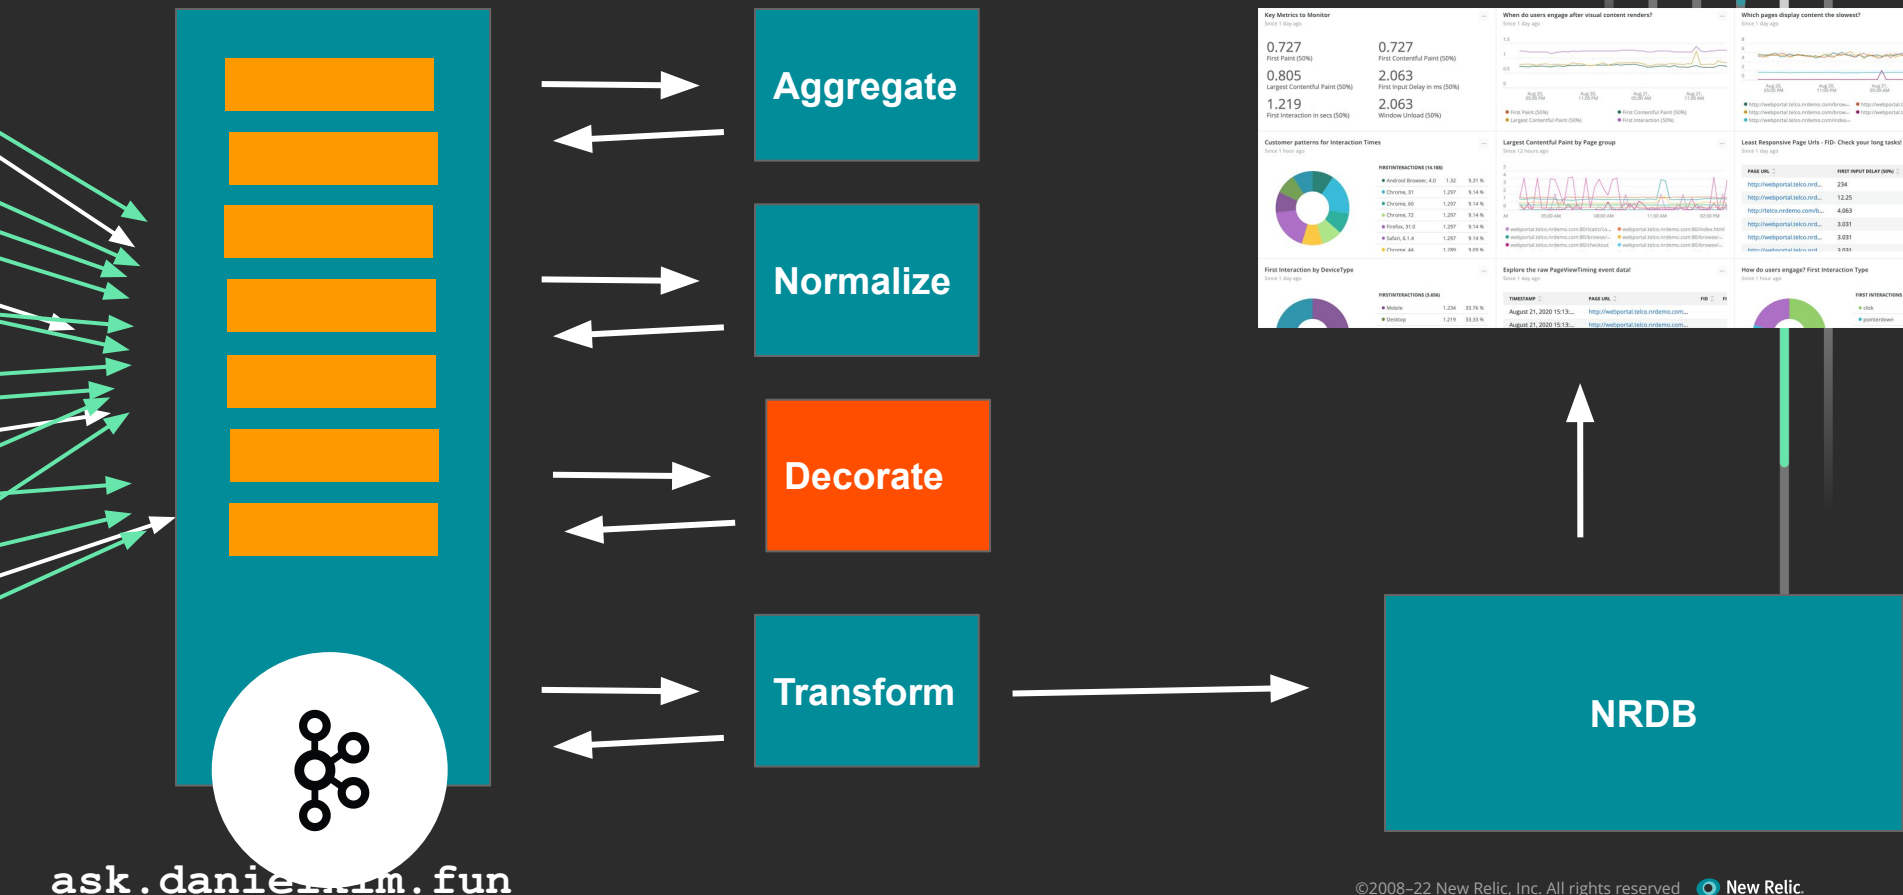




HTTP Requests

`ask.danielkim.fun`





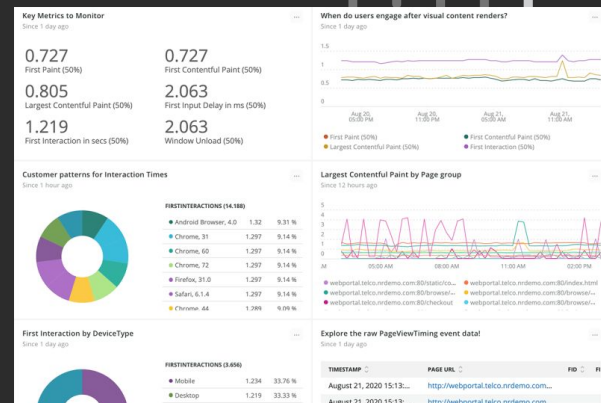




Kafka retains messages for a set amount of time so consumers can potentially go down or lag behind incoming data and recover without data loss

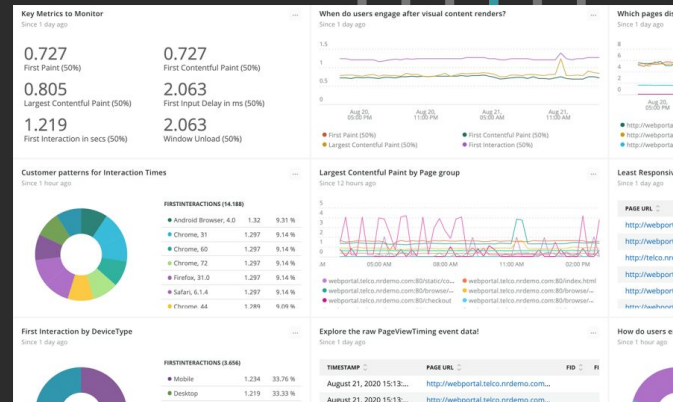
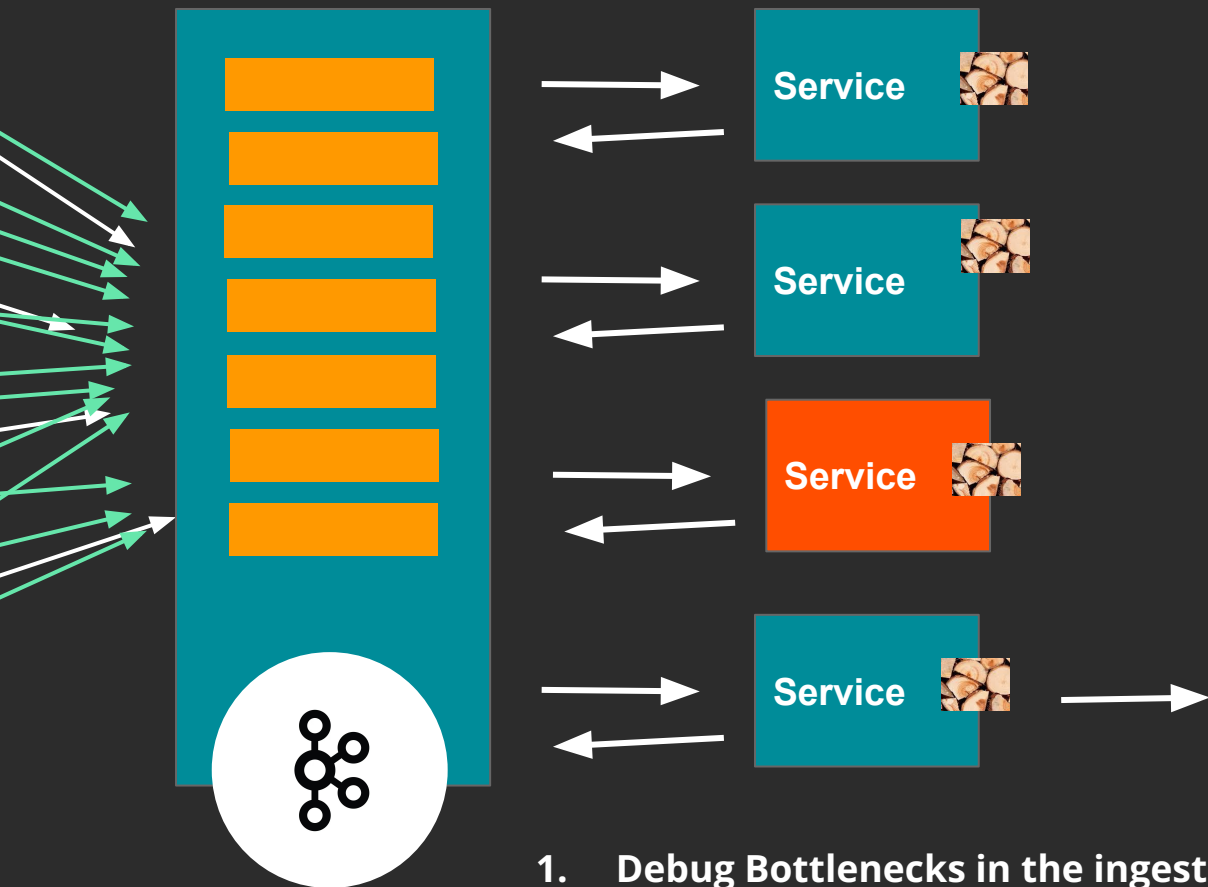
Why do you need to **trace** Kafka?

Agent

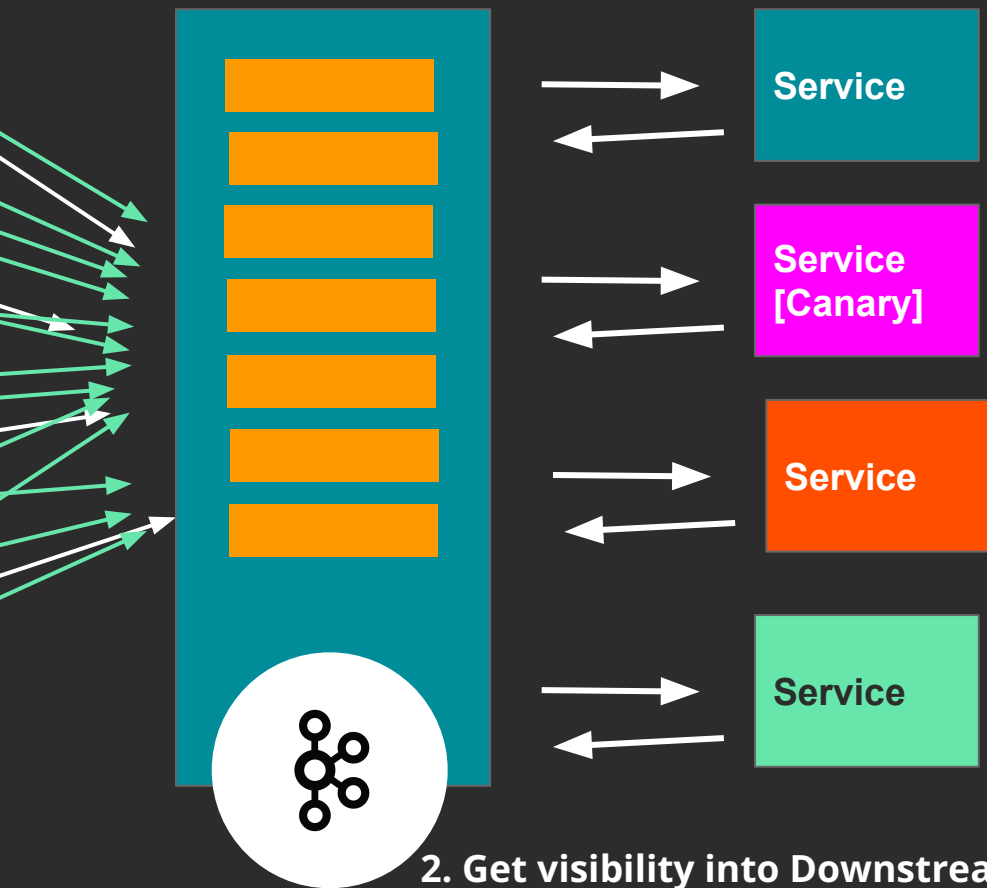


How do we optimize “time to glass”?


ask.danielkim.fun

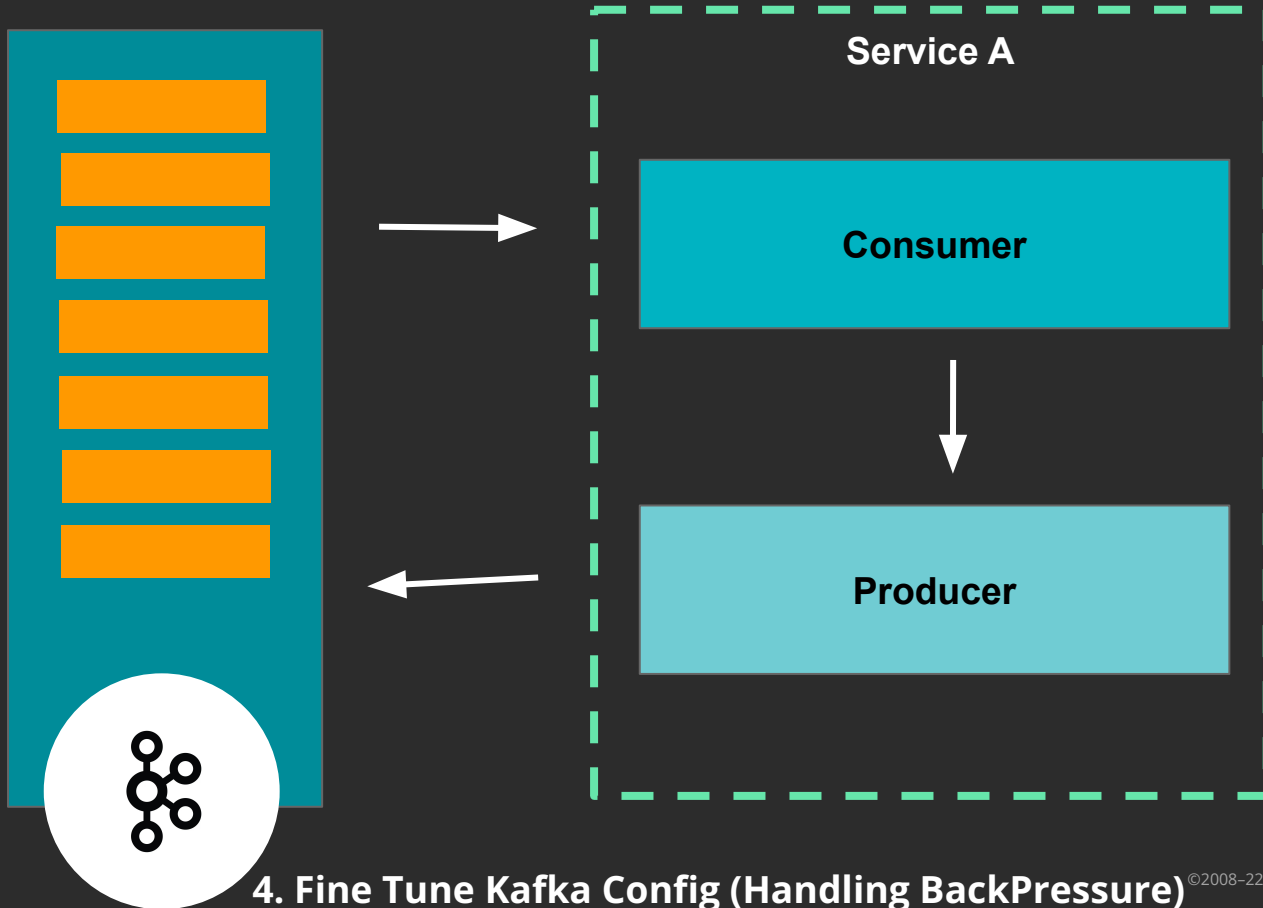


1. Debug Bottlenecks in the ingest pipeline

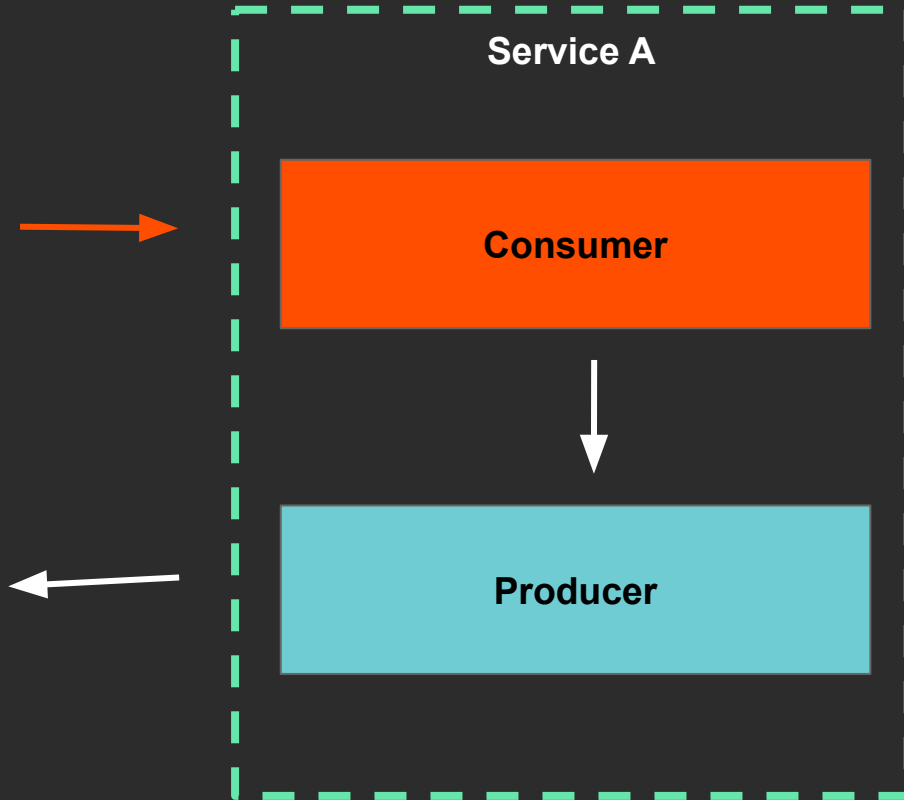



2. Get visibility into Downstream impacts of your service






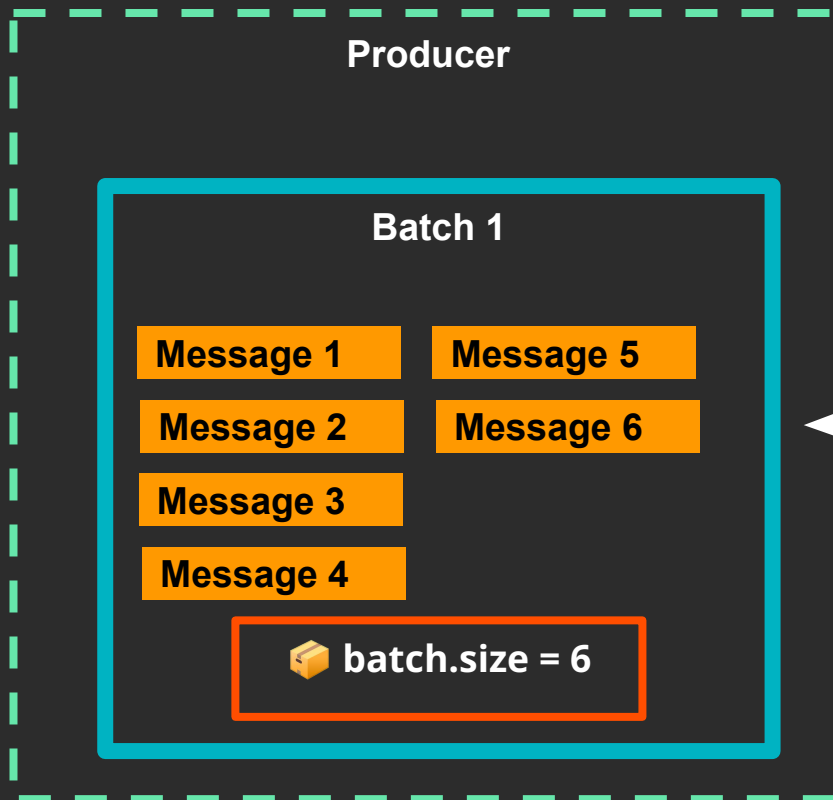
4. Fine Tune Kafka Config (Handling BackPressure)




4. Fine Tune Kafka Config (Handling BackPressure)



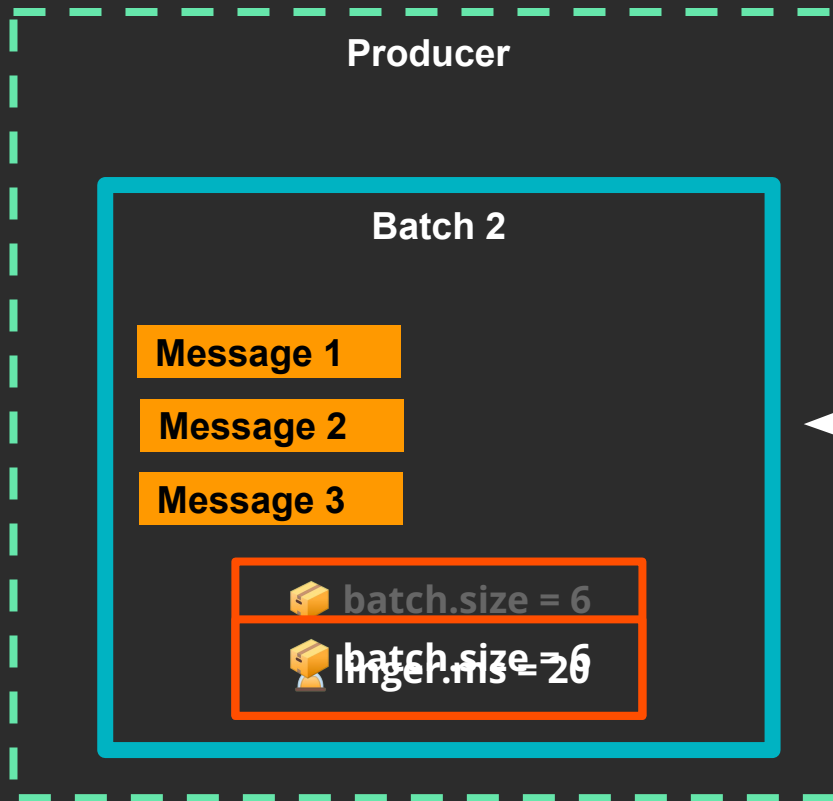
Compression



5. Fine Tune Kafka Config (Batching/Compressing Messages)



Compression





5. Fine Tune Kafka Config (Batching/Compressing Messages)

What



is

Distributed Tracing?

What is Distributed Tracing?



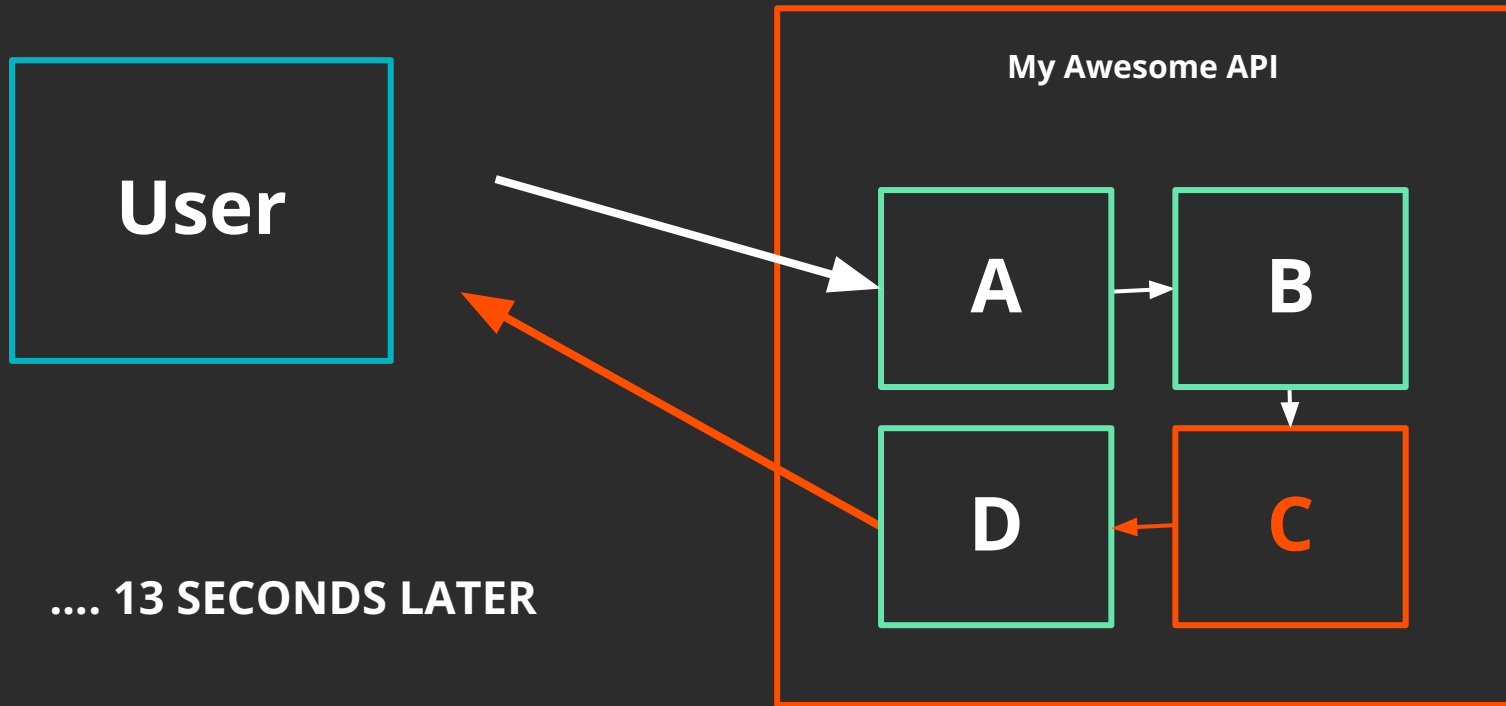
What is Distributed Tracing?



**My Awesome
API**

.... 13 SECONDS LATER

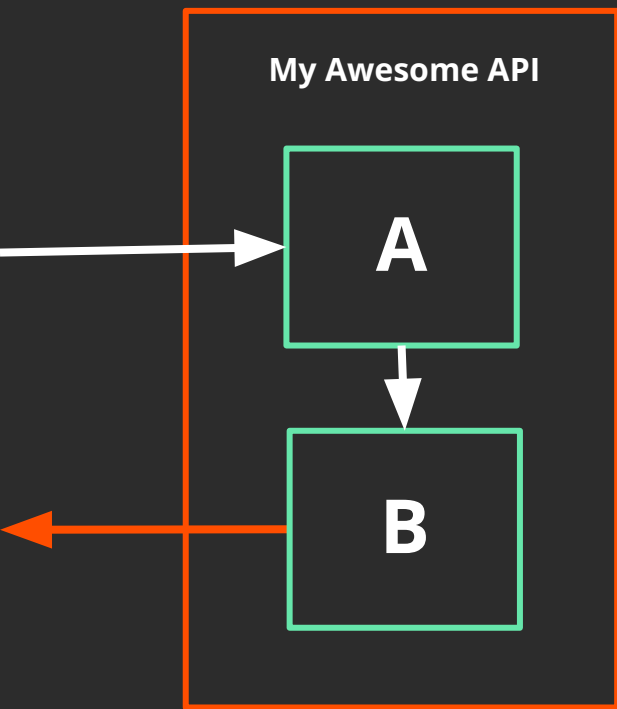
What is Distributed Tracing?



How do we connect the dots?

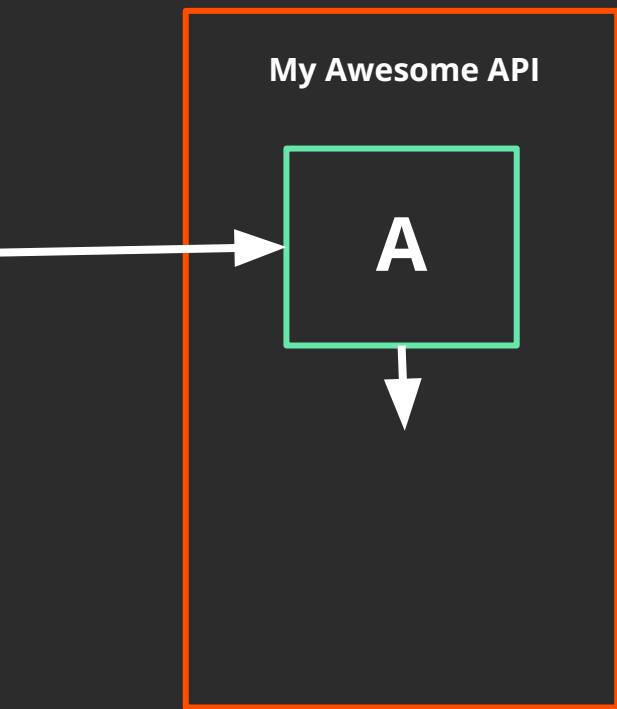
Context

Context in action



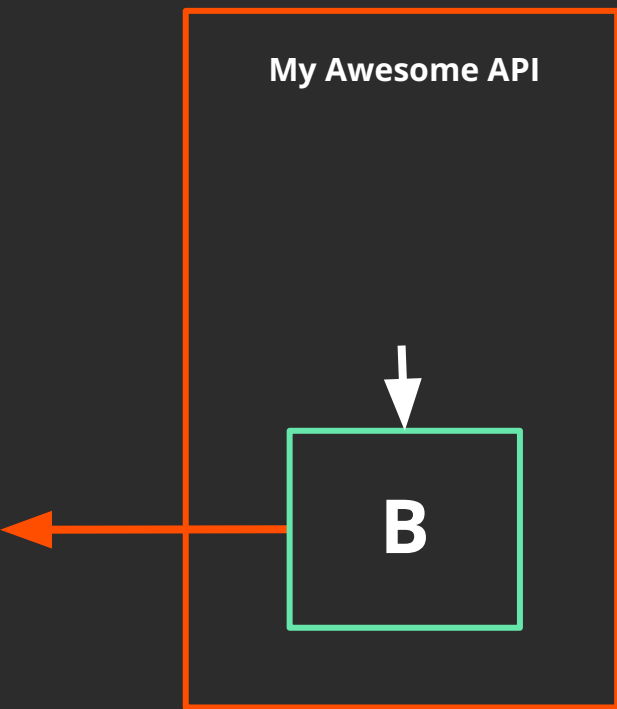
We only see one side of the transaction.

Context in action



.Here A is sending a payload to somewhere.

Context in action




**Here B is receiving a payload
from somewhere**

`ask.danielkim.fun`

Context in action


My Awesome API



Each Individual Span gets
its own id

Context in action

My Awesome API




**Each Individual Span gets
its own id (transaction)**

**And the whole trace gets its
own id (parentid)**

Context in action

My Awesome API




Each Individual Span gets
its own id (transaction)

And the whole trace gets its
own id (parentid)

**Both data points get
encoded as a set of HTTP
headers**

Context in action




Context

transaction=12H4
parentid=x1Fd

Overall Trace 12H4
And the parent span is x1FD


Context in action




Received by A

transaction=12H4
parentid=x1Fd

Overall Trace 12H4
And the parent span is x1FD







What if we have microservices with different languages or frameworks?

What if microservices are using different observability systems?


What if we have decoupled microservices via Kafka?

W3C standard





OSS Standard for Context Propagation




**W3C Trace Context
enables cross-vendor
interoperation of traces**

Adoption of W3C Tracing standard



OpenTelemetry

Vendors adopting W3C



<https://github.com/w3c/trace-context/blob/main/implementations.md>


What

is

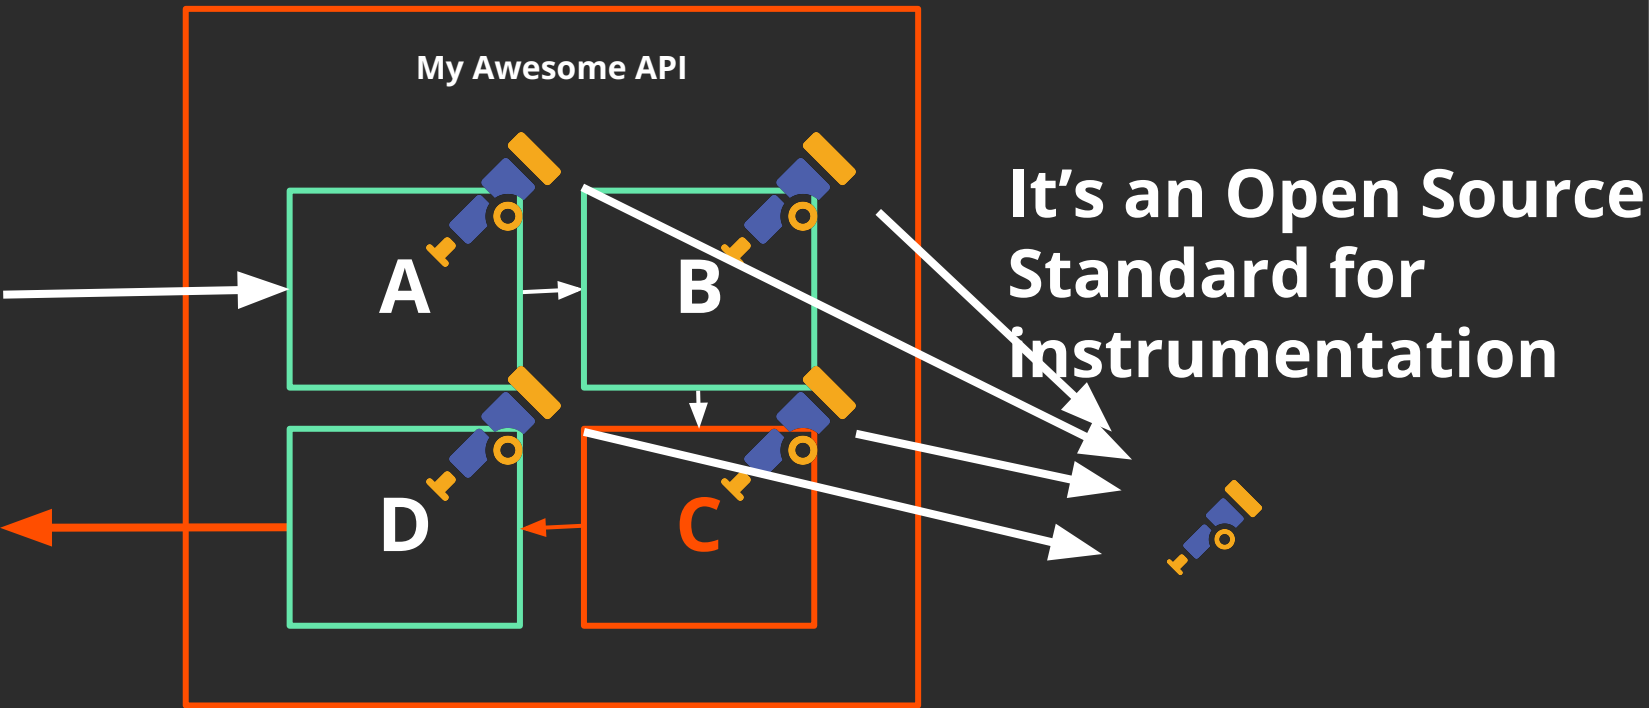


OpenTelemetry

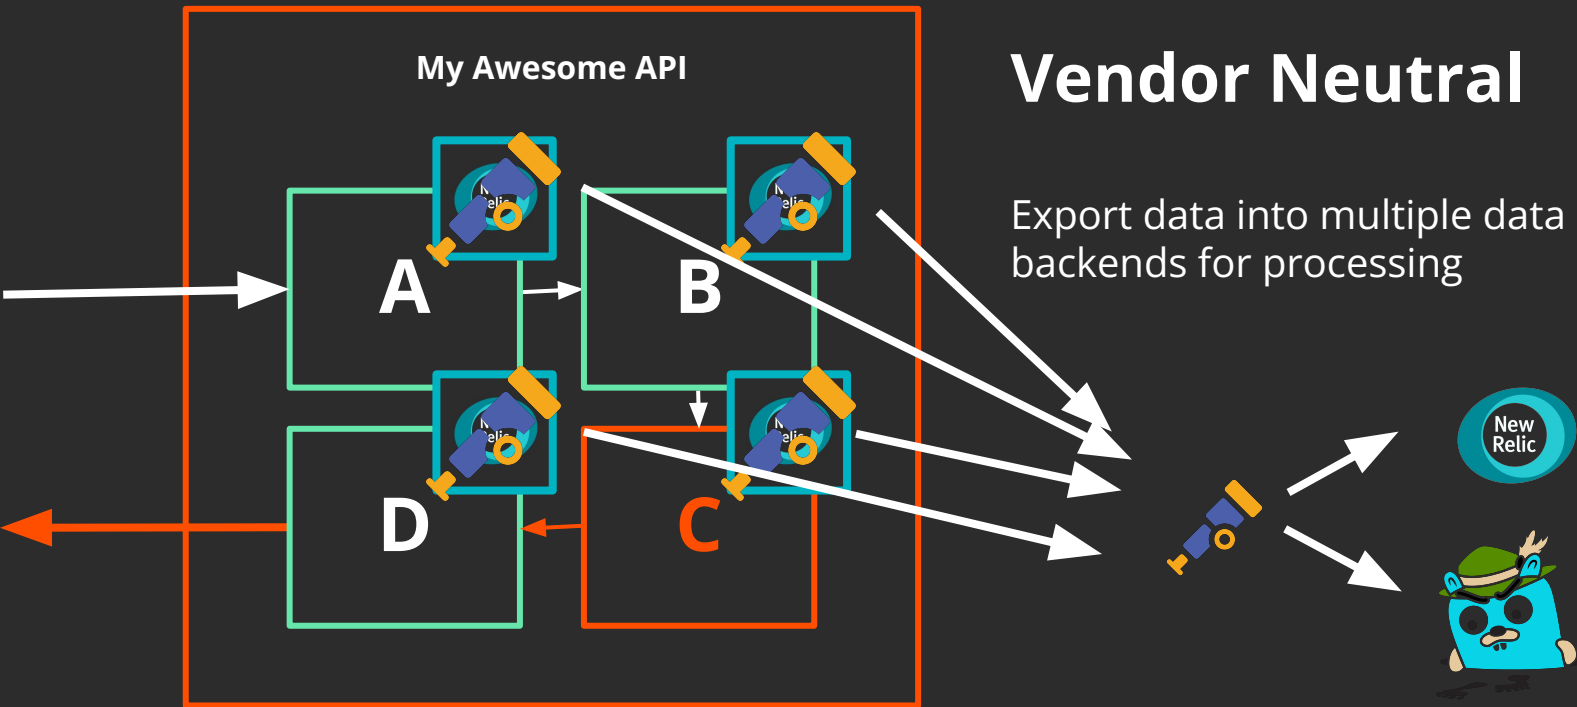
What is OpenTelemetry?



What is OpenTelemetry?



What is OpenTelemetry?

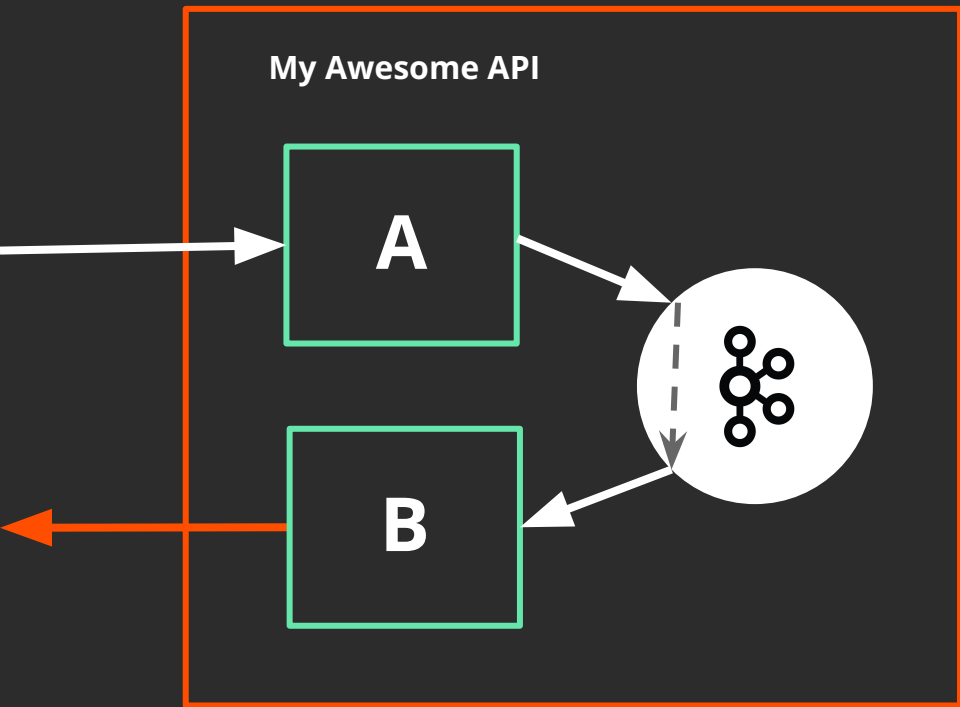


Implementing

Distributed Tracing

For Kafka

What happens with Kafka?



**Context is lost and
transaction is broken**

What happens with Kafka?

Consumer

Kafka decouples producers and consumers

Producer

`ask.danielkim.fun`

Adding

CONTEXT

to traces

Otelsarama injects context into Kafka messages, allowing you to trace asynchronous messages

Shopify/sarama

Sarama is a Go library for Apache Kafka.

216

Contributors

6k


Used by


8k

Stars

1k

Forks





ProductTeamEnterpriseExploreMarketplacePricing

Search

Sign inSign up

open-telemetry / opentelemetry-go-contribPublic

NotificationsFork224Star381

<>CodeIssues68Pull requests13ActionsProjects1WikiSecurity

main

opentelemetry-go-contrib / instrumentation / github.com / Shopify / sarama / otelsarama /

Go to file

MrAlias and Aneurysm9 Release v1.6.0/v0.31.0 (#2134)

56111417 hours agoHistory

..		
example	Release v1.6.0/v0.31.0 (#2134)	7 hours ago
test	Release v1.6.0/v0.31.0 (#2134)	7 hours ago
consumer.go	Build tooling refactor (#1974)	3 days ago
consumer_group.go	Build tooling refactor (#1974)	3 days ago
consumer_group_test.go	Update instrumentation naming (#359)	2 years ago
consumer_test.go	Remove use of oteltest in otelsarama (#1007)	7 months ago
dispatcher.go	Build tooling refactor (#1974)	3 days ago
doc.go	Update instrumentation naming (#359)	2 years ago
go.mod	Release v1.6.0/v0.31.0 (#2134)	7 hours ago
go.sum	Release v1.6.0/v0.31.0 (#2134)	7 hours ago
message.go	Build tooling refactor (#1974)	3 days ago
message_test.go	Bumping otel version to v0.18.0. Prepare for releasing v0.18.0 (#600)	13 months ago
option.go	Build tooling refactor (#1974)	3 days ago
option_test.go	Prerelease v0.24.0 (#1125)	6 months ago
github.com/open-telemetry/...sarama	Build tooling refactor (#1974)	3 days ago

Kafka Headers

HEADER

Message

**Add metadata about your record without
changing the content of your message**

Kafka Headers


TraceState

TraceParent

Message

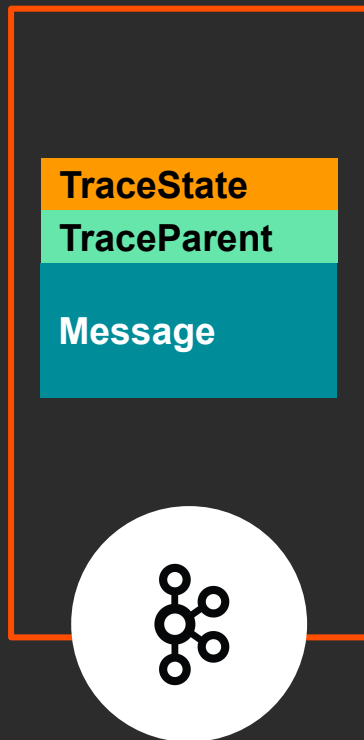
Let's inject context (W3C Tracing Headers)

Kafka Headers

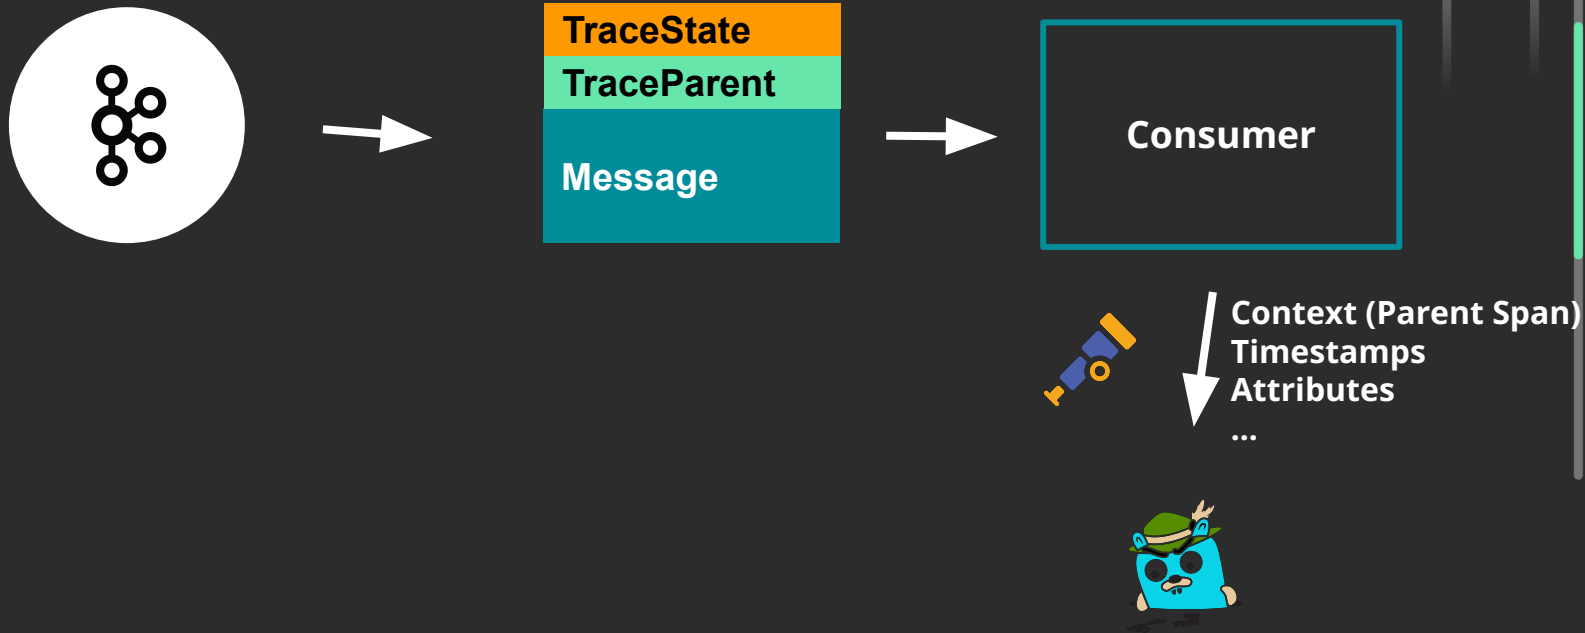


**Producer needs to
Inject context into
Kafka Header**

Kafka Headers



Kafka Headers



Let's print out what is injected in the header



kcat

<https://github.com/edenhill/kcat>

Key (8 bytes): question

Value (15 bytes): 🖐 Daniel: food

Timestamp: 1649611412653

Partition: 0

Offset: 7

→ Headers: tracestate=,traceparent=00-6626f84bd20a9a171a699a560387102b-10ba2f2a74c884bf-01

Reading the trace headers

00-6626f84bd20a9a171a699a560387102b-10ba2f2a74c884bf-01



**Tracing Standard Version
number (Base16)**

Reading the trace headers

00-6626f84bd20a9a171a699a560387102b-10ba2f2a74c884bf-01



Trace ID (Base16)

This is the ID of the whole trace and is used to uniquely identify a distributed trace through a system.

Reading the trace headers

00-6626f84bd20a9a171a699a560387102b-10ba2f2a74c884bf-01

Trace
Parent(Base16)



Reading the trace headers

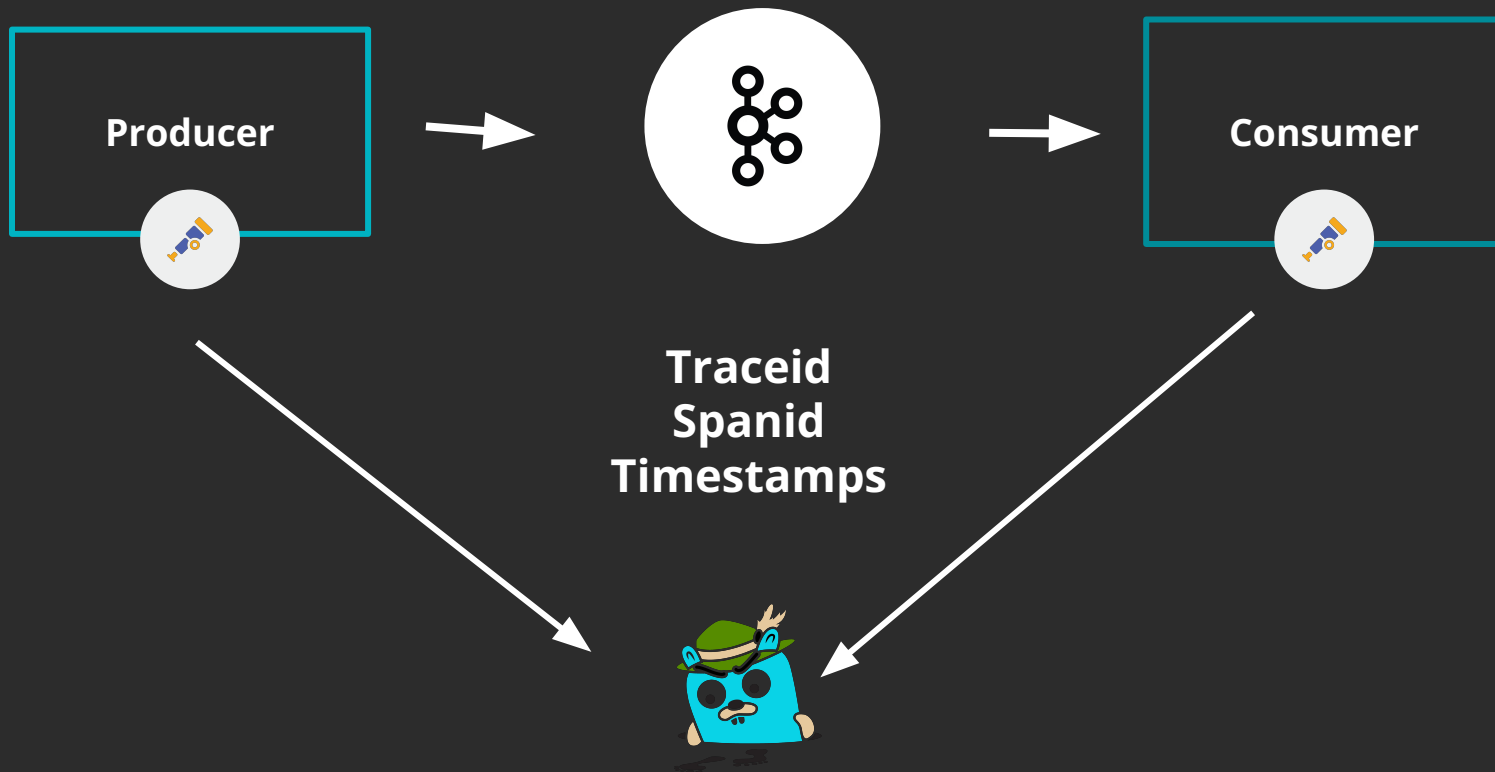
00-6626f84bd20a9a171a699a560387102b-10ba2f2a74c884bf-01

This flag indicates that the span
has been sampled and will be
exported

Tracing Flags



OTel in Action





Live Demo

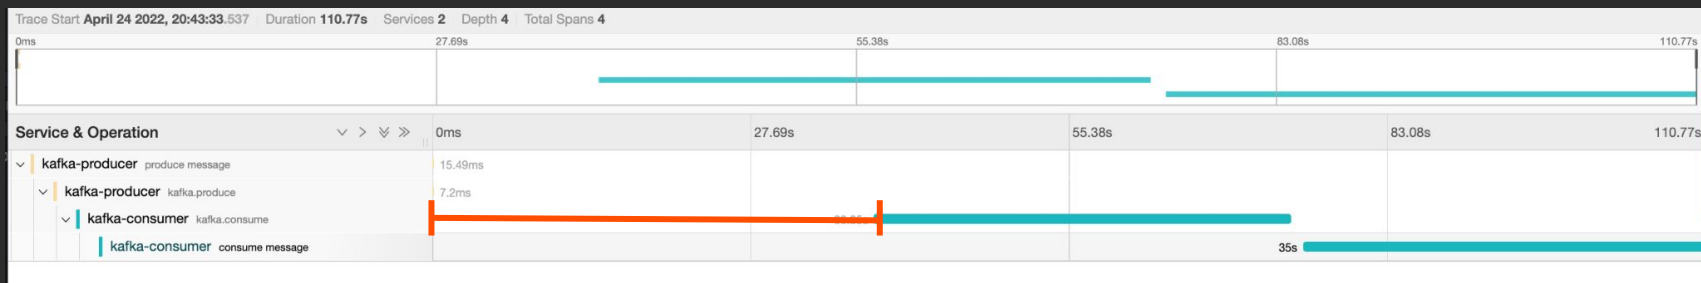
`github.com/lazyplatypus/kafka-otel`

<input type="checkbox"/>	kafka-producer: produce message a7d900f	110.77s
4 Spans	<div><div>kafka-consumer (2)</div><div>kafka-producer (2)</div></div>	Today 8:43:33 pm 26 minutes ago
<input type="checkbox"/>	kafka-producer: produce message 282c9fd	82.03s
4 Spans	<div><div>kafka-consumer (2)</div><div>kafka-producer (2)</div></div>	Today 8:43:25 pm 26 minutes ago
<input type="checkbox"/>	kafka-producer: produce message f938c84	52.12s
4 Spans	<div><div>kafka-consumer (2)</div><div>kafka-producer (2)</div></div>	Today 8:43:19 pm 26 minutes ago
<input type="checkbox"/>	kafka-producer: produce message a64377b	36.02s
4 Spans	<div><div>kafka-consumer (2)</div><div>kafka-producer (2)</div></div>	Today 8:42:59 pm 26 minutes ago

What's Happening?



The consumer was busy handling other requests so the message was sitting in the kafka queue



What's Happening?

ConsumeClaim

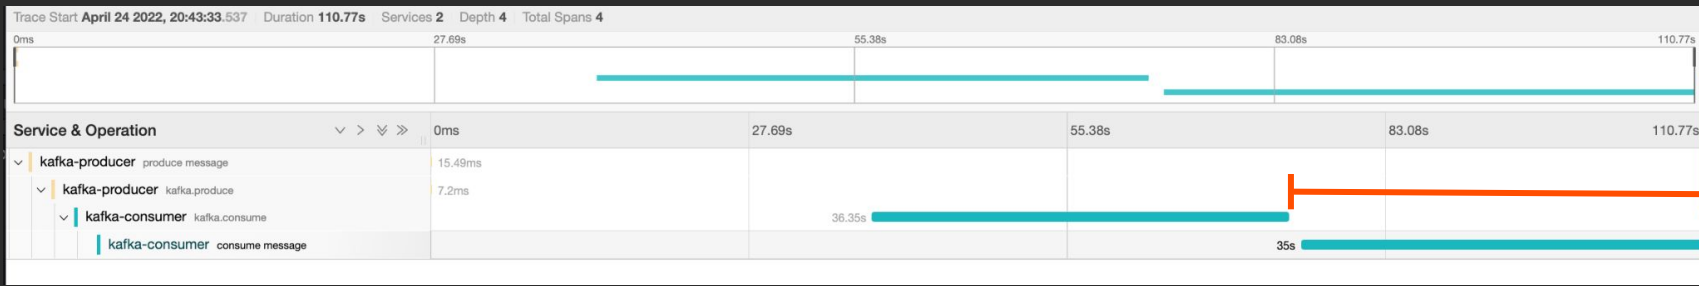
time.Sleep

printMessage

MarkMessage

tweetMessage

```
// Emulate work loads (or any...)
time.Sleep(35 * time.Second)
```



What have

WE

Learned?

Caveat:

Implementing Distributed Tracing is **NOT EASY**

Especially at Scale, a lot of Cross Functional work to make sure tracing is implemented in every microservice in your set up.

It is gonna be a lot of data, even when you sample every 10k, 50k, 100k traces.

Distributed Tracing is AWESOME!

- It gets you better visibility into the WHY on things are behaving funky and helps you optimize your data ingest pipeline
- It helps you be a better neighbor
- Helps you tweak config variables to make your system more performant.
- OSS Instrumentation/standards is the future of observability!

Thank you!

twitter.com/learnwdaniel

ask.danielkim.fun