

Developer's Guide to Contributing Code to Kafka

Mickael Maison
Tom Bentley

Who are we?

Mickael

Contributor since 2016

Committer since 2019

PMC member since 2020

Tom

Contributor since 2017

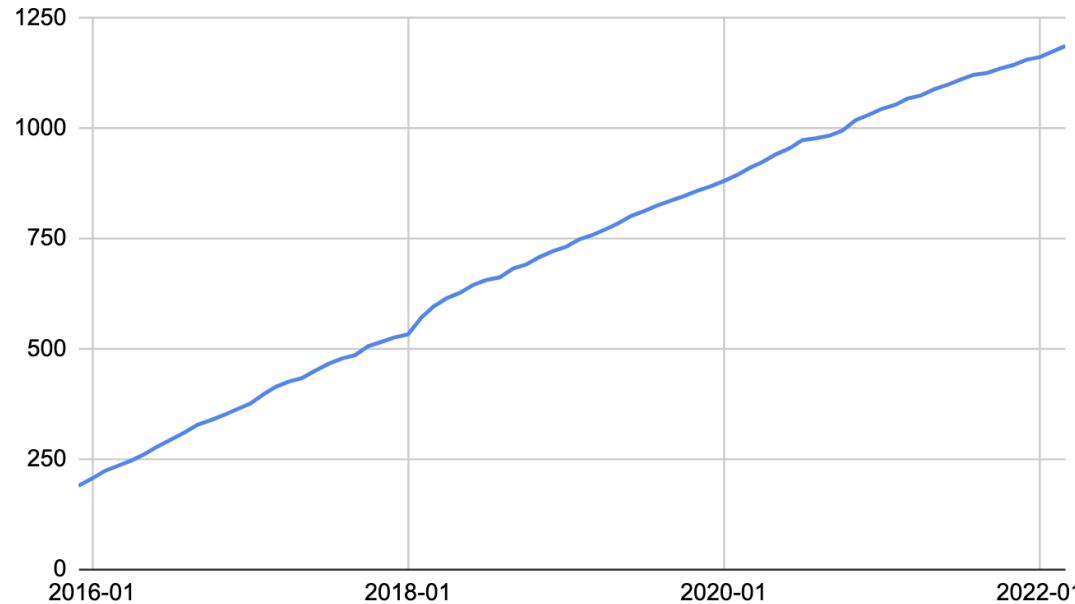
Committer since 2021

PMC member since 2021

Apache Kafka: An Open Source Project



Become a contributor!



<https://github.com/apache/kafka/blob/trunk/CONTRIBUTING.md>

Collaboration



Dev Mailing List: <https://kafka.apache.org/contact>



Confluence

Wiki: <https://cwiki.apache.org/confluence/display/KAFKA/Index>



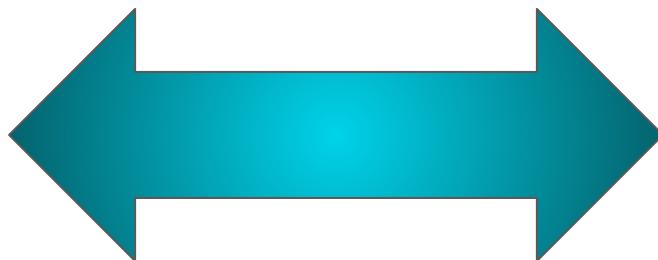
Bug tracker: <https://issues.apache.org/jira/projects/KAFKA/issues>



Code: <https://github.com/apache/kafka>

Finding something to work on

“I just need to fix
this one bug”

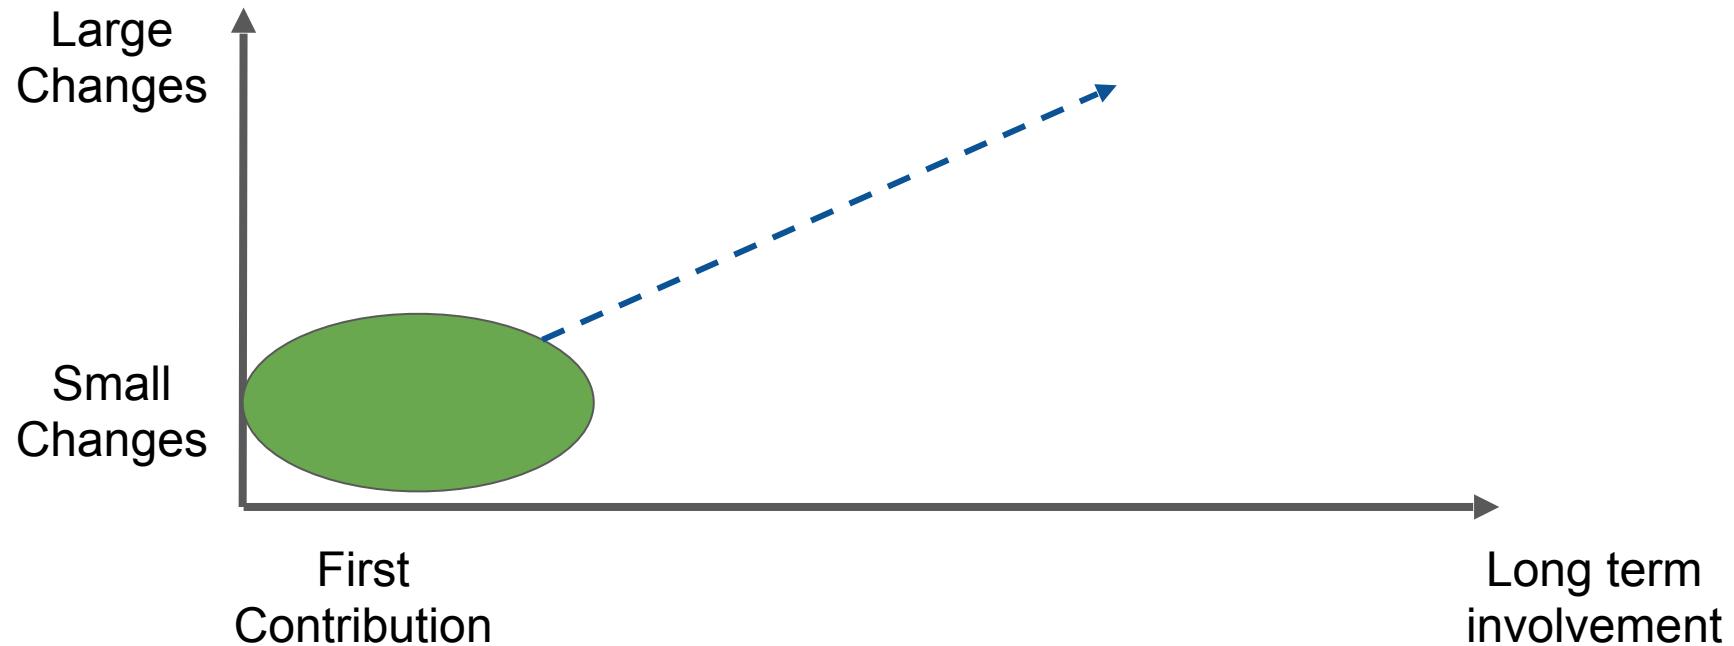


“I just want to
contribute”

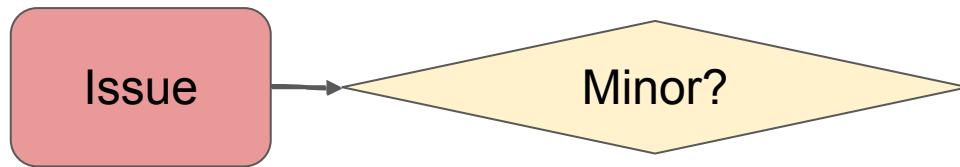
Finding something to work on



Finding something to work on



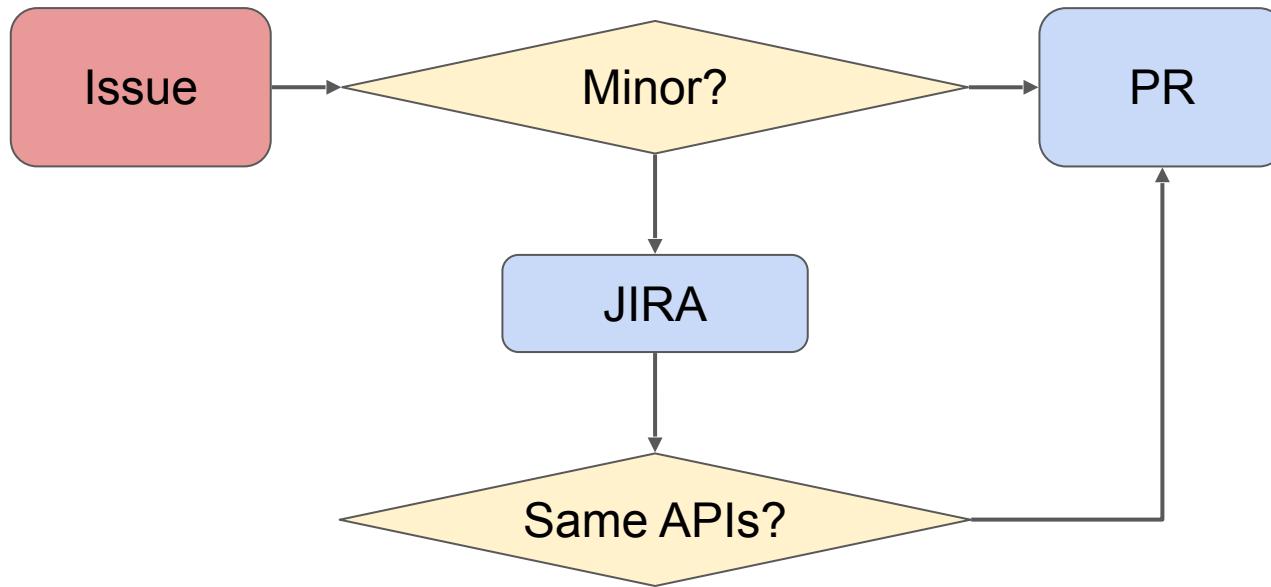
Contributing Code Changes



Contributing Code Changes



Contributing Code Changes



Writing code *and* tests



<https://github.com/apache/kafka/blob/trunk/README.md>

Gradle Tasks

Compile: assemble



2 minutes

Build + run test: build



20-30 minutes

Run test on specific project: clients:test



variable

Run specific tests: core:test --tests "*PartitionTest"

Opening a Pull Request

MINOR: Refactor `kafka.cluster.Replica` #12081

 Open

dajac wants to merge 1 commit into `apache:trunk` from `dajac:KAFKA-13790-2` 

Conversation 0

Commits 1

Checks 11

Files changed 7



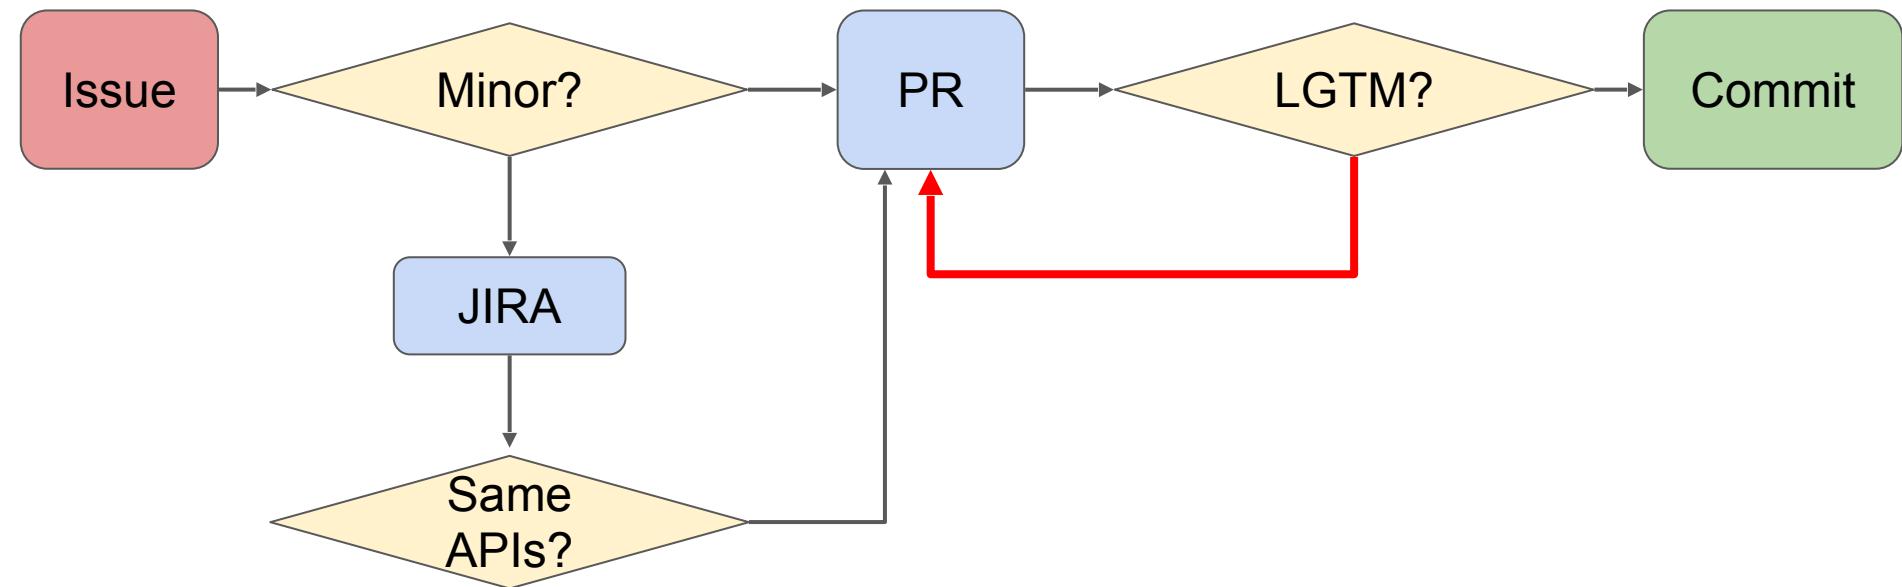
dajac commented 2 hours ago • edited 

Member 

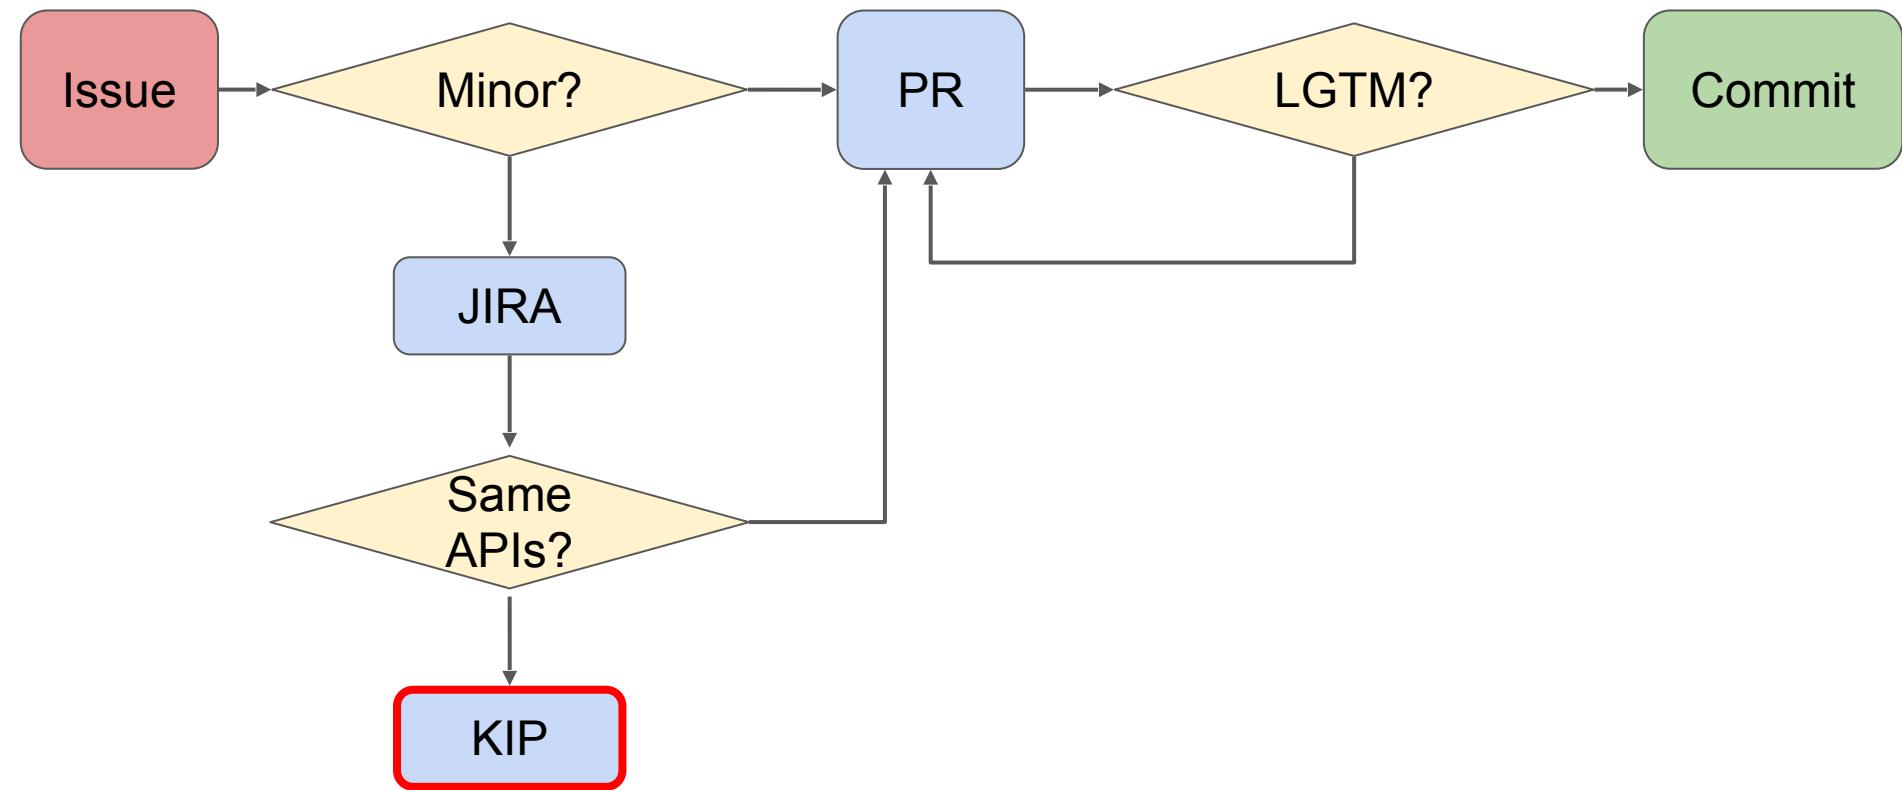
This patch refactors `kafka.cluster.Replica`, its usages and tests. This is part of the work in [KAFKA-13790](#).

Reviewers: <https://kafka.apache.org/committers>

Making a Change



Kafka Improvement Proposals



ONE DOES NOT SIMPLY



CHANGE A PUBLIC KAFKA API

Configs
Interfaces
Public Classes
Network Protocol
Log Format
Behaviour
Metrics
Tools

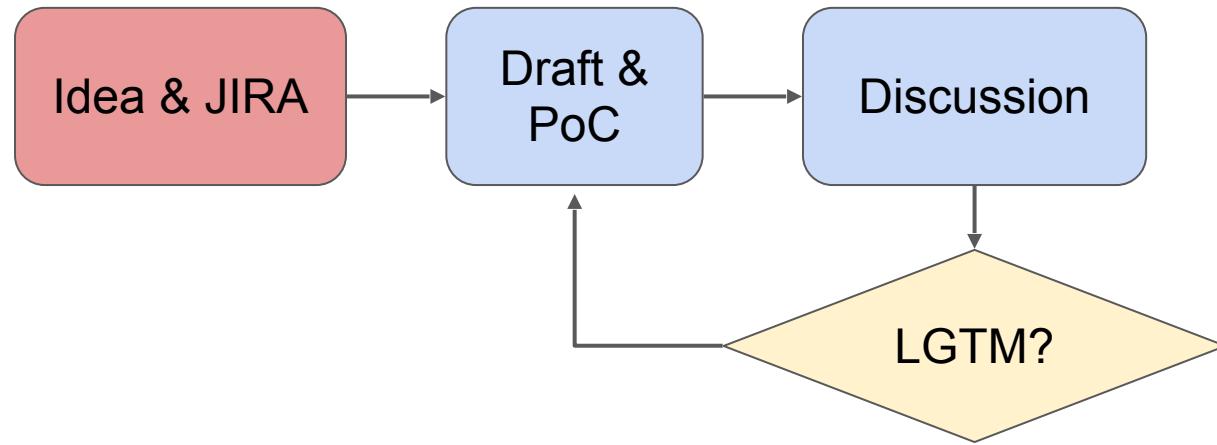
Idea & JIRA

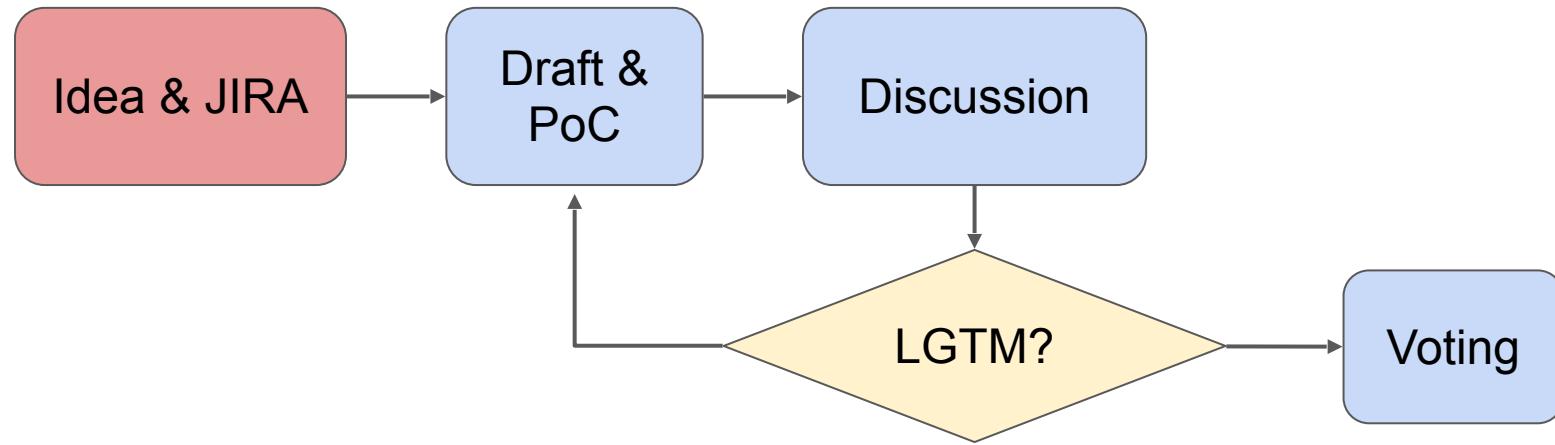
Idea & JIRA

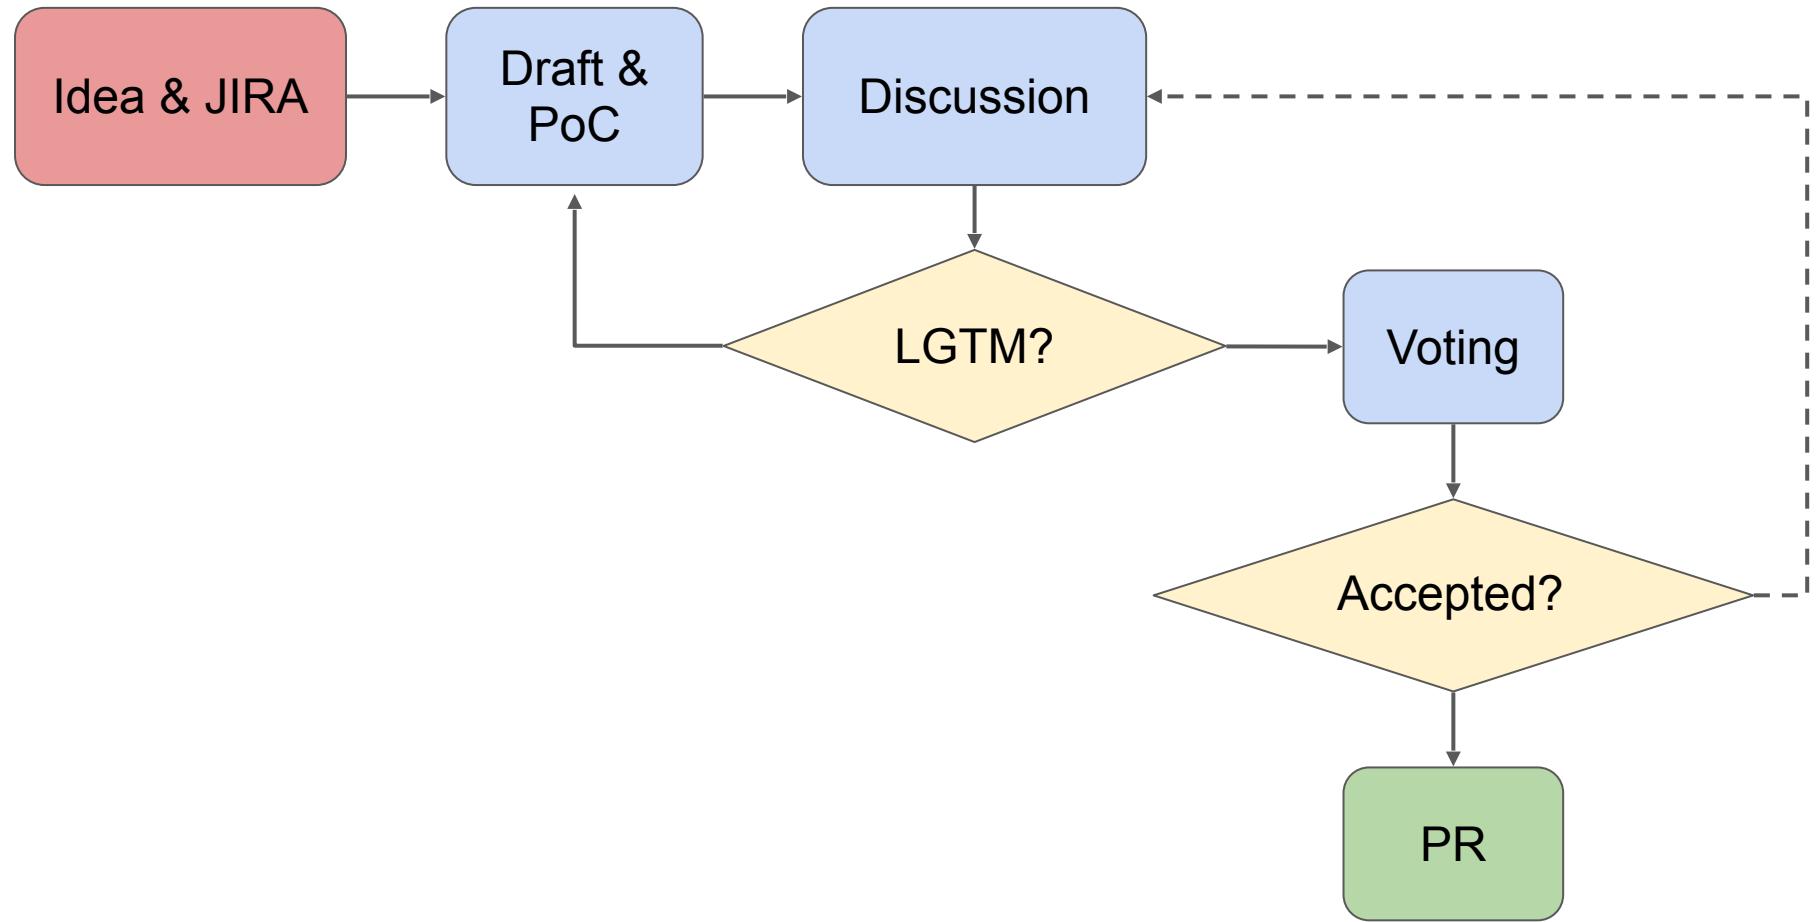
Draft &
PoC











Draft &
PoC

Anatomy of a KIP

[KIP Index](#) lists all the KIPs, and their current state

Standard sections:

Status

Motivation

Public interfaces

Proposed changes

Compatibility, Deprecation, Migration

Test Plan

Rejected alternatives

Anatomy of a KIP - Status section

Status

Current state: [*One of "Under Discussion", "Accepted", "Rejected"*]

Discussion thread: [here](#)

JIRA: [here](#)

Released: <Kafka Version>

Please keep the discussion on the mailing list rather than commenting on the wiki (wiki discussions get unwieldy fast).

Anatomy of a KIP - Motivation section

Motivation

Describe the problems you are trying to solve.

Example:

[KIP-730: Producer ID generation in KRaft mode](#)

Anatomy of a KIP - Public interfaces

Public Interfaces

Briefly list any new interfaces that will be introduced as part of this proposal or any existing interfaces that will be removed or changed. The purpose of this section is to concisely call out the public contract that will come along with this feature.

Example:

KIP-664 Provide tooling to detect and abort hanging transactions

Anatomy of a KIP - Proposed changes

Proposed Changes

Describe the new thing you want to do in appropriate detail. This may be fairly extensive and have large subsections of its own. Or it may be a few sentences. Use judgement based on the scope of the change.

Example:

KIP-98: Exactly Once Delivery and Transactional Messaging

Anatomy of a KIP - Compatibility, Deprecation, Migration

Compatibility, Deprecation, and Migration Plan

- *What impact (if any) will there be on existing users?*
- *If we are changing behavior how will we phase out the older behavior?*
- *If we need special migration tools, describe them here.*
- *When will we remove the existing behavior?*

Example:

[KIP-802 Validation Support for Kafka Connect SMT and Converter Options](#)

Anatomy of a KIP - Test plan

Test Plan

Describe in few sentences how the KIP will be tested. We are mostly interested in system tests (since unit-tests are specific to implementation details). How will we know that the implementation works as expected? How will we know nothing broke?

Example:

None! This is new

Anatomy of a KIP - Rejected alternatives

Rejected Alternatives

If there are alternative ways of accomplishing the same thing, what were they? The purpose of this section is to motivate why the design is the way it is and not some other way.

Example:

KIP-618: Exactly Once support for Source Connectors

Discussion

To: dev@kafka.apache.org
Subject: [DISCUSS] KIP-XXX: My great idea



Voting

To: dev@kafka.apache.org
Subject: [VOTE] KIP-XXX: My great idea

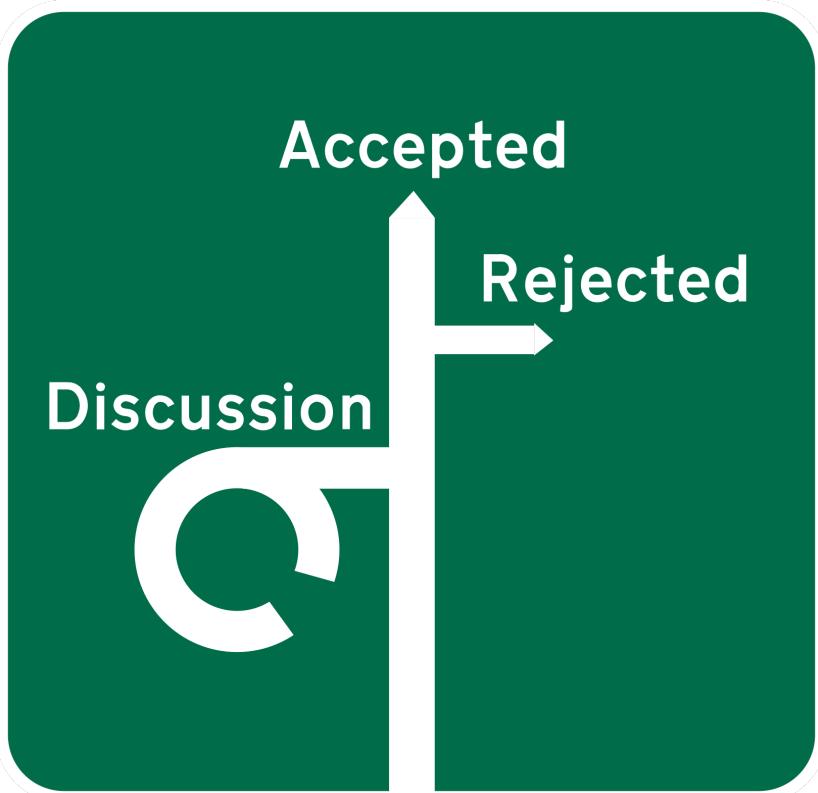




IT'S FINALLY OVER

WHEN YOUR PR IS MERGED

Interminable discussion



- Real problem?
- Clear + strong motivation?
- Too narrow?
- Too broad?
- Listened to feedback?

Contributing regularly



Thank you!



Key Resources:

<https://kafka.apache.org/contributing>

<https://cwiki.apache.org/confluence/display/KAFKA/Kafka+Improvement+Proposals>

<https://kafka.apache.org/committers>

<https://github.com/apache/kafka>