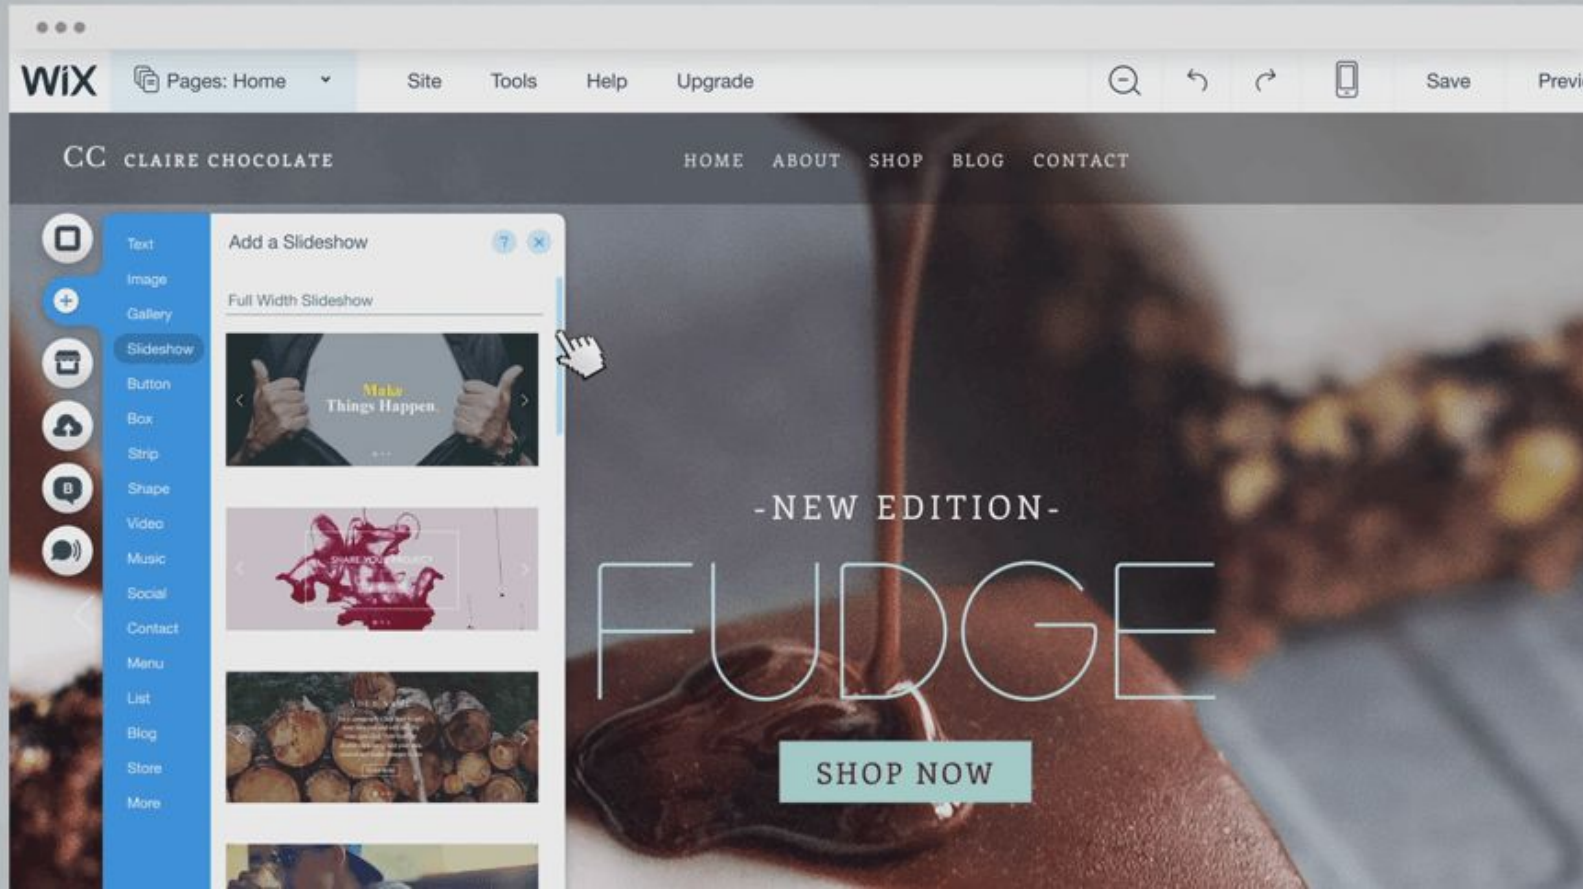


Kafka-based Global Data Mesh at Wix

Natan Silnitsky Backend Infra TL, Wix.com

natansil.com [twitter@NSilnitsky](https://twitter.com/NSilnitsky) [linkedin/natansilnitsky](https://linkedin.com/natansilnitsky) github.com/natansil



Our mission

Enable any business,
community or person on earth
to make their **online dream
come true.**

At Wix



±1B
Unique
visitors

use Wix
platform
every month

>100M
Websites

Published by
Wix users
(5-7% of the
entire internet
websites)

±5000
People

Work at Wix

Scaling out

500B

Daily HTTP
Transactions

>600

GAs Per Day

~1500

Developers


18

Data
Centers &
Pops

2

Cloud Providers
[Google/AWS]

Wix has over 2200 microservices in production



15B daily business **Events**


4 **Regions** @ GCP/AWS

~50k **Topics**

> 500k **Partitions**

45% of the data is **replicated**

+ We assume
x10 scale factor
during 12-24 next
months



A Fully Distributed Environment

Growing Pains



**Event-driven
architecture**

**Large scale
traffic**

**Cluster
overload**

**Hundreds of
developers**

**DC
awareness**

**Diverse
tech stack**

Growing Pains



Event-driven
architecture

Large scale
traffic

Cluster
overload

Hundreds of
developers

DC
awareness

Diverse
tech stack

Growing Pains



Event-driven
architecture

Large scale
traffic

**Cluster
overload**

Hundreds of
developers

**DC
awareness**

Diverse
tech stack

Growing Pains



Kafka

Event-driven
architecture

Large scale
traffic

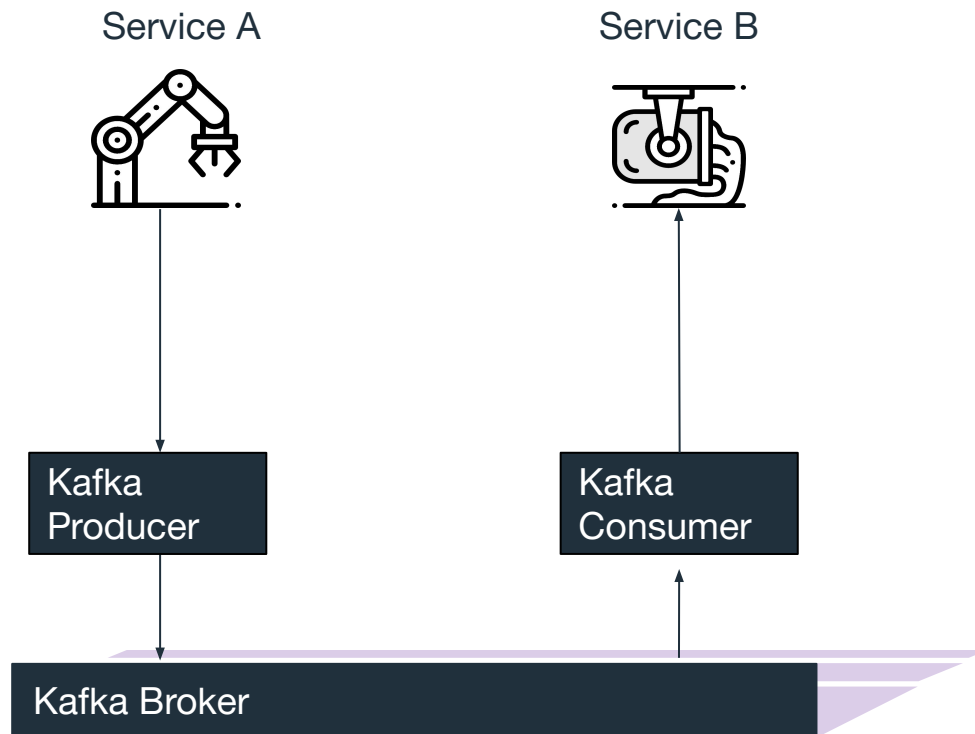
Hundreds of
developers

Diverse
tech stack

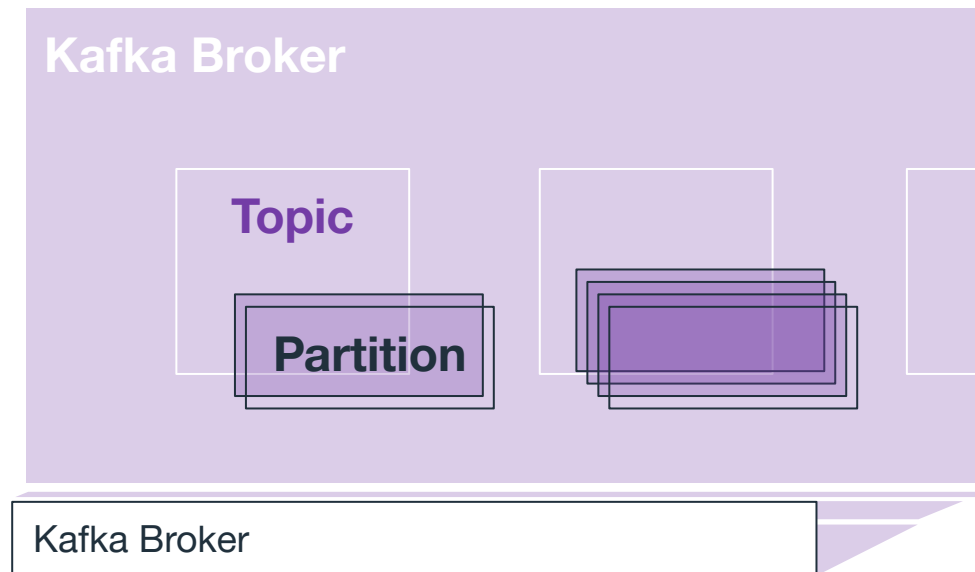
Cluster
overload

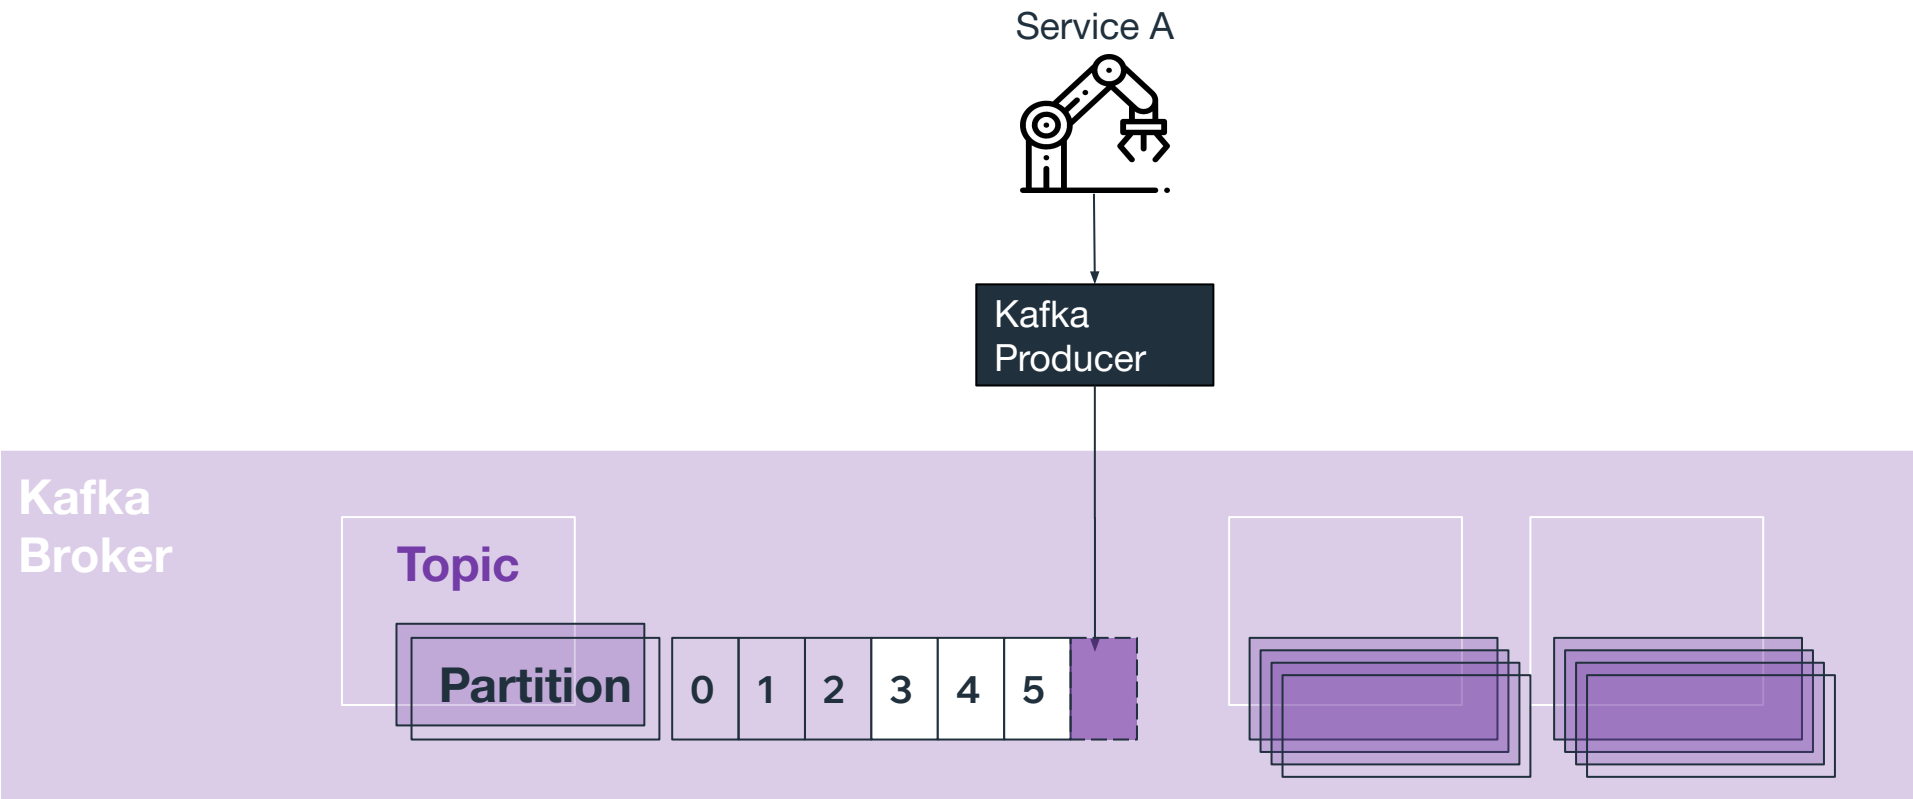
DC
awareness

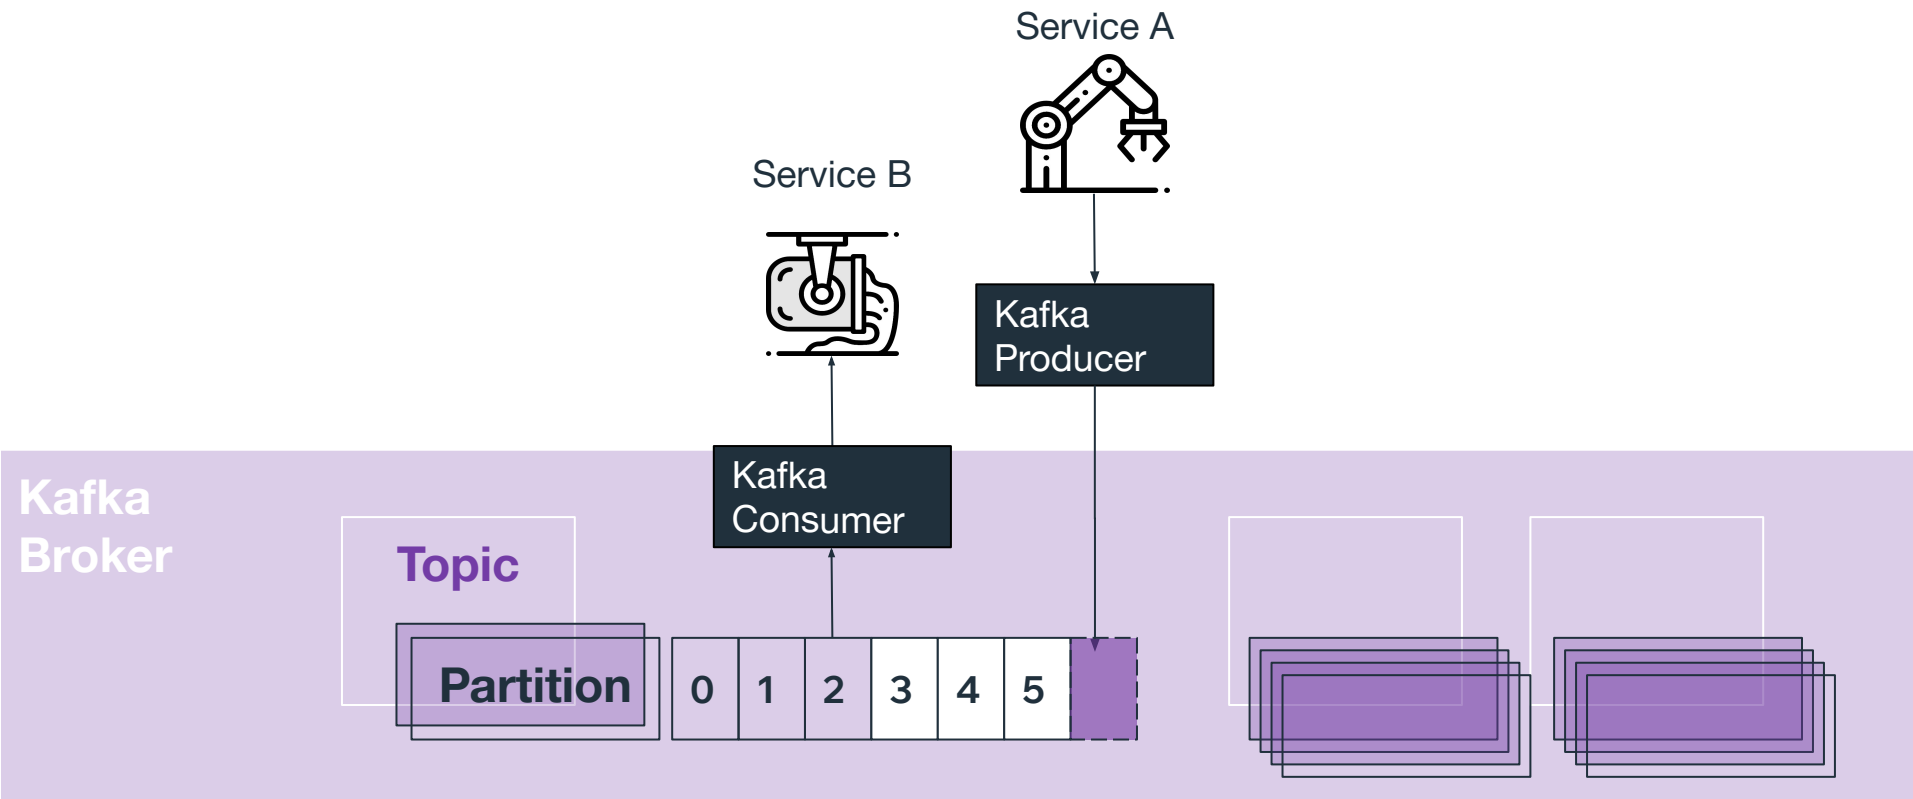
You Know Kafka



If You Don't Know Kafka







Growing Pains



Kafka

Event-driven
architecture

Large scale
traffic

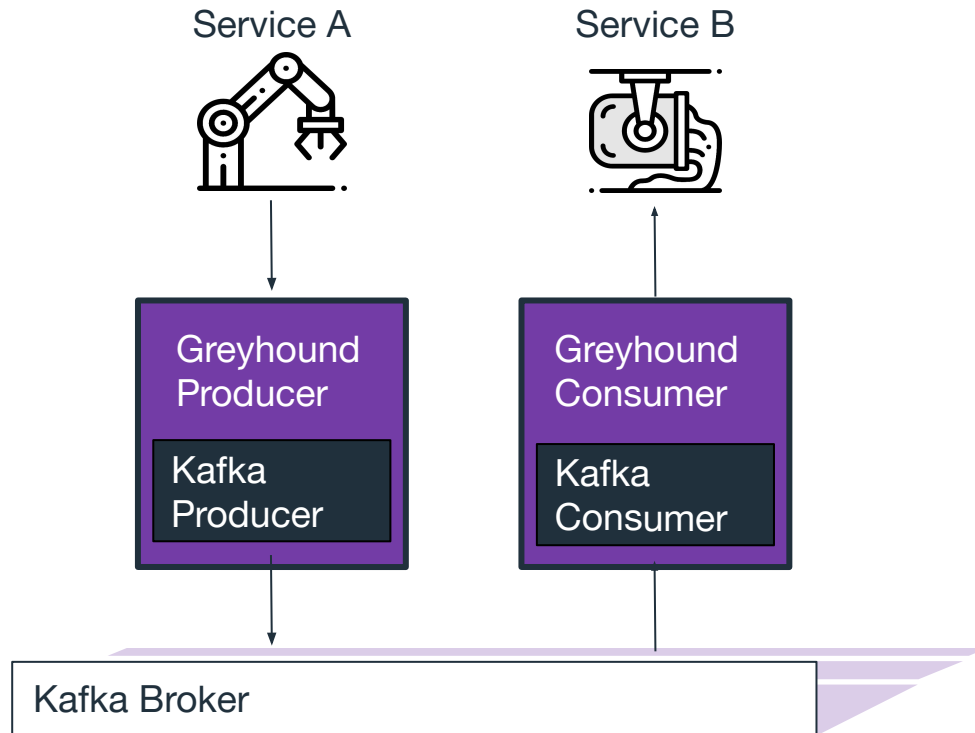
Hundreds of
developers

Diverse
tech stack

Cluster
overload

DC
awareness

Greyhound* Wraps Kafka






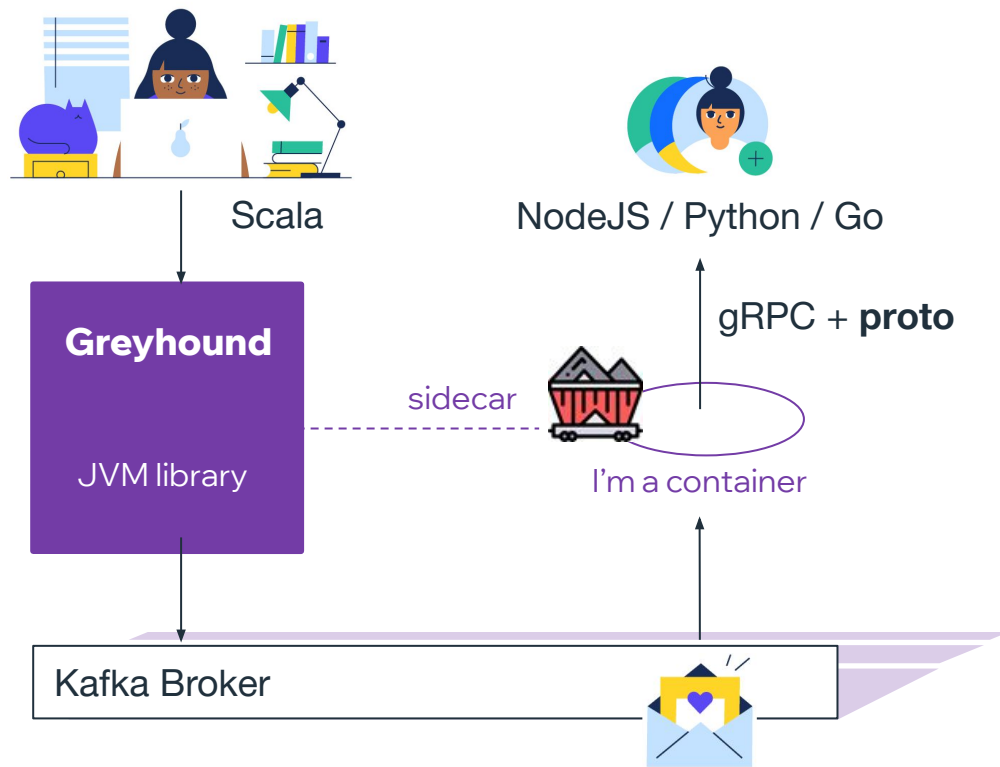

```
aGreyhoundConsumerSpec(groupId, messageHandler, topicName)
```

```
val producer = GreyhoundBuilder.resilientProducer(topicName)
```

Self Service

- ✓ Parallel Consumption
- ✓ Batch Consumer
- ✓ Consumer Retry Strategies
- ✓ Resilient Producer
- ✓ Context Propagation
- ✓ Metrics reporting

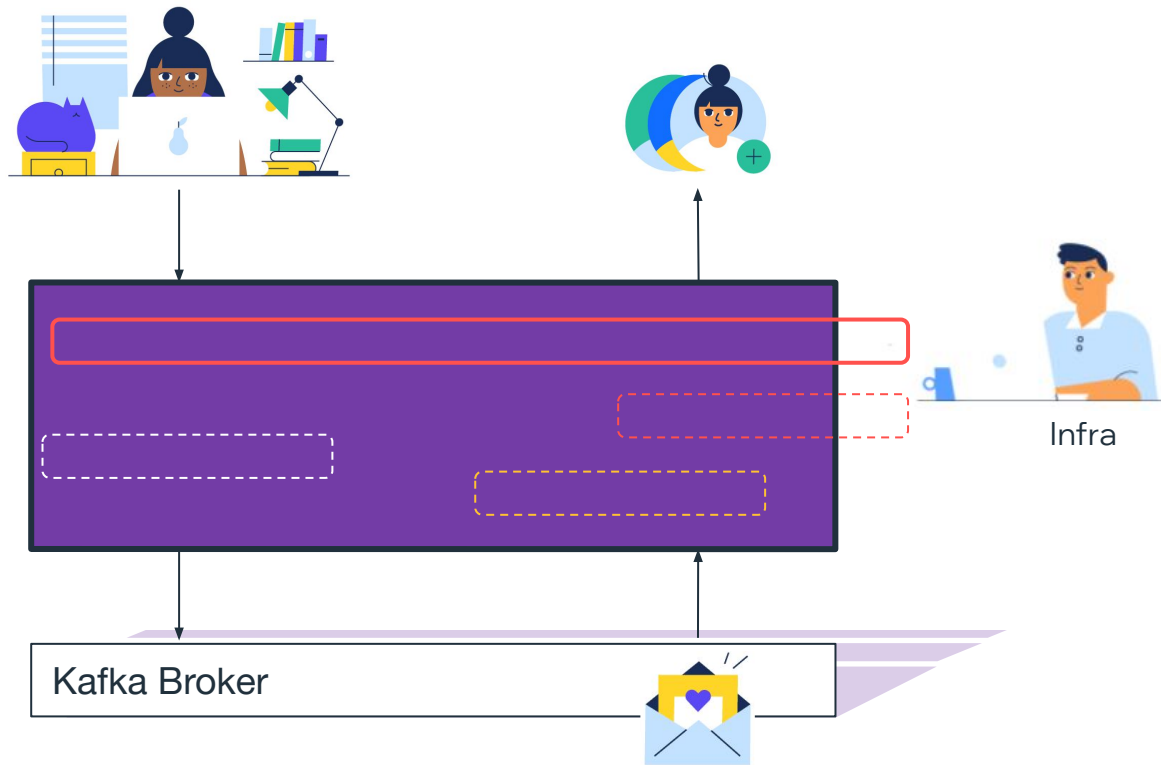




We use it for...



Domain Events
ETL to DWH
Task Scheduling
Async Queue
...



Growing Pains



Kafka

Event-driven
architecture

+ Greyhound

Large scale
traffic

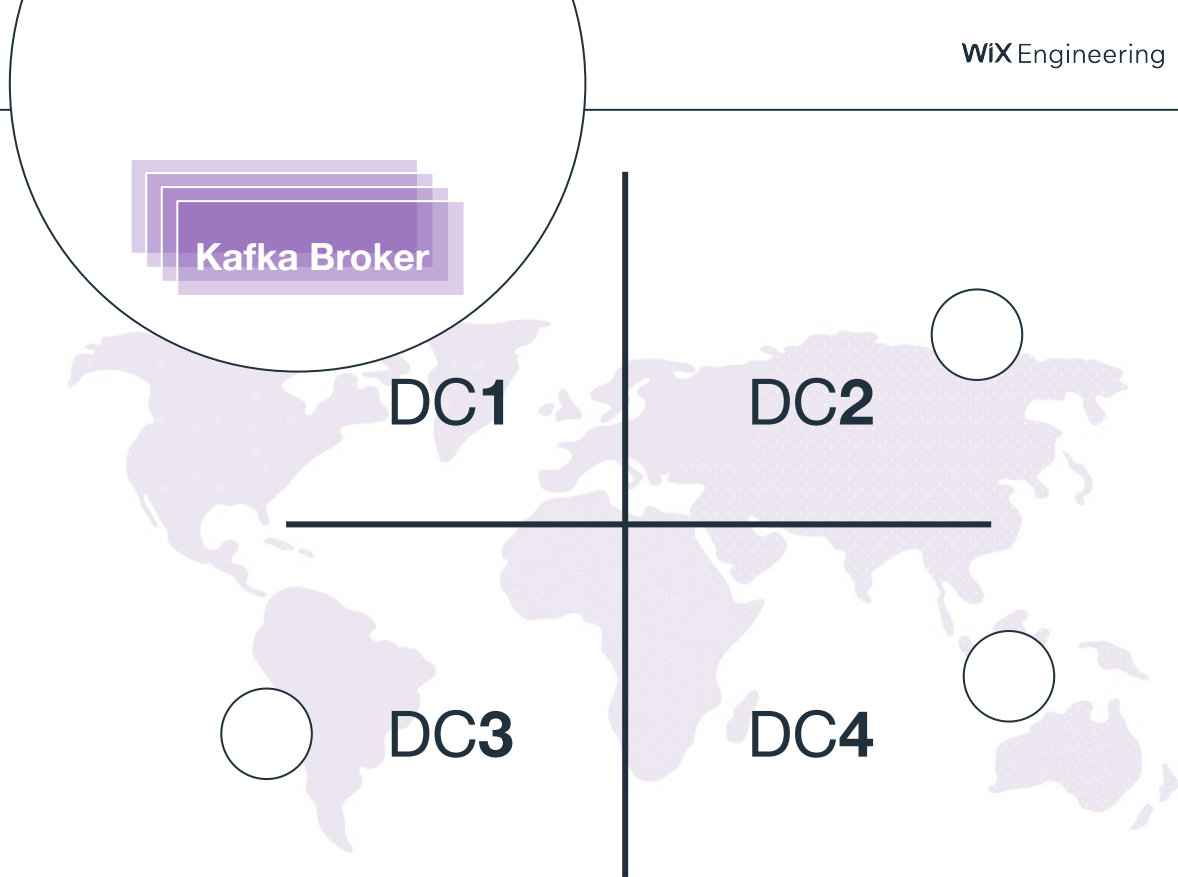
Hundreds of
developers

Diverse
tech stack

**Cluster
Overload**

DC
awareness

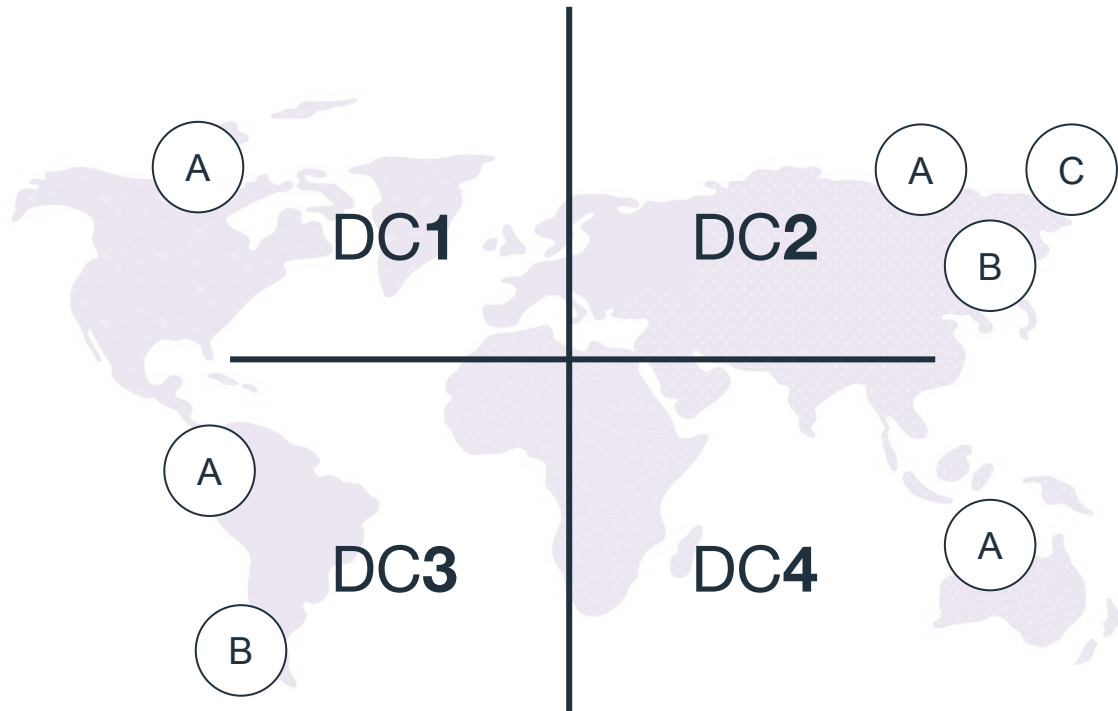
Single Cluster Overload



Single
Cluster
Overload




Multi
Cluster

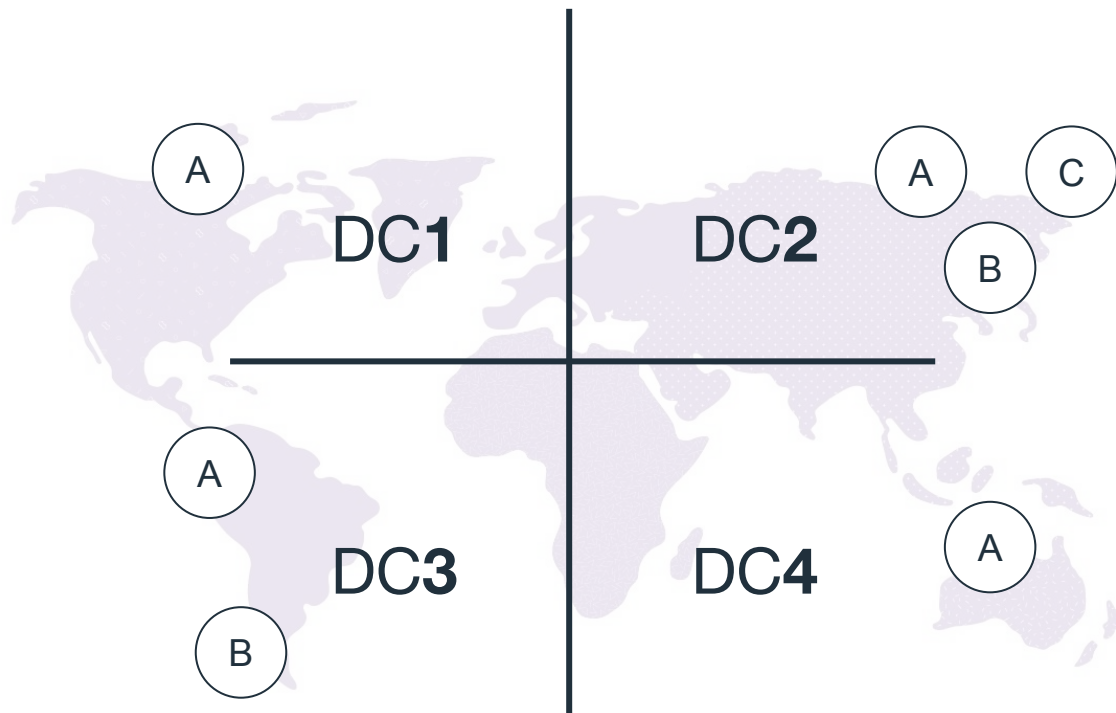


Kafka Cluster (by SLAs)

I want to produce
a domain event.



To which
Cluster?



Kafka Cluster (by SLAs)

```
aGreyhoundConsumerSpec(groupName, messageHandler, topicName)  
  .ClusterA
```

```
val producer = GreyhoundBuilder.resilientProducer(topicName,  
  _.ClusterB
```



Wix Greyhound Admin

Topics Configuration

natans@wix.com



Throughput and Lags

Live Stream

Topics Configuration

Topics Replication

Key-Value Stores

'[redacted]' Configuration last updated: 10/23/2021, 8:03:...

Refresh

Select a topic

+ New Topic

Cancel

Save

Topic: [redacted]

Config Key	Cluster A + DC1	Cluster B + DC1	Add DC +
partitions	[redacted]	[redacted]	
cleanup.policy ⓘ	[redacted]	[redacted]	...
delete.retention.ms ⓘ	[redacted]	[redacted]	...
min.insync.replicas ⓘ	[redacted]	[redacted]	...
retention.ms ⓘ	[redacted]	[redacted]	...
segment.bytes ⓘ	[redacted]	[redacted]	...
segment.ms ⓘ	[redacted]	[redacted]	...

Add new config entry

Growing Pains



Kafka

**Event-driven
architecture**

+ Multi Cluster

+ Greyhound

**Large scale
traffic**

**Hundreds of
developers**

**Diverse
tech stack**

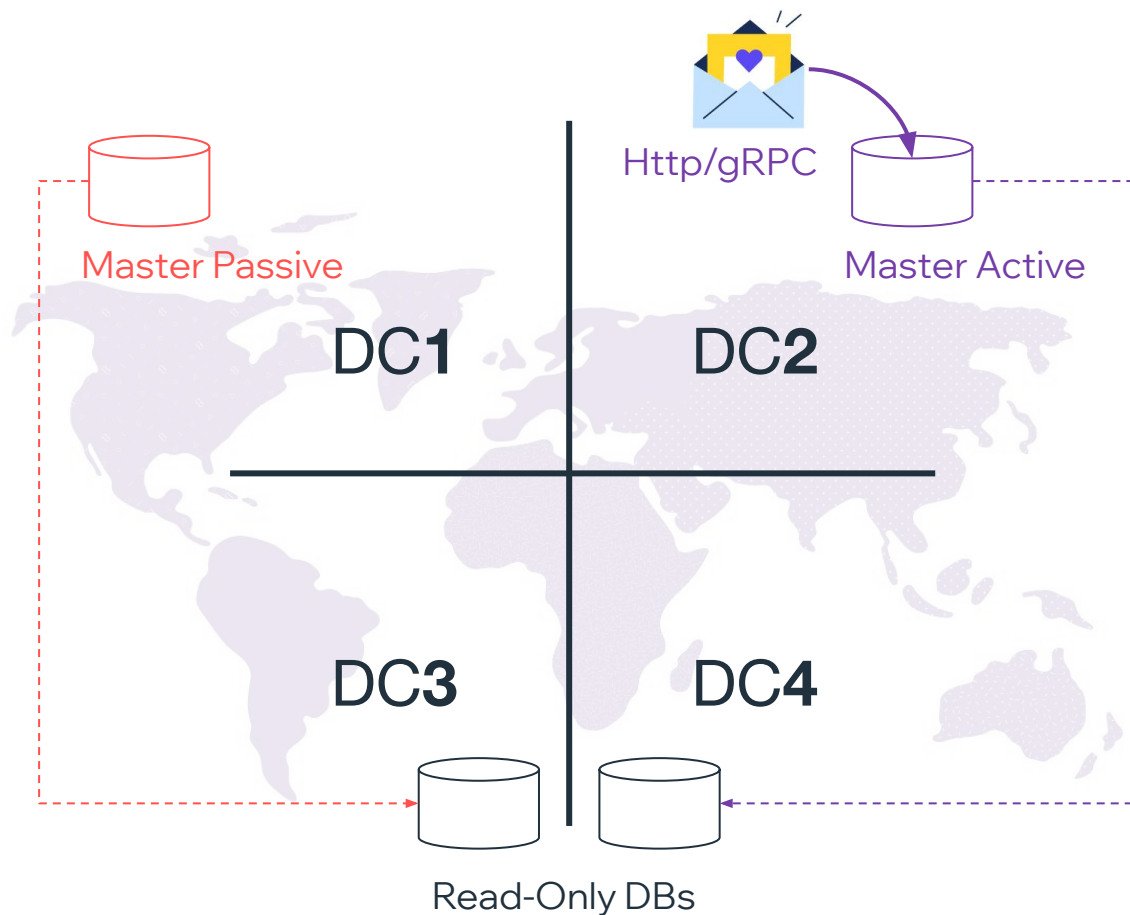
**Cluster
Overload**

**DC
awareness**

I want to
process the event
and save the data.



To which
DC?



Growing Pains



Kafka

**Event-driven
architecture**

+ Multi Cluster

+ Greyhound

**Large scale
traffic**


**Hundreds of
developers**

**Diverse
tech stack**

+ Cross DC services

**Cluster
Overload**

**DC
awareness**



Greyhound

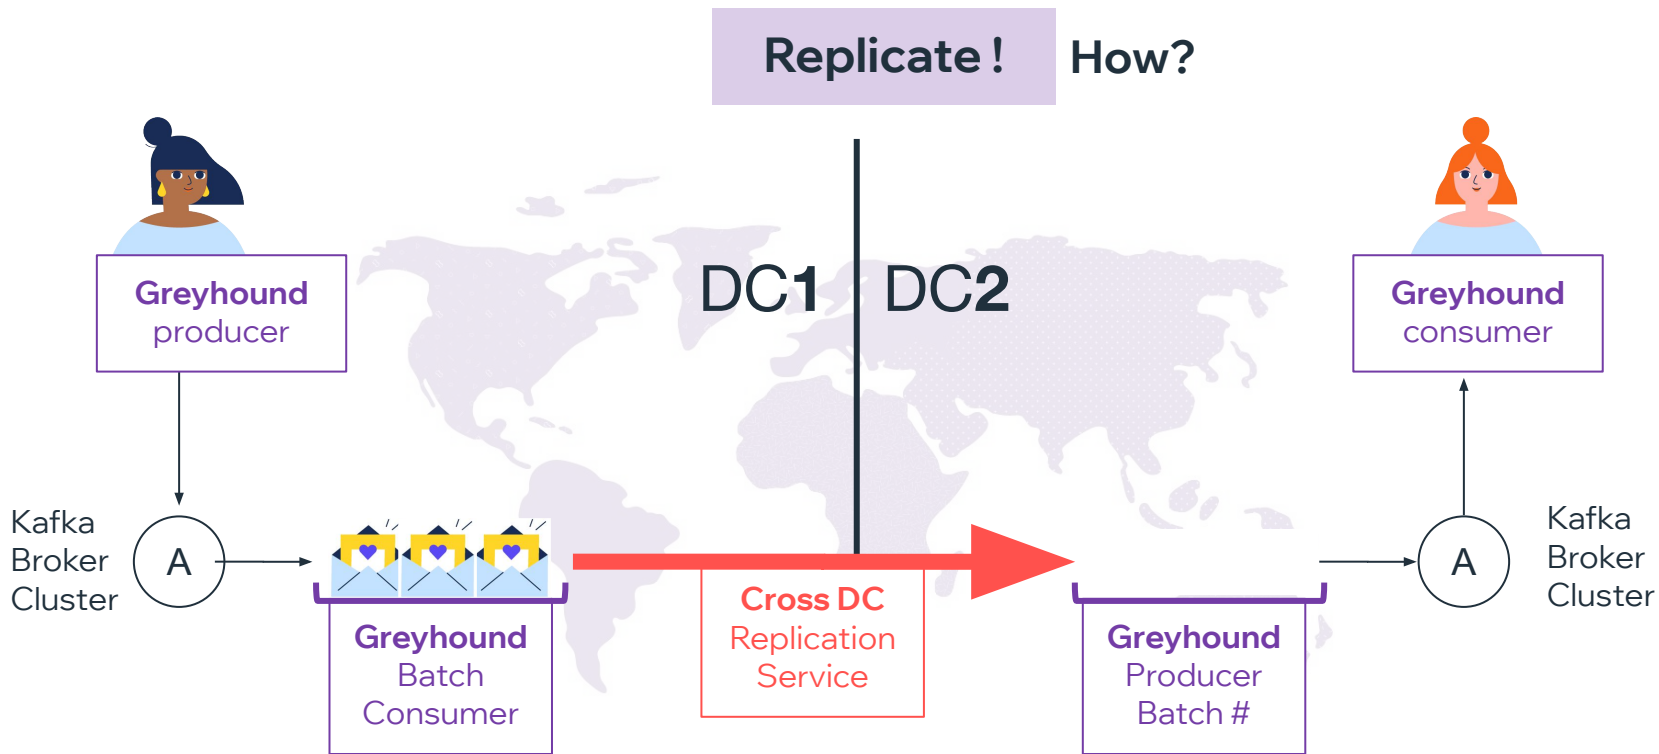
Kafka

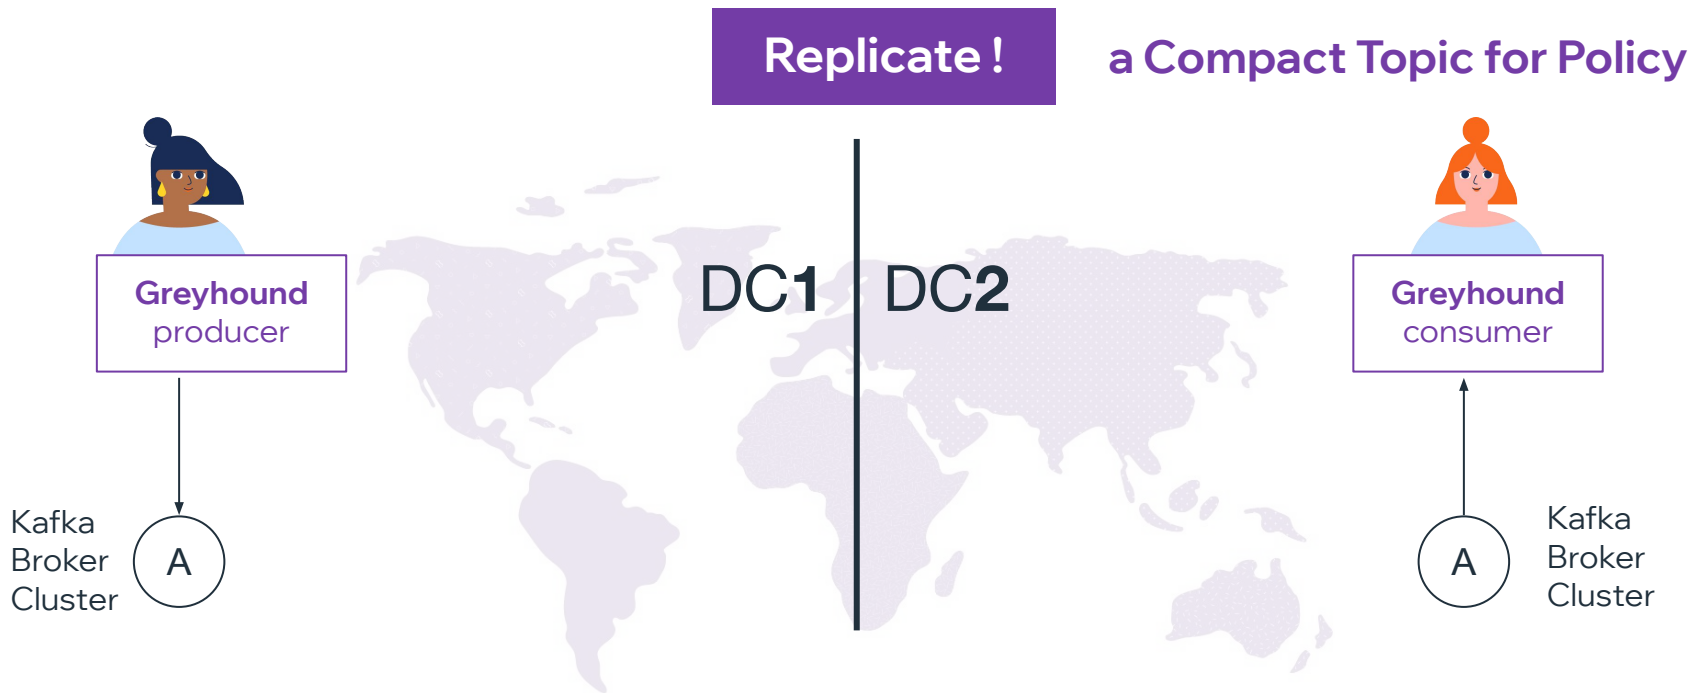
Cross-DC Self Services

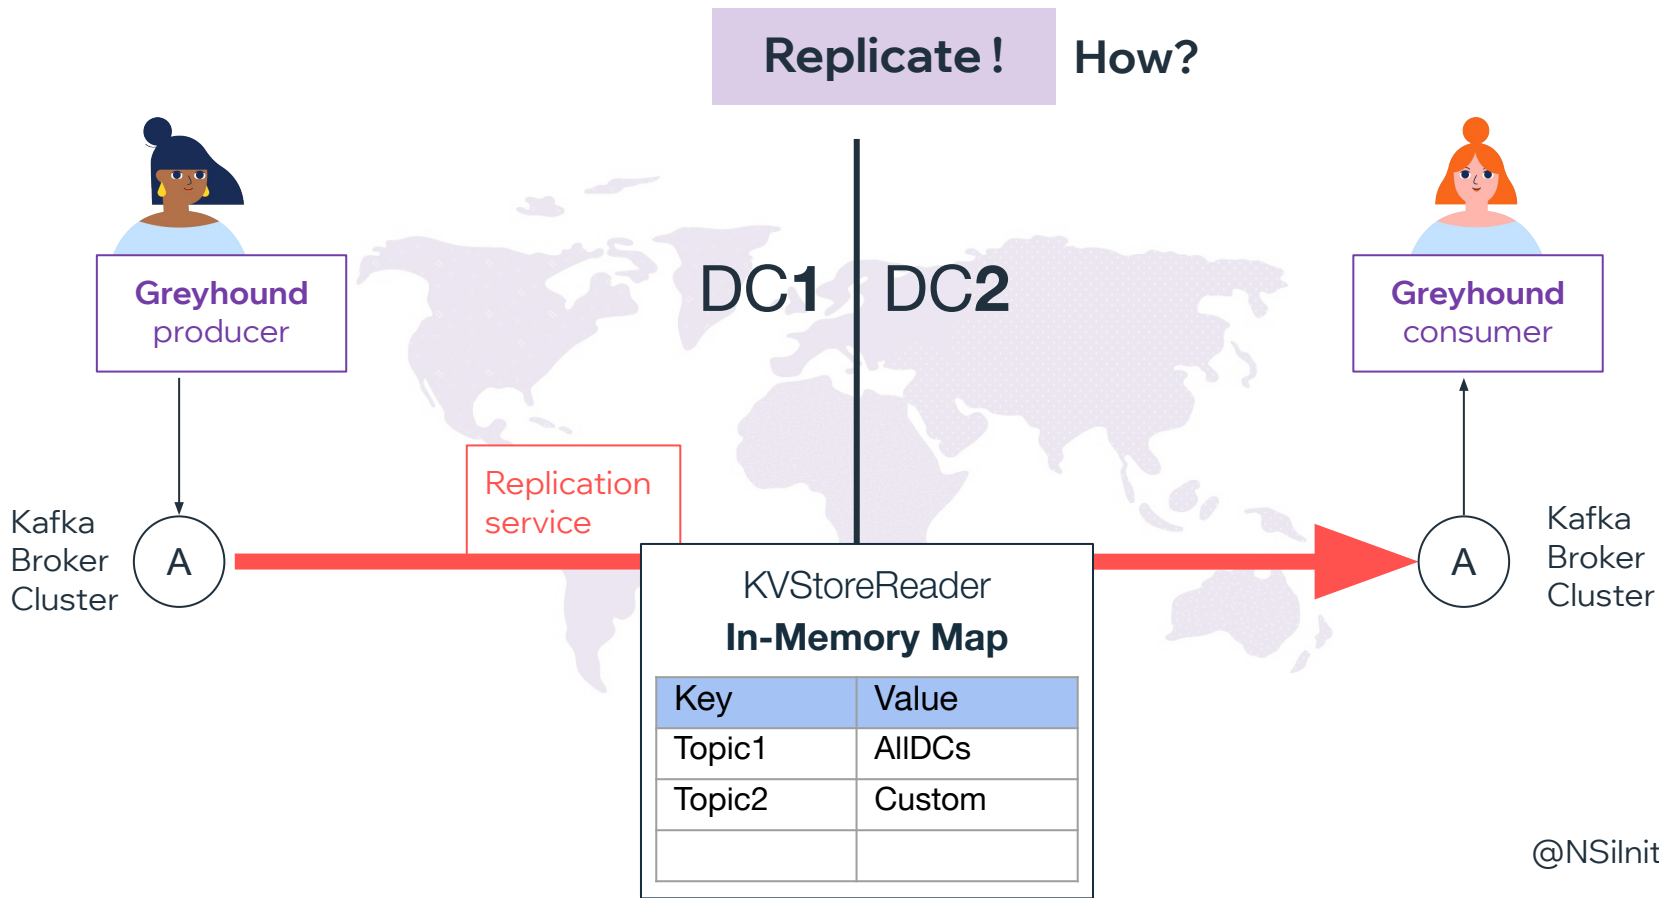
1. Replicate across DCs
2. Produce to Main DC

[WIP] Just Produce 











Wix Greyhound Admin

Topics Replication

natans@wix.com



Throughput and Lags

Live Stream

Topics Configuration

Topics Replication

Key-Value Stores

Select a topic

Topic:

DCs: Cluster: Legacy

Policy:

All DCs

Update

Refresh

Cluster:

ANY



ALL



ALL



ALL



ALL



in	⇒		⇒	
	Lag	Msg/min	Lag	Msg/min
116.1	0 ▲33.76	▼33.76	0 ▲33.71	▼33.71
	0 ▲43.94	▼43.94	0 ▲39.05	▼39.05
28.97	X		0 ▲0	▼0
51.34	0 ▲55.24	▼55.24	X	



Wix Greyhound Admin

Topics Replication

natans@wix.com



Throughput and Lags

Live Stream

Topics Configuration

Topics Replication

Key-Value Stores

Select a topic

Topic: DCs: Cluster: Legacy

Policy:

All DCs



Unreplicated

All DCs


Main DCs

Custom Rules

Update

Refresh

	→ <input type="text"/>			→ <input type="text"/>	
	Lag	Msg/min		Lag	Msg/min
1	0	▲33.76 ▼33.76	0	▲33.71 ▼33.71	
	0	▲43.94 ▼43.94	0	▲39.05 ▼39.05	
7	X		0	▲0 ▼0	
4	0	▲55.24 ▼55.24	X		



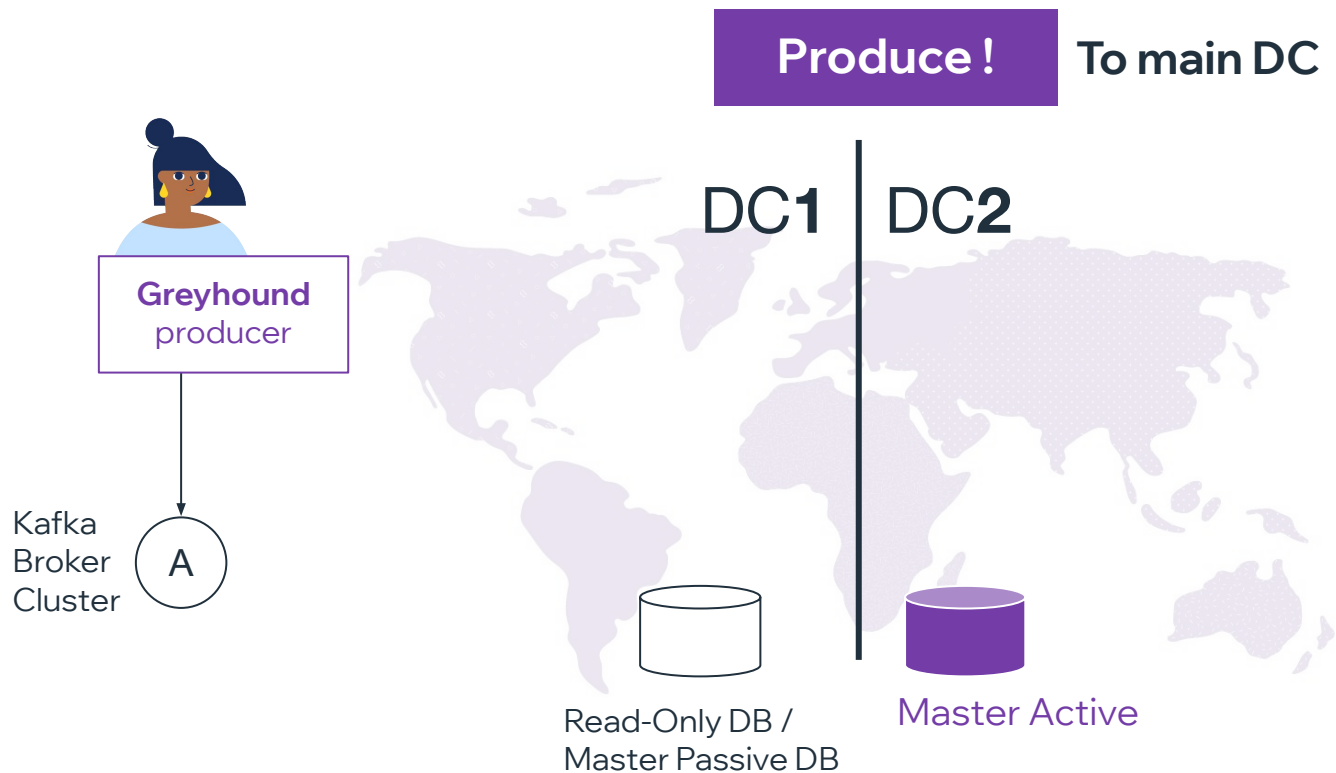
Greyhound

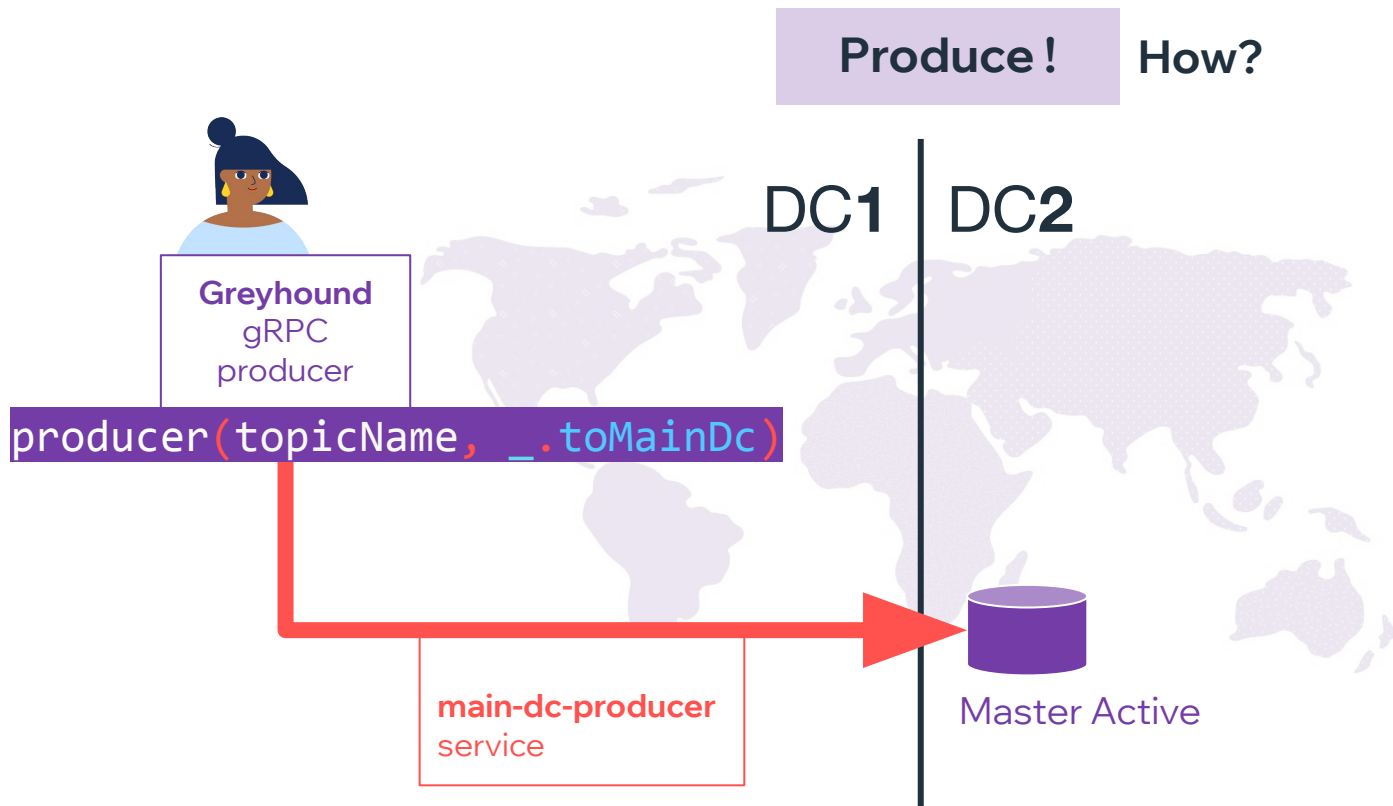
Kafka


Cross-DC Self Services ☐

1. ✓ Replicate across DCs
2. **Produce to Main DC**

[WIP] Just Produce ❤️







Greyhound

Kafka

Cross-DC Self Services ☐

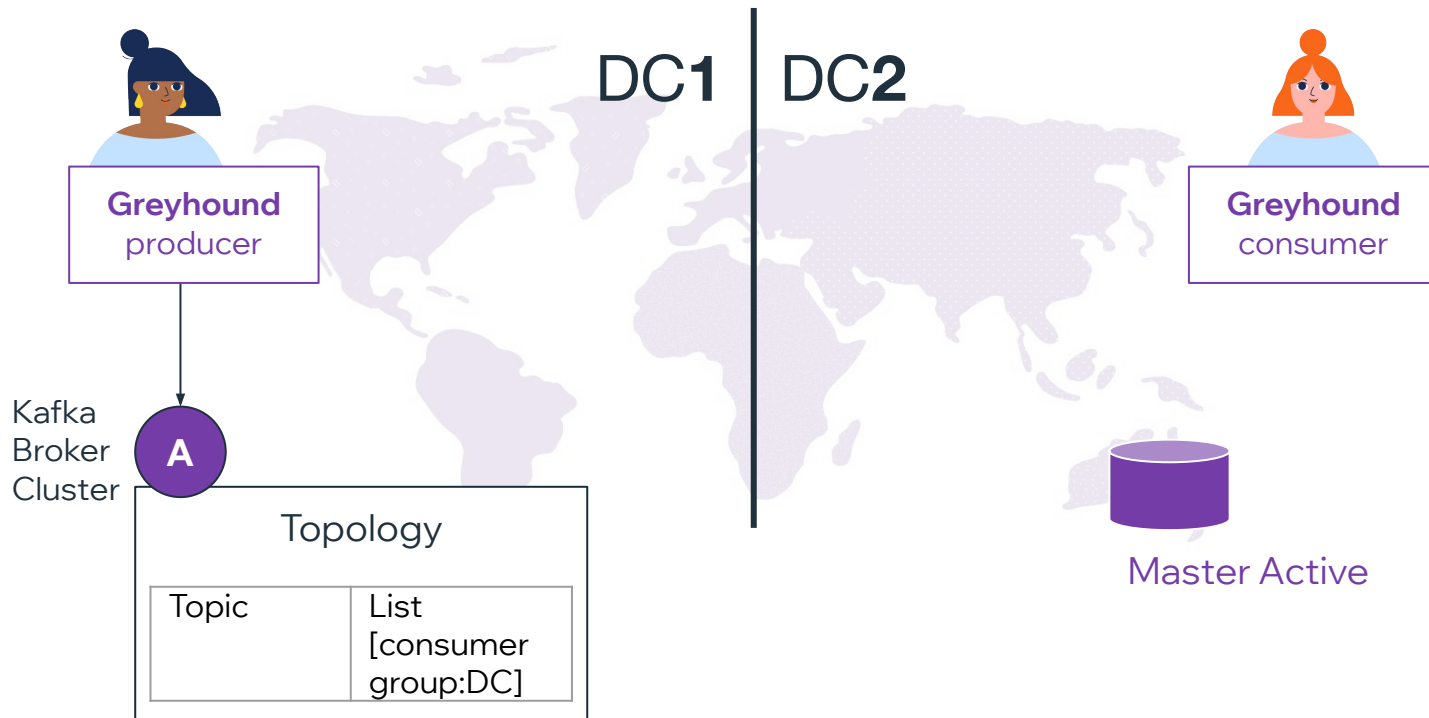
1. ✓ Replicate across DCs
2. ✓ Produce to Main DC

[WIP] Just Produce ❤️

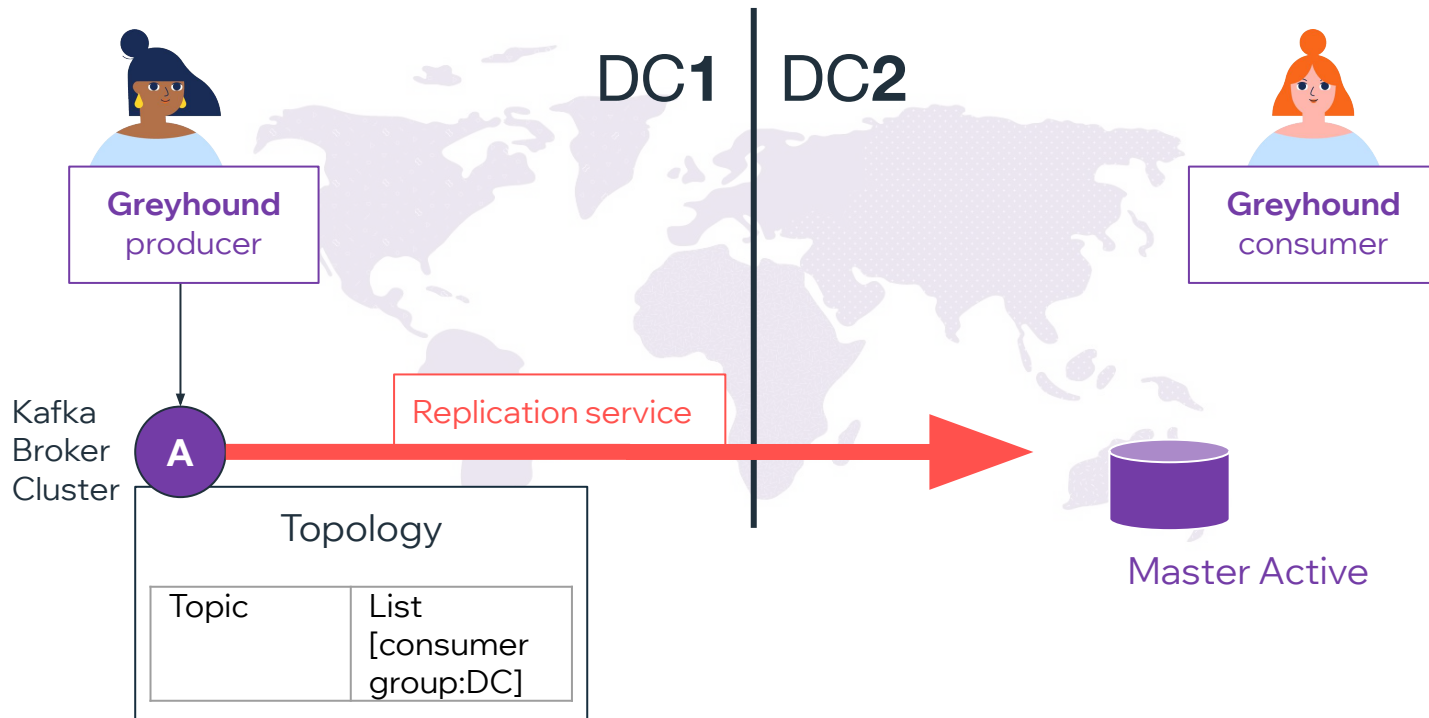
[WIP] Just Produce ❤️ To where needed




[WIP] Just Produce ❤️ To where needed



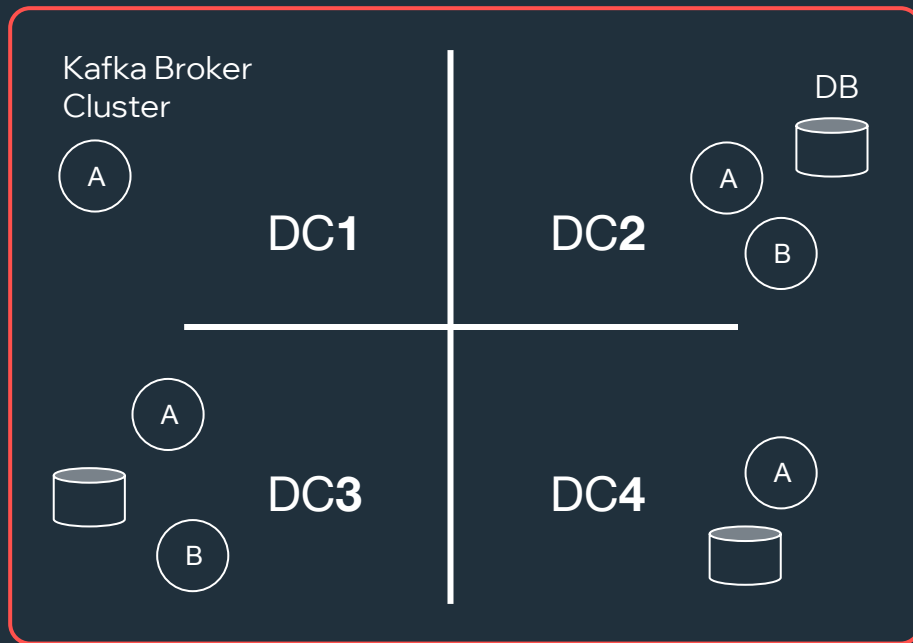

[WIP] Just Produce ❤️ To where needed




* reintroduce decoupling



Just Produce ❤️



To a Single Logical
Kafka Cluster.




developer

Any
Tech Stack



Anywhere
You Need

github.com/wix/greyhound




developer



☰ README.md ✎

Greyhound

maven central 0.2.0 CI passing Scala/Java High-level SDK for [Apache Kafka](#).



Why Greyhound?

Kafka is shipped with a Java SDK which allows developers to interact with a Kafka cluster. However, this SDK consists of somewhat low-level APIs which are difficult to use correctly. Greyhound seeks to provide a higher-level interface to Kafka and to express richer semantics such as parallel message handling or retry policies with ease.

You can read more about it on this blog post - "[Building a High-level SDK for Kafka: Greyhound Unleashed](#)"

The Open-sourced version of Greyhound is still in stages of initial rollout, so APIs might not be fully stable yet.

Thank You!

👉 slideshare.net/NatanSilnitsky

Any questions?

Q&A

👉 slideshare.net/NatanSilnitsky

Any questions?