



Evolving from Messaging to Data in Motion

Confluent for Messaging Modernization

Perry Krol, Manager Solutions Engineering CEMEA

Every Business is Becoming Software



Then

Banking



Mortgage



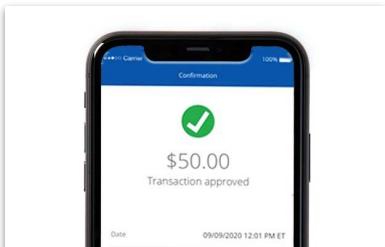
Taxi



Grocery



Now



Example: Loan Application Process



Software-using



1-2 Weeks

Software-defined



Seconds



Unlocking Value from Data, in Real-time, is Imperative to Transformation



Purchases



Workflows



App data



Security events



Interactions



Search data



IoT and so much more



Employees



Customers

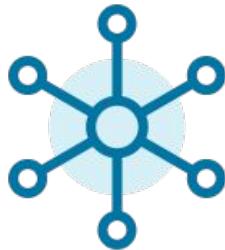
Enterprises are Undertaking Multiple Initiatives to Enable Transformation



Microservices



Cloud



Automation



Machine Learning



The Problem: your use of data has changed
but the supporting infrastructure has not!

Paradigm for Data-at-Rest: Databases





Paradigm for Data Movement: Messaging Middleware



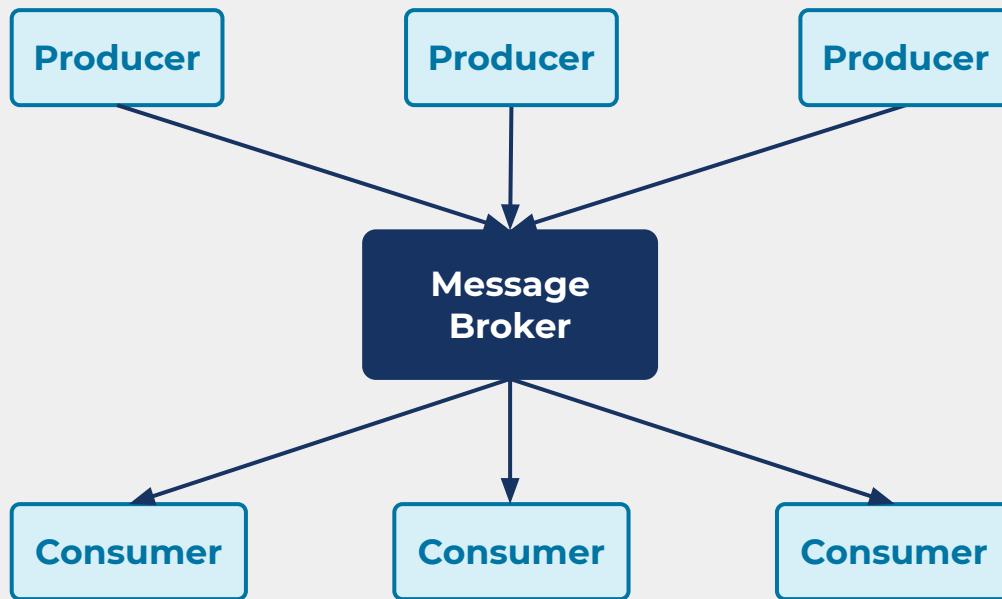


“Message Exchange”



Message brokers were originally architected as centralised systems

Traditional Hub & Spoke Message Broker

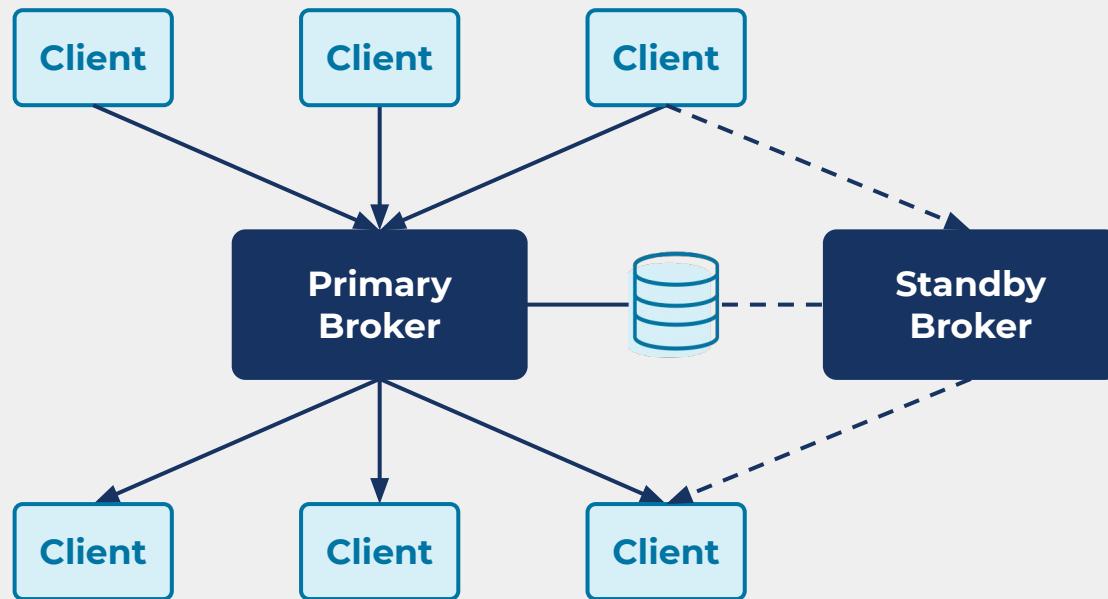


Single point of failure,
resulting in **low fault-tolerance and high downtime**



For improved fault-tolerance, brokers are often deployed in high availability pairs

High Availability Pairs of Brokers

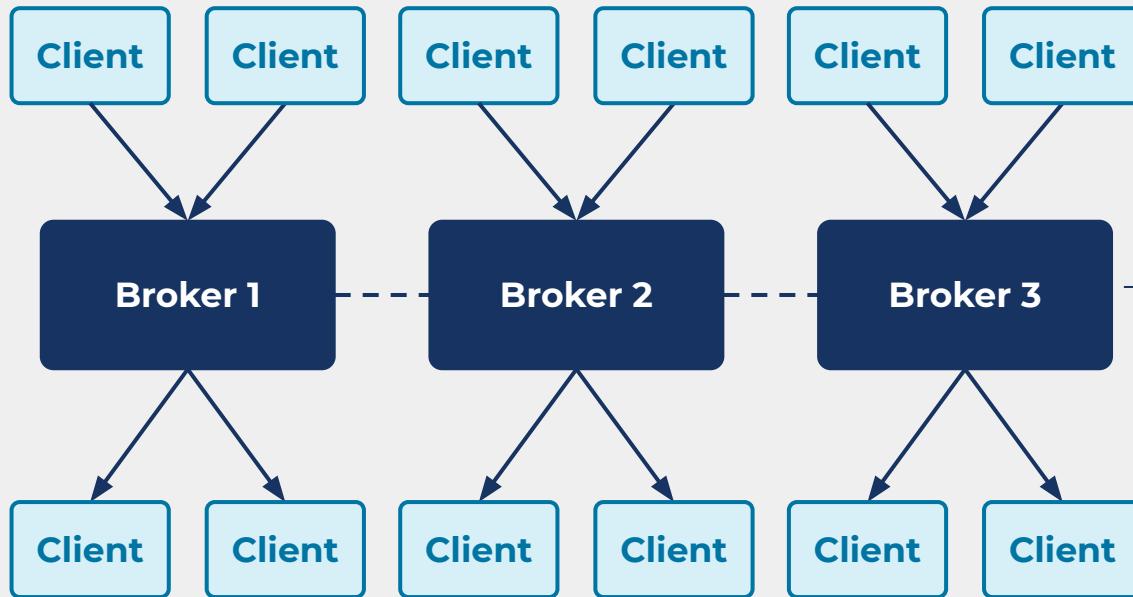


Clients still **only connect to the active brokers** though, so this solution **lacks scalability...**



Over time, some brokers evolved to multi-node networks, but clients still connect to one broker

Interconnected group of Message Brokers



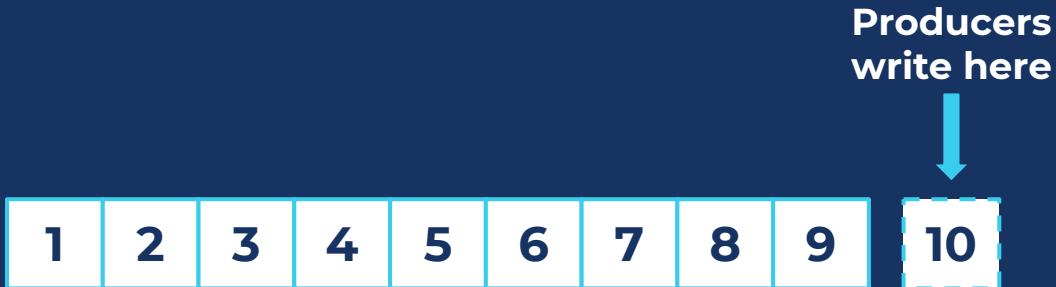
Because clients still **only connect to one broker**, this solution **still does not provide horizontal scalability**

Managed independently



The world is not Ephemeral!

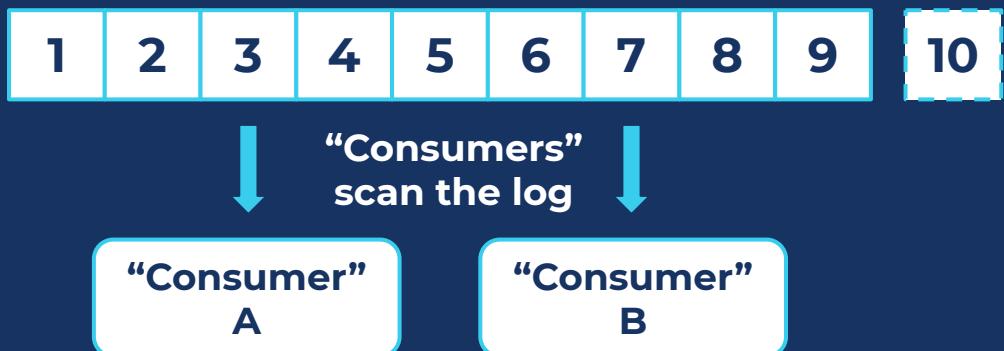
Let's use an immutable log to share data!



Kafka producers write to an append-only, immutable, ordered sequence of messages, which is always ordered by time

- Sequential writes only
- No random disk access
- All operations are $O(1)$
- Highly efficient

A log is like a queue, but re-readable :-D



“Better than a queue”-like behavior as Kafka consumer groups allows for parallel in-order consumption of data, which is something that shared queues in traditional message brokers do not support.

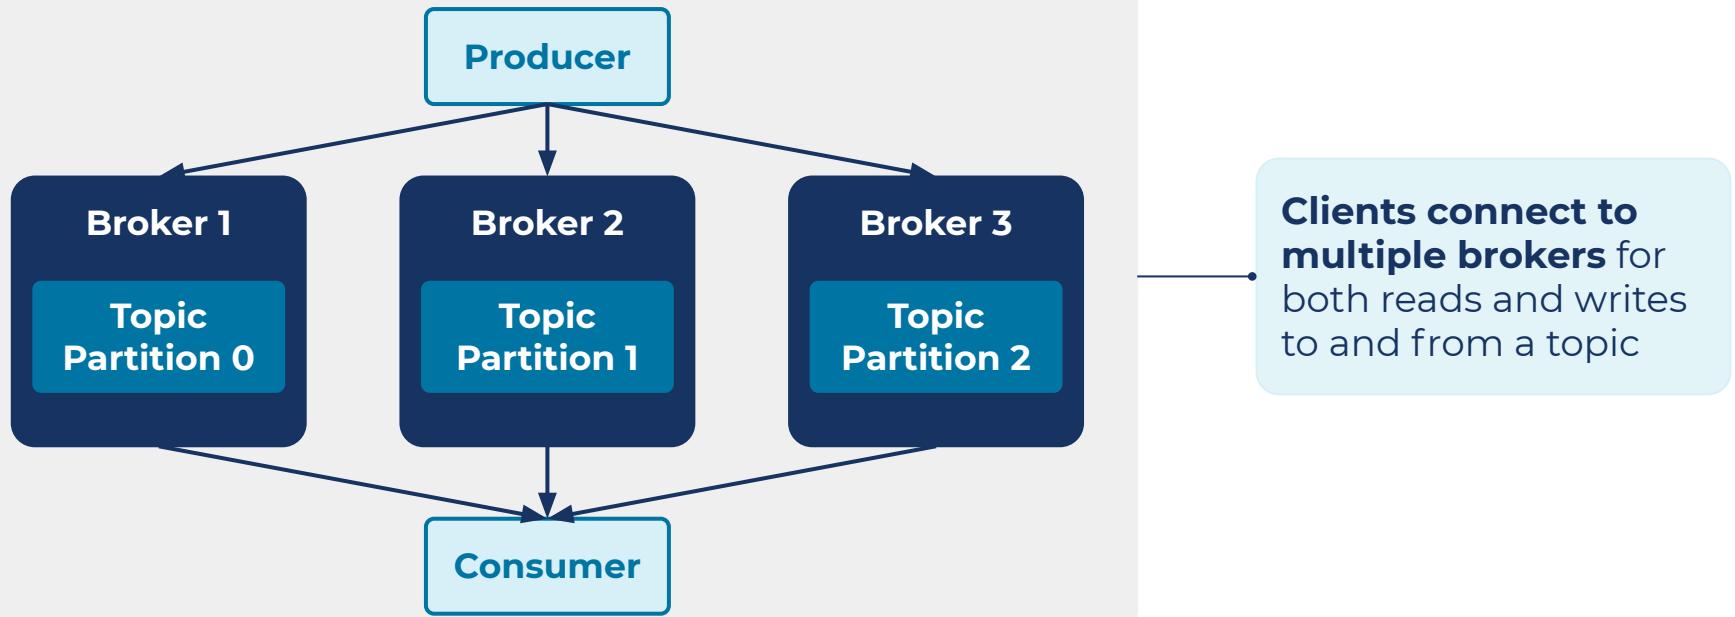
- Sequential reads only
- Start at any offset
- All operations are $O(1)$
- Highly efficient

Slow consumers don't back up the broker: **THE STREAM GOES ON.**



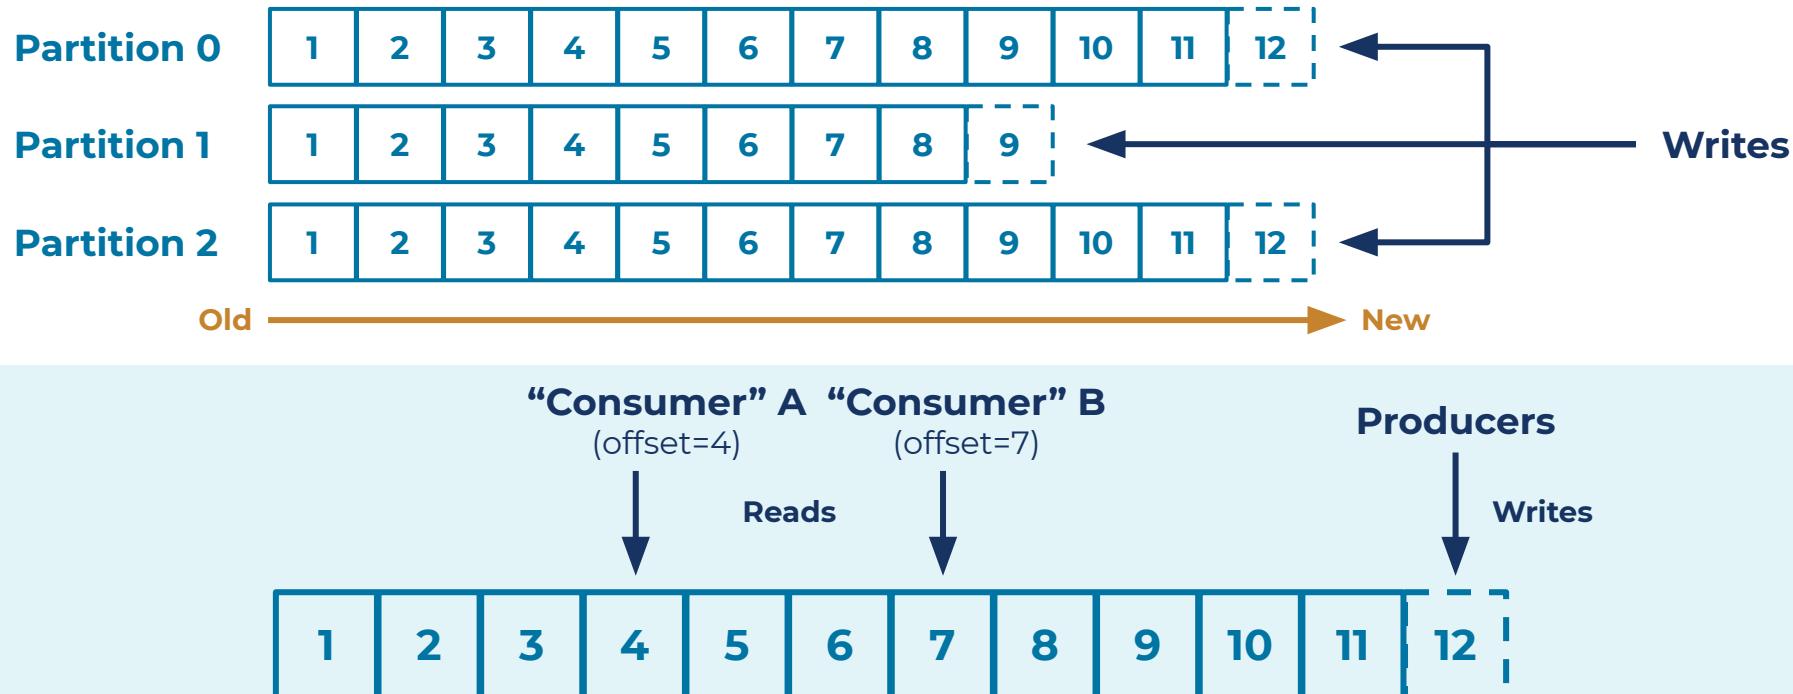
Kafka takes a different approach by partitioning topics across the brokers in a cluster

Kafka Cluster with Topic Partitions & Multiple Client Connections



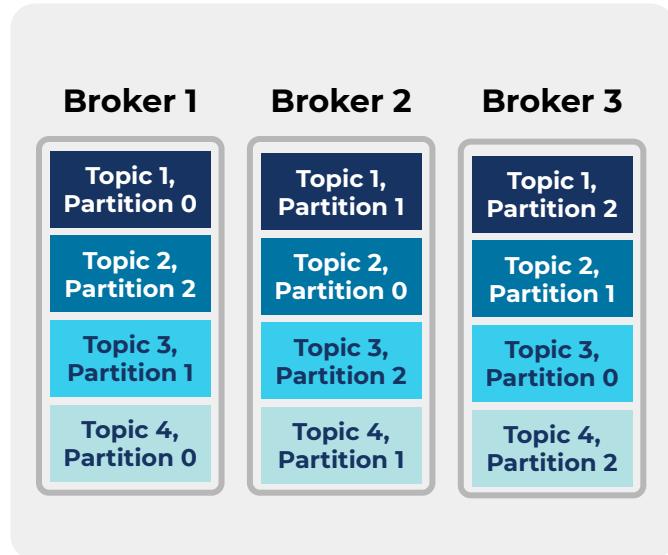


Kafka topics are designed as a commit log that captures events in a durable, scalable way

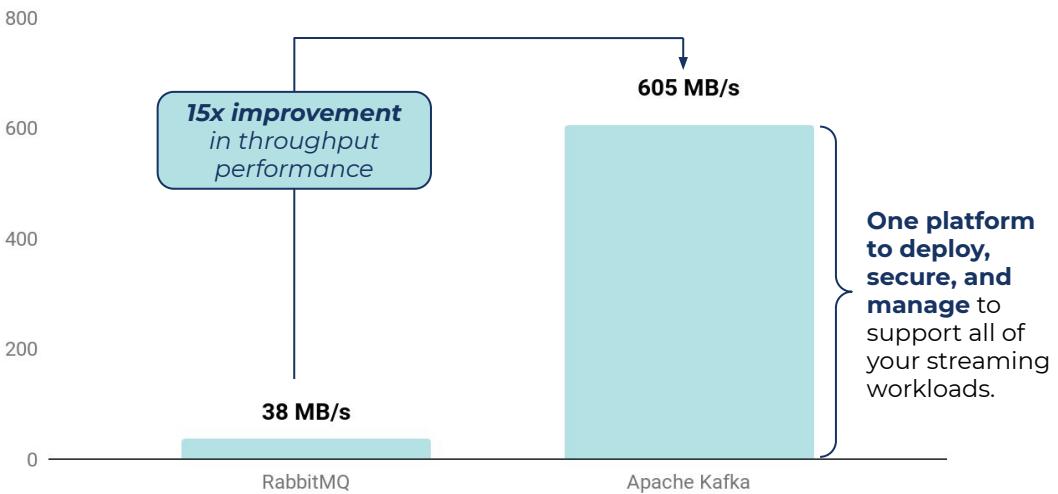




Partitioning topics enables greater horizontal scalability and enterprise-scale throughput



Peak throughput (MB/s): Kafka vs RabbitMQ





How else is Kafka different from traditional messaging queues?



Topic partitions are replicated to maximize fault-tolerance

In addition to partitioning topics, each partition can be replicated across multiple brokers to **ensure high uptime even if a broker is lost.**



Producers and consumers scale independently from brokers

Production and consumption rates (e.g. spike or slow consumer issue) have no effect on the broker. **THE STREAM GOES ON.**



Event streams can be enriched in real-time with stream processing

ksqlDB and Kafka Streams **enable event streams to be processed “in-flight”** rather than with a separate batch solution

Message System Benefits

+

Real-time (low latency)

Broad Adoption (familiar technology)

Message System Drawbacks

-

Lacks Common Structure

No Stream Processing

No Persistence After Consumption

Low Fault Tolerance at Scale

Slow Consumers Drag Performance

Complexity and Technical Debt

Message Systems



Real-time (low latency)

Broad Adoption (Apache)

Confluent



Real-time
Broad Adoption



Schema Registry
Stream Processing
Durable & Persistent
Elastically Scalable
Reliable

Common Drawbacks



Common Structure

Stream Processing

After Consumption

Tolerance at Scale

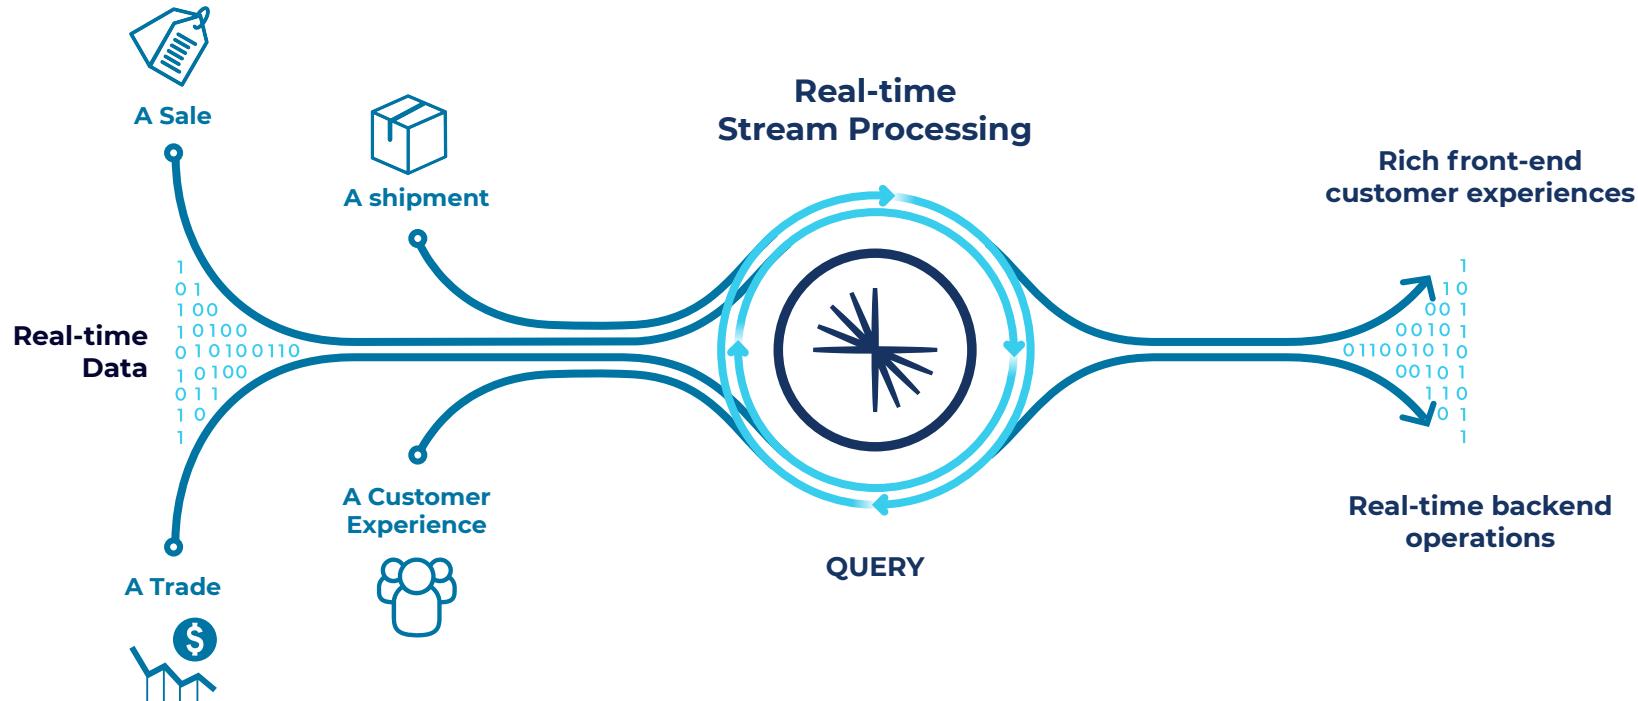
Drag Performance

Complexity and Technical Debt

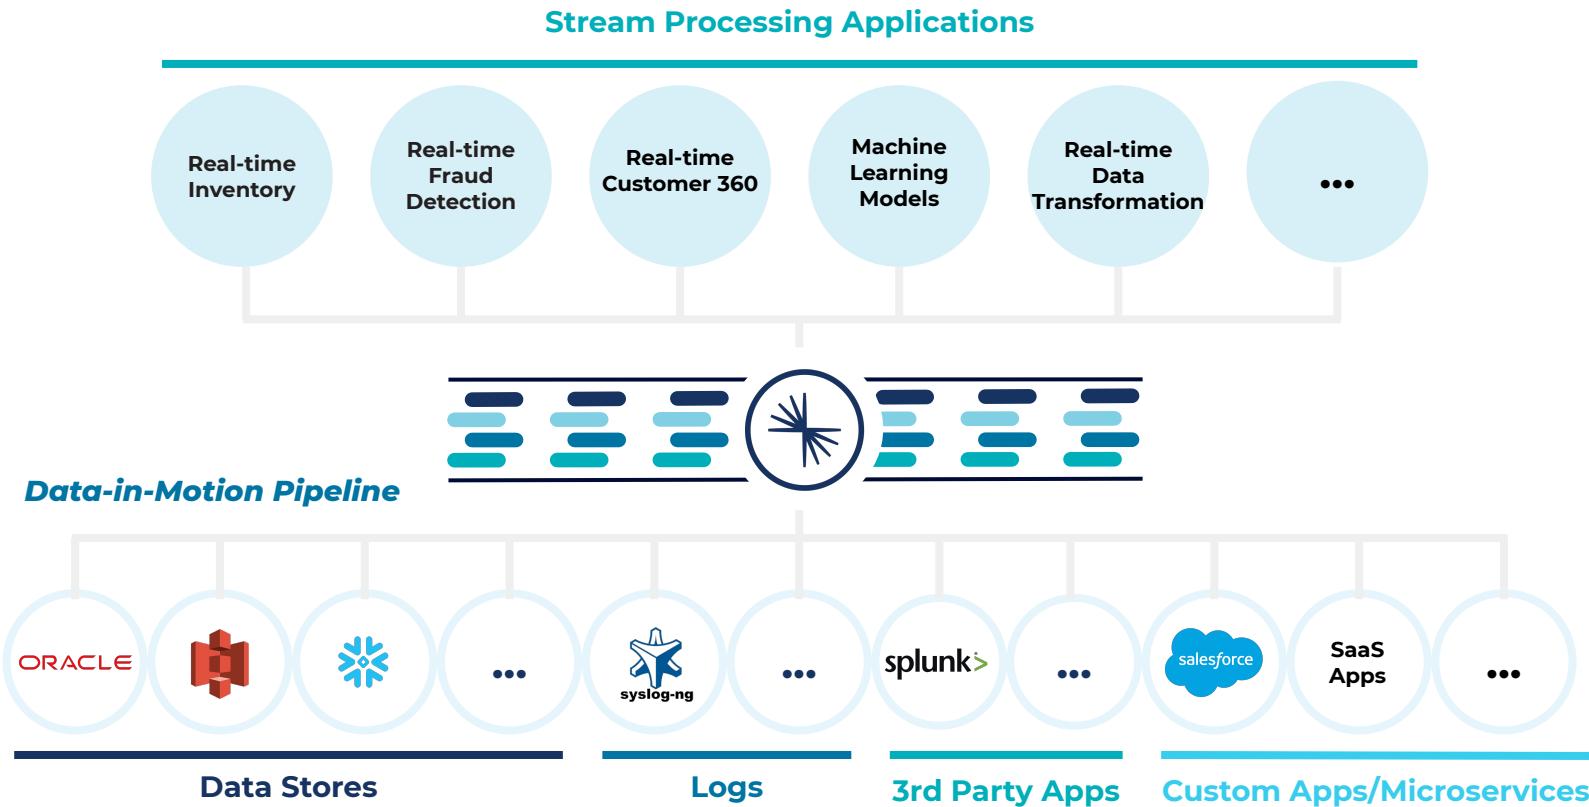


Introduction to Confluent

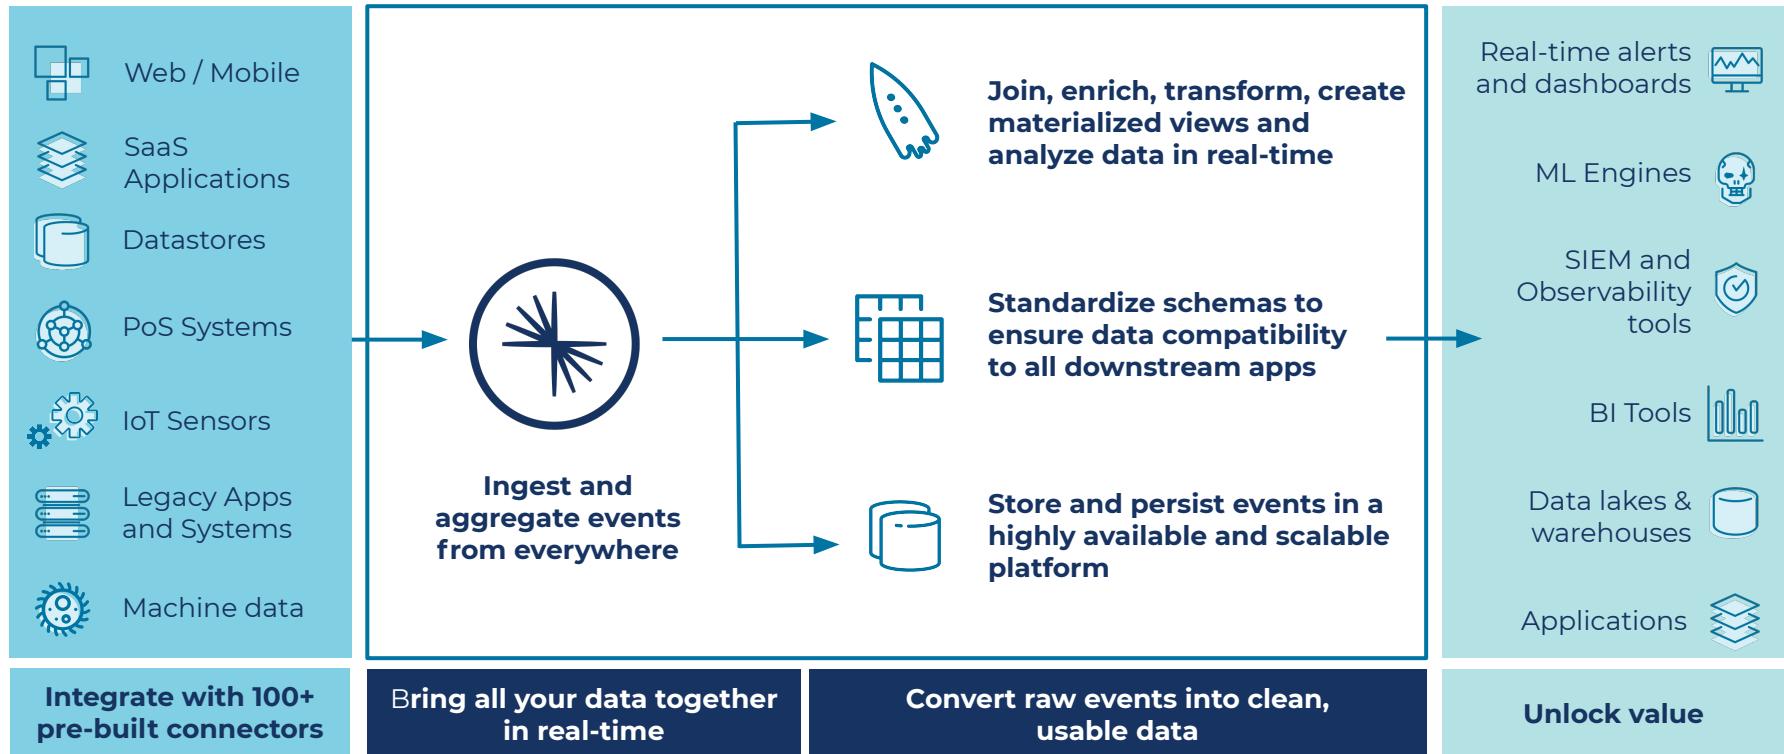
Confluent: A New Paradigm for Data-in-Motion



Implication: Central Nervous System for Enterprise



How Confluent Works



Confluent Modernizes Enterprise Messaging Infrastructure



Build a modern architecture

Enable real-time workloads and applications



Achieve any scale

Future-proof data architecture and achieve high performance and availability at scale



Migrate with ease

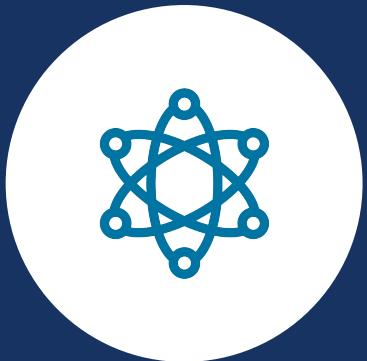
Easily augment or migrate apps to Confluent



Reduce TCO and tech debt

Future proof cloud-ready architecture

Build a modern architecture



Confluent offers a way to modernize your data architecture and accelerates your developer velocity



Cloud-native design

Embrace modern DevOps and easily deploy, manage, and scale your infrastructure, whether on your private cloud or a public cloud



Simple integrations

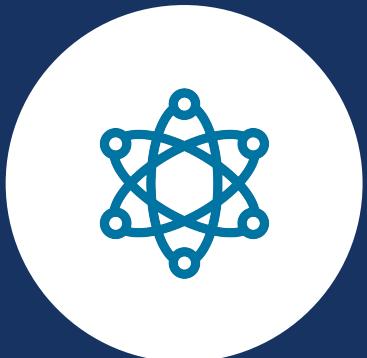
Use best-of-breed solutions across your architecture and easily integrate them with Confluent connectors, reducing your reliance on a single infrastructure vendor



Message replay

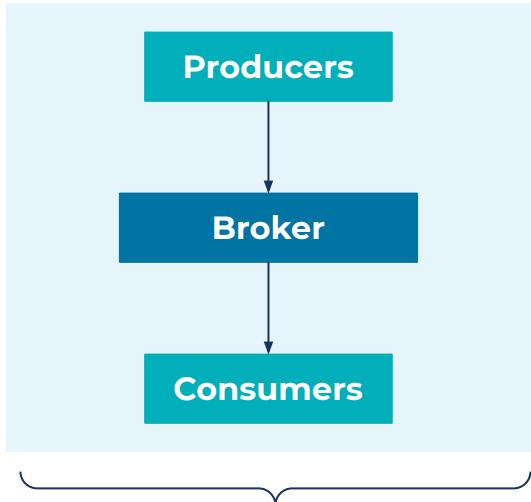
Easily replay messages, provide event sourcing and command sourcing for error handling and eliminate complexity for resending messages

Build a modern architecture



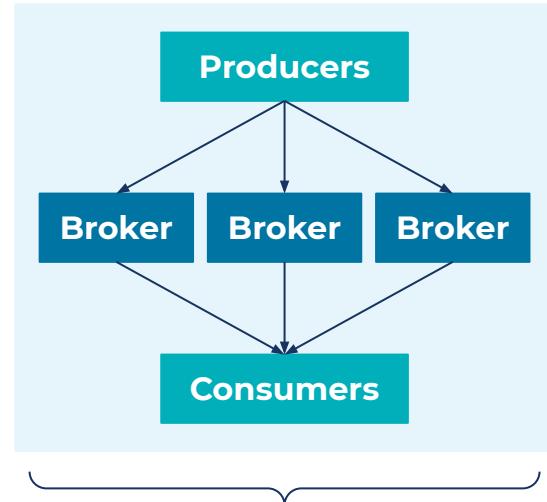
Confluent offers a way to modernize your data architecture and accelerates your developer velocity

Legacy messaging brokers are monolithic



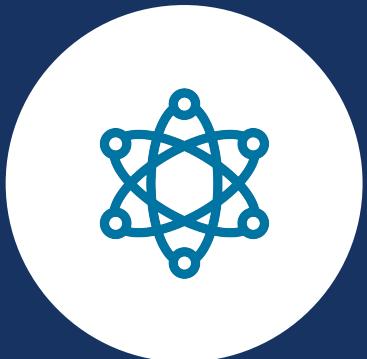
More point of failure, lack of horizontal scalability & higher risk of downtime: poor fit for mission-critical use cases

Kafka supports containerization and multi-DC deployments



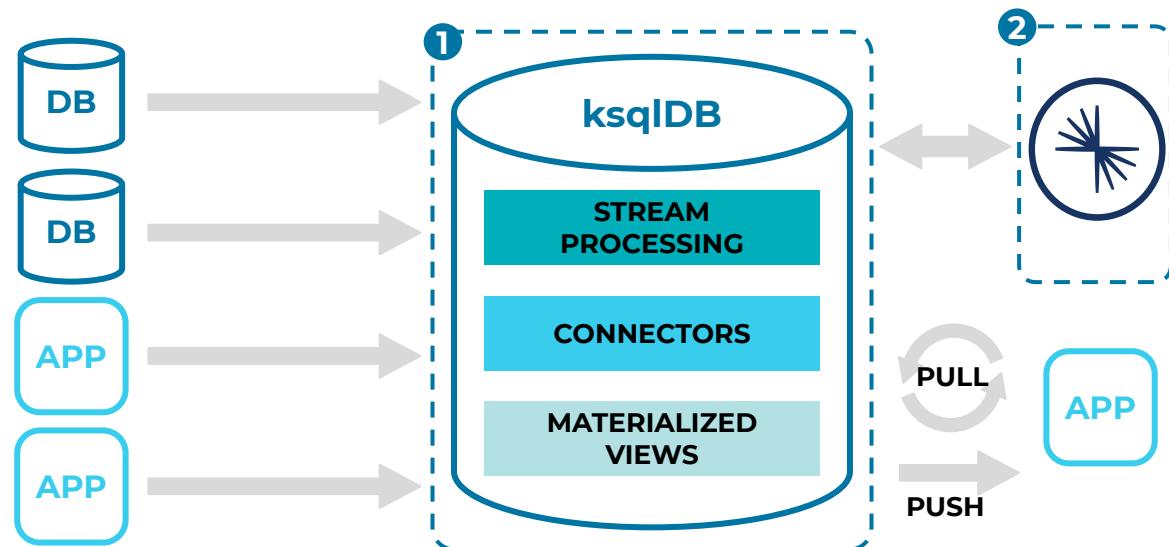
Fault tolerant and horizontally scalable: ensures high availability for mission-critical apps and more cloud-native experience

Build a modern architecture

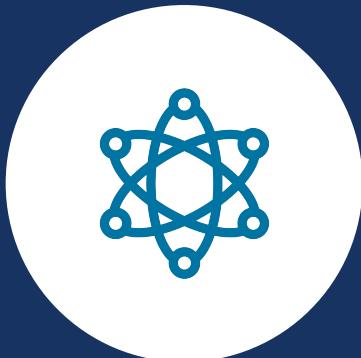


Confluent offers a way to modernize your data architecture and accelerates your developer velocity

ksqlDB provides everything you need to build a complete real-time application entirely with SQL syntax



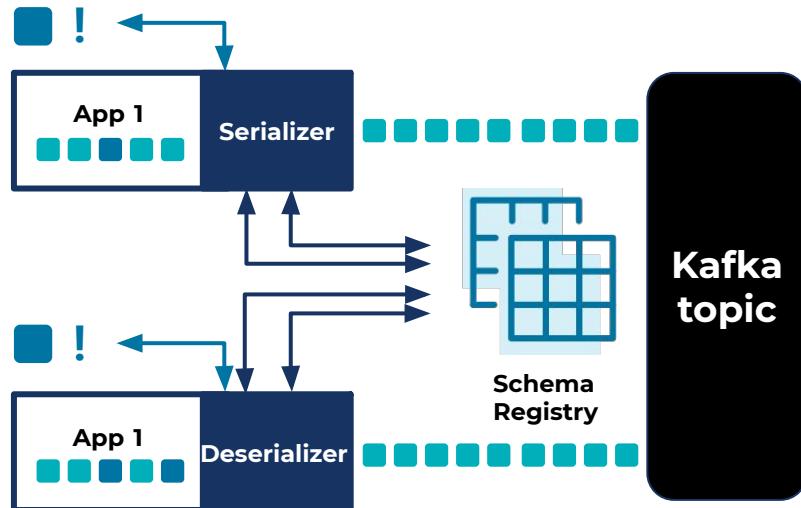
Build a modern architecture



Confluent offers a way to modernize your data architecture and accelerates your developer velocity

Data discovery

- Validate data compatibility and get warnings
- Let developers focus on deploying apps



Scale with confidence

- Store a versioned history of all schemas
- Enable evolution of schemas while preserving backwards compatibility for existing consumers

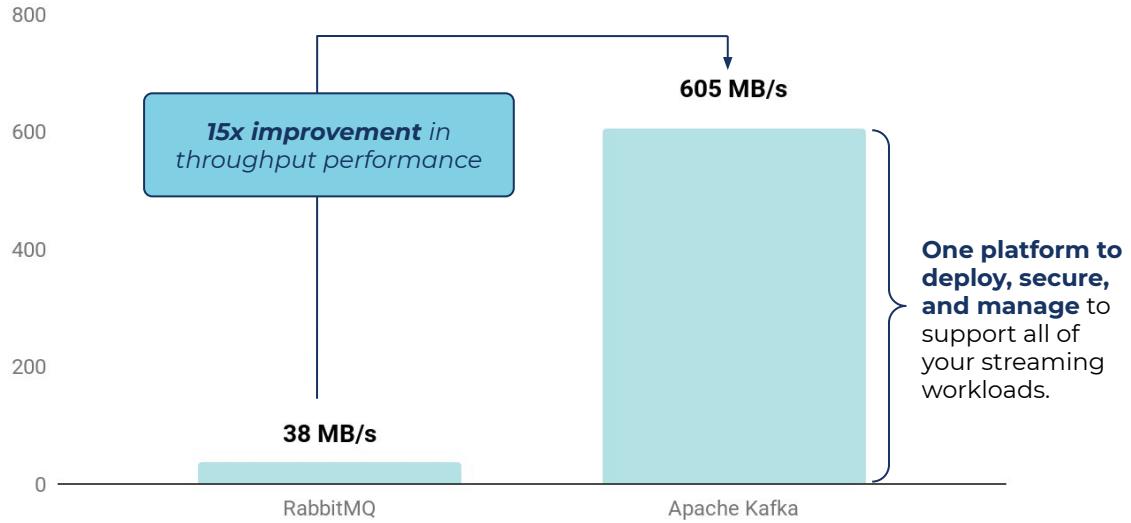
Achieve any scale



Confluent was designed for high performance and high availability at massive scale to support your mission-critical use cases

Confluent can support **all of your enterprise's streaming use cases** by achieving a dramatically higher throughput

Peak throughput (MB/s): Kafka vs RabbitMQ



[Read more about our internal performance benchmarking](#)

Achieve any scale



Confluent was designed for high performance at massive scale to support your mission-critical use cases

Confluent was designed for high performance at massive scale to support your mission-critical use cases

5ms

Confluent achieves <5ms latency at massive throughput



Synchronize data across your organization in real-time



Take action on insights from your data immediately



Remove data silos by moving from batch to data streaming

Migrate with ease



Confluent provides the tools and services required to effectively migrate from your messaging queue

We offer a robust set of connectors to pull data from your MQs into Confluent...



TIBCO®

solace•

JMS

RabbitMQ

Apache ACTIVEMQ

...and connectors to push data from Confluent into your modern, cloud-native sinks



mongoDB®



Google Cloud Storage



Amazon S3



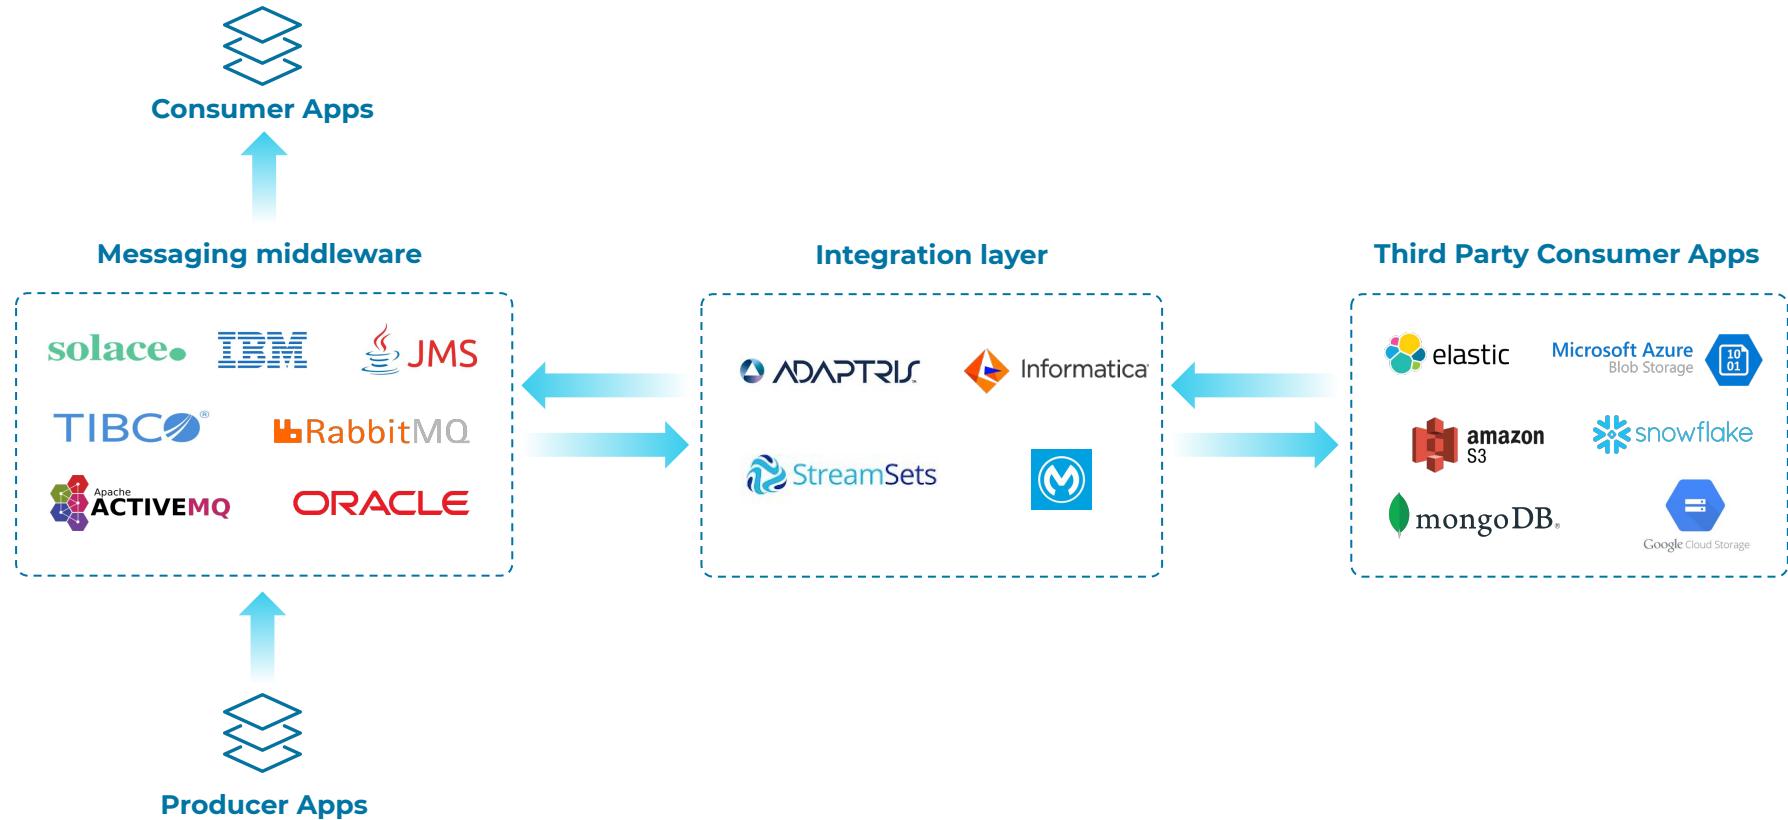
Microsoft Azure Blob Storage

snowflake



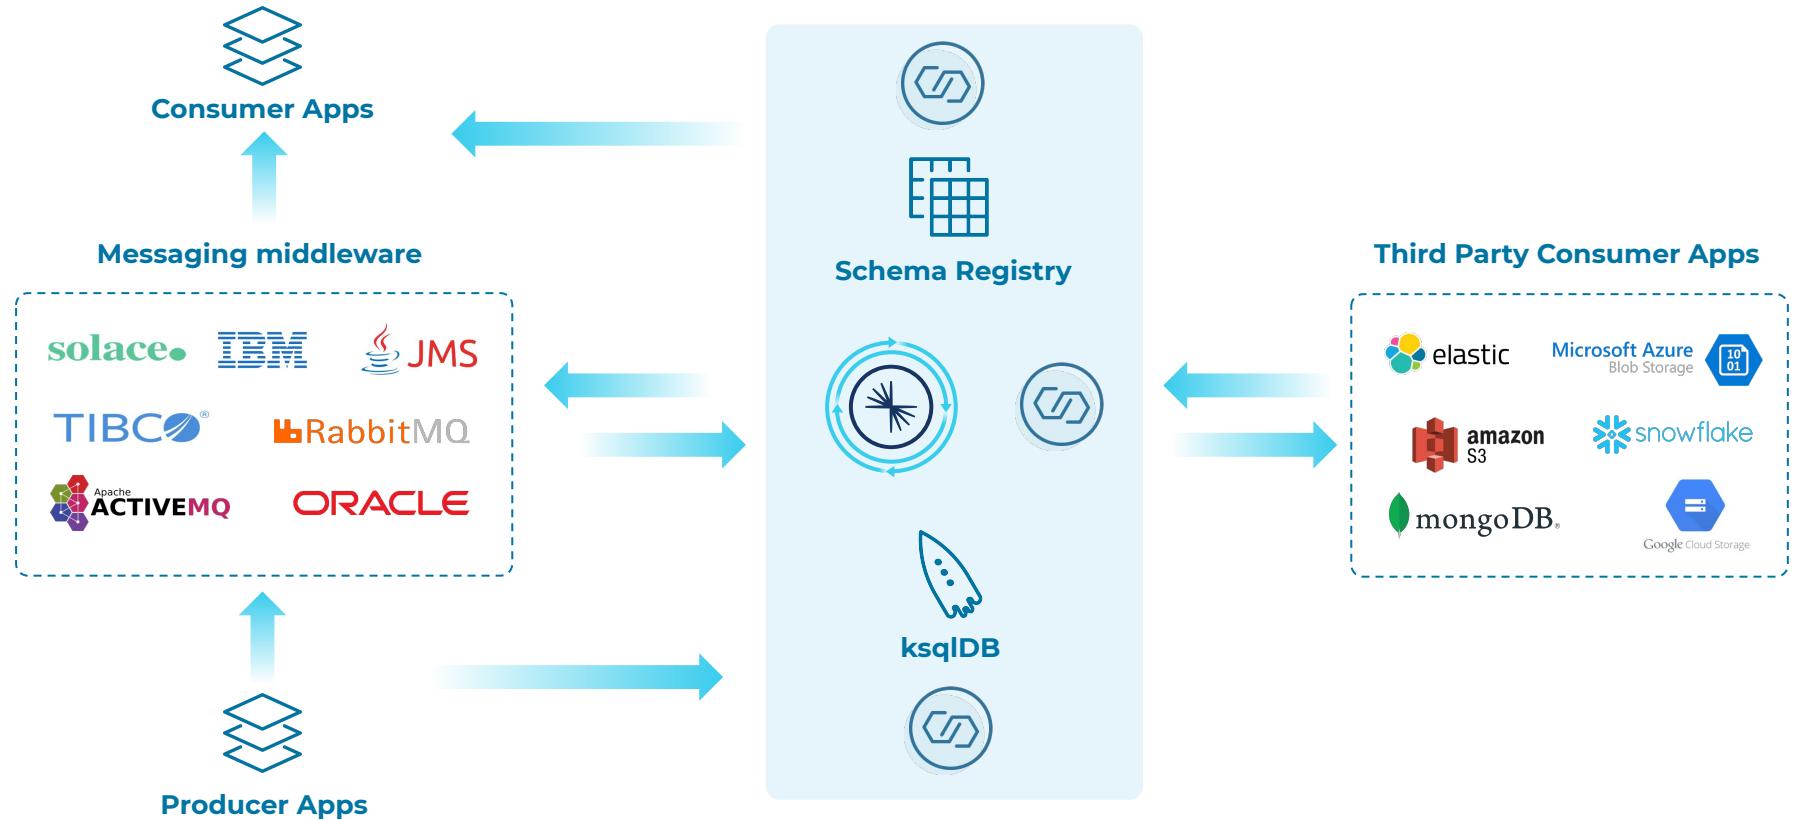
How We Do It

Current approach to Messaging

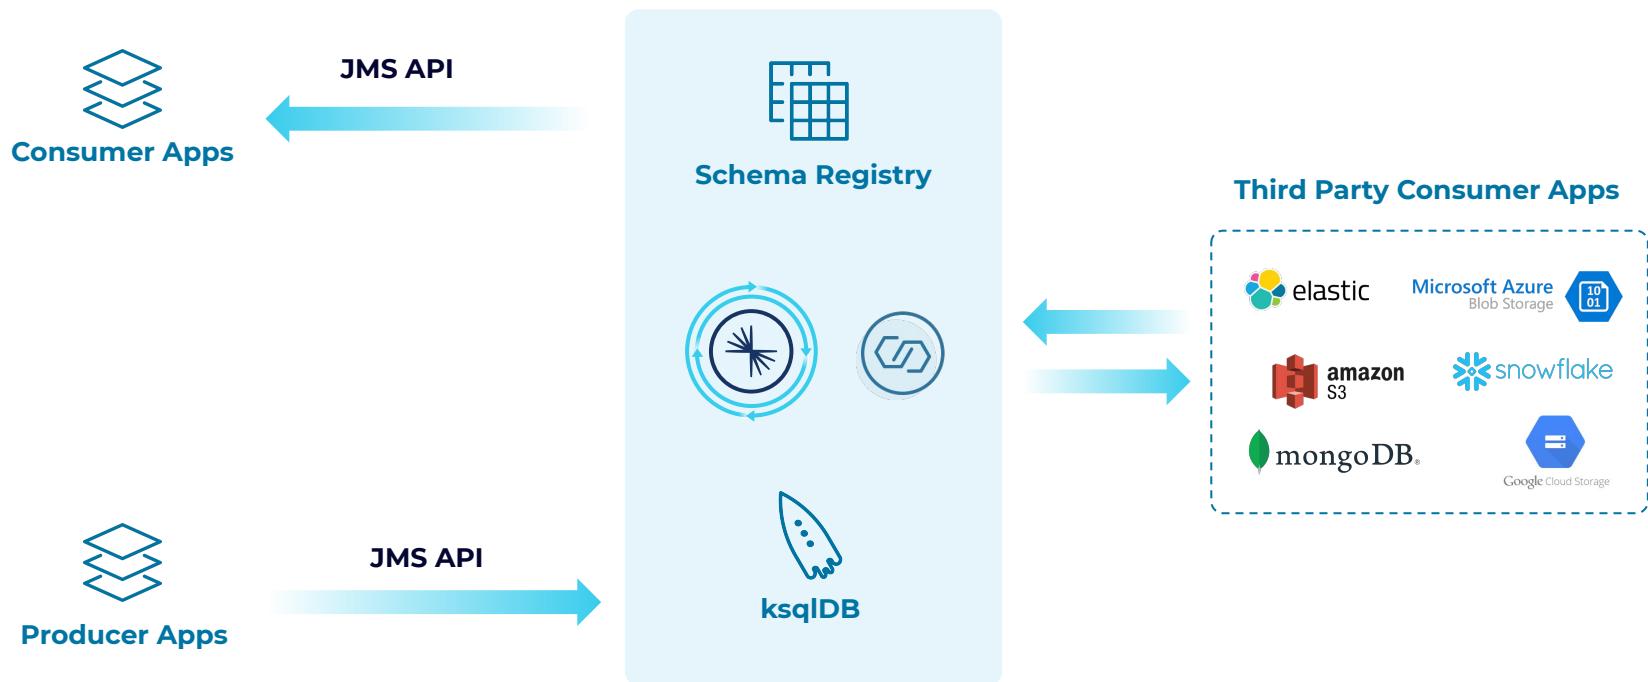




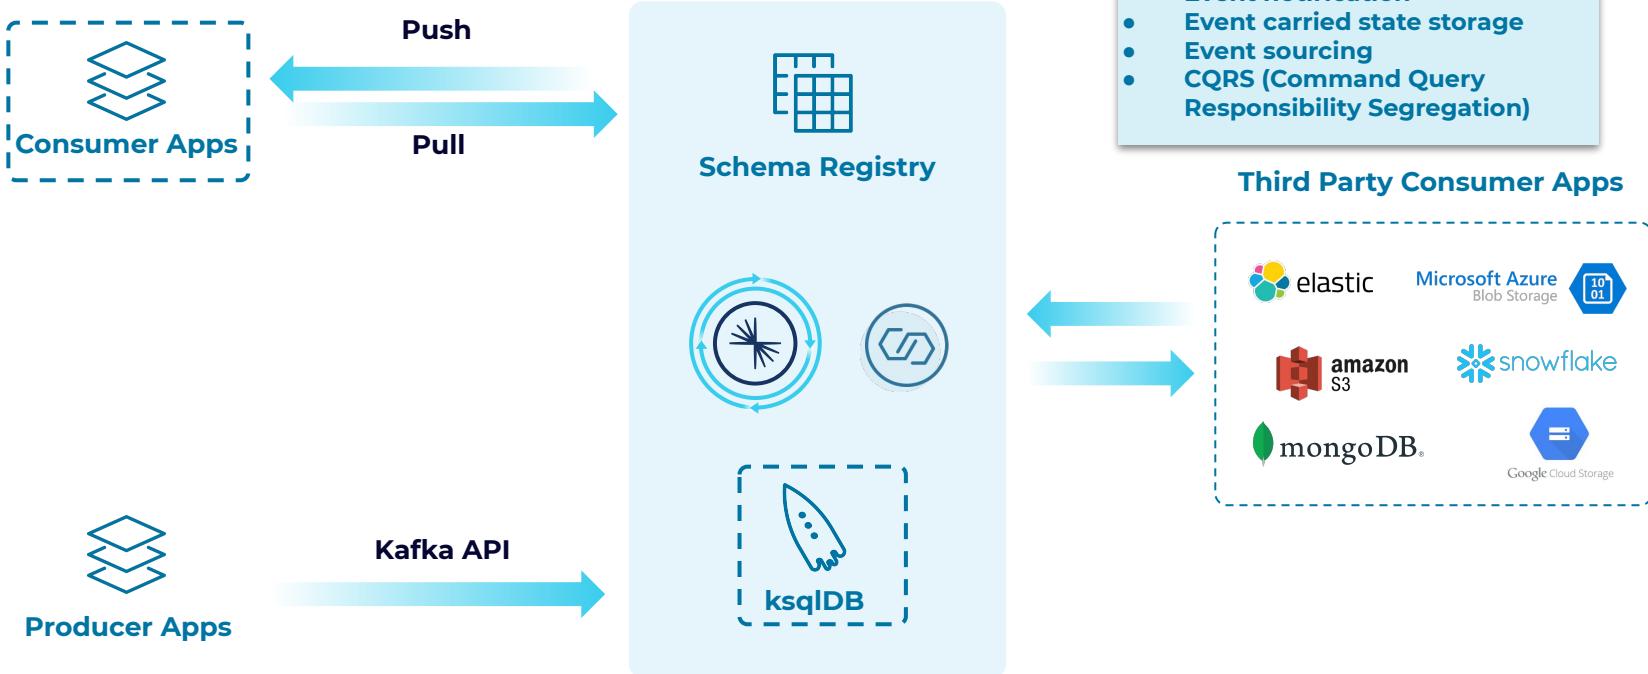
Step 1: Use pre-built connectors to start integrating data from your applications



Step 2: Refactor applications to use JMS APIs and replace messaging middleware over time



Step 3: Move off JMS and build streaming applications with new event-driven patterns



Unify and simplify



Lower TCO



Future-proof data architecture



CONFLUENT