Availability of Kafka - Beyond the Brokers

Andrew Borley

Emma Humber

Apache Kafka

Kafka Summit Europe 2021

High availability

Kafka has guarantees around the number of server failures a cluster can tolerate

What if the environment becomes unavailable?

Many layers: Applications, Kafka, OS, Network, Storage

Design into the system

Constraints

What guarantees do you need

- Consistency, availability, performance, manual intervention

What resources do you have

- Datacenters, network, cost

Consistency OR availability

Stretch Clusters

Availability zones

Set of isolated infrastructure

- Compute, storage and network connectivity + associated power and cooling
- Limits the blast radius of an infrastructure problem

A datacenter, or a failure domain within a datacenter

Geographic regions (eg Central Europe) often support multiple availability zones

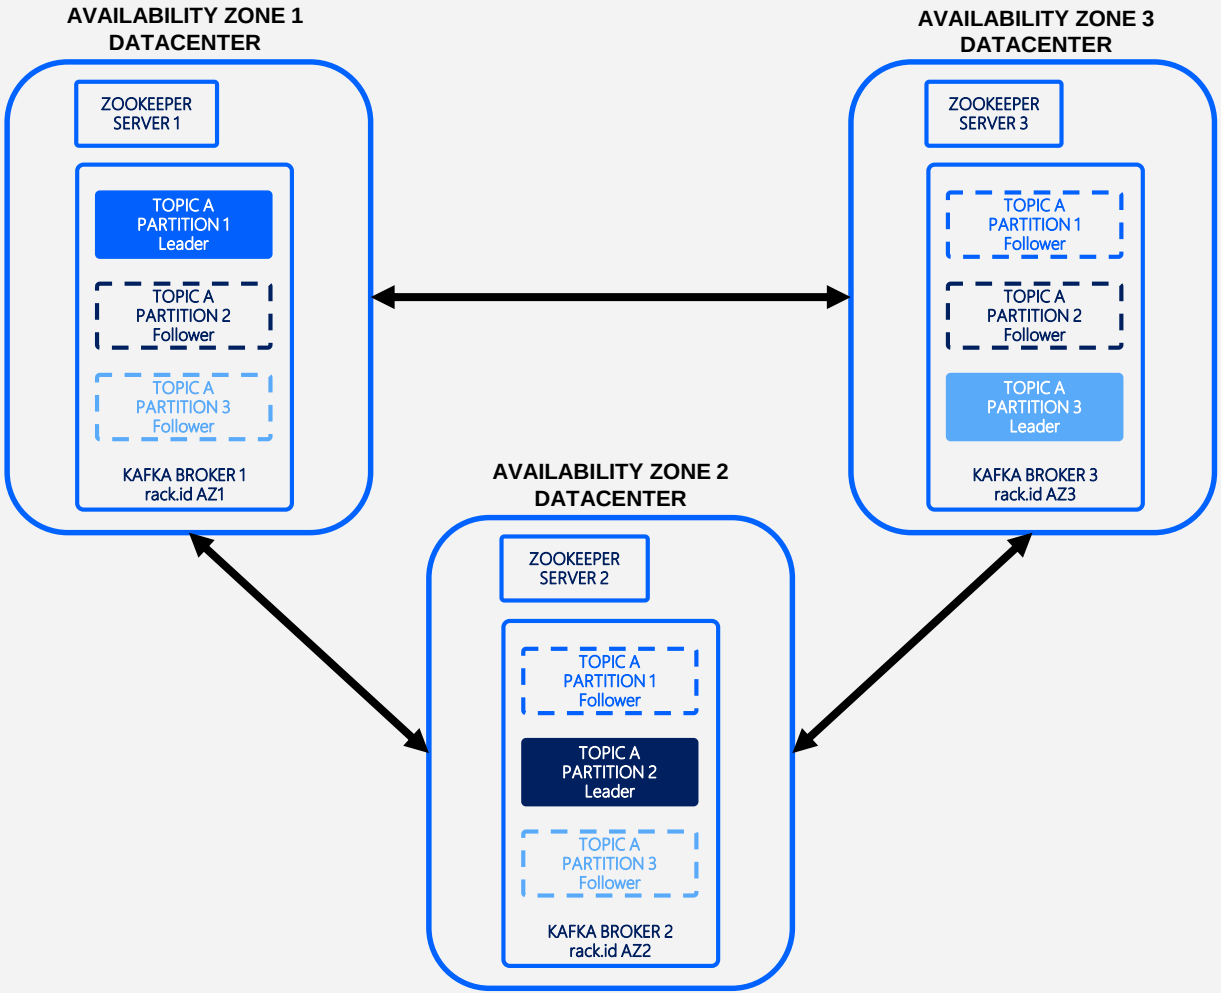

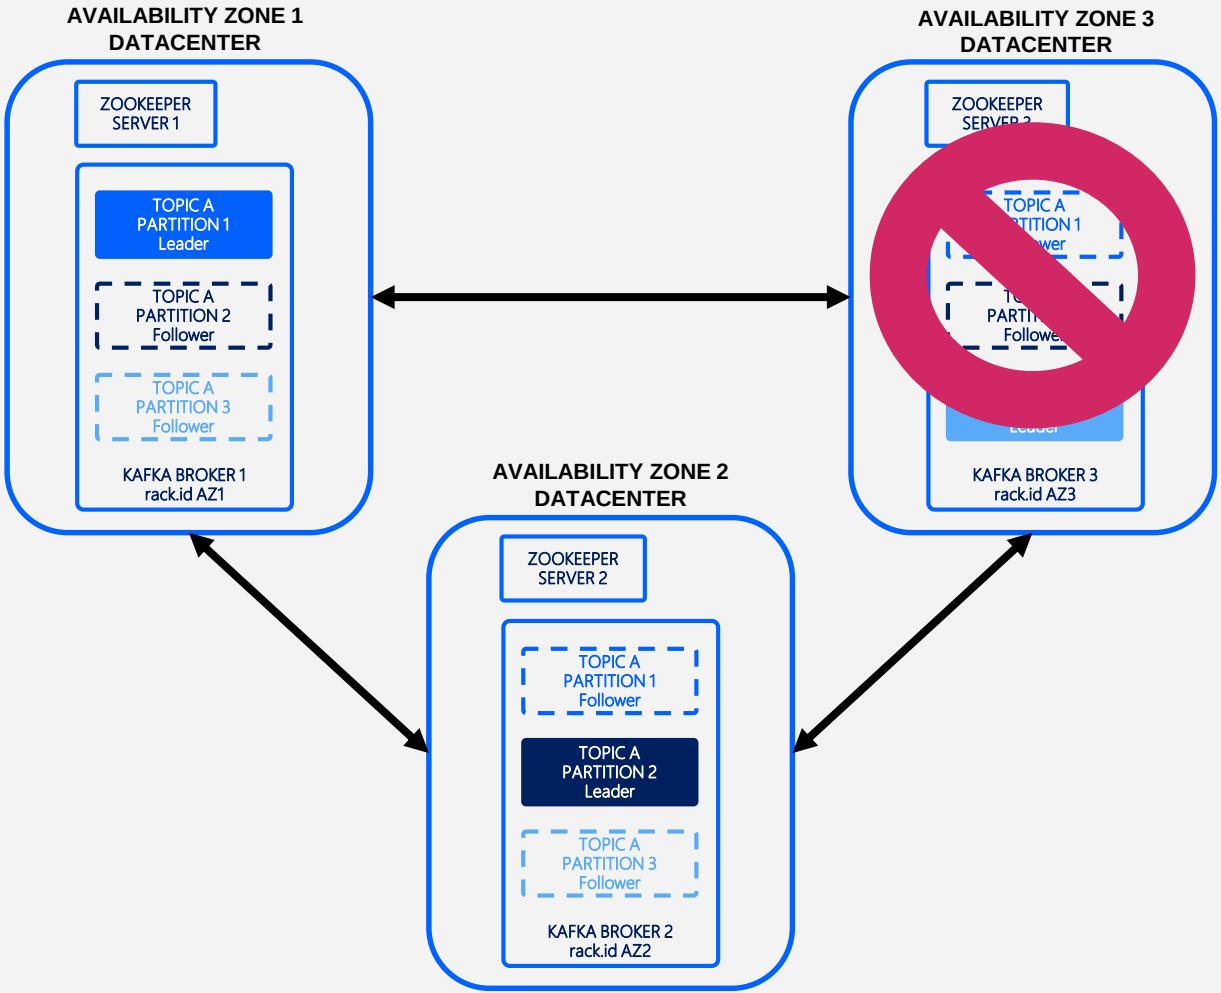

Configuration

Zone aware platforms automatically set up clusters in multiple availability zones

- Each Kubernetes node assigned a zone and labelled with `failure-domain.beta.kubernetes.io/zone`

Kafka's `rack.id : zone`

Low latency a MUST

Look at timeouts and client configuration

`zookeeper.connection.timeout.ms`

`replica.lag.time.max.ms`

`acks=all`

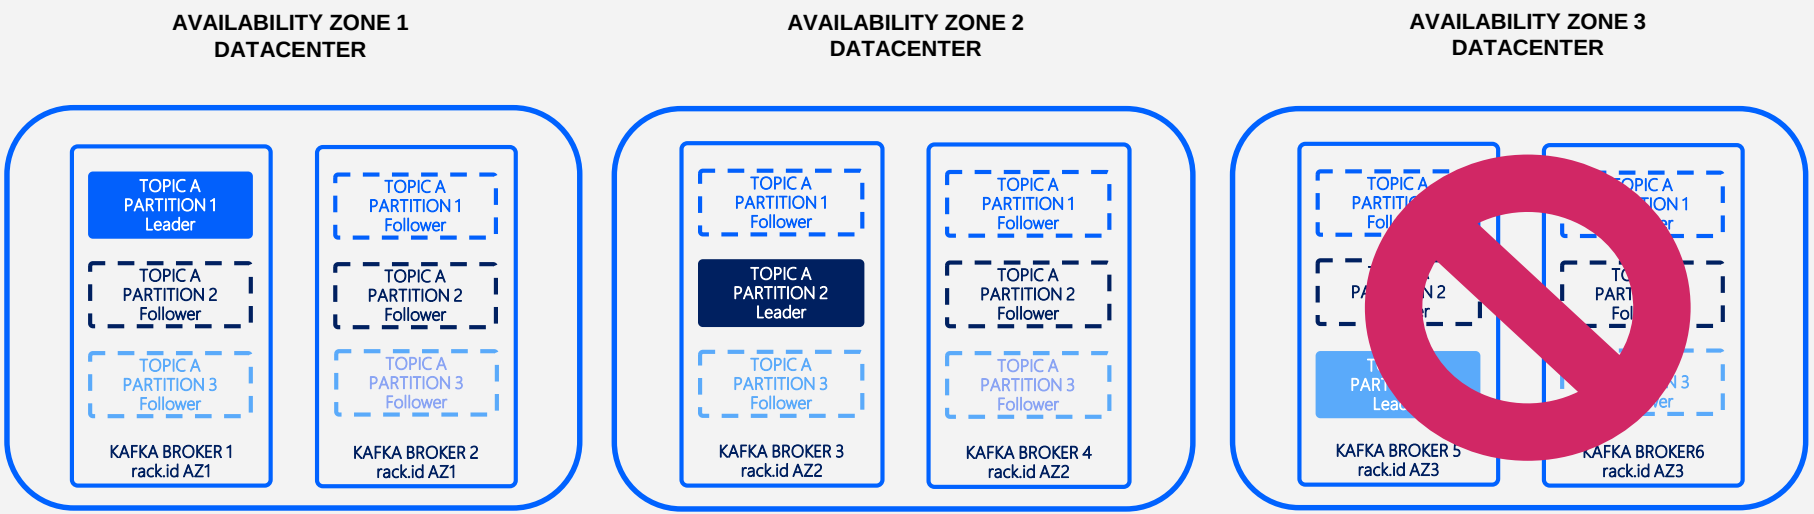

`Replication factor` > `min.insync.replicas`

Ensure there are sufficient replicas to cover all zones

`client.rack` tags which node the client application is running on

- Consumers to fetch from the closest replica

Zookeeper

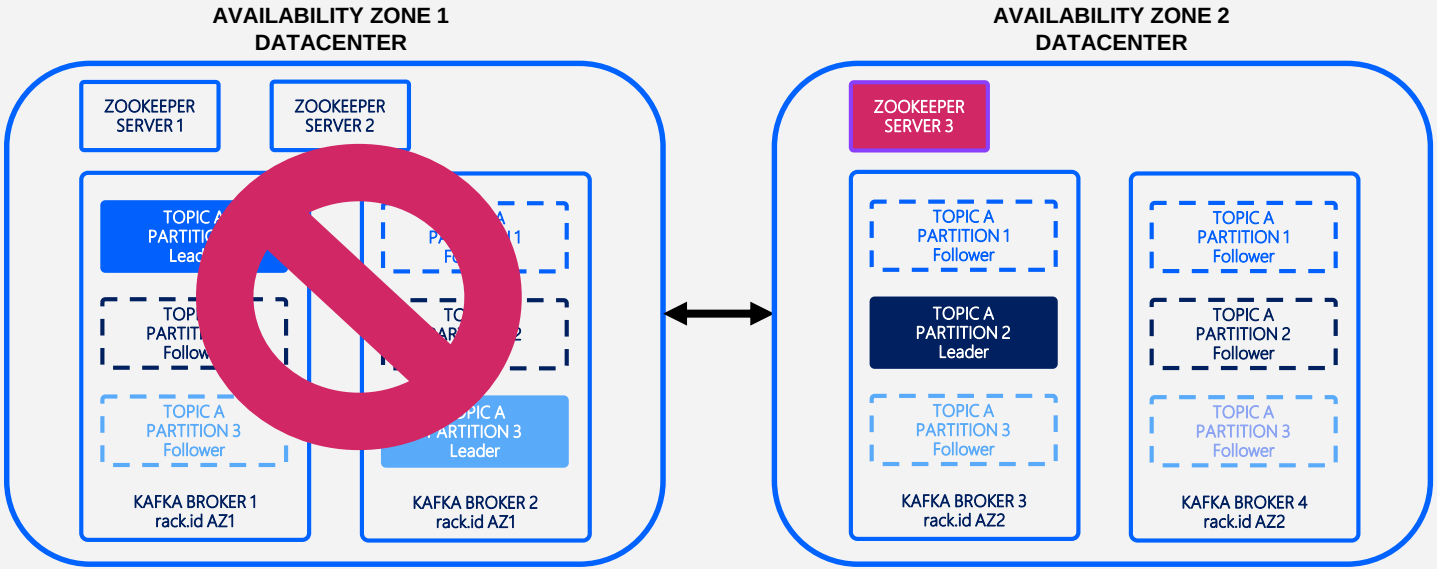

Stretch clusters

Kafka and Zookeeper replication ensures data is highly available

Data consistency

Exactly once processing possible

Event order can be preserved

Guards against the loss of a data center

Simple client configuration

No offset lag or offset translation

Utilizes all brokers

Considerations

Stable, low latency, high-bandwidth connection is a must : care if crossing regions

No o downtime upgrades

Doesn't protect against whole cluster failure

Cross-availability zone data transfer fees will apply : especially ingress/egress

Third datacenter required

Complexity of configuration

Multiple Clusters

Multiple clusters

Multiple independent Kafka clusters in different regions. Topic data is mirrored across clusters

- **Active/passive**: Produce to primary cluster, consume from any
- **Active/active**: Produce to and consume from any cluster
- **Federated**: Central cluster with multiple regional clusters

A **mirror making** technology mirrors the data between clusters

- Typically runs at the target cluster
- Data is consumed remotely and produced locally

Active/passive clusters

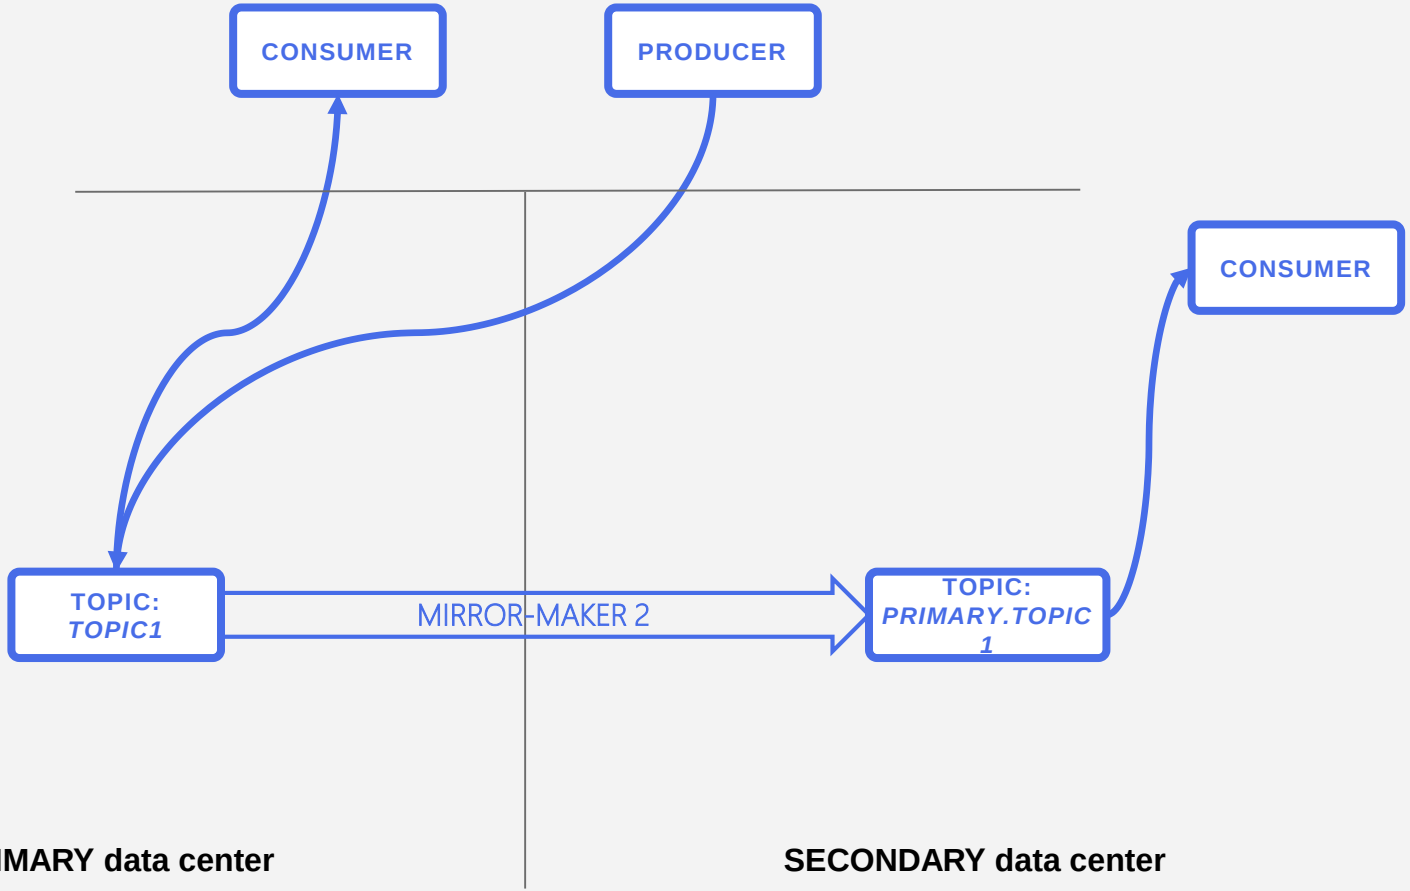

Active/passive clusters

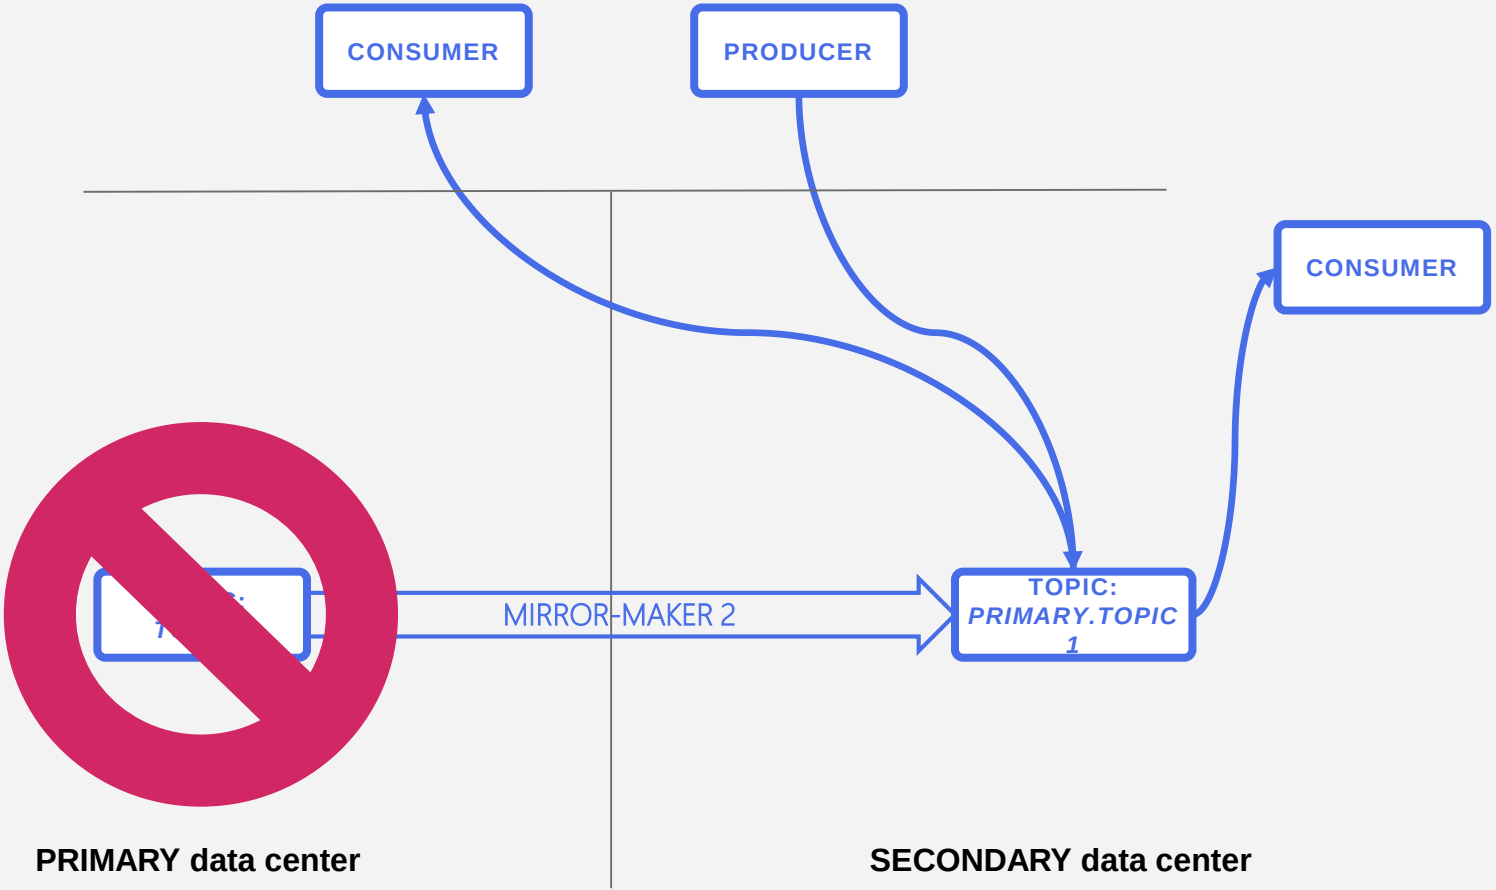

Active/passive clusters

Data back-up to multiple destinations

Disaster recovery fail-over after loss of infrastructure

Data migration:

- Moving to a new cluster
- Moving from a staging environment to a production environment

Clusters are independent

Message keys used for partitioning, so order is preserved on a per-key basis

Source cluster must be the 'owner' of the data. Target cluster is essentially read-only

Considerations

Target cluster will lag behind source cluster - mirroring is asynchronous, so monitor that it's not too far behind

- High traffic topics could mean 1000's of un-mirrored messages, even if there is only a few milliseconds of lag

Message offsets in the source and target topics may not match up

- On fail-over consumers need to be updated to start at the correct offset
- Returning to the source cluster will require a similar process

Configuration changes on the source may need to be propagated to the target

Target cluster may be unused (until it's needed)

Active/active clusters

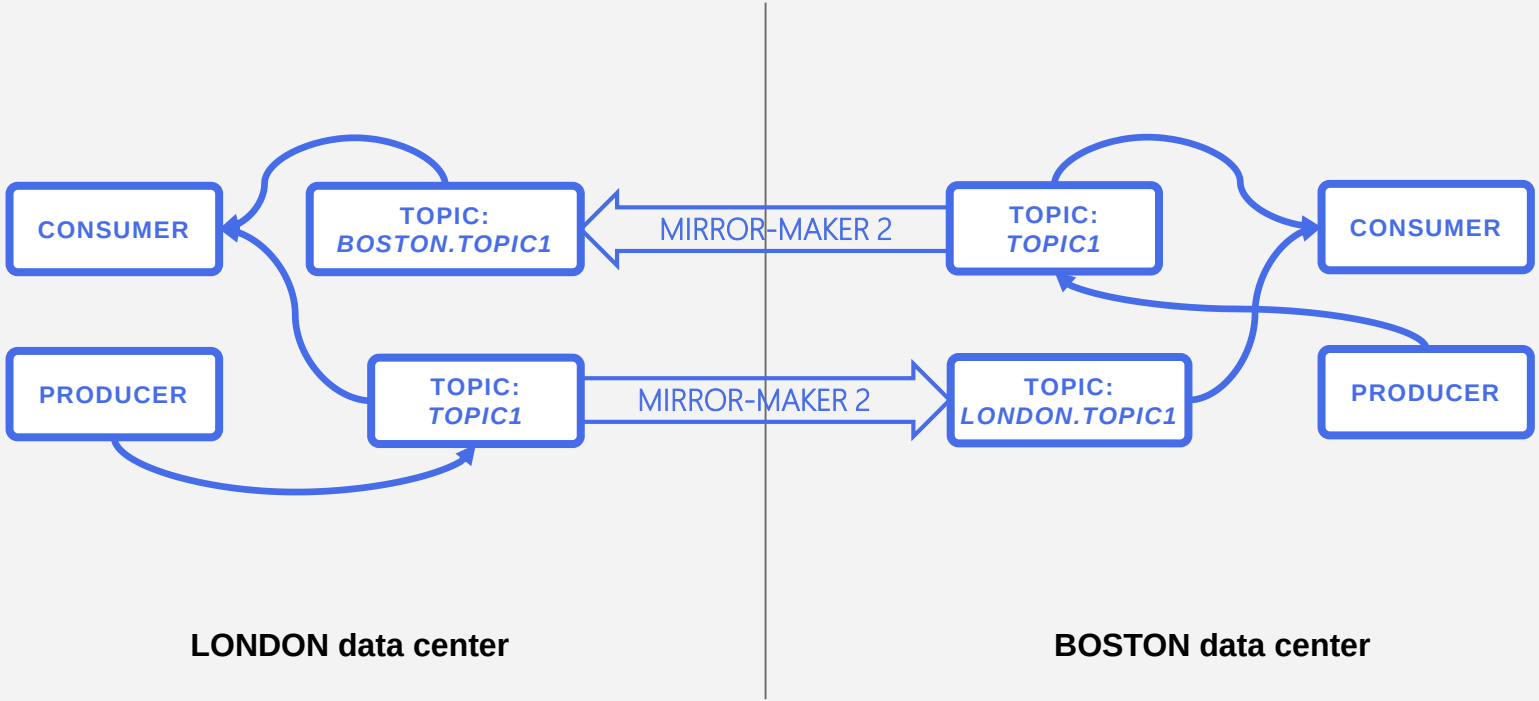

Active/active clusters

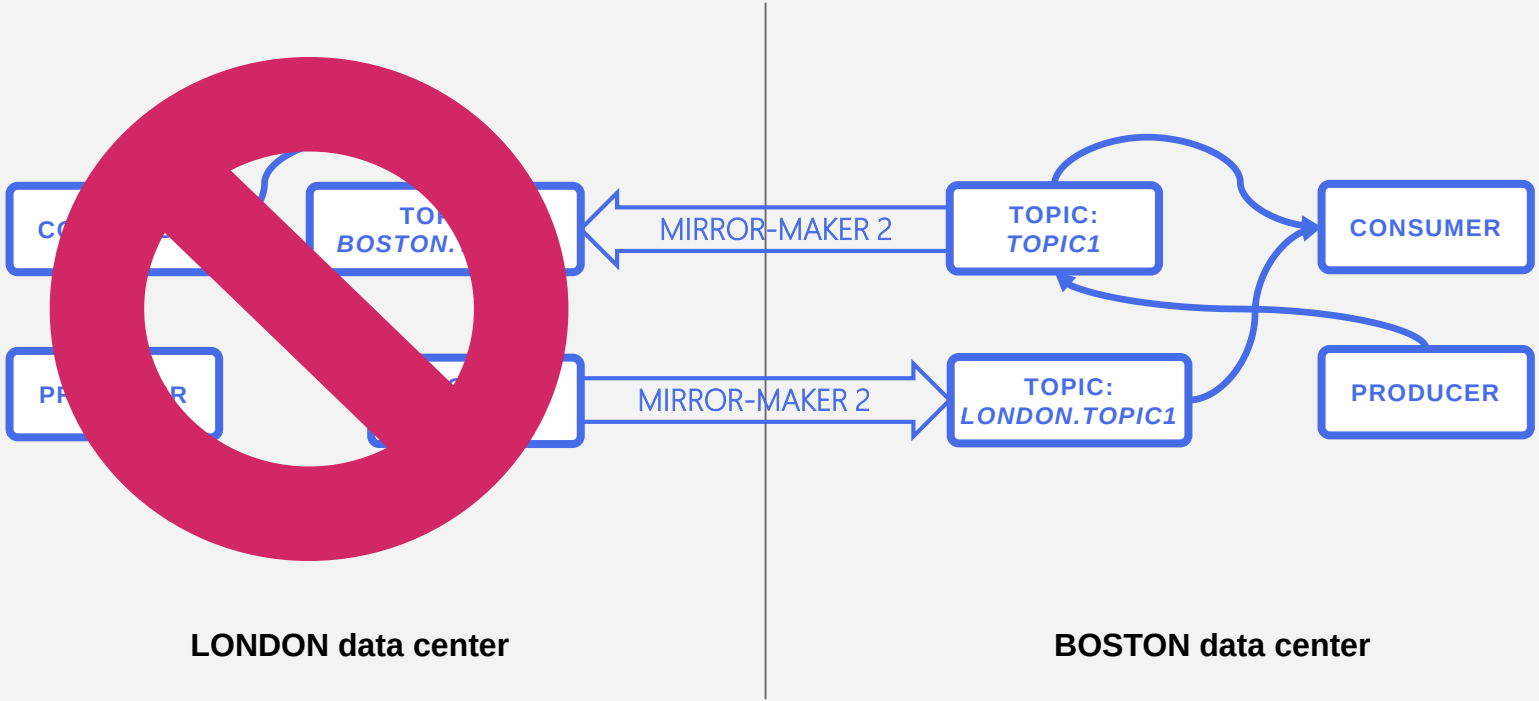

Active/active clusters

Implicit disaster recovery – data centers can operate independently

Data back-up to multiple destinations

A ‘virtual’ topic shared across geographies

Serve users from a nearby data center to increase performance

Redundancy and resilience - can easily do a network redirect on failure to route all traffic to surviving cluster

Considerations

Possible to accidentally configure loops of data

Lag and ordering issues may arise if trying to consume from a different data center to where the data was produced

Federated (hub and spoke) clusters

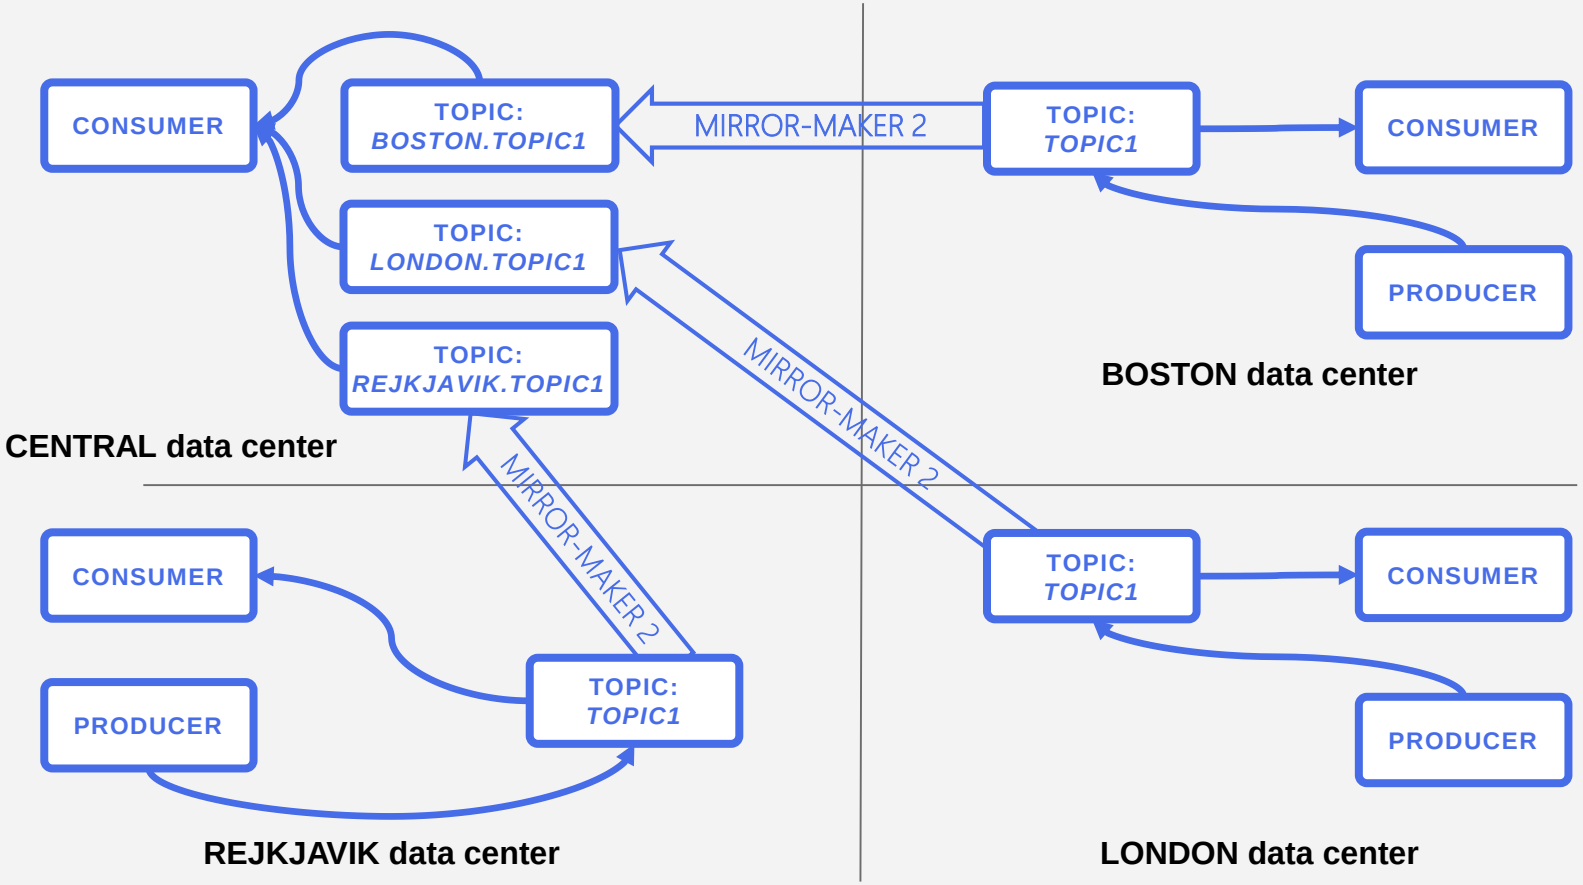

Federated (hub and spoke) clusters

Each region has a Kafka cluster to handle data for the region

There is no requirement for one region to know about data for other regions

Serve users from a nearby data center to increase performance

Central cluster consolidates data from each regional center

- Can be used when central processing requires access to the full set of data

Mirroring is single direction

- Makes it simple to configure, deploy and monitor

Be aware

Not useful if all regions need access to the data

Central cluster is for secondary processing or back-up

Lag and ordering issues at the central cluster may arise if regional data is co-dependent

Planning multi-cluster architectures for fail-over

Consider how applications will fail-over to different clusters

- Bootstrap address / DNS
- Certificates, user credentials, permissions
- Consumer offsets
- Topic subscriptions
- Message ordering

Consider effect of duplicates/lost messages

- Idempotent writes

Monitor that data is arriving at remote clusters as well as your primary system health

Practice your fail-over and fail-back

Summary

Kafka provides a good degree of high availability

Some applications require additional guarantees

Dependent on your requirements and your infrastructure

Thank you

Andrew Borley
Emma Humber

IBM Event Streams
—
borley@uk.ibm.com
emma.humber@uk.ibm.com