

Agenda

- What is a Service Mesh
- How we got here: A story
- Architecture and details
- Q & A





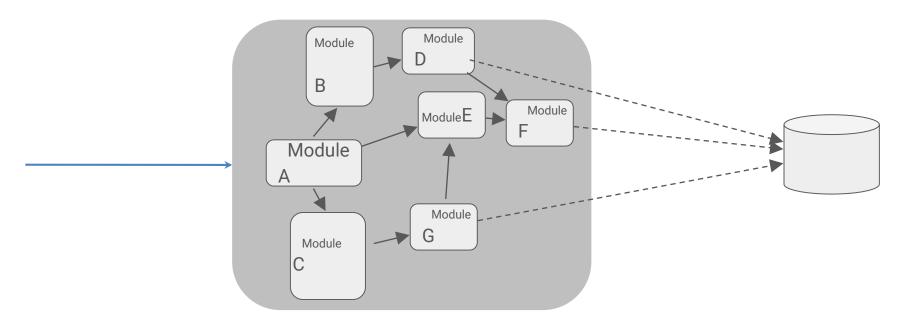
What is a 'Service Mesh'?

- Traffic Control
- Visibility
- Resiliency & Efficiency
- Security
- Policy Enforcement



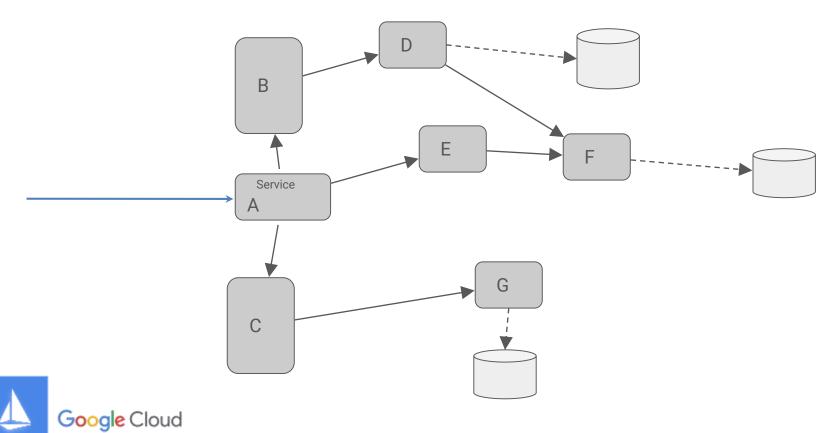


Remember the Monolith?





Micro Services



Micro Services FTW

- 1. Gained development velocity!
- 2. Easy testing because of abstractions.
- 3. Scale services independently.

Problem Solved!



Yay!

1. Gained development velocity!

2. Easy testing because of abstractions.

3. Scale services independently.

Problem Solved!





What have we lost?

- 1. I replaced a **reliable** in-process call with an **unreliable rpc**.
- 2. Trivial single stepping replaced by ...?
- 3. Secure in-process communication is replaced by insecure network.
- 4. Access control within process was a NOOP
- 5. Latency went up

That abstraction was leaky ...



Can we fix it?

- 1. Add retry logic to the application code.
- 2. Add entry-exit traces.
- 3. Secure inter service connections with strong authentication.

Now that we are adding code ... choose the rpc endpoint intelligently

- a. Endpoints with low latency.
- b. Endpoints with warm caches.



Service Mesh

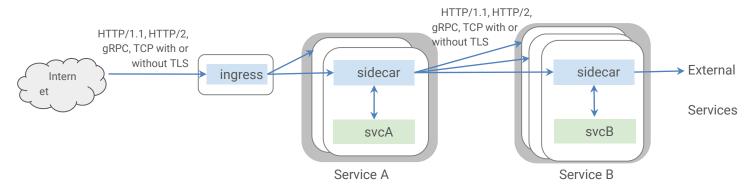
Address service level concerns.

• Unlock the full power of microservices.





Weaving the mesh





Outbound features:

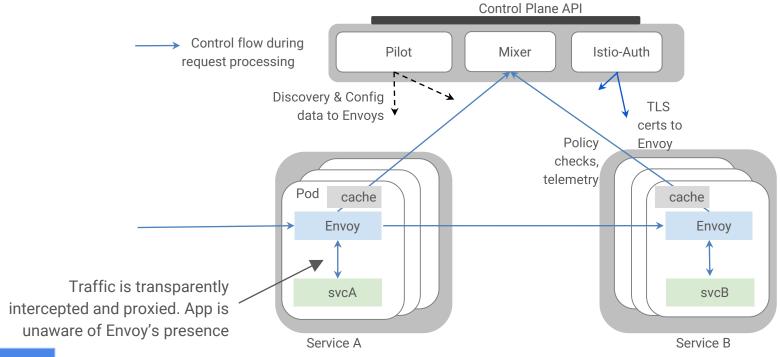
- Service authentication
- Load balancing
- Retry and circuit breaker
- Fine-grained routing
- Telemetry
- Request Tracing
- Fault Injection

Inbound features:

- Service authentication
- Authorization
- Rate limits
- Load shedding
- Telemetry
- Request Tracing
- Fault Injection



Architecture





istio Service Mesh

- Traffic Control
- Visibility
- Resiliency & Efficiency
- Security



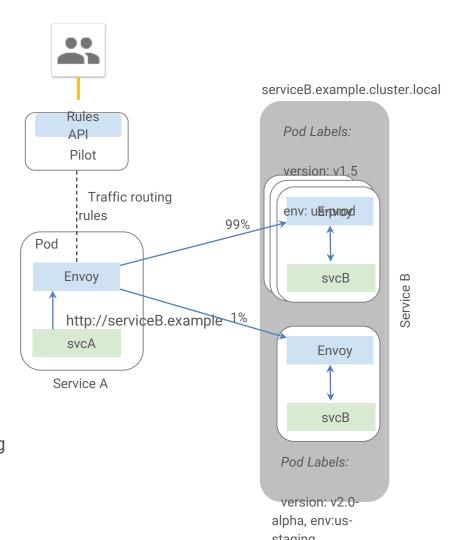


Application Rollout

```
// A simple traffic splitting rule
destination: serviceB.example.cluster.local
match:
 source: serviceA.example.cluster.local
route:
- tags:
    version: v1.5
    env: us-prod
 weight: 99
- tags:
    version: v2.0-alpha
    env: us-staging
 weight: 1
```

Traffic control is decoupled from infrastructure scaling



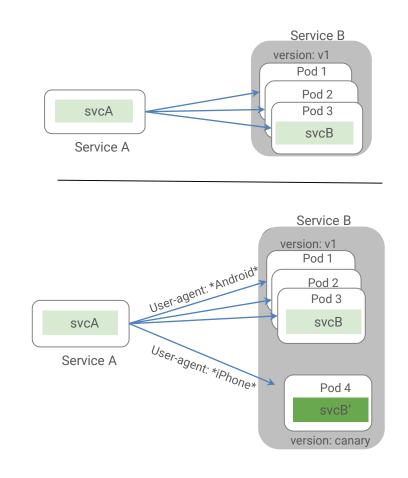


Traffic Steering

```
// Content-based traffic steering rule

destination: serviceB.example.cluster.local
match:
   httpHeaders:
       user-agent:
       regex: ^(.*?;)?(iPhone)(;.*)?$
precedence: 2
route:
- tags:
   version: canary
```

Content-based traffic steering

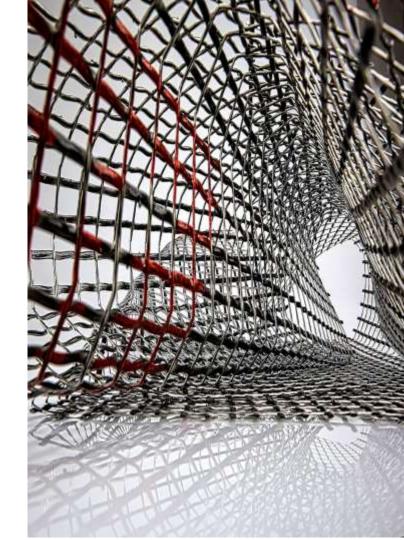




Istio Service Mesh

- Traffic Control
- Visibility
- Resiliency & Efficiency
- Security





Visibility

Monitoring & tracing should not be an afterthought in the infrastructure

Goals

- Metrics without instrumenting apps
- Consistent metrics across fleet
- Trace flow of requests across services
- Portable across metric backend providers



Istio - Grafana dashboard w/ Prometheus backend

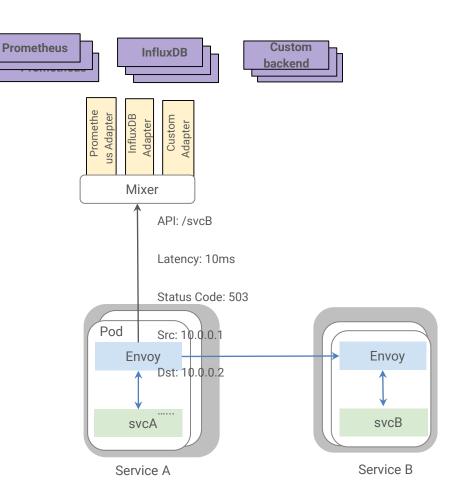


Istio Zipkin tracing dashboard



Metrics flow

- Mixer collects metrics emitted by Envoys
- Adapters in the Mixer normalize and forward to monitoring backends
- Metrics backend can be swapped at runtime





Visibility: Tracing

Application do not have to deal with generating spans or correlating causality

Envoys generate spans

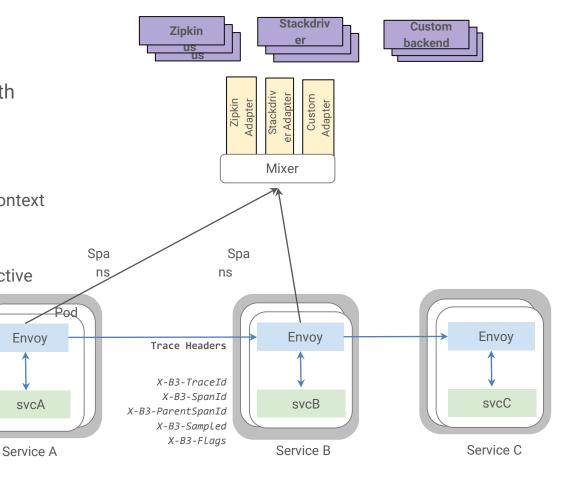
Applications need to *forward* context headers on outbound calls

svcA

Envoys send traces to Mixer

Adapters at Mixer send traces to respective

backends

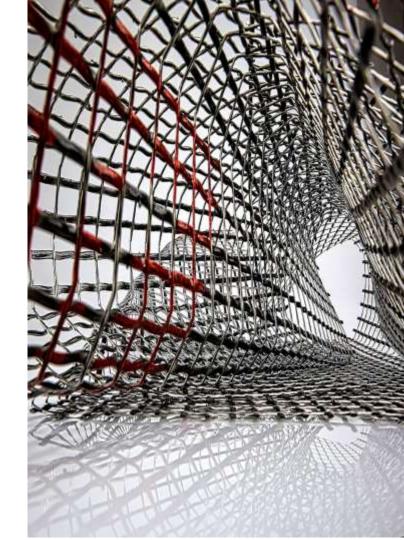




Istio Service Mesh

- Traffic Control
- Visibility
- Resiliency & Efficiency
- Security





Resiliency

Istio adds fault tolerance to your application without any changes to code

```
// Circuit breakers
 destination: serviceB.example.cluster.local
policy:
- tags:
    version: v1
  circuitBreaker:
    simpleCb:
      httpConsecutiveErrors: 7
      sleepWindow: 5m
      httpDetectionInterval: 1m
```

Resilience features

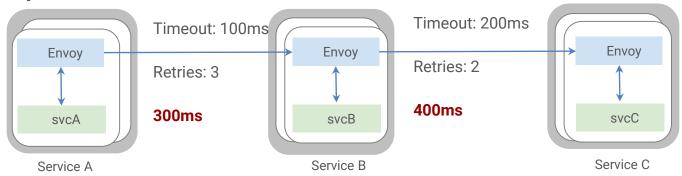
- Timeouts
- Retries with timeout budget
- Circuit breakers
- Health checks
- AZ-aware load balancing w/ automatic failover
- Control connection pool size and request load
- Systematic fault injection



Resiliency Testing

Systematic fault injection to identify weaknesses in failure recovery policies

- HTTP/gRPC error codes
- Delay injection



Efficiency

- L7 load balancing
 - Passive/Active health checks, circuit breaks
 - Backend subsets
 - Affinity
- TLS offload
 - No more JSSE or stale SSL versions.
- HTTP/2 and gRPC proxying



Istio Service Mesh

- Traffic Control
- Visibility
- Resiliency & Efficiency
- Security





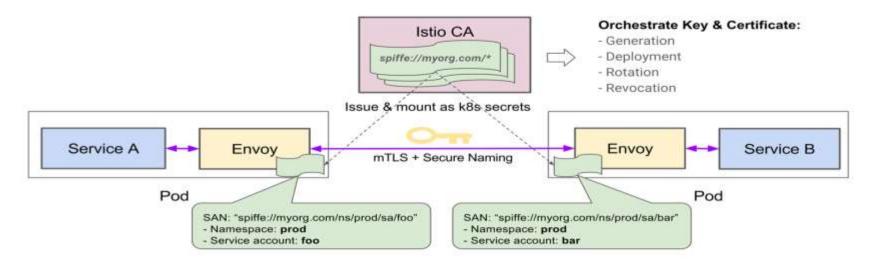
Securing Microservices

- Verifiable identity
- Secure naming / addressing
- Traffic encryption
- Revocation





Istio - Security at Scale

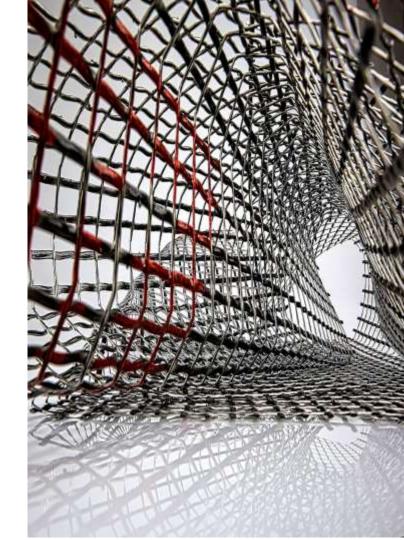




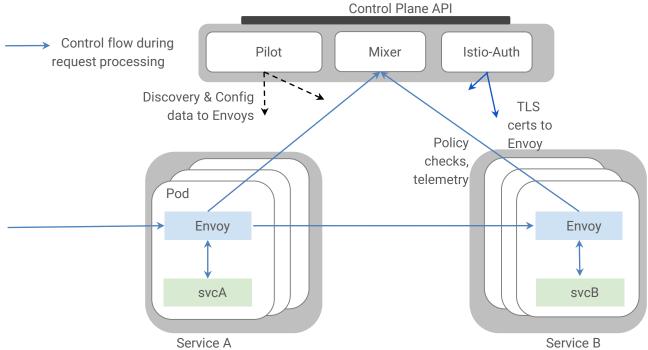
Istio Service Mesh

- Traffic Control
- Visibility
- Resiliency & Efficiency
- Security





Putting it all together



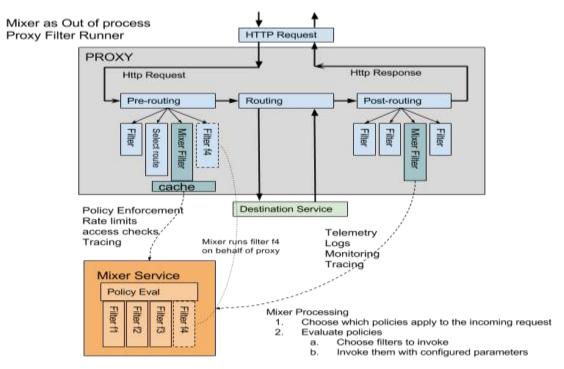


What does Mixer do?

- Check()
 - Precondition checking
 - Quotas & Rate Limiting
- Report()
 - telemetry reporting
- Primary point of extensibility
- Enabler for platform mobility
- Operator-focused configuration model



Mixer - Proxy Filter Extension





Mixer Metrics

```
kind: metrics
metadata:
   name: requestcount
spec:
   value: "1"
   dimensions:
      Destination_service: destination.service
      Source_service: source.service
      Response_code: response.code
```

```
kind: rule
metadata:
  name: prometheus
spec:
  match: request.headers["x-user"] != "admin"
  actions:
  - handler: handler.prometheus
  instances:
  - requestcount.metric
```



