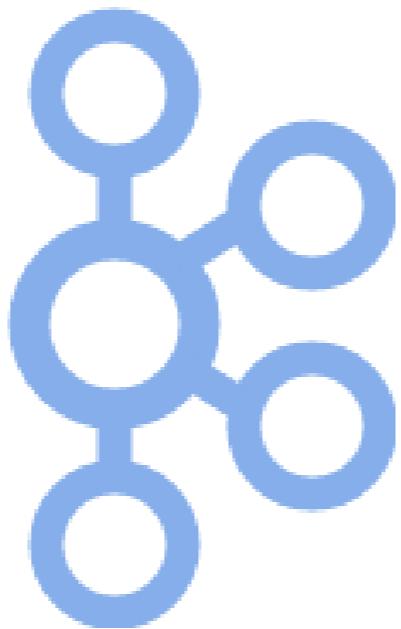


Tutorial on Kafka Streaming Platform

Part 2

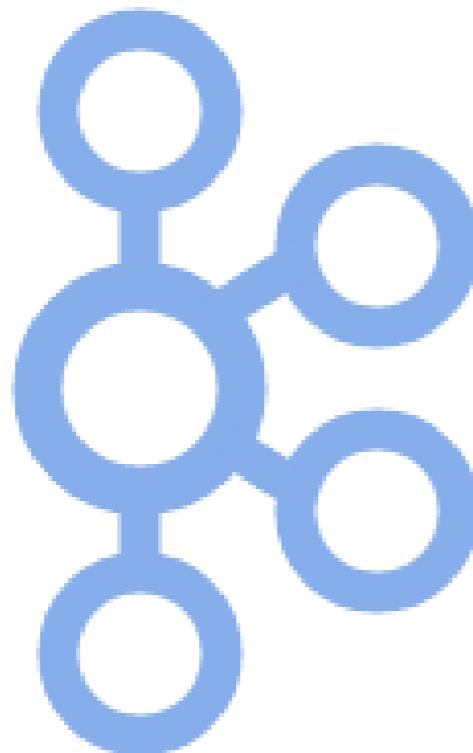


Cloudurable provides Cassandra and Kafka Support on AWS/EC2

Kafka Tutorial (Part 2)

Java Producer and Consumer examples

Schema Registry





StockPrice Producer Java Example

Lab StockPrice Producer Java Example

StockPrice App to demo Advanced Producer



- ❖ **StockPrice** - holds a stock price has a name, dollar, and cents
- ❖ **StockPriceKafkaProducer** - Configures and creates **KafkaProducer<String, StockPrice>**, **StockSender** list, ThreadPool (ExecutorService), starts **StockSender** runnable into thread pool
- ❖ **StockAppConstants** - holds topic and broker list
- ❖ **StockPriceSerializer** - can serialize a **StockPrice** into **byte[]**
- ❖ **StockSender** - generates somewhat random stock prices for a given **StockPrice** name, Runnable, 1 thread per StockSender
 - ❖ Shows using **KafkaProducer** from many threads



StockPrice domain object

```
c StockPrice.java x
StockPrice
1 package com.cloudurable.kafka.producer.model;
2
3 import io.advantageous.boon.json.JsonFactory;
4
5 public class StockPrice {
6
7     private final int dollars;
8     private final int cents;
9     private final String name;
10
11    public String toJson() {
12        return "{" +
13            "\"dollars\": " + dollars +
14            ", \"cents\": " + cents +
15            ", \"name\": \"\" + name + '\"' +
16            "}";
17    }
18}
```

- ❖ has name
- ❖ dollars
- ❖ cents
- ❖ converts itself to JSON



StockPriceKafkaProducer

- ❖ Import classes and setup logger
- ❖ Create ***createProducer*** method to create ***KafkaProducer*** instance
- ❖ Create ***setupBootstrapAndSerializers*** to initialize bootstrap servers, client id, key serializer and custom serializer (***StockPriceSerializer***)
- ❖ Write ***main()*** method - creates ***producer***, create ***StockSender*** list passing each instance a ***producer***, creates a thread pool so every stock sender gets its own thread, runs each ***stockSender*** in its own thread



StockPriceKafkaProducer imports, createProducer

StockPriceKafkaProducer main()

```
1 package com.cloudurable.kafka.producer;
2
3 import com.cloudurable.kafka.StockAppConstants;
4 import com.cloudurable.kafka.producer.model.StockPrice;
5 import io.advantageous.boon.core.Lists;
6 import org.apache.kafka.clients.producer.*;
7 import org.apache.kafka.common.serialization.StringSerializer;
8 import org.slf4j.Logger;
9 import org.slf4j.LoggerFactory;
10
11 import java.util.List;
12 import java.util.Properties;
13 import java.util.concurrent.ExecutorService;
14 import java.util.concurrent.Executors;
15 import java.util.concurrent.TimeUnit;
16
17 public class StockPriceKafkaProducer {
18     private static final Logger logger =
19         LoggerFactory.getLogger(StockPriceKafkaProducer.class);
20
21     private static Producer<String, StockPrice> createProducer() {
22         final Properties props = new Properties();
23         setupBootstrapAndSerializers(props);
24         return new KafkaProducer<>(props);
25     }
}
```

Import classes and
setup logger

createProducer
used to create a
KafkaProducer

createProducer()
calls
setupBootstrapAnd
Serializers()

Configure Producer Bootstrap and Serializer



StockPriceKafkaProducer

```
27     private static void setupBootstrapAndSerializers(Properties props) {  
28         props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,  
29                     StockAppConstants.BOOTSTRAP_SERVERS);  
30         props.put(ProducerConfig.CLIENT_ID_CONFIG, "StockPriceKafkaProducer");  
31         props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,  
32                     StringSerializer.class.getName());  
33  
34         //Custom Serializer - config "value.serializer"  
35         props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,  
36                     StockPriceSerializer.class.getName());  
37     }
```

- ❖ Create **setupBootstrapAndSerializers** to initialize bootstrap servers, client id, key serializer and custom serializer (**StockPriceSerializer**)
- ❖ **StockPriceSerializer** will serialize **StockPrice** into bytes

StockPriceKafkaProducer.main()



```
c StockPriceKafkaProducer.java x
StockPriceKafkaProducer main()
39 >     public static void main(String... args)
40             throws Exception {
41
42             //Create Kafka Producer
43             final Producer<String, StockPrice> producer = createProducer();
44
45             //Create StockSender list
46             final List<StockSender> stockSenders = getStockSenderList(producer);
47
48             //Create a thread pool so every stock sender gets its own.
49             final ExecutorService executorService =
50                 Executors.newFixedThreadPool(stockSenders.size());
51
52             //Run each stock sender in its own thread.
53             stockSenders.forEach(executorService::submit);
54
55 }
```

- ❖ **main** method - creates **producer**,
- ❖ create **StockSender** list passing each instance a **producer**
- ❖ creates a thread pool (executorService)
- ❖ every StockSender runs in its own thread



StockAppConstants

```
1 package com.cloudurable.kafka;  
2  
3 public class StockAppConstants {  
4     public final static String TOPIC = "stock-prices";  
5     public final static String BOOTSTRAP_SERVERS =  
6         "localhost:9092,localhost:9093,localhost:9094";  
7  
8 }
```

- ❖ topic name for Producer example
- ❖ List of bootstrap servers

StockPriceKafkaProducer.getStockSenderList



```
StockPriceKafkaProducer getStockSenderList()
57 @ private static List<StockSender> getStockSenderList(
58     final Producer<String, StockPrice> producer) {
59     return Lists.list(
60         new StockSender(StockAppConstants.TOPIC,
61             new StockPrice( name: "IBM", dollars: 100, cents: 99),
62             new StockPrice( name: "IBM", dollars: 50, cents: 10),
63             producer,
64             delayMinMs: 1, delayMaxMs: 10
65         ),
66         new StockSender(
67             StockAppConstants.TOPIC,
68             new StockPrice( name: "SUN", dollars: 100, cents: 99),
69             new StockPrice( name: "SUN", dollars: 50, cents: 10),
70             producer,
71             delayMinMs: 1, delayMaxMs: 10
72         ),
73         new StockSender(
74             StockAppConstants.TOPIC,
75             new StockPrice( name: "GOOG", dollars: 500, cents: 99),
76             new StockPrice( name: "GOOG", dollars: 400, cents: 10),
77             producer,
78             delayMinMs: 1, delayMaxMs: 10
79         ),
80         new StockSender(
81             StockAppConstants.TOPIC,
82             new StockPrice( name: "INEL", dollars: 100, cents: 99),
83             new StockPrice( name: "INEL", dollars: 50, cents: 10),
84             producer,
85             delayMinMs: 1, delayMaxMs: 10
86     );
87 }
```



StockPriceSerializer

```
1 package com.cloudurable.kafka.producer;
2 import com.cloudurable.kafka.producer.model.StockPrice;
3 import org.apache.kafka.common.serialization.Serializer;
4 import java.nio.charset.StandardCharsets;
5 import java.util.Map;
6
7 public class StockPriceSerializer implements Serializer<StockPrice> {
8
9     @Override
10    public byte[] serialize(String topic, StockPrice data) {
11        return data.toJson().getBytes(StandardCharsets.UTF_8);
12    }
13
14    @Override
15    public void configure(Map<String, ?> configs, boolean isKey) {
16    }
17
18    @Override
19    public void close() {
20    }
21}
```

- ❖ Converts **StockPrice** into byte array



StockSender

- ❖ Generates random stock prices for a given ***StockPrice*** name,
- ❖ StockSender is Runnable
- ❖ 1 thread per StockSender
- ❖ Shows using ***KafkaProducer*** from many threads
- ❖ Delays random time between delayMin and delayMax,
 - ❖ then sends random StockPrice between ***stockPriceHigh*** and ***stockPriceLow***



StockSender imports, Runnable

```
StockSender
1 package com.cloudurable.kafka.producer;
2
3 import com.cloudurable.kafka.producer.model.StockPrice;
4 import org.apache.kafka.clients.producer.Producer;
5 import org.apache.kafka.clients.producer.ProducerRecord;
6 import org.apache.kafka.clients.producer.RecordMetadata;
7 import org.slf4j.Logger;
8 import org.slf4j.LoggerFactory;
9
10 import java.util.Date;
11 import java.util.Random;
12 import java.util.concurrent.ExecutionException;
13 import java.util.concurrent.Future;
14
15 public class StockSender implements Runnable{
```

- ❖ Imports Kafka **Producer**, **ProducerRecord**, **RecordMetadata**, **StockPrice**
- ❖ Implements Runnable, can be submitted to **ExecutionService**



StockSender constructor

StockSender

```
17     private final StockPrice stockPriceHigh;
18     private final StockPrice stockPriceLow;
19     private final Producer<String, StockPrice> producer;
20     private final int delayMinMs;
21     private final int delayMaxMs;
22     private final Logger logger = LoggerFactory.getLogger(StockSender.class);
23     private final String topic;
24
25     public StockSender(final String topic, final StockPrice stockPriceHigh,
26                         final StockPrice stockPriceLow,
27                         final Producer<String, StockPrice> producer,
28                         final int delayMinMs,
29                         final int delayMaxMs) {
30         this.stockPriceHigh = stockPriceHigh;
31         this.stockPriceLow = stockPriceLow;
32         this.producer = producer;
33         this.delayMinMs = delayMinMs;
34         this.delayMaxMs = delayMaxMs;
35         this.topic = topic;
36     }
```

- ❖ takes a topic, high & low stockPrice, producer, delay min & max



StockSender run()

```
StockSender run()

38
39 ①↑ public void run() {
40     final Random random = new Random(System.currentTimeMillis());
41     int sentCount = 0;
42
43     while (true) {
44         sentCount++;
45         final ProducerRecord<String, StockPrice> record =
46                         createRandomRecord(random);
47         final int delay = randomIntBetween(random, delayMaxMs, delayMinMs);
48
49         try {
50             final Future<RecordMetadata> future = producer.send(record);
51             if (sentCount % 100 == 0) {displayRecordMetaData(record, future);}
52             Thread.sleep(delay);
53         } catch (InterruptedException e) {
54             if (Thread.interrupted()) {
55                 break;
56             }
57         } catch (ExecutionException e) {
58             logger.error("problem sending record to producer", e);
59         }
60     }
61 }
```

- ❖ In loop, creates random record, **send** record, waits random time

StockSender

createRandomRecord



StockSender createRandomRecord()

```
77
78     private final int randomIntBetween(final Random random,
79                                         final int max,
80                                         final int min) {
81         return random.nextInt( bound: max - min + 1 ) + min;
82     }
83
84     private ProducerRecord<String, StockPrice> createRandomRecord(
85         final Random random) {
86
87         final int dollarAmount = randomIntBetween(random,
88             stockPriceHigh.getDollars(), stockPriceLow.getDollars());
89
90         final int centAmount = randomIntBetween(random,
91             stockPriceHigh.getCents(), stockPriceLow.getCents());
92
93         final StockPrice stockPrice = new StockPrice(
94             stockPriceHigh.getName(), dollarAmount, centAmount);
95
96         return new ProducerRecord<>(topic, stockPrice.getName(),
97             stockPrice);
98     }
```

- ❖ *createRandomRecord* uses *randomIntBetween*
- ❖ creates **StockPrice** and then wraps **StockPrice** in **ProducerRecord**

StockSender

displayRecordMetaData



StockSender displayRecordMetaData()

```
62
63     private void displayRecordMetaData(final ProducerRecord<String, StockPrice> record,
64                                         final Future<RecordMetadata> future)
65                                         throws InterruptedException, ExecutionException {
66         final RecordMetadata recordMetadata = future.get();
67         logger.info(String.format("\n\t\tkey=%s, value=%s " +
68                         "\n\t\tsent to topic=%s part=%d off=%d at time=%s",
69                         record.key(),
70                         record.value().toJson(),
71                         recordMetadata.topic(),
72                         recordMetadata.partition(),
73                         recordMetadata.offset(),
74                         new Date(recordMetadata.timestamp())
75                     ));
76     }
```

- ❖ Every 100 records displayRecordMetaData gets called
- ❖ Prints out record info, and recordMetadata info:
 - ❖ key, JSON value, topic, partition, offset, time
 - ❖ uses Future from call to **producer.send()**

Run it

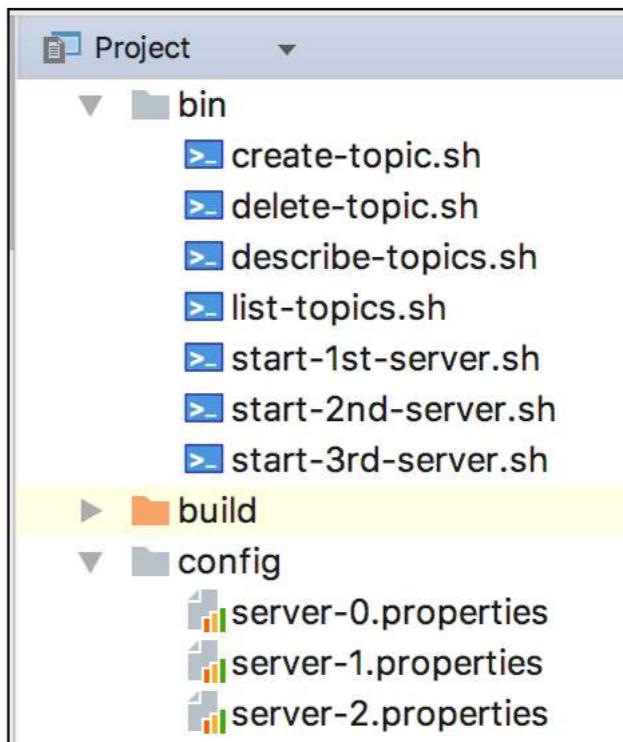


- ❖ Run ZooKeeper
- ❖ Run three Brokers
- ❖ run `create-topic.sh`
- ❖ Run **`StockPriceKafkaProducer`**



Run scripts

- ❖ run ZooKeeper from ~/kafka-training
- ❖ use ***bin/create-topic.sh*** to create topic
- ❖ use ***bin/delete-topic.sh*** to delete topic
- ❖ use ***bin/start-1st-server.sh*** to run Kafka Broker 0
- ❖ use ***bin/start-2nd-server.sh*** to run Kafka Broker 1
- ❖ use ***bin/start-3rd-server.sh*** to run Kafka Broker 2



Config is under directory called config

server-0.properties is for Kafka Broker 0
server-1.properties is for Kafka Broker 1
server-2.properties is for Kafka Broker 2



Run All 3 Brokers

```

1 #!/usr/bin/env bash
2 CONFIG=`pwd`/config
3 cd ~/kafka-training
4 ## Run Kafka
5 kafka/bin/kafka-server-start.sh \
6   "$CONFIG/server-0.properties"
7

```

```

start-2nd-server.sh x

1 #!/usr/bin/env bash
2 CONFIG=`pwd`/config
3 cd ~/kafka-training
4 ## Run Kafka
5 kafka/bin/kafka-server-start.sh \
6   "$CONFIG/server-1.properties"
7

```

```

start-3rd-server.sh x

1 #!/usr/bin/env bash
2 CONFIG=`pwd`/config
3 cd ~/kafka-training
4 ## Run Kafka
5 kafka/bin/kafka-server-start.sh \
6   "$CONFIG/server-2.properties"
7

```

server-0.properties x

```

1 broker.id=0
2 port=9092
3 log.dirs=./logs/kafka-0
4 min.insync.replicas=3
5 compression.type=producer
6 auto.create.topics.enable=false
7 message.max.bytes=65536
8 replica.lag.time.max.ms=5000
9 delete.topic.enable=true

```

server-1.properties x

```

1 broker.id=1
2 port=9093
3 log.dirs=./logs/kafka-1
4 min.insync.replicas=3
5 compression.type=producer

```

server-2.properties x

```

1 broker.id=2
2 port=9094
3 log.dirs=./logs/kafka-2
4 min.insync.replicas=3
5 compression.type=producer
6 auto.create.topics.enable=false
7 message.max.bytes=65536
8 replica.lag.time.max.ms=5000
9 delete.topic.enable=true

```



Run create-topic.sh script

```
create-topic.sh x
1 #!/usr/bin/env bash
2
3 cd ~/kafka-training
4
5 kafka/bin/kafka-topics.sh \
6   --create \
7   --zookeeper localhost:2181 \
8   --replication-factor 3 \
9   --partitions 3 \
10  --topic stock-prices
11
```

Name of the topic
is stock-prices

Three partitions

Replication factor
of three

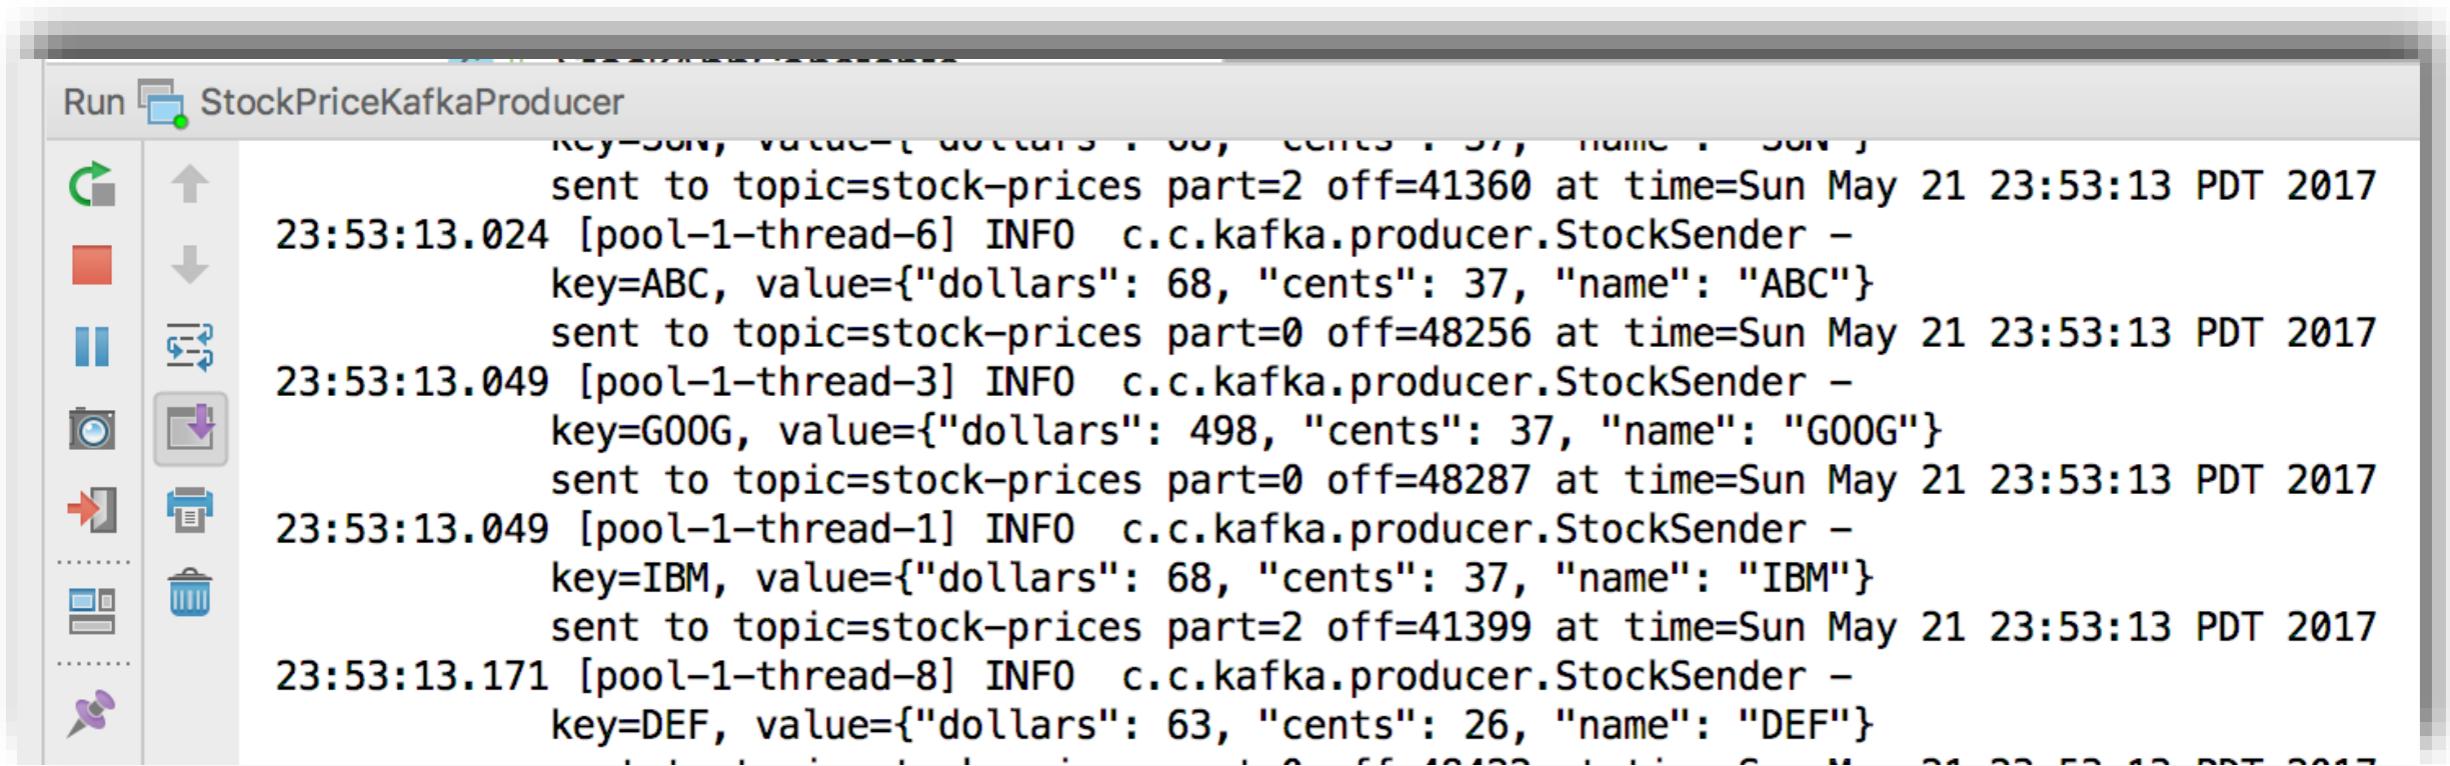
Terminal

+	Local	Local (1)	Local (2)	Local (3)	Local (4)
---	-------	-----------	-----------	-----------	-----------

\$./bin/create-topic.sh
Created topic "stock-prices".



Run StockPriceKafkaProducer



```
Run StockPriceKafkaProducer
23:53:13.024 [pool-1-thread-6] INFO c.c.kafka.producer.StockSender - key=ABC, value={"dollars": 68, "cents": 37, "name": "ABC"} sent to topic=stock-prices part=2 off=41360 at time=Sun May 21 23:53:13 PDT 2017
23:53:13.049 [pool-1-thread-3] INFO c.c.kafka.producer.StockSender - key=G00G, value={"dollars": 498, "cents": 37, "name": "G00G"} sent to topic=stock-prices part=0 off=48256 at time=Sun May 21 23:53:13 PDT 2017
23:53:13.049 [pool-1-thread-1] INFO c.c.kafka.producer.StockSender - key=IBM, value={"dollars": 68, "cents": 37, "name": "IBM"} sent to topic=stock-prices part=0 off=48287 at time=Sun May 21 23:53:13 PDT 2017
23:53:13.171 [pool-1-thread-8] INFO c.c.kafka.producer.StockSender - key=DEF, value={"dollars": 63, "cents": 26, "name": "DEF"} sent to topic=stock-prices part=2 off=41399 at time=Sun May 21 23:53:13 PDT 2017
```

- ❖ Run **StockPriceKafkaProducer** from the IDE
- ❖ You should see log messages from **StockSender(s)** with **StockPrice** name, JSON value, partition, offset, and time



Lab Adding an orderly shutdown flush and close

Using flush and close



Shutdown Producer nicely

- ❖ Handle ctrl-C shutdown from Java
- ❖ Shutdown thread pool and wait
- ❖ Flush producer to send any outstanding batches if using batches (***producer.flush()***)
- ❖ Close Producer (***producer.close()***) and wait



Nice Shutdown producer.close()

StockPriceKafkaProducer.java ×

```
StockPriceKafkaProducer main()
41 > 41  public static void main(String... args)
42      throws Exception {
43          //Create Kafka Producer
44          final Producer<String, StockPrice> producer = createProducer();
45          //Create StockSender list
46          final List<StockSender> stockSenders = getStockSenderList(producer);
47          //Create a thread pool so every stock sender gets its own.
48          final ExecutorService executorService =
49              Executors.newFixedThreadPool(stockSenders.size());
50          //Run each stock sender in its own thread.
51          stockSenders.forEach(executorService::submit);
52
53          //Register nice shutdown of thread pool, then flush and close producer.
54          Runtime.getRuntime().addShutdownHook(new Thread(() -> {
55              executorService.shutdown();
56              try {
57                  executorService.awaitTermination( timeout: 200, TimeUnit.MILLISECONDS );
58                  logger.info("Flushing and closing producer");
59                  producer.flush();
60                  producer.close( timeout: 10_000, TimeUnit.MILLISECONDS );
61              } catch (InterruptedException e) {
62                  logger.warn("shutting down", e);
63              }
64          }));
65      }
```

Restart Producer then shut it down



- ❖ Add shutdown hook
- ❖ Start StockPriceKafkaProducer
- ❖ Now stop it (CTRL-C or hit stop button in IDE)

The screenshot shows the Eclipse IDE interface with the 'External Terminals' view open. The terminal window title is 'Run StockPriceKafkaProducer'. The output log shows the producer sending four stock price messages to the 'stock-prices' topic and then shutting down:

```
00:06:52.709 [pool-1-thread-4] INFO c.c.kafka.producer.StockSender - key=INEL, value={"dollars": 73, "cents": 10, "name": "INEL"} sent to topic=stock-prices part=0 off=93522 at time=Mon May 22 00:06:52 PDT 2017
00:06:52.710 [pool-1-thread-5] INFO c.c.kafka.producer.StockSender - key=UBER, value={"dollars": 547, "cents": 90, "name": "UBER"} sent to topic=stock-prices part=2 off=80143 at time=Mon May 22 00:06:52 PDT 2017
00:06:52.717 [pool-1-thread-2] INFO c.c.kafka.producer.StockSender - key=SUN, value={"dollars": 73, "cents": 10, "name": "SUN"} sent to topic=stock-prices part=2 off=80146 at time=Mon May 22 00:06:52 PDT 2017
00:06:52.718 [pool-1-thread-1] INFO c.c.kafka.producer.StockSender - key=IBM, value={"dollars": 73, "cents": 10, "name": "IBM"} sent to topic=stock-prices part=2 off=80148 at time=Mon May 22 00:06:52 PDT 2017
00:06:52.967 [Thread-1] INFO c.c.k.p.StockPriceKafkaProducer - Flushing and closing producer
00:06:52.968 [Thread-1] INFO o.a.k.clients.producer.KafkaProducer - Closing the Kafka producer
Process finished with exit code 130 (interrupted by signal 2: SIGINT)
```



Lab Configuring Producer Durability



Lab Batching Records



Lab Adding Retries and Timeouts



Kafka Advanced Consumers Currently Not added to sample



Working with Avro Serialization and the Schema Registry

Schema Management in Kafka

Understanding Avro and the Schema Registry for Kafka



Overview

- ❖ First we cover Avro
 - ❖ Avro Schema
 - ❖ Using a Schema to write data
 - ❖ Using a Schema to read data
- ❖ Then we cover the Schema Registry
 - ❖ Managing Schemas
 - ❖ Kafka Avro Serialization



Avro

Introduction to Avro for Big Data and Streaming Architecture

Apache Avro Data
Serialization for Hadoop and
Kafka



Avro Overview

- ❖ Why Avro?
- ❖ What is Avro?
- ❖ Why are schemas important
- ❖ Examples reading and writing with Avro
- ❖ Working with Code generation
- ❖ Working with Avro Generic Record



Why Avro?

- ❖ Direct mapping to JSON
- ❖ Compact binary format
- ❖ Very Fast
- ❖ Polyglot bindings
- ❖ Code generation for static languages
- ❖ Code generation not needed
- ❖ Supports Schema
- ❖ Supports compatibility checks
- ❖ Allows evolving your data over time



Apache Avro

- ❖ Data serialization system
- ❖ Data structures
- ❖ Binary data format
- ❖ Container file format to store persistent data
- ❖ RPC capabilities
- ❖ Does not require code generation to use



Avro Schemas

- ❖ Supports schemas for defining data structure
- ❖ Serializing and deserializing data, uses schema
- ❖ File schema
 - ❖ Avro files store data with its schema
- ❖ RPC Schema
 - ❖ RPC protocol exchanges schemas as part of the handshake
- ❖ Schemas written in JSON



Avro compared to...

- ❖ Similar to Thrift, Protocol Buffers, JSON, etc.
- ❖ Does not require code generation
- ❖ Avro needs less encoding as part of the data since it stores names and types in the schema
- ❖ It supports evolution of schemas.



Why schemas?

- ❖ Multiple Producers and Consumers evolve over time
- ❖ Schemas keep your data clean
 - ❖ NoSQL was no schema but not as much: Cassandra has schemas
 - ❖ REST/JSON was schema-less and IDL-less but not anymore Swagger and RAML

Schema provides Robustness



- ❖ Streaming architecture support
 - ❖ **Challenging:** Consumers and Producers **evolve** on different timelines.
 - ❖ Producers send streams of records that 0 to * Consumers read
 - ❖ All consumers or use cases are **not known ahead of time**
 - ❖ Data could go to Hadoop or some store used years later **in use cases you can't imagine**
 - ❖ Schemas help **future proof data**; makes data more robust.
 - ❖ Avro Schema with support for evolution is important
 - ❖ **Schema provides robustness:** provides meta-data about data stored with data



Future Compatibility

- ❖ Data record format compatibility is a hard problem to solve
- ❖ Schema capture a point in time of what your data when it was stored
- ❖ Data will evolve. New fields are added.
- ❖ Kafka streams are often recorded in data lakes - needs to be less rigid than schema for operational data
- ❖ Producers can have Consumers that they never know about
- ❖ You can't test a Consumer that you don't know about
- ❖ Avro Schemas support agility sakes
- ❖ 80% of data science is data cleanup, Schemas help



Avro Schema

The screenshot shows the IntelliJ IDEA IDE interface. The left sidebar displays the project structure under the 'avro' module:

- .gradle
- .idea
- build
 - classes
 - dependency-cache
- generated-main-avro-java [main] s
 - com.cloudurable.phonebook
 - c Employee
- libs
- tmp
- gradle
- src
 - main
 - avro
 - com.cloudurable.phonebook

The 'Employee.avsc' file is open in the main editor window. The code is a JSON schema for an 'Employee' record:

```
1 {"namespace": "com.cloudurable.phonebook",
2   "type": "record",
3   "name": "Employee",
4   "fields": [
5     {"name": "firstName", "type": "string"},
6     {"name": "nickName", "type": ["string", "null"]},
7     {"name": "lastName", "type": "string"},
8     {"name": "age", "type": "int"},
9     {"name": "phoneNumber", "type": "string"}]
```

Avro schema stored in `src/main/avro` by default.



Code Generation

The screenshot shows a Java project structure in an IDE. The project root is named 'avro' and contains files like .gradle, .idea, build, gradle, and src. The src directory has main and test sub-directories. The main/java directory contains an 'employees.avro' file. The test/java/com.cloudurable directory contains an 'EmployeeTest.java' file. The build.gradle file is open in the code editor, defining a Gradle build script. It includes dependencies on 'com.commercehub.gradle.plugin.avro' and 'org.apache.avro:avro:1.8.1', and repositories for jcenter() and mavenCentral(). The avro plugin configuration specifies createSetters = false and fieldVisibility = "PRIVATE".

```
plugins {
    id "com.commercehub.gradle.plugin.avro" version "0.9.0"
}

group 'cloudurable'
version '1.0-SNAPSHOT'
apply plugin: 'java'
sourceCompatibility = 1.8

dependencies {
    compile "org.apache.avro:avro:1.8.1"
    testCompile group: 'junit', name: 'junit', version: '4.11'
}

repositories {
    jcenter()
    mavenCentral()
}

avro {
    createSetters = false
    fieldVisibility = "PRIVATE"
}
```

Employee Code Generation



The screenshot shows an IDE interface with the following details:

- Project Structure:** The left pane displays the project structure under "generated-main-avro-java". It includes folders for **avro**, **build**, **gradle**, **src**, and **test**. In the **src/main/java/com.cloudurable.phonebook** directory, there is an **Employee.java** file and an **Employee.avsc** schema file.
- Code Editor:** The right pane shows the generated **Employee.java** code. The code is annotated with **Employee** and **Employee()**. It includes imports for **com.cloudurable.phonebook** and **org.apache.avro.specific.SpecificData**. It uses annotations **@SuppressWarnings("all")** and **@org.apache.avro.specific.AvroGenerated**. The class **Employee** extends **org.apache.avro.specific.SpecificRecord**. It has private static final fields for **serialVersionUID** and **SCHEMA\$**. A static method **getClassSchema()** returns **SCHEMA\$**. The class contains four private fields: **firstName**, **lastName**, **age**, and **phoneNumber**. A detailed comment at the bottom explains the default constructor and its behavior regarding field initialization.



Using Generated Avro class

```
Employee bob = Employee.newBuilder().setAge(35)
    .setFirstName("Bob")
    .setLastName("Jones")
    .setPhoneNumber("555-555-1212")
    .build();

assertEquals(expected: "Bob", bob.getFirstName());
```



Writing employees to an Avro File

```
final DatumWriter<Employee> datumWriter = new SpecificDatumWriter<>(Employee.class);
final DataFileWriter<Employee> dataFileWriter = new DataFileWriter<>(datumWriter);

try {
    dataFileWriter.create(employeeList.get(0).getSchema(),
        new File( pathname: "employees.avro"));
    employeeList.forEach(employee -> {
        try {
            dataFileWriter.append(employee);
        } catch (IOException e) {
            throw new RuntimeException(e);
        }
    });
} finally {
    dataFileWriter.close();
}
```

Reading employees From a File



```
final File file = new File( pathname: "employees.avro");
final List<Employee> employeeList = new ArrayList<>();
final DatumReader<Employee> empReader = new SpecificDatumReader<>(Employee.class);
final DataFileReader<Employee> dataFileReader = new DataFileReader<>(file, empReader);

while (dataFileReader.hasNext()) {
    employeeList.add(dataFileReader.next(new Employee()));
}

employeeList.forEach(System.out::println);
```

All 2 tests passed – 333ms

```
/Library/Java/JavaVirtualMachines/jdk1.8.0_66.jdk/Contents/Home/bin/java ...
{"firstName": "Bob0", "lastName": "Jones0", "age": 25, "phoneNumber": "555-555-1212"}
{"firstName": "Bob1", "lastName": "Jones1", "age": 26, "phoneNumber": "555-555-1212"}
{"firstName": "Bob2", "lastName": "Jones2", "age": 27, "phoneNumber": "555-555-1212"}
{"firstName": "Bob3", "lastName": "Jones3", "age": 28, "phoneNumber": "555-555-1212"}
{"firstName": "Bob4", "lastName": "Jones4", "age": 29, "phoneNumber": "555-555-1212"}
 {"firstName": "Bob5", "lastName": "Jones5", "age": 30, "phoneNumber": "555-555-1212"}
 {"firstName": "Bob6", "lastName": "Jones6", "age": 31, "phoneNumber": "555-555-1212"}
 {"firstName": "Bob7", "lastName": "Jones7", "age": 32, "phoneNumber": "555-555-1212"}
 {"firstName": "Bob8", "lastName": "Jones8", "age": 33, "phoneNumber": "555-555-1212"}
 {"firstName": "Bob9", "lastName": "Jones9", "age": 34, "phoneNumber": "555-555-1212"}
 {"firstName": "Bob10", "lastName": "Jones10", "age": 35, "phoneNumber": "555-555-1212"}
 {"firstName": "Bob11", "lastName": "Jones11", "age": 36, "phoneNumber": "555-555-1212"}
 {"firstName": "Bob12", "lastName": "Jones12", "age": 37, "phoneNumber": "555-555-1212"}  
...
```



Using GenericRecord

```
final String schemaLoc = "src/main/avro/com/cloudurable/phonebook/Employee.avsc";
final File schemaFile = new File(schemaLoc);
final Schema schema = new Schema.Parser().parse(schemaFile);

GenericRecord bob = new GenericData.Record(schema);
bob.put(key: "firstName", v: "Bob");
bob.put(key: "lastName", v: "Smith");
bob.put(key: "phoneNumber", v: "555-555-1212");
bob.put(key: "age", v: 35);

assertEquals(expected: "Bob", bob.get("firstName"));
```



Writing Generic Records

```
final List<GenericRecord> employeeList = new ArrayList<>();
```

```
final DatumWriter<GenericRecord> datumWriter = new GenericDatumWriter<>(schema);
final DataFileWriter<GenericRecord> dataFileWriter = new DataFileWriter<>(datumWriter);

try {
    dataFileWriter.create(employeeList.get(0).getSchema(),
        new File( pathname: "employees2.avro"));
    employeeList.forEach(employee -> {
        try {
            dataFileWriter.append(employee);
        } catch (IOException e) {
            throw new RuntimeException(e);
        }
    });
} finally {
    dataFileWriter.close();
}
```

Reading using Generic Records



```
final File file = new File( pathname: "employees2.avro");
final List<GenericRecord> employeeList = new ArrayList<>();
final DatumReader<GenericRecord> empReader = new GenericDatumReader<>();
final DataFileReader<GenericRecord> dataFileReader = new DataFileReader<>(file, empReader);

while (dataFileReader.hasNext()) {
    employeeList.add(dataFileReader.next( reuse: null));
}

employeeList.forEach(System.out::println);
```

All 2 tests passed – 307ms

```
/Library/Java/JavaVirtualMachines/jdk1.8.0_66.jdk/Contents/Home/bin/java ...
{"firstName": "Bob0", "nickName": null, "lastName": "Smith0", "age": 25, "phoneNumber": "555-555-1212"}
{"firstName": "Bob1", "nickName": null, "lastName": "Smith1", "age": 26, "phoneNumber": "555-555-1212"}
{"firstName": "Bob2", "nickName": null, "lastName": "Smith2", "age": 27, "phoneNumber": "555-555-1212"}
{"firstName": "Bob3", "nickName": null, "lastName": "Smith3", "age": 28, "phoneNumber": "555-555-1212"}
{"firstName": "Bob4", "nickName": null, "lastName": "Smith4", "age": 29, "phoneNumber": "555-555-1212"}
{"firstName": "Bob5", "nickName": null, "lastName": "Smith5", "age": 30, "phoneNumber": "555-555-1212"}
{"firstName": "Bob6", "nickName": null, "lastName": "Smith6", "age": 31, "phoneNumber": "555-555-1212"}
{"firstName": "Bob7", "nickName": null, "lastName": "Smith7", "age": 32, "phoneNumber": "555-555-1212"}
{"firstName": "Bob8", "nickName": null, "lastName": "Smith8", "age": 33, "phoneNumber": "555-555-1212"}
{"firstName": "Bob9", "nickName": null, "lastName": "Smith9", "age": 34, "phoneNumber": "555-555-1212"}
```



Avro Schema Validation

```
GenericRecord employee = new GenericData.Record(schema);
employee.put( key: "firstName", v: "Bob" + index);
employee.put( key: "lastName", v: "Smith"+ index);
employee.put( key: "phoneNumber", v: "555-555-1212");
//employee.put("age", index % 35 + 25);
employee.put( key: "age", v: "old");
employeeList.add(employee);
```

2 tests done: 1 failed – 395ms

```
org.apache.avro.file.DataFileWriter$AppendWriteException: java.lang.ClassCastException: java.lang.String cannot be cast to java.lang.Number
        at org.apache.avro.file.DataFileWriter.append(DataFileWriter.java:308)
        at com.cloudurable.phonebook.EmployeeTestNoGen.lambda$testWrite$0(EmployeeTestNoGen.java:71)
        at java.util.ArrayList.forEach(ArrayList.java:1249)
        at com.cloudurable.phonebook.EmployeeTestNoGen.testWrite(EmployeeTestNoGen.java:69) <27 internal calls>
Caused by: java.lang.ClassCastException: java.lang.String cannot be cast to java.lang.Number
        at org.apache.avro.generic.GenericDatumWriter.writeWithoutConversion(GenericDatumWriter.java:117)
        at org.apache.avro.generic.GenericDatumWriter.write(GenericDatumWriter.java:73)
        at org.apache.avro.generic.GenericDatumWriter.writeField(GenericDatumWriter.java:153)
        at org.apache.avro.generic.GenericDatumWriter.writeRecord(GenericDatumWriter.java:143)
        at org.apache.avro.generic.GenericDatumWriter.writeWithoutConversion(GenericDatumWriter.java:105)
```



Avro supported types

- ❖ Records
- ❖ Arrays
- ❖ Enums
- ❖ Unions
- ❖ Maps
- ❖ String, Int, Boolean, Decimal, Timestamp, Date



Avro Types 1

avro type	json type	example	description
null	null	null	no value
boolean	boolean	true	true or false
int	integer	1	32 bit signed int
long			64 bit signed int
float	number	1.1	single precision (32-bit) floating-point number
double	number	1.2	double precision (64-bit) floating-point number
bytes	string	"\u00FF"	
string	string	"foo"	
record	object	{"a": 1}	complex type / object
enum	string	"Hi Mom"	Enumerated type
array	array	[1]	array of any value
union			can be used to define nullable fields
map	object	{"a": 1}	
fixed	string	"\u00ff"	fixed length string



Avro Types 2

avro type	json type	example	description
date	int	2000	stores the number of days from the unix epoch, 1 January 1970
time-millis	int	10800000	stores the number of milliseconds after midnight
time-micros	long	1000000	microseconds after midnight
timestamp-milis	long	1494437796880	milliseconds since epoch
decimal	complex/object	22.555	specifies scale and precision like BigDecimal in Java
duration	complex/object		Specifies duration in months, days and milliseconds

Fuller example Avro Schema



Project Tree:

- avro ~/kafka-training/avro
 - .gradle
 - .idea
 - build
 - classes
 - dependency-cache
 - generated-main-avro-java [main]
 - com.cloudurable.phonebook
 - Employee
 - PhoneNumber
 - Status
 - libs
 - reports
 - test-results
 - tmp
- gradle
- src
 - main
 - avro
 - com.cloudurable.phonebook
 - Employee.avsc
 - java
- test
 - java
 - com.cloudurable.phonebook
 - EmployeeTest
 - EmployeeTestNoGen

- build.gradle

```

1 {"namespace": "com.cloudurable.phonebook",
2  "type": "record",
3  "name": "Employee",
4  "fields": [
5    {"name": "firstName", "type": "string"},
6    {"name": "nickName", "type": ["null", "string"], "default": null},
7    {"name": "lastName", "type": "string"},
8    {"name": "age", "type": "int"},
9    {"name": "emails", "default": [], "type": {"type": "array", "items": "string"}},
10   {"name": "phoneNumber", "type":
11     [
12       "null",
13       {
14         "type": "record", "name": "PhoneNumber",
15         "fields": [
16           {"name": "areaCode", "type": "string"},
17           {"name": "countryCode", "type": "string", "default": ""},
18           {"name": "prefix", "type": "string"},
19           {"name": "number", "type": "string"}
20         ]
21       }
22     ],
23   {"name": "status", "default": "SALARY", "type": {
24     "type": "enum", "name": "Status",
25     "symbols": ["RETIRED", "SALARY", "HOURLY", "PART_TIME"]
26   }
27 }

```

```

public class PhoneNumber extends org.apache.avro.specific.SpecificRecord<org.apache.avro.Schema>
{
    private static final long serialVersionUID = -313877L;
    public static final org.apache.avro.Schema SCHEMA$ = ...
    public static org.apache.avro.Schema getClassSchema()
    {
        ...
    }
    private java.lang.String areaCode;
    private java.lang.String countryCode;
    private java.lang.String prefix;
    private java.lang.String number;
}

```

```

package com.cloudurable.phonebook;
@SuppressWarnings("all")
@org.apache.avro.specific.AvroGenerated
public enum Status {
    RETIRED, SALARY, HOURLY, PART_TIME ;
    public static final org.apache.avro.Schema SCHEMA$ = ...
}

```



Avro Record Attributes

- ❖ ***name***: name of the record (required).
- ❖ ***namespace***: equates to packages or modules
- ❖ ***doc***: documentation for future user of this schema
- ❖ ***aliases***: array aliases (alias names)
- ❖ ***fields***: an array of fields



Avro Field Attributes

- ❖ ***name***: name of the field (required)
- ❖ ***doc***: description of field (important for future usage)
- ❖ ***type***: JSON object defining a schema, or a JSON string naming a record definition (required)
- ❖ ***default***: Default value for this field
- ❖ ***order***: specifies sort ordering of record (optional).
 - ❖ "ascending" (the default), "descending", "ignore"
- ❖ ***aliases***: array of alternate names



Tips for Using Avro

- ❖ Avoid advance features which are not supported by polyglot language mappings
 - ❖ Think simple data objects
- ❖ Instead of magic strings use enums (better validation)
- ❖ Document all records and fields in the schema
- ❖ Avoid complex union types (use Unions for nullable fields only)
- ❖ Avoid recursive types
- ❖ Use reasonable field names
 - ❖ employee_id instead of id and then use employee_id in all other records that have a field that refer to the employee_id



Avro

- ❖ Fast data serialization
- ❖ Supports data structures
- ❖ Supports Records, Maps, Array, and basic types
- ❖ You can use it direct or use Code Generation
- ❖ [Read more.](#)



Avro

Confluent Schema Registry and Kafka Avro Serialization

Managing Record Schema in
Kafka



Overview

- ❖ What is the Schema Registry?
- ❖ Why do you want to use it?
- ❖ Understanding Avro Schema Evolution
- ❖ Setting up and using Schema Registry
- ❖ Managing Schemas with REST interface
- ❖ Writing Avro Serializer based Producers
- ❖ Writing Avro Deserializer based Consumers



Confluent Schema Registry

- ❖ Confluent Schema Registry manage Avro Schemas for Kafka clients
- ❖ Provides REST interface for putting and getting Avro schemas
- ❖ Stores a history of schemas
 - ❖ versioned
 - ❖ allows you to configure compatibility setting
 - ❖ supports evolution of schemas
- ❖ Provides serializers used by Kafka clients which handles schema storage and serialization of records using Avro
- ❖ Don't have to send schema just the schema uuid, schema can be looked up in Registry. Helps with Speed!

Avro Producer Schema Process



- ❖ Producer creates a record/message, which is an Avro record
- ❖ Record contains schema and data
- ❖ Schema Registry Avro Serializer serializes the data and schema id (just id)
 - ❖ Keeps a cache of registered schemas from Schema Registry to ids
- ❖ Consumer receives payload and deserializes it with Kafka Avro Deserializers which use Schema Registry
- ❖ Deserializer looks up the full schema from cache or Schema Registry based on id



Why Schema Registry?

- ❖ Consumer has its schema, one it is expecting record/message to conform to
 - ❖ Compatibility check is performed or two schemas
 - ❖ if no match, but are compatible, then payload transformation happens aka Schema Evolution
 - ❖ if not failure
- ❖ Kafka records have Key and Value and schema can be done on both



Schema Registry Actions

- ❖ Register schemas for key and values of Kafka records
- ❖ List schemas (subjects); List all versions of a subject (schema)
- ❖ Retrieve a schema by version or id; get latest version of schema
- ❖ Check to see if schema is compatible with a certain version
- ❖ Get the compatibility level setting of the Schema Registry
 - ❖ BACKWARDS, FORWARDS, FULL, NONE
- ❖ Add compatibility settings to a subject/schema



Schema Compatibility

- ❖ ***Backward compatibility***
 - ❖ data written with older schema can be read with newer schema
- ❖ ***Forward compatibility***
 - ❖ data written with newer schema can be read with an older schema
- ❖ ***Full compatibility***
 - ❖ New version of a schema is backward and forward compatible
- ❖ ***None***
 - ❖ Schema will not be validated for compatibility at all



Schema Registry Config

- ❖ Compatibility can be configured globally or per schema
- ❖ Options are:
 - ❖ NONE - don't check for schema compatibility
 - ❖ FORWARD - check to make sure last schema version is forward compatible with new schemas
 - ❖ BACKWARDS (default) - make sure new schema is backwards compatible with latest
 - ❖ FULL - make sure new schema is forwards and backwards compatible from latest to new and from new to latest



Schema Evolution

- ❖ Avro schema is changed after data has been written to store using an older version of that schema, then Avro might do a Schema Evolution
- ❖ Schema evolution happens only during deserialization at Consumer
- ❖ If Consumer's schema is different from Producer's schema, then value or key is automatically modified during deserialization to conform to consumers reader schema



Schema Evolution 2

- ❖ Schema evolution is automatic transformation of Avro schema
 - ❖ transformation is between version of consumer schema and what the producer put into the Kafka log
 - ❖ When Consumer schema is not identical to the Producer schema used to serialize the Kafka Record, then a data transformation is performed on the Kafka record (key or value)
 - ❖ If the schemas match then no need to do a transformation



Allowed Schema Modifications

- ❖ Add a field with a default
- ❖ Remove a field that had a default value
- ❖ Change a field's order attribute
- ❖ Change a field's default value
- ❖ Remove or add a field alias
- ❖ Remove or add a type alias
- ❖ Change a type to a union that contains original type

Rules of the Road for modifying Schema



- ❖ Provide a default value for fields in your schema
 - ❖ Allows you to delete the field later
- ❖ Don't change a field's data type
- ❖ When adding a new field to your schema, you have to provide a default value for the field
- ❖ Don't rename an existing field
 - ❖ You can add an alias



Remember our example Employee

The screenshot shows a file browser on the left and a code editor on the right. The file browser displays a project structure under the 'avro' folder, including '.gradle', '.idea', 'build', 'classes', 'dependency-cache', 'generated-main-avro-java [main]', 'com.cloudurable.phonebook', 'Employee', 'libs', 'tmp', 'gradle', 'src', 'main', 'avro', and 'com.cloudurable.phonebook'. The 'Employee.avsc' file is selected in the code editor. The code editor displays the following JSON-based Avro schema:

```
1 {"namespace": "com.cloudurable.phonebook",
2  "type": "record",
3  "name": "Employee",
4  "fields": [
5    {"name": "firstName", "type": "string"},
6    {"name": "nickName", "type": ["string", "null"]},
7    {"name": "lastName", "type": "string"},
8    {"name": "age", "type": "int"},
9    {"name": "phoneNumber", "type": "string"}]
```

Avro covered in [Avro/Kafka Tutorial](#)



Let's say

- ❖ Employee did not have an age in version 1 of the schema
- ❖ Later we decided to add an age field with a default value of -1
- ❖ Now let's say we have a Producer using version 2, and a Consumer using version 1



Scenario adding new field age w/ default value

- ❖ Producer uses version 2 of the Employee schema and creates a com.cloudurable.Employee record, and sets age field to 42, then sends it to Kafka topic new-employees
- ❖ Consumer consumes records from new-employees using version 1 of the Employee Schema
 - ❖ Since Consumer is using version 1 of schema, age field is removed during deserialization

Scenario adding new field age w/ default value 2



- ❖ Same consumer modifies some fields, then writes record to a NoSQL store
 - ❖ When it does this, the age field is missing from record that it writes to NoSQL store
- ❖ Another client using version 2 reads the record from the NoSQL store
 - ❖ Age field is missing from the record (because the Consumer wrote it with version 1), age is set to default value of -1

Schema Registry Actions

REST



- ❖ Register schemas for key and values of Kafka records
- ❖ List schemas (subjects)
- ❖ List all versions of a subject (schema)
- ❖ Retrieve a schema by version or id
 - ❖ get latest version of schema

Schema Registry Actions 2



- ❖ Check to see if schema is compatible with a certain version
- ❖ Get the compatibility level setting of the Schema Registry
 - ❖ BACKWARDS, FORWARD, FULL, NONE
- ❖ Add compatibility settings to a subject/schema



Register a Schema

```
private final static MediaType SCHEMA_CONTENT =
    MediaType.parse("application/vnd.schemaregistry.v1+json");

private final static String EMPLOYEE_SCHEMA = "{\n" +
    "  \"schema\": \"\" +\n    "{" +\n      "  \"namespace\": \"com.cloudurable.phonebook\", " +\n      "  \"type\": \"record\", " +\n      "  \"name\": \"Employee\", " +\n      "  \"fields\": [ " +\n        "    {\"name\": \"firstName\", \"type\": \"string\"}, " +\n        "    {\"name\": \"lastName\", \"type\": \"string\"}, " +\n        "    {\"name\": \"age\", \"type\": \"int\"}, " +\n        "    {\"name\": \"phoneNumber\", \"type\": \"string\"} " +\n      "  ]" +\n    " }\" +\n  \"\";\n}\n";
```



Register a Schema

```
final OkHttpClient client = new OkHttpClient();

//POST A NEW SCHEMA
Request request = new Request.Builder()
    .post(RequestBody.create(SCHEMA_CONTENT, EMPLOYEE_SCHEMA))
    .url("http://localhost:8081/subjects/Employee/versions")
    .build();

String output = client.newCall(request).execute().body().string();
System.out.println(output);
```

```
curl -X POST -H "Content-Type: application/vnd.schemaregistry.v1+json"
--data '{"schema": "{\"type\": ...}"' \
http://localhost:8081/subjects/Employee/versions
```

{"id":2}



List All Schema

```
//LIST ALL SCHEMAS
request = new Request.Builder()
    .url("http://localhost:8081/subjects")
    .build();

output = client.newCall(request).execute().body().string();
System.out.println(output);
```

```
curl -X GET http://localhost:8081/subjects
```

["Employee","Employee2","FooBar"]



Working with versions

```
//SHOW ALL VERSIONS OF EMPLOYEE
request = new Request.Builder()
    .url("http://localhost:8081/subjects/Employee/versions/")
    .build(); [1,2,3,4,5]

output = client.newCall(request).execute().body().string();
System.out.println(output);

//SHOW VERSION 2 OF EMPLOYEE
request = new Request.Builder()
    .url("http://localhost:8081/subjects/Employee/versions/2")
    .build();
output = client.newCall(request).execute().body().string();
System.out.println(output); {"subject":"Employee","version":2,"id":4,"schema": "{\"type\":\"record\",\"name\":\"Employee\", \"namespace\":\"com.cloudurable.phonebook\", ...}

//SHOW THE SCHEMA WITH ID 3
request = new Request.Builder()
    .url("http://localhost:8081/schemas/ids/3")
    .build();
output = client.newCall(request).execute().body().string(); {"subject":"Employee","version":1,"id":3,"schema": "{\"type\":\"record\",\"name\":\"Employee\", \"namespace\":\"com.cloudurable.phonebook\", ...}
System.out.println(output);
```



Working with Schemas

```
//SHOW THE LATEST VERSION OF EMPLOYEE 2
request = new Request.Builder()
    .url("http://localhost:8081/subjects/Employee/versions/latest")
    .build();

output = client.newCall(request).execute().body().string();
System.out.println(output);
```

```
//CHECK IF SCHEMA IS REGISTERED
request = new Request.Builder()
    .post(RequestBody.create(SCHEMA_CONTENT, EMPLOYEE_SCHEMA))
    .url("http://localhost:8081/subjects/Employee")
    .build();
```

```
output = client.newCall(request).execute().body().string();
System.out.println(output);
```

```
//TEST COMPATIBILITY
request = new Request.Builder()
    .post(RequestBody.create(SCHEMA_CONTENT, EMPLOYEE_SCHEMA))
    .url("http://localhost:8081/compatibility/subjects/Employee/versions/latest")
    .build();
```

Changing Compatibility Checks



```
// SET TOP LEVEL CONFIG
// VALUES are none, backward, forward and full
request = new Request.Builder()
    .put(RequestBody.create(SCHEMA_CONTENT, "{\"compatibility\": \"none\"}"))
    .url("http://localhost:8081/config")
    .build();

output = client.newCall(request).execute().body().string();
System.out.println(output);

// SET CONFIG FOR EMPLOYEE
// VALUES are none, backward, forward and full
request = new Request.Builder()
    .put(RequestBody.create(SCHEMA_CONTENT, "{\"compatibility\": \"backward\"}"))
    .url("http://localhost:8081/config/Employee")
    .build();

output = client.newCall(request).execute().body().string();
System.out.println(output);
```



Incompatible Change

```

private final static String EMPLOYEE_SCHEMA = "{\n" +
    "  \"schema\": \"\"\" +\n    \"{\" +\n      \"\\\"namespace\\\": \\\\"com.cloudurable.phonebook\\\\",\" +\n      \"\\\"type\\\": \\\\"record\\\\",\" +\n      \"\\\"name\\\": \\\\"Employee\\\\",\" +\n      \"\\\"fields\\\": [\" +\n        {\\\"name\\\": \\\\"fName\\\\", \\\"type\\\": \\\\"string\\\\"},\" +\n        {\\\"name\\\": \\\\"lName\\\\", \\\"type\\\": \\\\"string\\\\"},\" +\n        {\\\"name\\\": \\\\"age\\\\", \\\"type\\\": \\\\"int\\\\"},\" +\n        {\\\"name\\\": \\\\"phoneNumber\\\\", \\\"type\\\": \\\\"string\\\\"}\"]\" +\n      \"}\" +\n    \"}\" +\n  \"\"\" +\n\"};"

```

```

//POST A NEW SCHEMA
Request request = new Request.Builder()
    .post(RequestBody.create(SCHEMA_CONTENT, EMPLOYEE_SCHEMA))
    .url("http://localhost:8081/subjects/Employee/versions")
    .build();

String output = client.newCall(request).execute().body().string();
System.out.println(output);

```

{"error_code":409,
 message:"Schema being registered is incompatible with an e



Incompatible Change

```
private final static String EMPLOYEE_SCHEMA = "{\n" +\n    "  \"schema\": \"\" +\n    "  {" +\n    "    \"namespace\": \"com.cloudurable.phonebook\", " +\n    "    \"type\": \"record\", " +\n    "    \"name\": \"Employee\", " +\n    "    \"fields\": [\" +\n    "      {\"name\": \"fName\", \"type\": \"string\"}, " +\n    "      {\"name\": \"lName\", \"type\": \"string\"}, " +\n    "      {\"name\": \"age\", \"type\": \"int\"}, " +\n    "      {\"name\": \"phoneNumber\", \"type\": \"string\"} " +\n    "    ]" +\n    "  }\" +\n"}";
```

```
//TEST COMPATIBILITY\nrequest = new Request.Builder()\n    .post(RequestBody.create(SCHEMA_CONTENT, EMPLOYEE_SCHEMA))\n    .url("http://localhost:8081/compatibility/subjects/Employee/versions/latest")\n    .build();
```

{"is_compatible":false}



Use Schema Registry

- ❖ Start up Schema Registry server pointing to Zookeeper cluster
- ❖ Import Kafka Avro Serializer and Avro Jars
- ❖ Configure Producer to use Schema Registry
- ❖ Use KafkaAvroSerializer from Producer
- ❖ Configure Consumer to use Schema Registry
- ❖ Use KafkaAvroDeserializer from Consumer



Start up Schema Registry Server

```
cat ~/tools/confluent-3.2.1/etc/schema-registry/schema-registry.properties
```

```
listeners=http://0.0.0.0:8081
kafkastore.connection.url=localhost:2181
kafkastore.topic=_schemas
debug=false
```

```
$ bin/schema-registry-start ~/tools/confluent-3.2.1/etc/schema-registry/schema-registry.properties
[2017-05-09 15:45:34,055] INFO SchemaRegistryConfig values:
  metric.reporters = []
  kafkastore.sasl.kerberos.kinit.cmd = /usr/bin/kinit
  response.mediatype.default = application/vnd.schemaregistry.v1+json
  kafkastore.ssl.trustmanager.algorithm = PKIX
  authentication.realm =
  ssl.keystore.type = JKS
  kafkastore.topic = _schemas
  metrics.jmx.prefix = kafka.schema.registry
```

Import Kafka Avro Serializer & Avro Jars



The screenshot shows a Java project structure in an IDE. The project is named "schema-registry". The build file, "build.gradle", is open in the editor. The code defines a group "cloudurable", version "1.0-SNAPSHOT", and applies the "java" plugin with source compatibility 1.8. It includes dependencies for org.apache.avro:avro:1.8.1, com.squareup.okhttp3:okhttp:3.7.0, junit:junit:4.11, org.apache.kafka:kafka-clients:0.10.2.0, and io.confluent:kafka-avro-serializer:3.2.1. It also configures repositories to use jcenter(), mavenCentral(), and a specific Maven repository at http://packages.confluent.io/maven/. The avro configuration section sets createSetters to false and fieldVisibility to PRIVATE.

```
group 'cloudurable'
version '1.0-SNAPSHOT'
apply plugin: 'java'
sourceCompatibility = 1.8

dependencies {
    compile "org.apache.avro:avro:1.8.1"
    compile 'com.squareup.okhttp3:okhttp:3.7.0'
    testCompile 'junit:junit:4.11'
    compile 'org.apache.kafka:kafka-clients:0.10.2.0'
    compile 'io.confluent:kafka-avro-serializer:3.2.1'
}

repositories {
    jcenter()
    mavenCentral()
    maven {
        url "http://packages.confluent.io/maven/"
    }
}

avro {
    createSetters = false
    fieldVisibility = "PRIVATE"
}
```

Configure Producer to use Schema Registry



The screenshot shows an IDE interface with the following details:

- Project Structure:** The project is named "schema-registry". It contains "src" and "main" directories. "src" includes "com.cloudurable.phonebook" (with "Employee", "PhoneNumber", and "Status" classes), "libs", "resources", "main", "tmp", and "gradle". "main" includes "avro" (with "employee.avsc" schema) and "java" (with "com.cloudurable.kafka" package containing "AvroConsumer", "AvroProducer", and "SchemaMain" classes).
- Code Editor:** The file "AvroProducer.java" is open. The code defines a class "AvroProducer" that creates a Kafka producer using Avro serialization. It sets properties for bootstrap servers, client ID, key serializer (LongSerializer), and value serializer (KafkaAvroSerializer). It also specifies the schema registry URL at "http://localhost:8081".

```
schema-registry > src > main > java > com > cloudurable > kafka > schema > AvroProducer.java

Project AvroProducer.java AbstractKafkaAvroSerDeConfig.class

com.cloudurable.phonebook Employee, PhoneNumber, Status
libs
resources
main
tmp
gradle
src
main
avro
employee.avsc
java
com.cloudurable.kafka
AvroConsumer, AvroProducer, SchemaMain
build.gradle
gradlew
gradlew.bat
settings.gradle
External Libraries
< 1.8 > /Library/Java/JavaVirtualMachines
Gradle: com.101tec:zkclient:0.10
Gradle: com.fasterxml.jackson.core:jackson-core:2.10.2
Gradle: com.fasterxml.jackson.core:jackson-databind:2.10.2
Gradle: com.fasterxml.jackson.core:jackson-annotations:2.10.2
Gradle: com.squareup.okhttp3:okhttp:4.3.1

AvroProducer
import org.apache.kafka.clients.producer.ProducerConfig;
import org.apache.kafka.clients.producer.ProducerRecord;
import org.apache.kafka.common.serialization.LongSerializer;
import io.confluent.kafka.serializers.KafkaAvroSerializer;

import java.util.Properties;
import java.util.stream.IntStream;

public class AvroProducer {

    private static Producer<Long, Employee> createProducer() {
        Properties props = new Properties();
        props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ProducerConfig.CLIENT_ID_CONFIG, "AvroProducer");
        props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
                  LongSerializer.class.getName());

        // Configure the KafkaAvroSerializer.
        props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
                  KafkaAvroSerializer.class.getName());

        // Schema Registry location.
        props.put(KafkaAvroSerializerConfig.SCHEMA_REGISTRY_URL_CONFIG,
                  "http://localhost:8081");

        return new KafkaProducer<>(props);
    }
}
```



Use KafkaAvroSerializer from Producer

The screenshot shows a Java code editor with the file `AvroProducer.java` open. The code demonstrates how to use the `KafkaAvroSerializer` from a producer. It includes imports for `Producer`, `Employee`, `PhoneNumber`, and `SchemaMain`. The code creates a producer for the topic "new-employees", builds an employee record (Bob Jones, age 35, phone number 301-555-1234), sends 100 such records to the topic, flushes the producer, and closes it. The code editor's sidebar shows the project structure with files like `employee.avsc` and `SchemaMain`.

```
private final static String TOPIC = "new-employees";

public static void main(String... args) {
    Producer<Long, Employee> producer = createProducer();

    Employee bob = Employee.newBuilder().setAge(35)
        .setFirstName("Bob")
        .setLastName("Jones")
        .setPhoneNumber(
            PhoneNumber.newBuilder()
                .setAreaCode("301")
                .setCountryCode("1")
                .setPrefix("555")
                .setNumber("1234")
                .build())
        .build();

    IntStream.range(1, 100).forEach(index->{
        producer.send(new ProducerRecord<>(TOPIC, key: 1L * index, bob));
    });

    producer.flush();
    producer.close();
}
```

Configure Consumer to use Schema Registry



```
AvroConsumer.java x
AvroConsumer
16 > public class AvroConsumer {
17
18     private final static String BOOTSTRAP_SERVERS = "localhost:9092";
19     private final static String TOPIC = "new-employees";
20
21     private static Consumer<Long, Employee> createConsumer() {
22         Properties props = new Properties();
23         props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, BOOTSTRAP_SERVERS);
24         props.put(ConsumerConfig.GROUP_ID_CONFIG, "KafkaExampleAvroConsumer");
25         props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
26                   LongDeserializer.class.getName());
27
28         //Use Kafka Avro Deserializer.
29         props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
30                   KafkaAvroDeserializer.class.getName()); //-----
31
32         //Use Specific Record or else you get Avro GenericRecord.
33         props.put(KafkaAvroDeserializerConfig.SPECIFIC_AVRO_READER_CONFIG, "true");
34
35
36         //Schema registry location.
37         props.put(KafkaAvroDeserializerConfig.SCHEMA_REGISTRY_URL_CONFIG,
38                   "http://localhost:8081"); //----- Run Schema Registry on 8081
39
40
41     return new KafkaConsumer<>(props);
42 }
43
```



Use KafkaAvroDeserializer from Consumer

```
public static void main(String... args) {  
  
    final Consumer<Long, Employee> consumer = createConsumer();  
    consumer.subscribe(Collections.singletonList(TOPIC));  
  
    IntStream.range(1, 100).forEach(index -> {  
  
        final ConsumerRecords<Long, Employee> records =  
            consumer.poll(timeout: 100);  
  
        if (records.count() == 0) {  
            System.out.println("None found");  
        } else records.forEach(record -> {  
  
            Employee employeeRecord = record.value();  
  
            System.out.printf("%s %d %d %s \n", record.topic(),  
                record.partition(), record.offset(), employeeRecord);  
        });  
    });  
}
```



Schema Registry

- ❖ Confluent provides Schema Registry to manage Avro Schemas for Kafka Consumers and Producers
- ❖ Avro provides Schema Migration
- ❖ Confluent uses Schema compatibility checks to see if Producer schema and Consumer schemas are compatible and to do Schema evolution if needed
- ❖ Use KafkaAvroSerializer from Producer
- ❖ Use KafkaAvroDeserializer from Consumer