


The Enterprise Guide to Building a Data Mesh

Introducing



Specification-Driven Data Mesh

Sion Smith CTO - oso.sh

Neil Avery CTO - liquidlabs.com

About us



Sion Smith

CTO, OSO

15 years consulting experience solving complex problems with various cloud and programming technologies



Neil Avery

CTO, Liquidlabs

Distributed systems, previously Confluent, Luxoft, startups and others



Emerging Technology



Agenda

Current **State of Play**

- The Gartner Hype Cycle
- Foundations of Data Mesh
- Evolution of central nervous system


The **Spec Mesh** Way

- Domain mapping out of the box
- Example specification
- Data Mesh lifecycle management

Developer Tooling

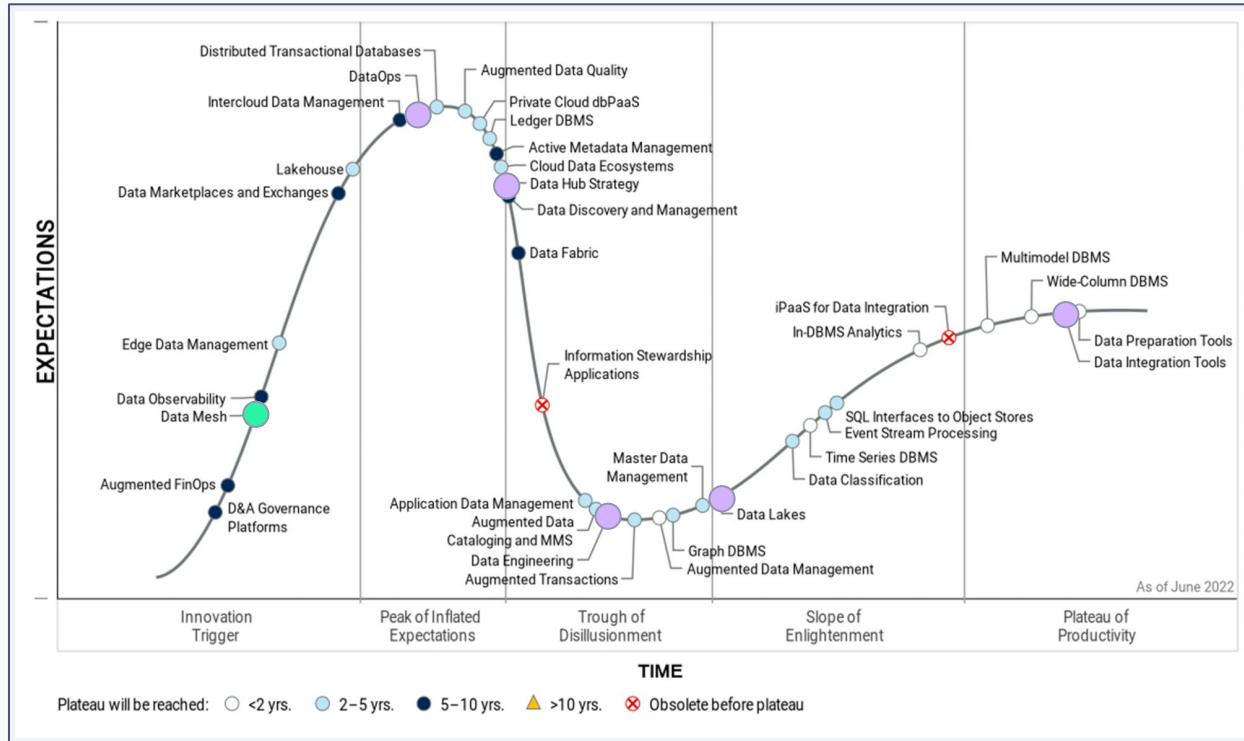
- Features of Spec Mesh
- Screenshots
- Developer roadmap

Our Data Mesh Journey



Is the Hype Real?


Gartner Hype Cycle for Data Management - 2022



Four pillars of Data Mesh

Domain-oriented
decentralised
data ownership &
architecture

Self-serve data
infrastructure as a platform



Data as a **Product**


Federated computational
Governance

Is Data Mesh *really* new?

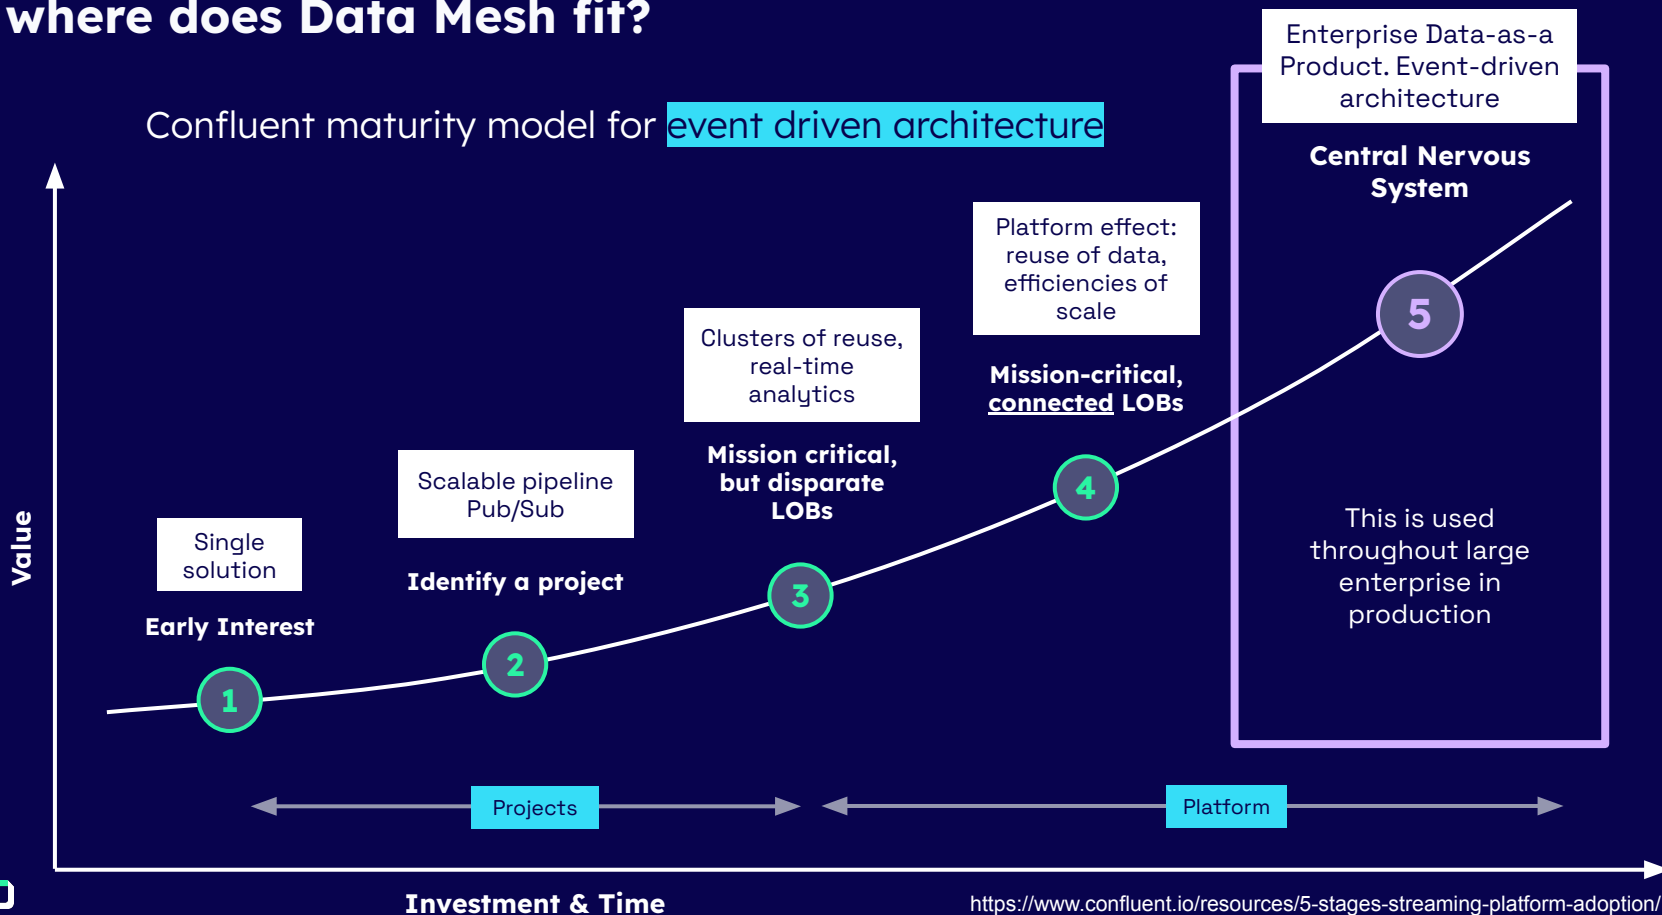
- + Data mesh **incremental evolution** of style of architecture we have been building for several years for **event streaming**
- + A mature data streaming system adopts a **central nervous system**
- + Can we build a data mesh around event streaming principles?
- + A central nervous system models topics with a domain structure and **federated computational governance**

Introducing:

- + An agreement / contract for data mesh using a specification



So where does Data Mesh fit?



Data Mesh should not boil the ocean



Introducing SpecMesh

Specification-Driven Data Mesh



CNS Patterns applied to Data Mesh

1. Events and storage comprise data platform fundamentals required to build almost anything
2. Events and storage **already exist** and that won't change
3. Organise data resources **hierarchically**
4. Enforcing a **domain model** is need to control complexity and cope with scale
5. Orgs have a **data quality** function
6. Orgs have a **data catalogue** function
7. **Registry models** are required to model data (and support data evolution etc)
8. Most organisations suffer from **inconsistent** mechanisms for data-as-an-api

Supporting the pillars of Data Mesh


Features	Domain ownership	Self-serve	Data as a Product	Federated Computational Governance
Spec driven (async-api spec)	✓		✓	
SDLC plugins (unit-test, integration-test)		✓		
SDLC plugins (provision - terraform)		✓	✓	✓
3rd party data catalogue	✓		✓	✓

Big Data London is our data product



Data Model

/ london / borough / venue / event
/ retail /
/ transport /
/ accommodation /
/ services /



/ london / hammersmith / olympia / bigdatalondon / public / attendee

/ london / hammersmith / transport / public / tube
/ london / heathrow / transport / public / airport

/ london / hammersmith / olympia / bigdatalondon / vendor / terra / public / visitor
/ london / hammersmith / olympia / bigdatalondon / retailer / subway / public / purchase

```
asyncapi: '2.4.0'
id: 'urn:london:hammersmith:olympia:bigdatalondon'
Info:
  title: BigDataLondon API
  version: '1.0.0'
  description: Simple model of BigDataLondon as a Data Product
```

Servers:

test:

```
url: test.mykafkacluster.org:8092
protocol: kafka-secure
```

channels:

public/attendee:

Publish:

```
summary: Humans arriving
```

Message:

```
name: Human
```

```
"tags": [
  "name": "human",
  "name": "big data london"]
```

Payload:

```
type: object
```

Properties:


Id:

```
type: integer
minimum: 0
description: Id of the human
```

Age:

```
type: integer
description: Age of the human
```

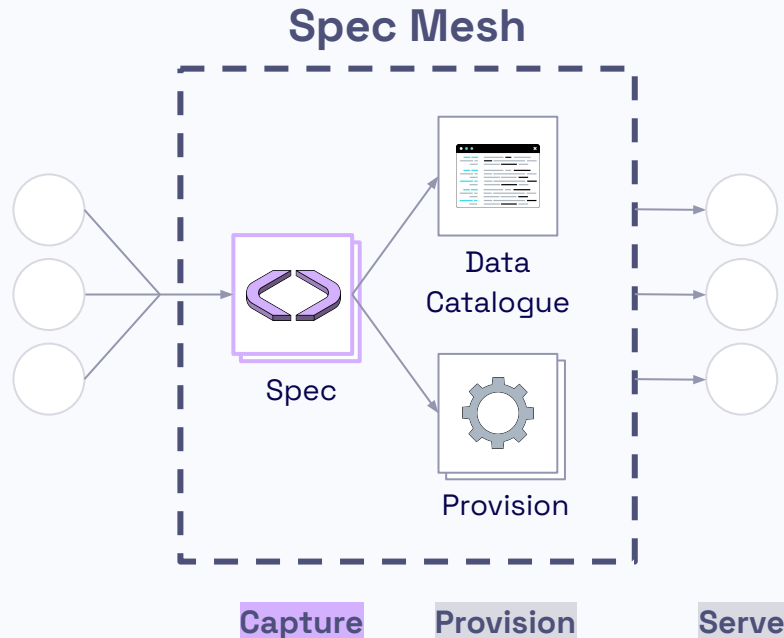
Data Mesh as Code: Spec Mesh



What is Spec Mesh?

Capture

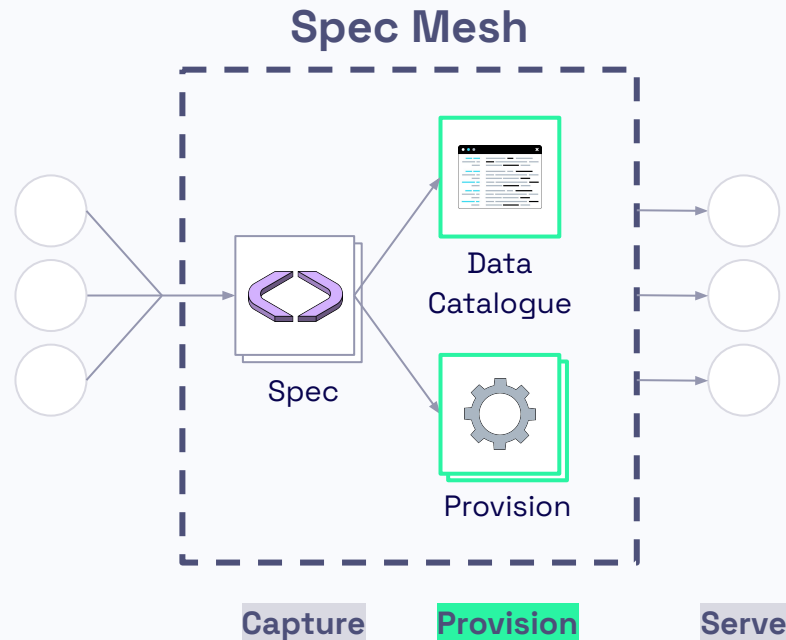
- Create repo from git template
- Async-api spec
- Apply the domain model/structure
- Specify data event streams & storage
- Model the data entities (schemas)
- Apply tags to data entities
- Permission data as public/protected/private
- Write tests and build your product



What is Spec Mesh?

Provision

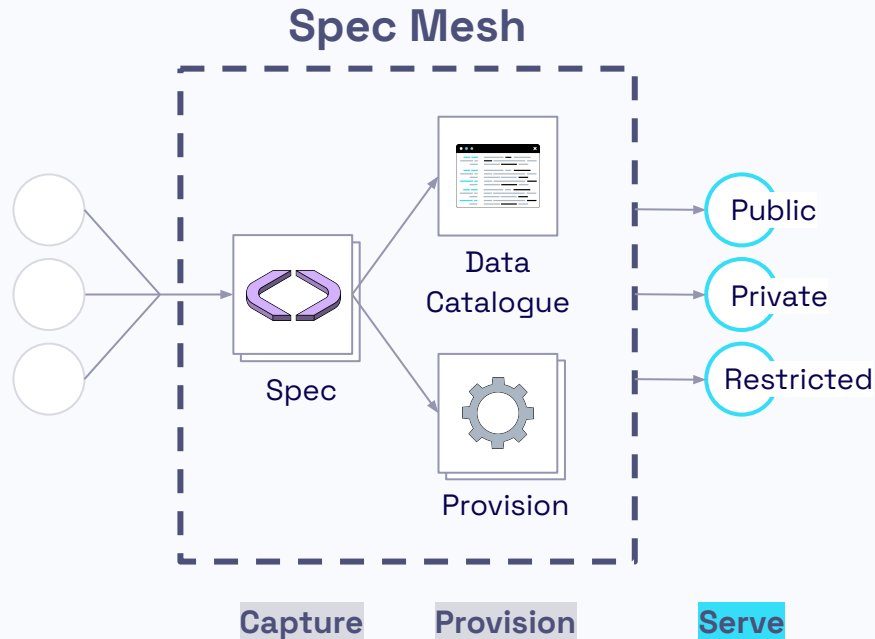
- Pipeline driven provisioning of data product based on Specification
- Includes domain structure, governance and permissions
- Data catalogue automatically updated (including tags)
- Data schema's published to registry



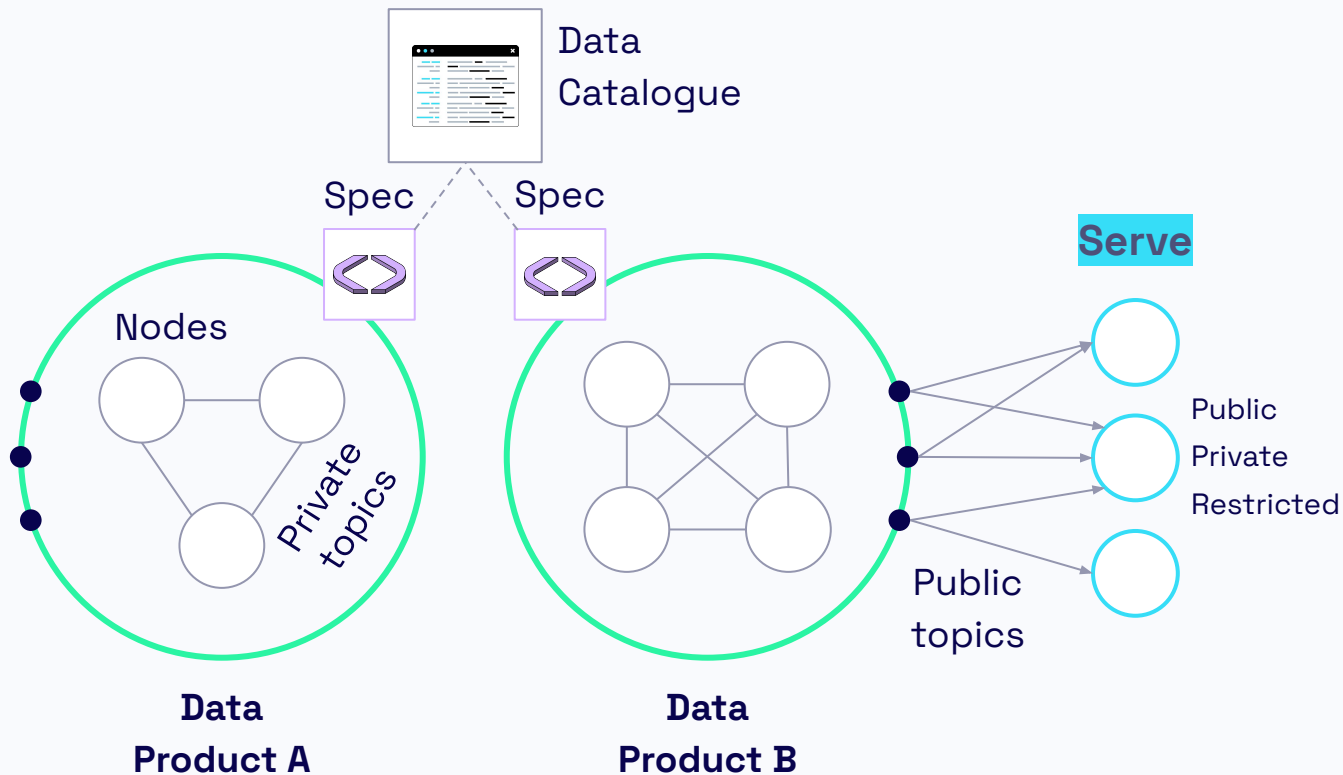
What is Spec Mesh?

Serve


- Governance integrated via underlying Data Platform. I.e. Kafka ACLs
- Data catalogue supports discovery (via tags), and access requests
- Data resources are structured using public, private and protected scope



Spec Mesh conceptual



Developer Tooling



Why use Spec Mesh?



Developer focused tooling: This is not some platform



Abstract away complexities: Leverage layers of abstraction to create repeatable patterns



Unit testing specifications: Java based tooling supports development lifecycle

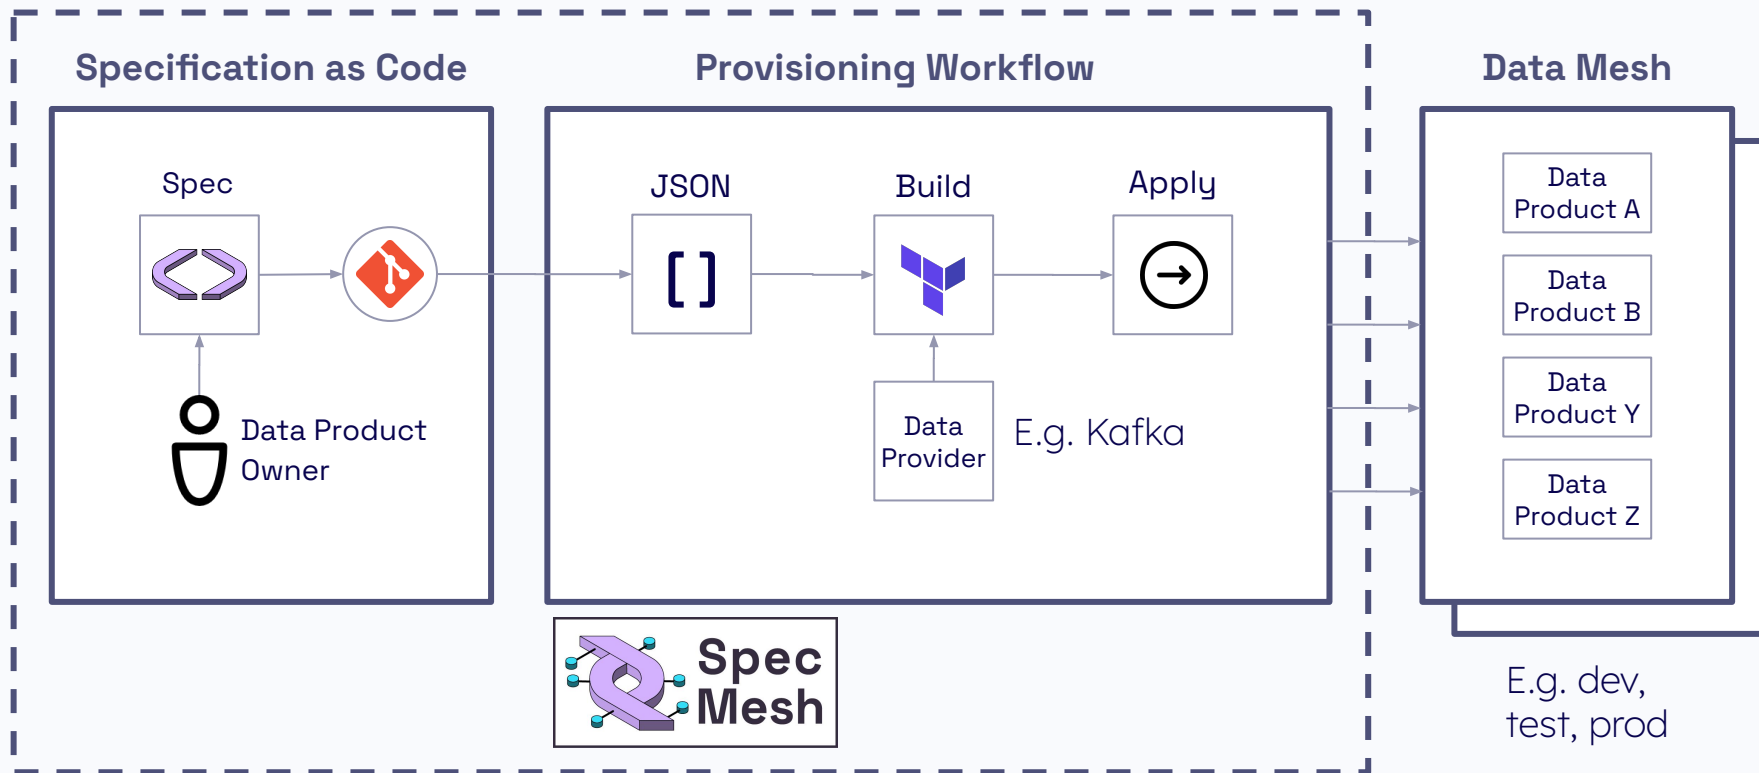


Tooling support flexibility: Using configurable guardrails



Increased extensibility: Interoperate with existing Data Management tooling and modules

What do you get out of the box?



Demo

Specification
output (JSON)

```

{
  "provider": {
    "kafka": [
      {
        "bootstrap_servers": [
          "localhost:29092"
        ],
        "sasl_password": "broker",
        "sasl_username": "broker",
        "tls_enabled": false
      }
    ]
  },
  "resource": {
    "kafka_topic": {
      "topic1": {
        "name": "sion1",
        "partitions": 3,
        "replication_factor": 1
      }
    }
  },
  "terraform": {
    "required_providers": {
      "kafka": {
        "source": "Mongey/kafka",
        "version": "0.5.1"
      }
    }
  }
}

```

Environment

Resources

Provisioning summary

Terraform has been successfully initialized!

No changes. Your infrastructure matches the configuration.

Terraform has compared your real infrastructure against your configuration and found no differences, so no changes are needed.

Apply complete! Resources: 0 added, 0 changed, 0 destroyed.

Kafka Infrastructure

```

➔ docker git:(main) x docker ps

```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS
21749a6e30a7	alpine/socat:1.7.4.3-r0	"/bin/sh -c 'socat T..."	3 seconds ago	Up 2 seconds	0.0.0.0:59109->2000/tcp
3b826c84d50b	confluentinc/cp-server:7.2.1	"/etc/confluent/dock..."	3 seconds ago	Up 2 seconds	0.0.0.0:9092->9092/tcp, 0.0.0.0:29092->29092/tcp
f3310329e408	confluentinc/cp-zookeeper:7.2.1	"/etc/confluent/dock..."	3 seconds ago	Up 2 seconds	2888/tcp, 0.0.0.0:2181->2181/tcp, 3888/tcp
820fdda7ebda	testcontainers/ryuk:0.3.3	"/app"	6 seconds ago	Up 5 seconds	0.0.0.0:59104->8080/tcp

```

➔ docker git:(main) x

```


Roadmap

Proof of Value - Q4 22

Developer Experience

- Local testing support (Java & Python) - Test containers via gradle plugin
- SpecMesh gradle plugin to use HCL Terraform & state mgmt

Domain model

- Domain modelling built into Spec.Id field

Build Pipeline

- Executions pushes the APISpec to apply orthogonal resource allocation

Governance

- Manual governance of 'public' resources through Credential creation (API Keys + ACLs)

Extensions

- Kafka support for Async API

Q1 - 2023

Extensions

- More Kafka extensions supported (incl provisioning support): Quotas, Storage & ACLs for Public, Private, Protected resources (basic governance)

Data Catalogue

- Initial integration to existing products to support discovery and basic governance

Storage

- Storage resource mapping

Governance

- Delegation to a Governance API where protected topics are specified and requested

Q2 - 2023

Data As A Product

- Data Product Open Registry; simple Data Catalogue with Discovery and tagging (Spec resources contain tags - show example)

Observability

- Topology view (derived from specs showing product/consume data flows and ownership visualisation)

Bolt-ons

- Multi-region
- Cross resource
- SMT
- DQ integration + other simple additions

Beyond

Community driven

Help us shape the roadmap by suggesting and contributing to features!

Thank you



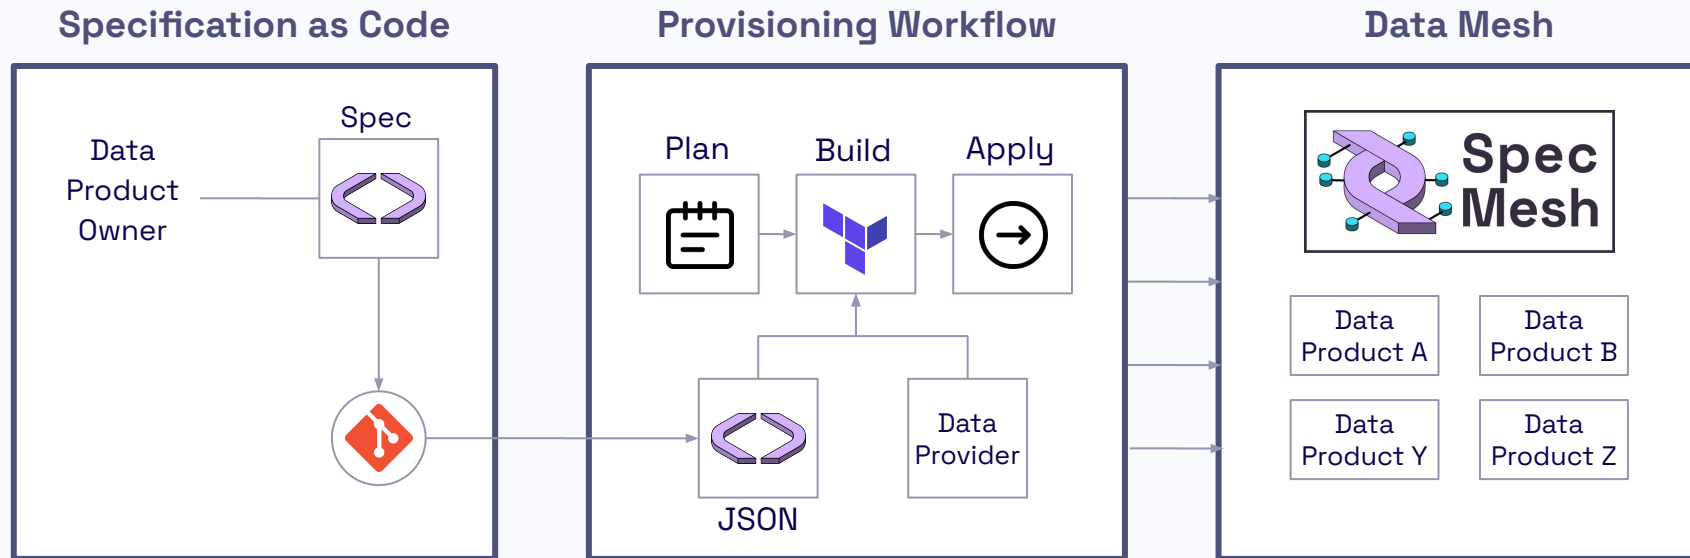
Get involved!
Scan & thumbs ups
on the Github issue

github.com/specmesh

@osodevops

@avery_neil

What do you get out of the box?



Roadmap - **THIS HAS BEEN MAPPED TO A SLIDE**

Proof of Value (0.1) (current status of Spec Mesh OS):

- Developer support for local testing (Java and Python) - i.e. TestContainers resources are provisioned using gradle plugin. The Spec Mesh Gradle plugin uses Terraform (making it very open to using existing operators, scripts, i.e. CDK)
- Domain modelling is built into the Spec.Id field - and upon provisioning via the Gradle/Terraform plugin applies correct domain knowledge to resources (see example)
- Build pipeline execution pushes the APISpec to apply orthogonal resource allocation (state mgmt via terraform) to provision pipe-line environment resources (build pipe example showing DEV, UAT and PROD)
- Manual governance of 'public' resources through Credential creation (API Keys + ACLs)
- Kafka extensions for Async API spec is being built

Roadmap after PoV (local testing and build pipeline)


1. More Kafka extensions supported (incl provisioning support): Quotas, Storage & ACLs for Public, Private, Protected resources (basic governance)
2. Storage resource mapping
3. Initial Catalogue integration to existing products to support discovery and basic governance
4. Delegation to a Governance API where protected topics are specified and requested
5. Data Product Open Registry: simple Data Catalogue with Discovery and tagging (Spec resources contain tags - show example)
6. Topology view (derived from specs showing product/consume data flows and ownership visualisation)
7. Bolt ons: multi-region, cross resource, SMT, DQ integration + other simple additions

Data Model

/ london / borough / venue / event



/ london / hammersmith / olympia / bigdataidn



Map domain model into spec

- Now that specification in my branch / git repo
- I need to build a test around this / init test my data product (include the topic/channel name)
- The unit test in this example is Big Data LDN (People coming to the event)
 - Test container example

Established patterns of Data Mesh and central nervous system

Data as a Product

Discoverability	Spec Mesh provides	
Addressability		
Trustworthy		
Self-describing semantics		
Interoperable / standards		
Security		


Build pipeline

- Leverage @IntegrationTest - gets run first and calls on the SpecMeshProvisioner (Java-> terraform) to parse the spec, build topic names and execute against target cluster location

DEMO MAYBE

- ALL IN JAVA
 - Drive TF from Java
 - Parse the spec from github location
 - Provision kafka topics from what is in the specification
-
- LETS SEE THIS BAD BOY IN ACTION

stuff



The Mission

To make Data Mesh simple, Open Source and available to all - without lockin, without complex tooling by using an approach centered around 'specifications', existing tools and baking in a 'domain' model.

<insert triangular image here>

Why we should build a mesh using spec's

Benefits

Organisation

- + Real life insights into technology problems
- + Uncover new opportunities
- + Builds trust and good working relationships
- + Brand awareness

Domain

- + Insights into emerging technologies
- + Access to technology experts
- + Technical focus
- + Professional development for staff
- + Builds team morale
- + Free lunch!

Set it up

Line up with client



Get it in the diary



Agree the theme and key questions



Arrange lunch




Run the session

Build using ASyncAPI specifications

Outline

Talk Structure (30 mins) 15 slides

- we believe data mesh should not be a large complex undertaken lasting multi years costing millions
- Target audience developers
- we believe data mesh should be useable and free for everyone. ITs not about complex custom built tooling
- foundations are opinionated on a specification way
- Problem statement (Why? Most Kafka environments we work in are shit, its a graveyard of topics - no uniform way of provisioning we cannot scale)
- Why do we think data mesh has anything todo with this. These are the principles: how we are using these principles to build Kafka at scale.
- Data ownership by domain
- your domain model is defined by people in your organisation, they should be able to look at it and know where they sit and where to discovery events they need.
- data governance / data as a product
- why we think the async api matter, data as a product is represented by the api / specification for data > that allows us to model the domain and thefore the
- Modelling London, here are the specs for parts of this
- data model Fulham
- you came in from as an event today
- Here is the spec for big data ldn : this is what it looks like when we use it > we are using this to test our code
- local test > we need to get away from the graveyard >
- Domain model for London (
 - private events buying a kebab
 - public events: me moving from Fulham to central

Data mesh, self service ubiquitous access to data

Goal: To make Data Mesh simple, Open Source and available to all - without lockin, without complex tooling by using an approach centered around 'specifications', existing tools and baking in a 'domain' model. (see image)

Scope - running OTS infrastructure like connectors and all things, webservices, microservices or processors.... They are the problem of the client, not the Mesh itself

Problem statement (what):

- Everyone is touting data mesh and it's getting confusing, many bespoke, many commercial solutions... it should be accessible & simple, without lockin
- We see Data Mesh as a natural, formalisation of building the central nervous system (everything via streams of topic data)
- Been building CNS for many years across some of the largest companies in the world
- CNS focuses on Event-Streams, but during implementation includes most of the guiding principles of Datamesh (ownership by domain, as a product, selfservice, governance)
- Avoid vendor lockin (doesn't need to be all encompassing and conflict with existing data infra)
- Most data-mesh solutions don't provide SDLC/development tools for test and infrastructure

Why:

- A specification based approach provides uniformity (part of the solution)
- Fixes Kafka resource mess/grave-yard
- Testing support (developer SDLC unlike most tooling)
- Repeatability -> provides guard-rails, reuse, reliable, simple and consistent approach for all
- Not just kafka, but eventually Kinesis, G PubSub, EventHubs, S3, from OnPrem to Cloud, multiple envs, clusters etc

How (Can we do this with minimal effort by leveraging existing tools and principles):

- Simple: Data Ubiquity should have 2 forms, Streams of Events (Kafka stream or others) + Storage (think S3) (because not everything fits in a stream). This simplifies everything - don't boil the ocean
- Simple: Specification based using AsyncAPI specs
- Simple and existing: Data as an API in the CNS maps into AsyncAPI Spec + Schemas
- Existing: by using the AsyncAPI Spec 'id' hierarchy we can express domain ownership (example)

Data Mesh Principle mapping

- Domain ownership is captured using the ID of the spec (example)
- Data as a product is reflected by the Spec itself, the spec is used to provision resources (example) (more later)
- Data available everywhere (discoverable) - built using existing tools such as AWS Glue Data Catalogue, Collibra and others
- Data governed everywhere - using streams and storage we build a integration to the data catalogue, and integrate access requests through the use of private, protected, public topics & storage while automating restrictive controls (i.e. ACLs in Kafka, principle mapping etc).

Worked example of the city of london → BDL

PoV (current status of Spec Mesh OS):

1. Developer support for local testing (Java and Python) - i.e. TestContainers resources are provisioned using gradle plugin. The Spec Mesh Gradle plugin uses Terraform (making it very open to using existing operators, scripts, i.e. CDK)
2. Domain modelling is built into the Spec.Id field - and upon provisioning via the Gradle/Terraform plugin applies correct domain knowledge to resources (see example)
3. Build pipeline execution pushes the APISpec to apply orthogonal resource allocation (state mgmt via terraform) to provision pipe-line environment resources (build pipe example showing DEV, UAT and PROD)
4. Manual governance of 'public' resources through Credential creation (API Keys + ACLs)
5. Kafka extensions for Async API spec is being built

Challenges! No one size fits all

There is not one tool to solve all the problems, a framework or suite of tools is needed.

- + Every company is different, there is **no prescribed solution**
- + **Scaling** your data products as more is captured
- + How to separate **signals** from **noise**?
- + **Complexity** of data is growing exponentially
- + Maintaining **data quality** whilst keeping consistent



Roadmap after PoV (local testing and build pipeline)

1. More Kafka extensions supported (incl provisioning support): Quotas, Storage & ACLs for Public, Private, Protected resources (basic governance)
2. Initial Catalogue integration to existing products to support discovery and basic governance
3. Delegation to a Governance API where protected topics are specified and requested
4. Data Product Open Registry: simple Data Catalogue with Discovery and tagging (Spec resources contain tags - show example)
5. Topology view (derived from specs showing product/consume data flows and ownership visualisation)
6. Bolt ons: multi-region, cross resource, SMT, DQ integration + other simple additions

Sion Notes

- People can work more collaboratively using a decentralised standardised approach
 - different levels of expertise of using data in each domain
- how do you mentor and empower each domain on contributing to the mesh
 - set some enterprise level standards so teams can educate themselves on things like the team structure etc.
- Decentralised doesn't mean free for all
- Federated governance is at the domain. Example the finance team needs to know what privacy
- NOT building a data silos, its more about people and processes not the technology
 - interoperability and easy to navigate
- Not every domain knows how to build and manage a scalable API.

Pillars of Modern Data Architecture

- Scalable on demand
- Purpose-build data services
- Seamless data movement
- unified governance
- Performant and cost-effective

Sion notes

Why Spec Mesh

- Encourage data-driven agility
- Support domain-local governance through lightweight specifications
- Isolate data resources with clear contracts
- Consider data-as-a-specification which can exist in any system

Introduction

Why title

+ aaa

Design and Build
title

+ aaa

Scale title

+ aaa